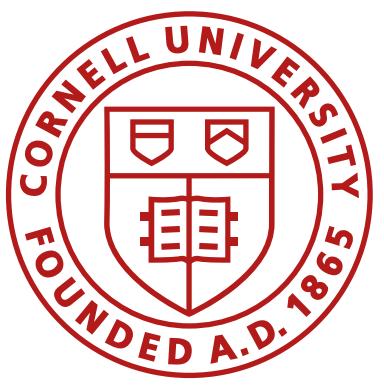


Maps and Graphs

Fast Robots, ECE4160/5160, MAE 4190/5190

E. Farrell Helbling, 3/13/25

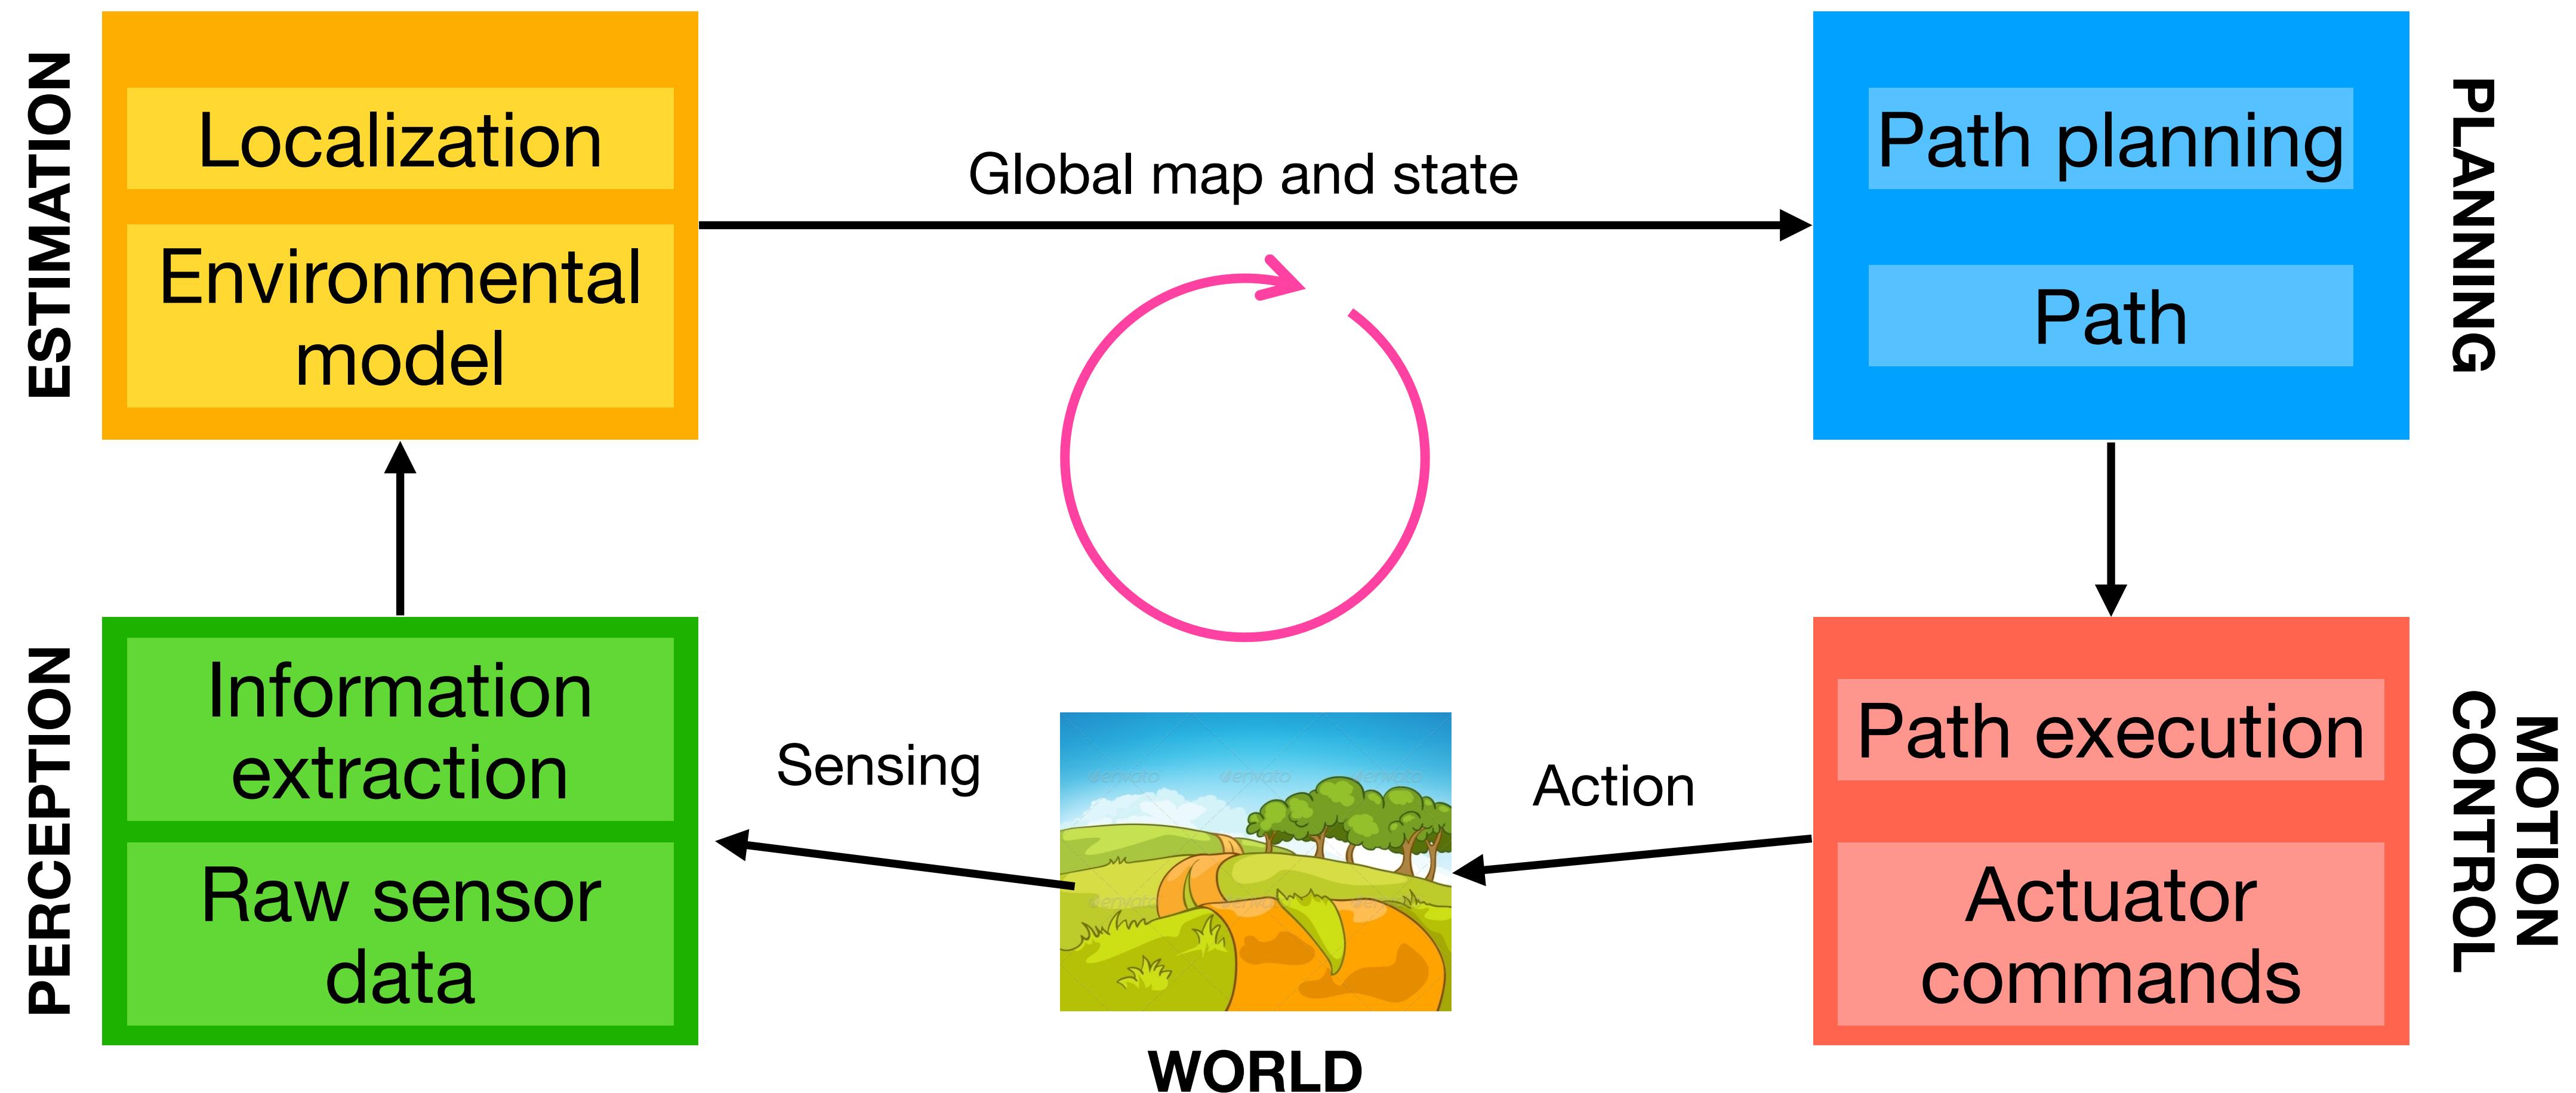


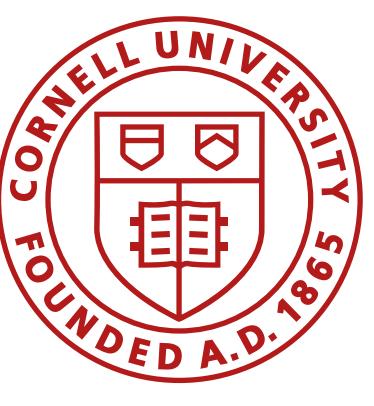
Class Action Items

- **Midsemester evaluations:** please fill them out! One point in your final grade is based on completing the midterm course evaluations.
Open until March 16th!
 - They remain anonymous, I get netIDs from MTEI
 - Lab 1 and Lab 2 regrade requests will close on **Sunday midnight**.
 - Lab 3 regrade requests will close on **Thursday midnight**.
 - Lab 6 how is it going?
 - Please start Lab 7 early!
 - If you get to the stunts lab next week, you definitely will not have to think about it over Spring Break!

Navigation

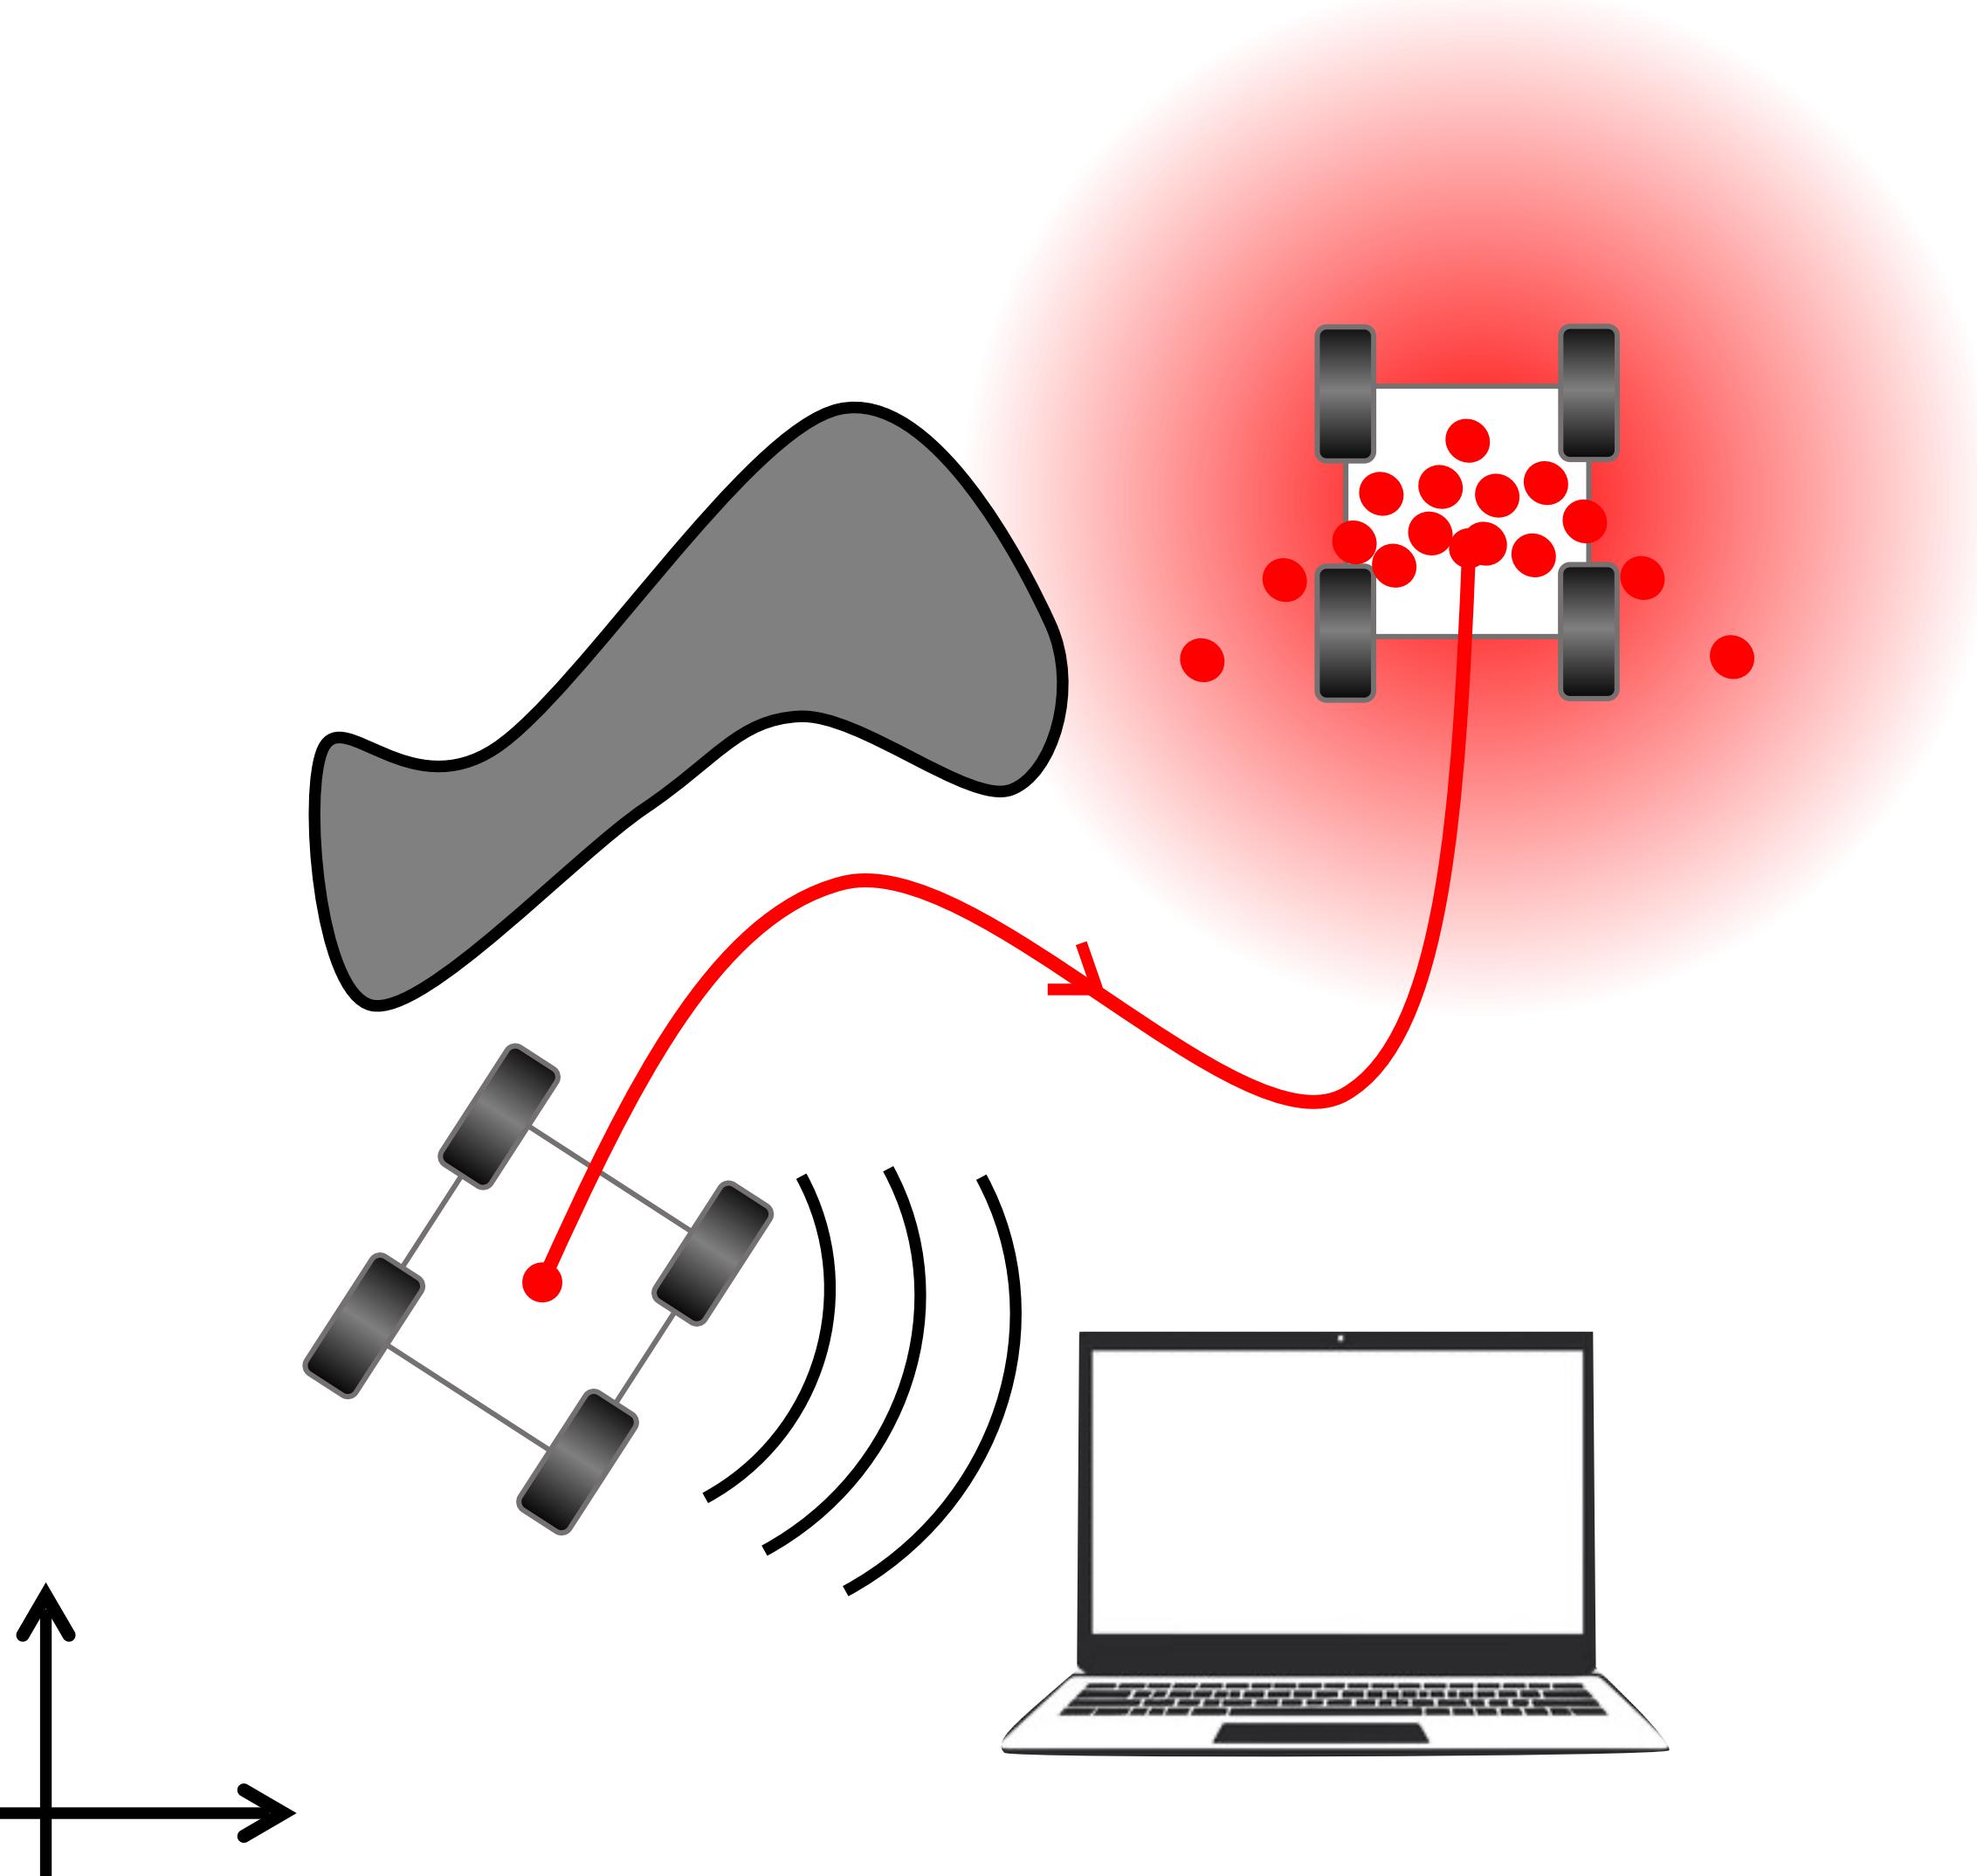
- Break the problem down: localization, map building, path planning

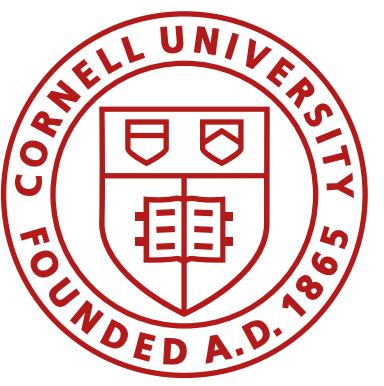




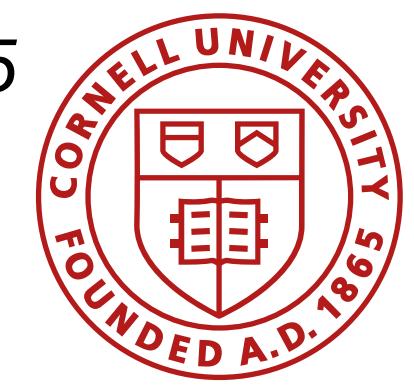
Navigation

- Local planners
- Global localization and planning
 - Map representations
 - Continuous
 - Discrete
 - Topological
 - Maps as graphs
 - Graph search algorithms
 - Breadth first search
 - Depth first search
 - Dijkstras
 - A*





Local Planners



Local path planning/ obstacle avoidance

- Use goal position, recent sensor readings, and relative position of robot to goal
 - Can be based on a local map
 - Often implemented as a separate task
 - Runs at a much faster rate than the global planner
- 3 examples:
 - Bug algorithms
 - Vector Field Histogram (VFH)
 - Dynamic Window Approach (DWA)

Wagner, ITS 2015

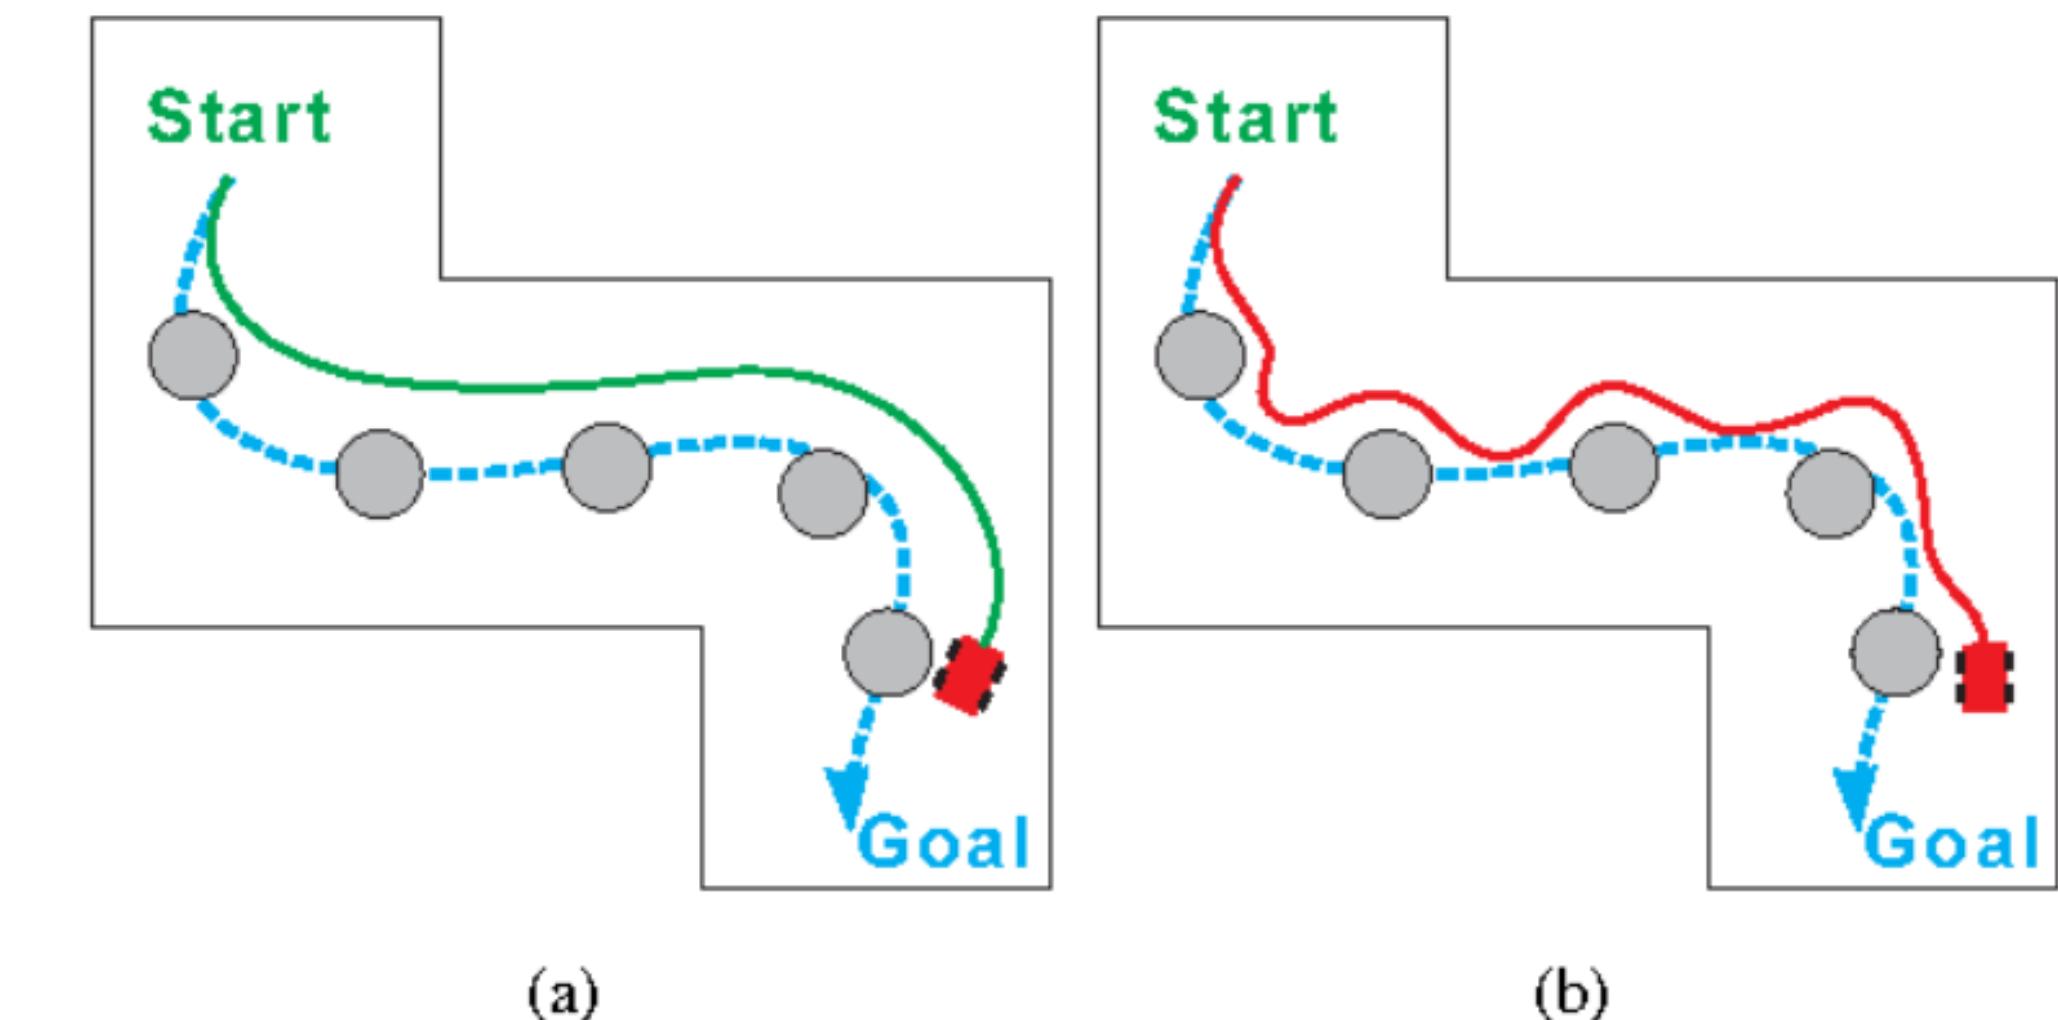
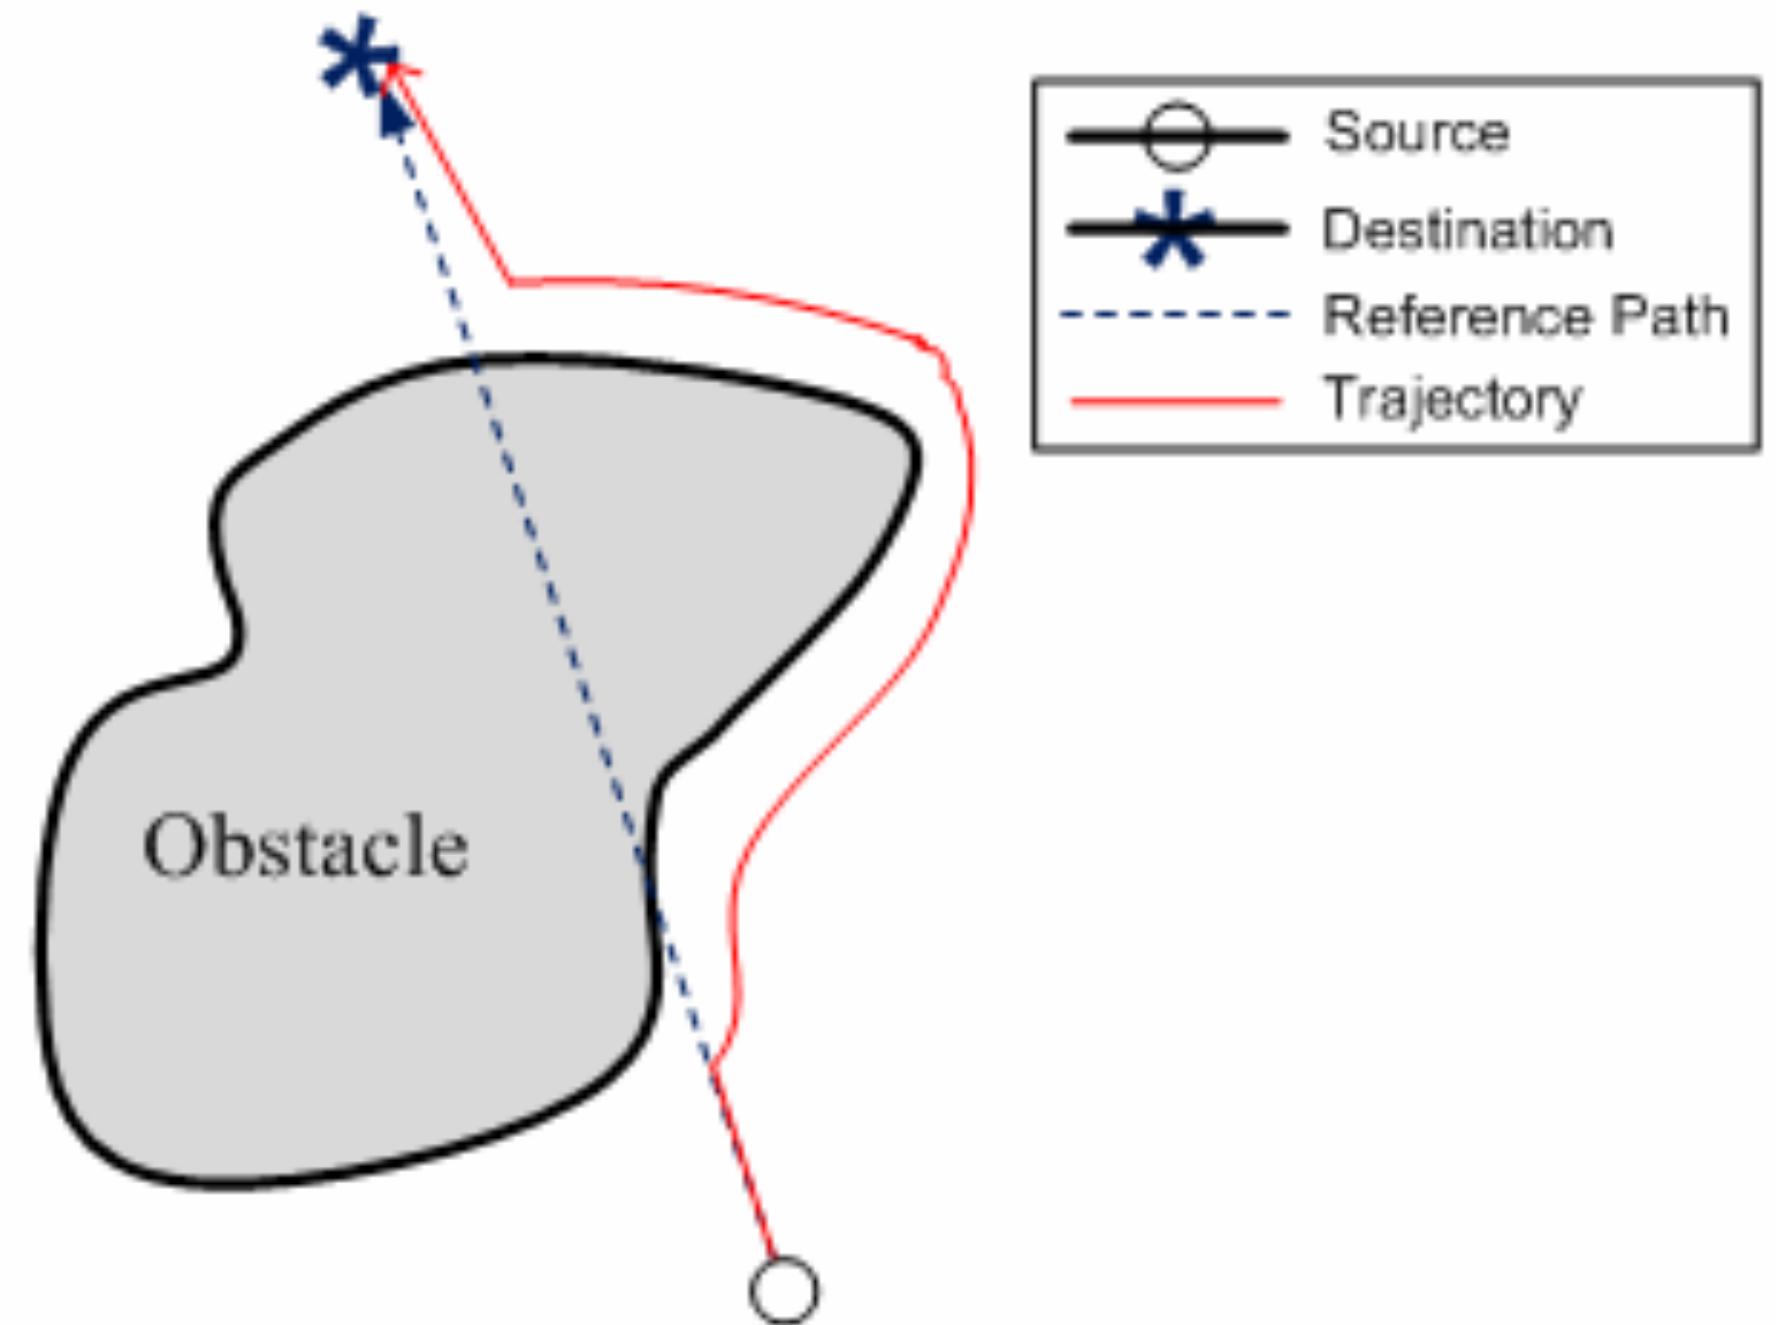


Fig. 1. Dashed blue spline is global path: a) Green spline is ideal local path; b) Red spline is actual local path

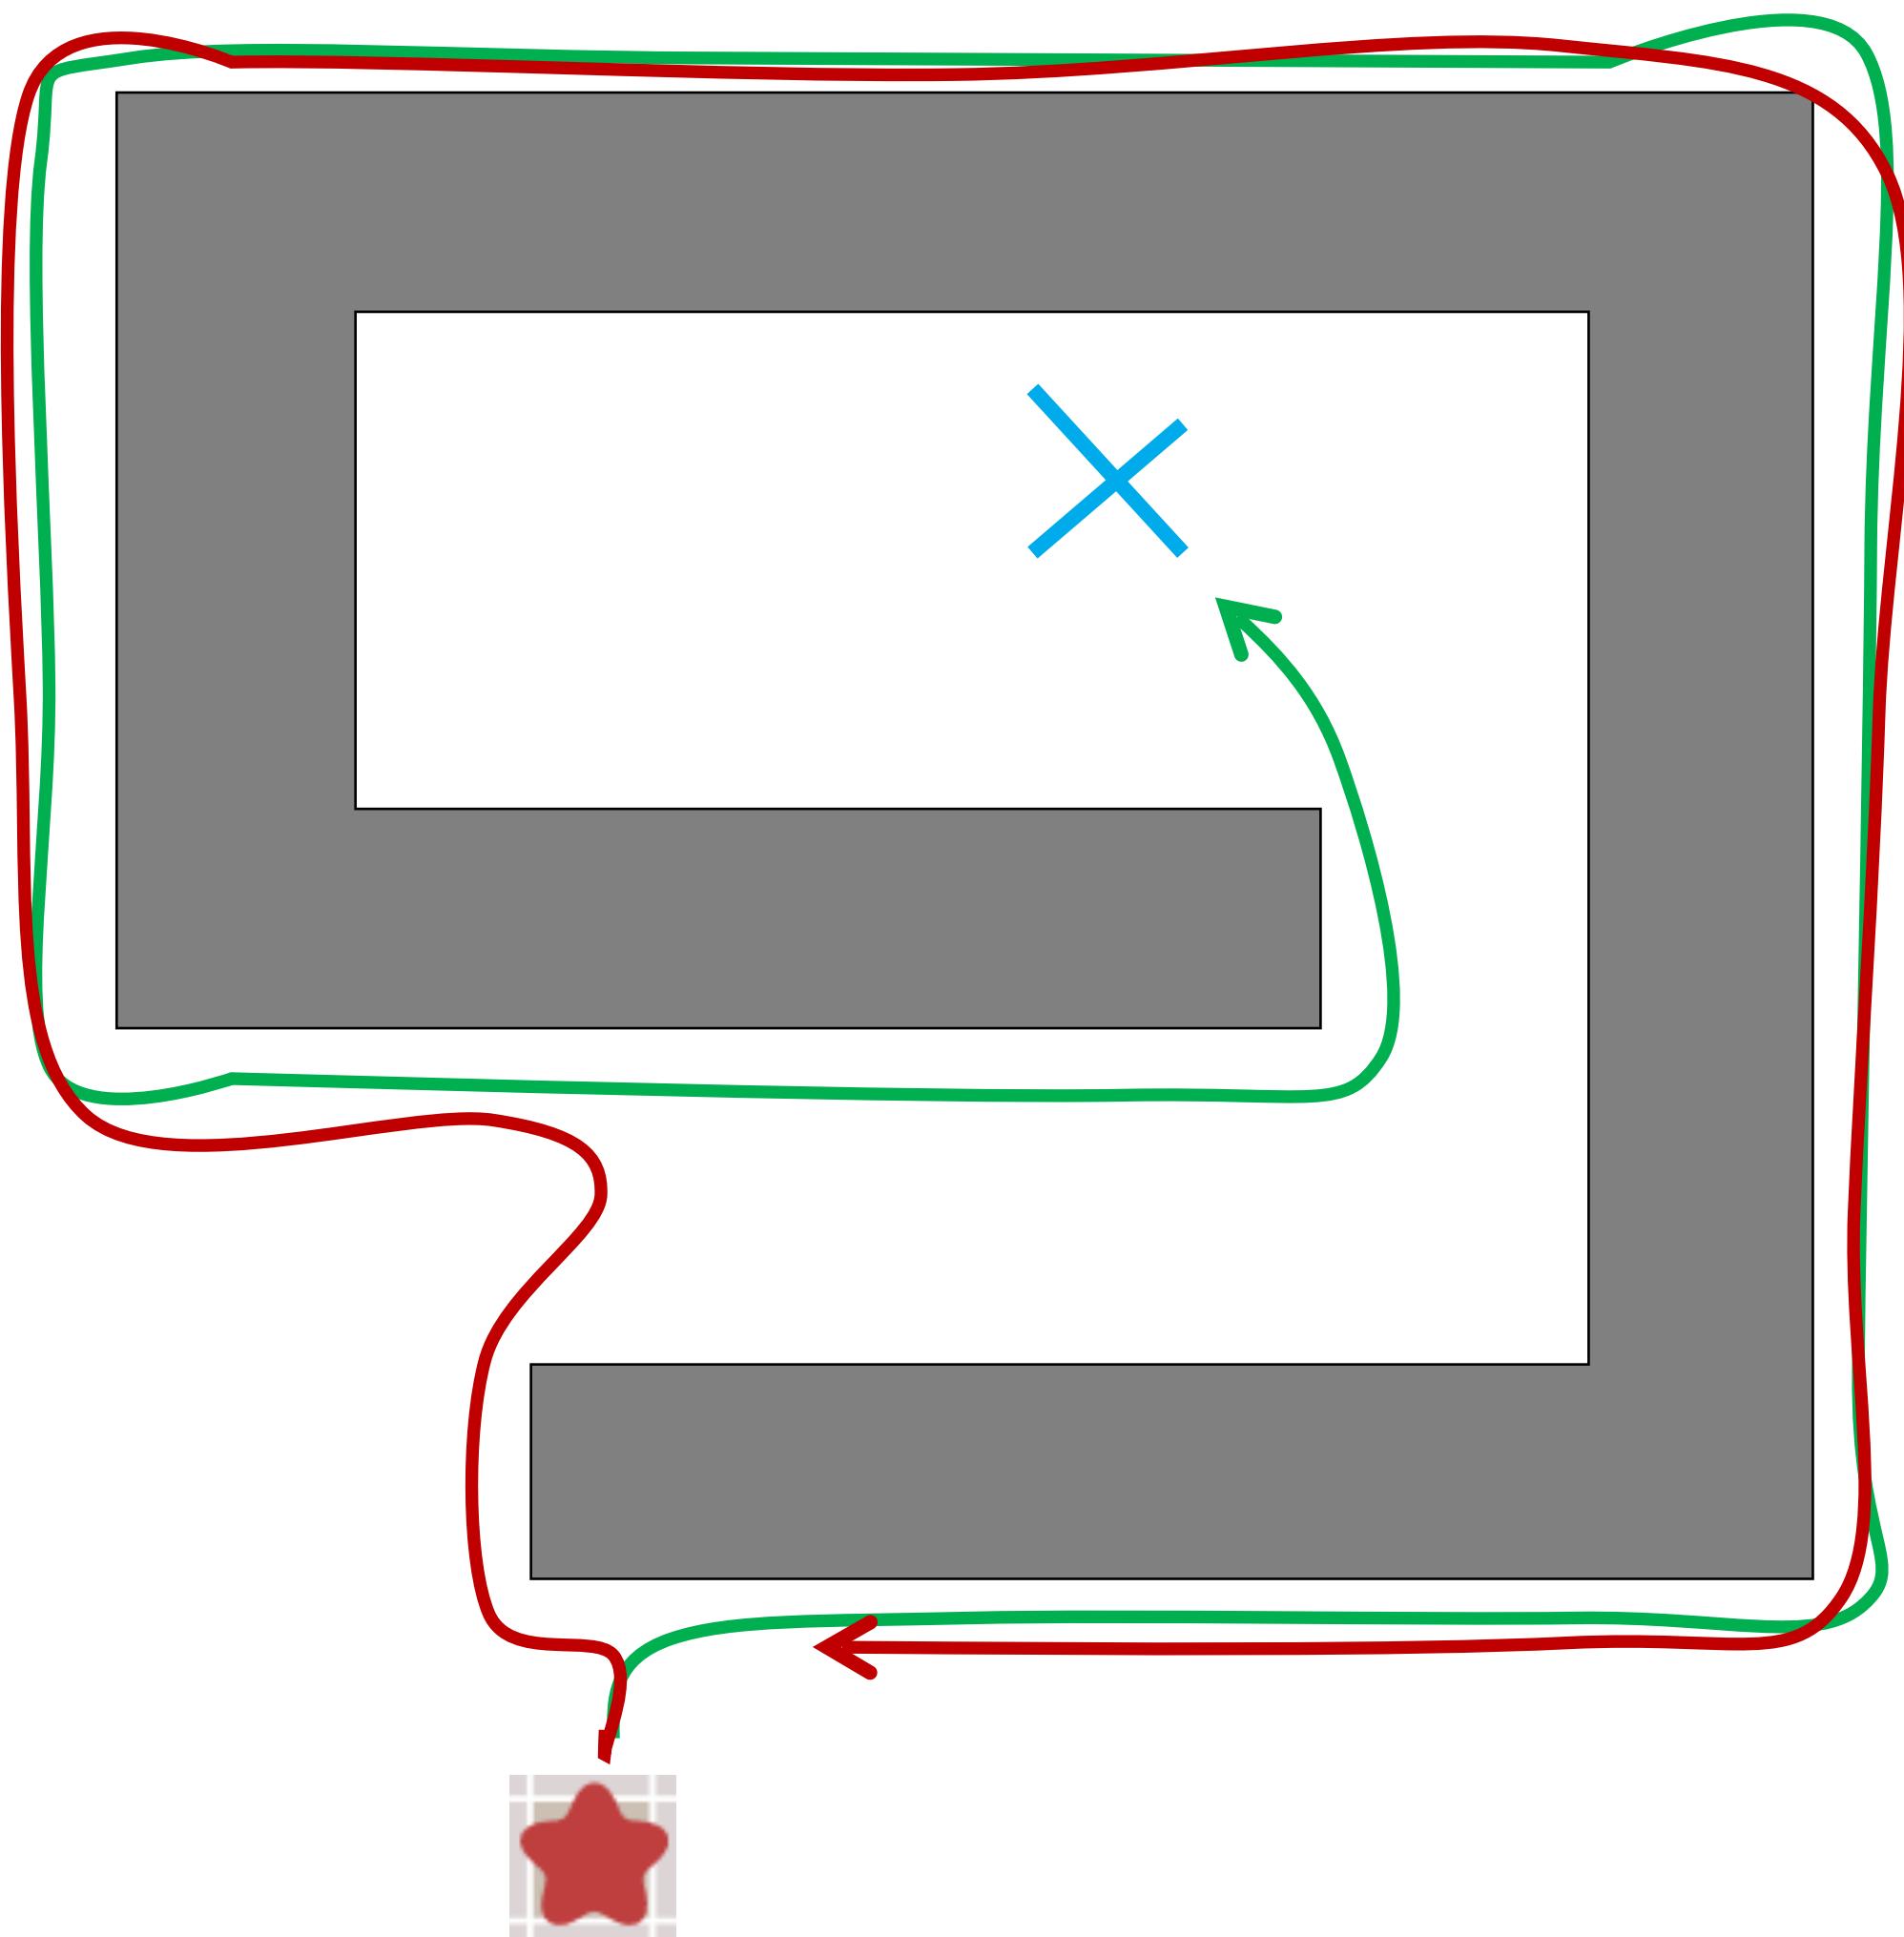
Bug algorithms

- Uses local knowledge and the direction and distance to the goal
- Basic idea
 - Follow the contour of obstacles until you see the goal
 - State 1: seek goal
 - State 2: follow wall
- Different Variants: Bug0, Bug1, Bug2
- Advantages
 - Super simple
 - No global map
 - Completeness
- Disadvantages
 - Suboptimal



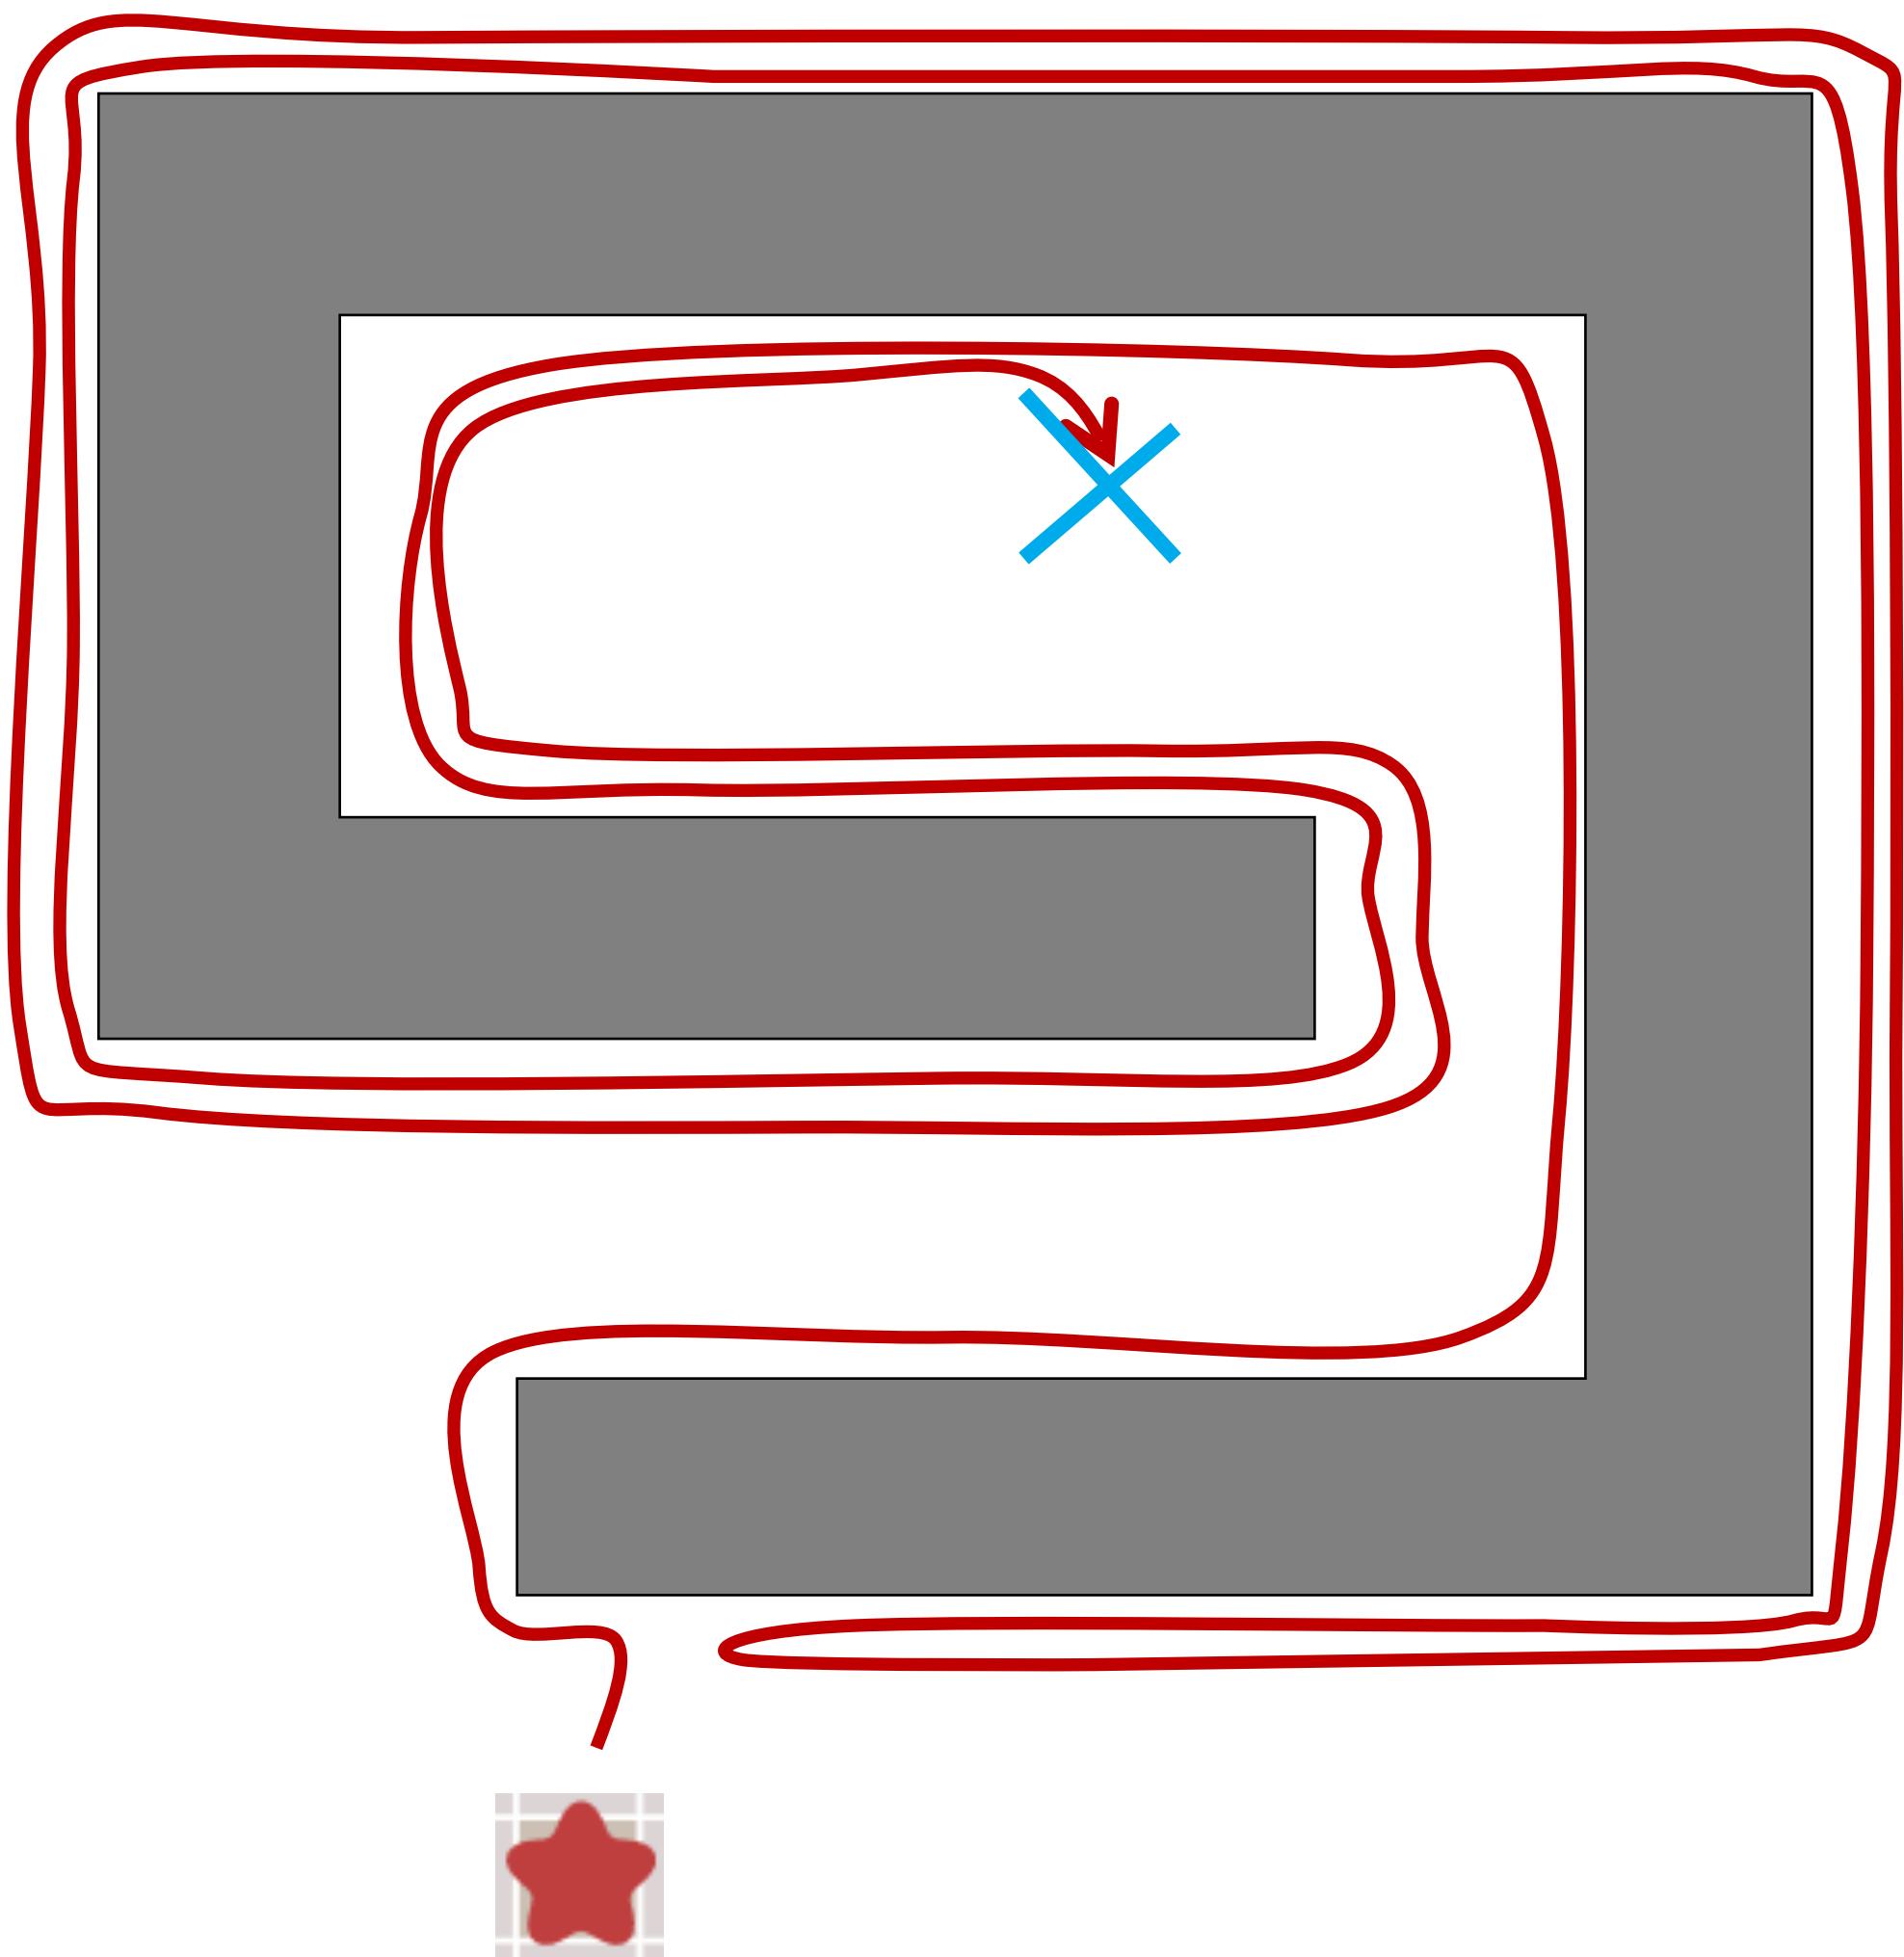
Bug 0

- Sensor Assumptions
 - Direction to the goal
 - Detect walls
- Algorithm
 - Go towards goal
 - Follow obstacles until you can go towards goal again
 - Loop



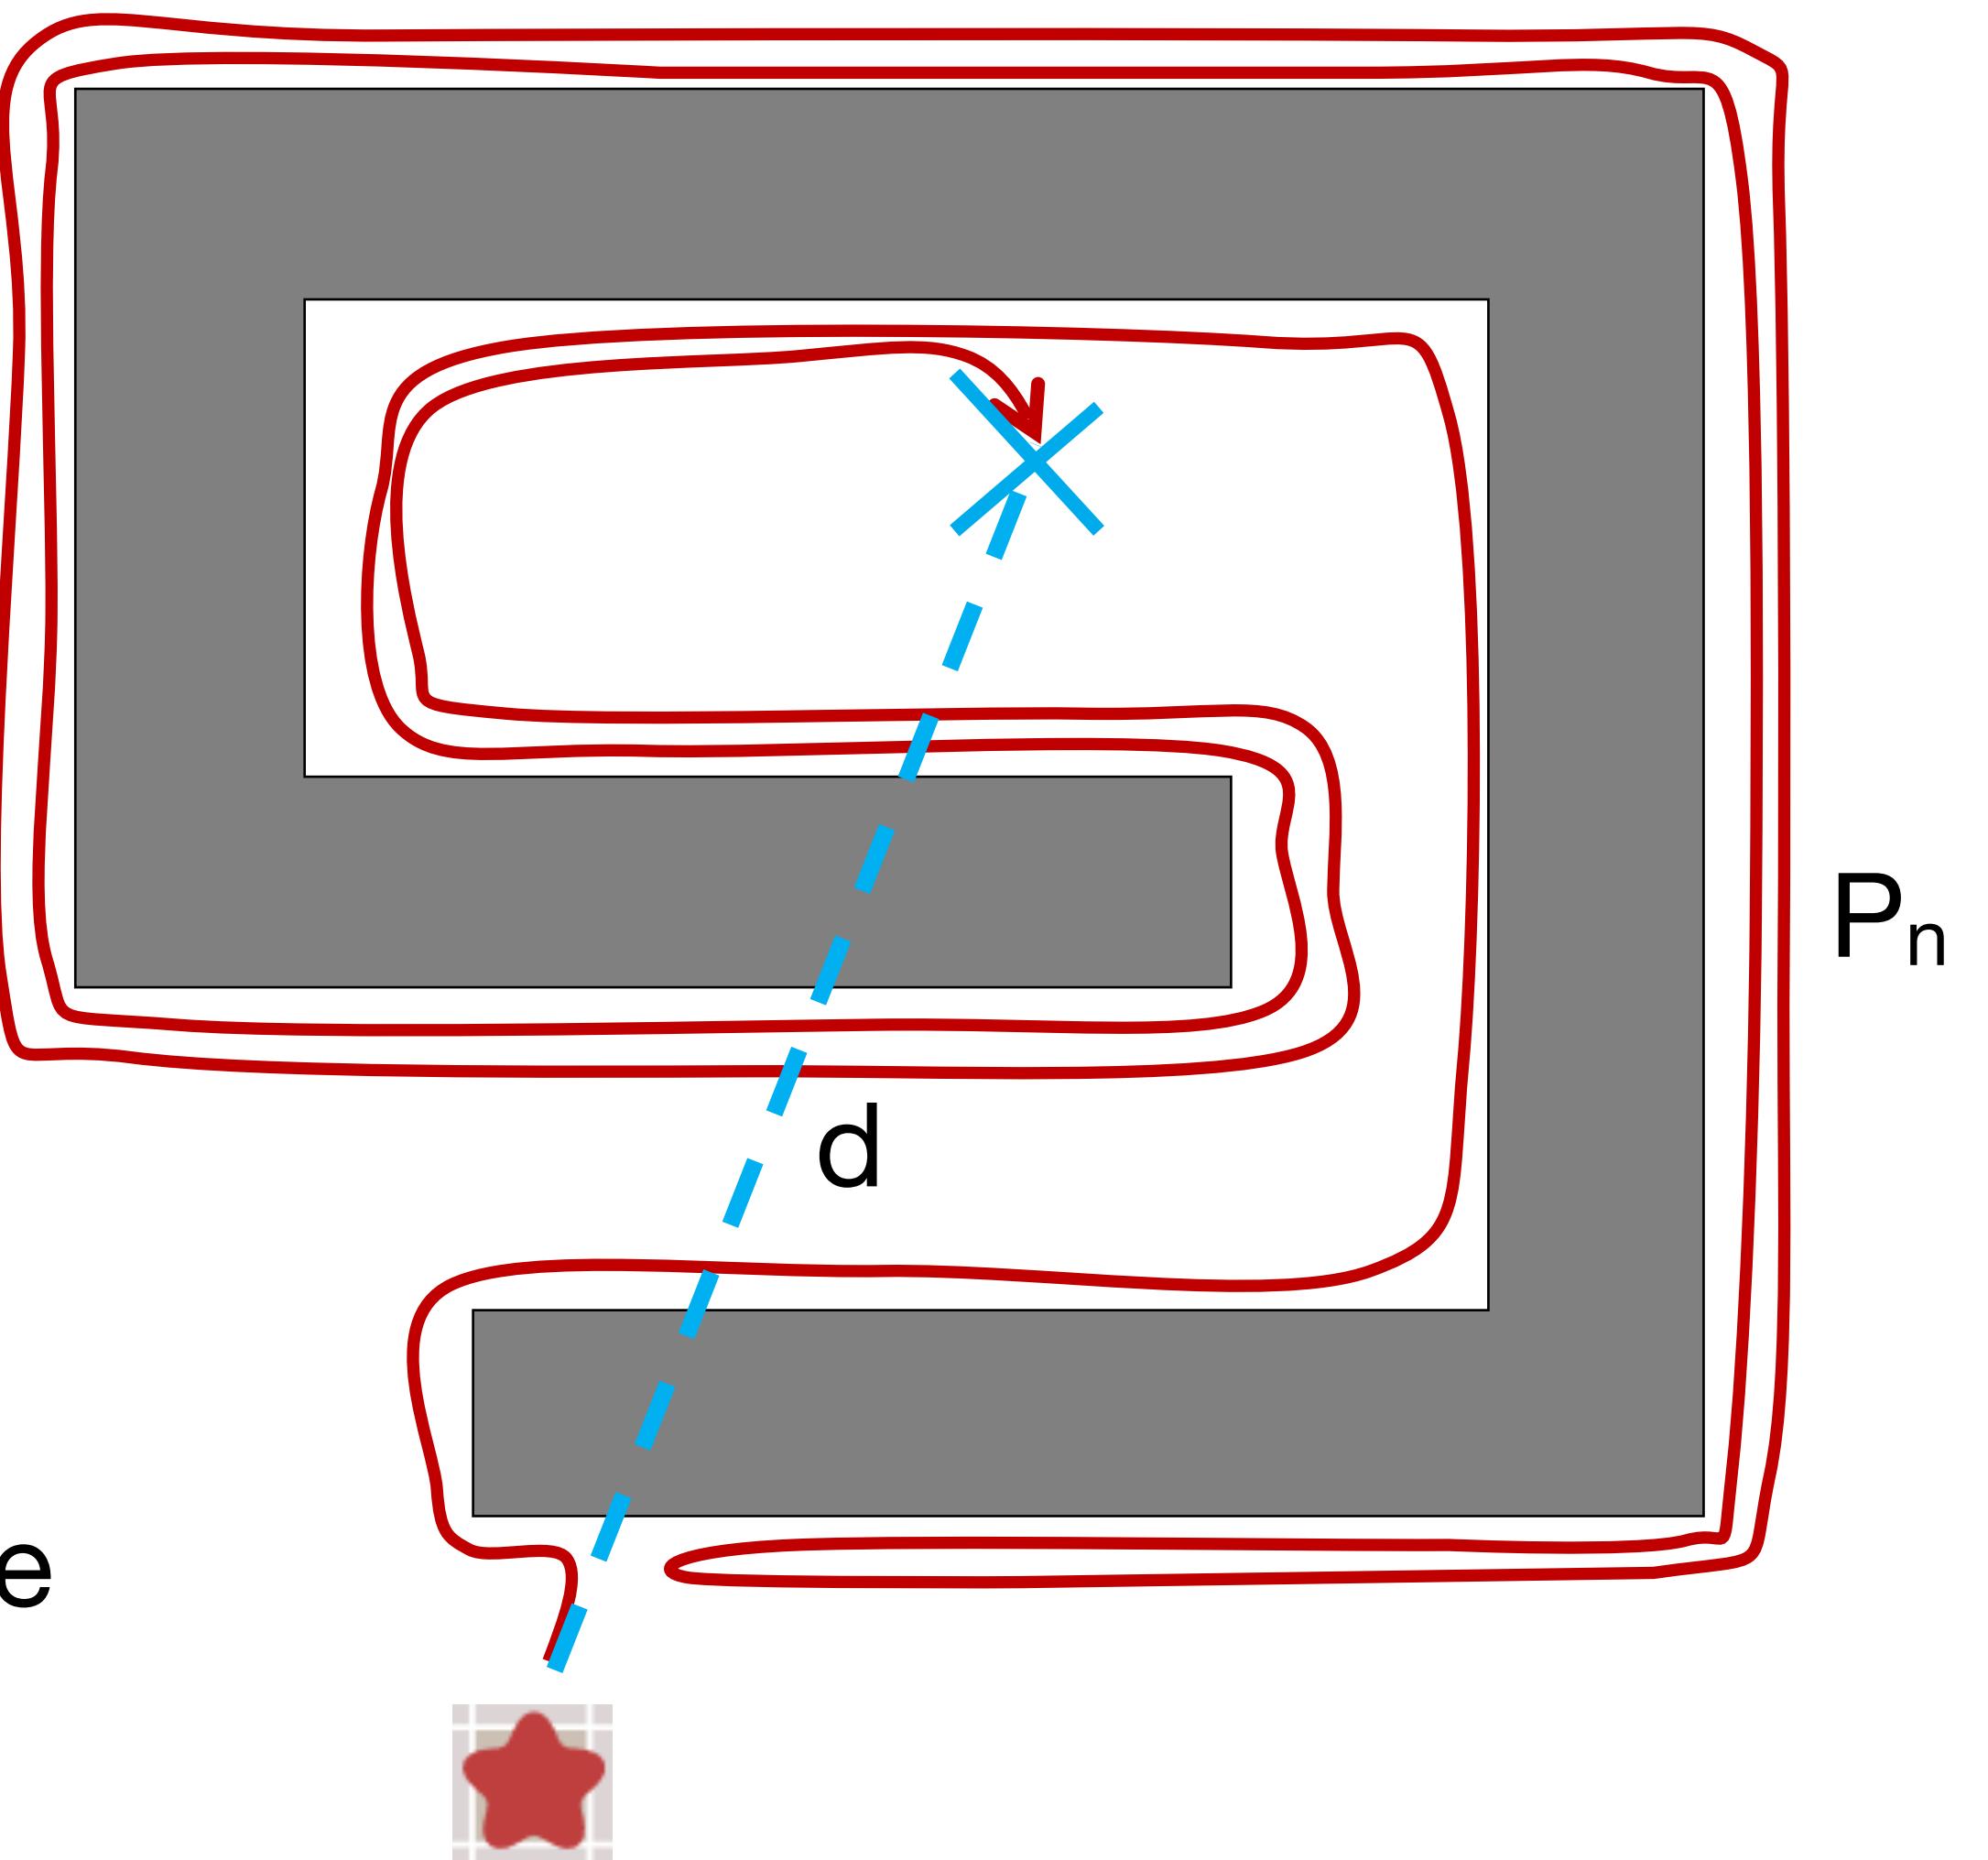
Bug 1

- Sensor Assumptions
 - Direction to the goal
 - Detect walls
 - Odometry
- Algorithm
 - Go towards goal
 - Follow obstacles and *remember how close you got to the goal*
 - Return to the closest point, loop



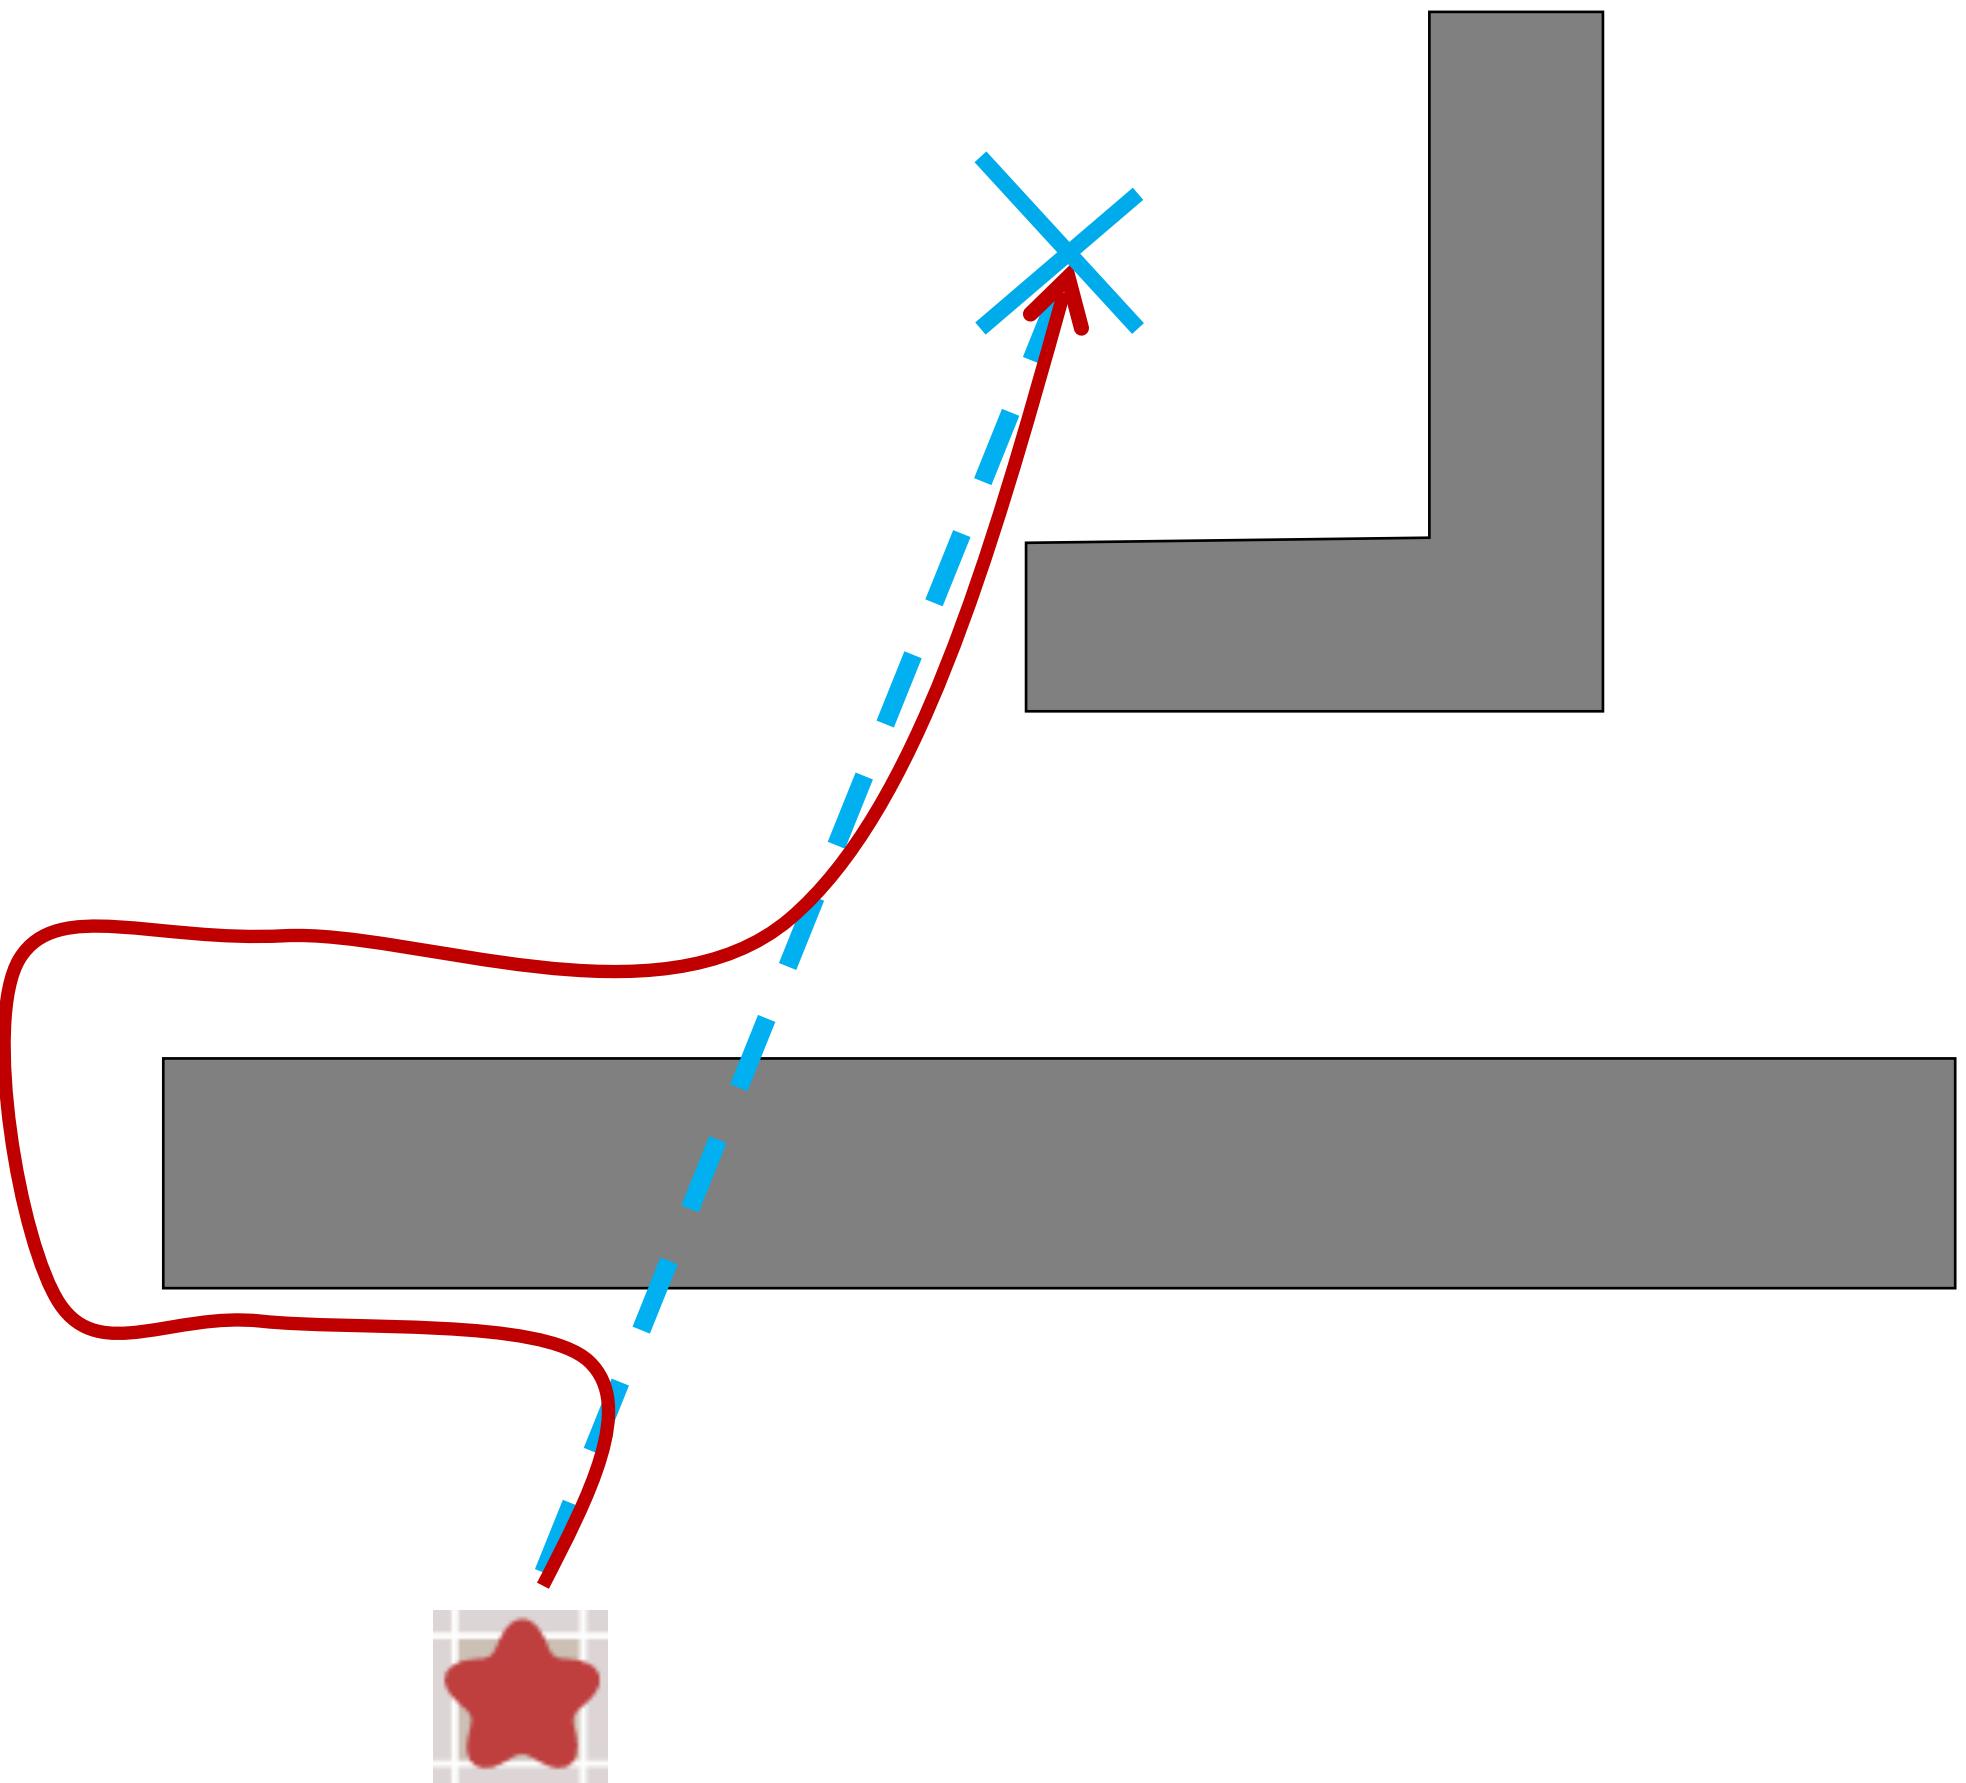
Bug 1 - formally

- Sensor Assumptions
 - Direction to the goal
 - Detect walls
 - Odometry
- Lower bound traversal? d
- Upper bound traversal? $d + 1.5\sum(P_n)$
- Pros?
 - If a path exists, it returns in finite time
 - It knows if none exist!



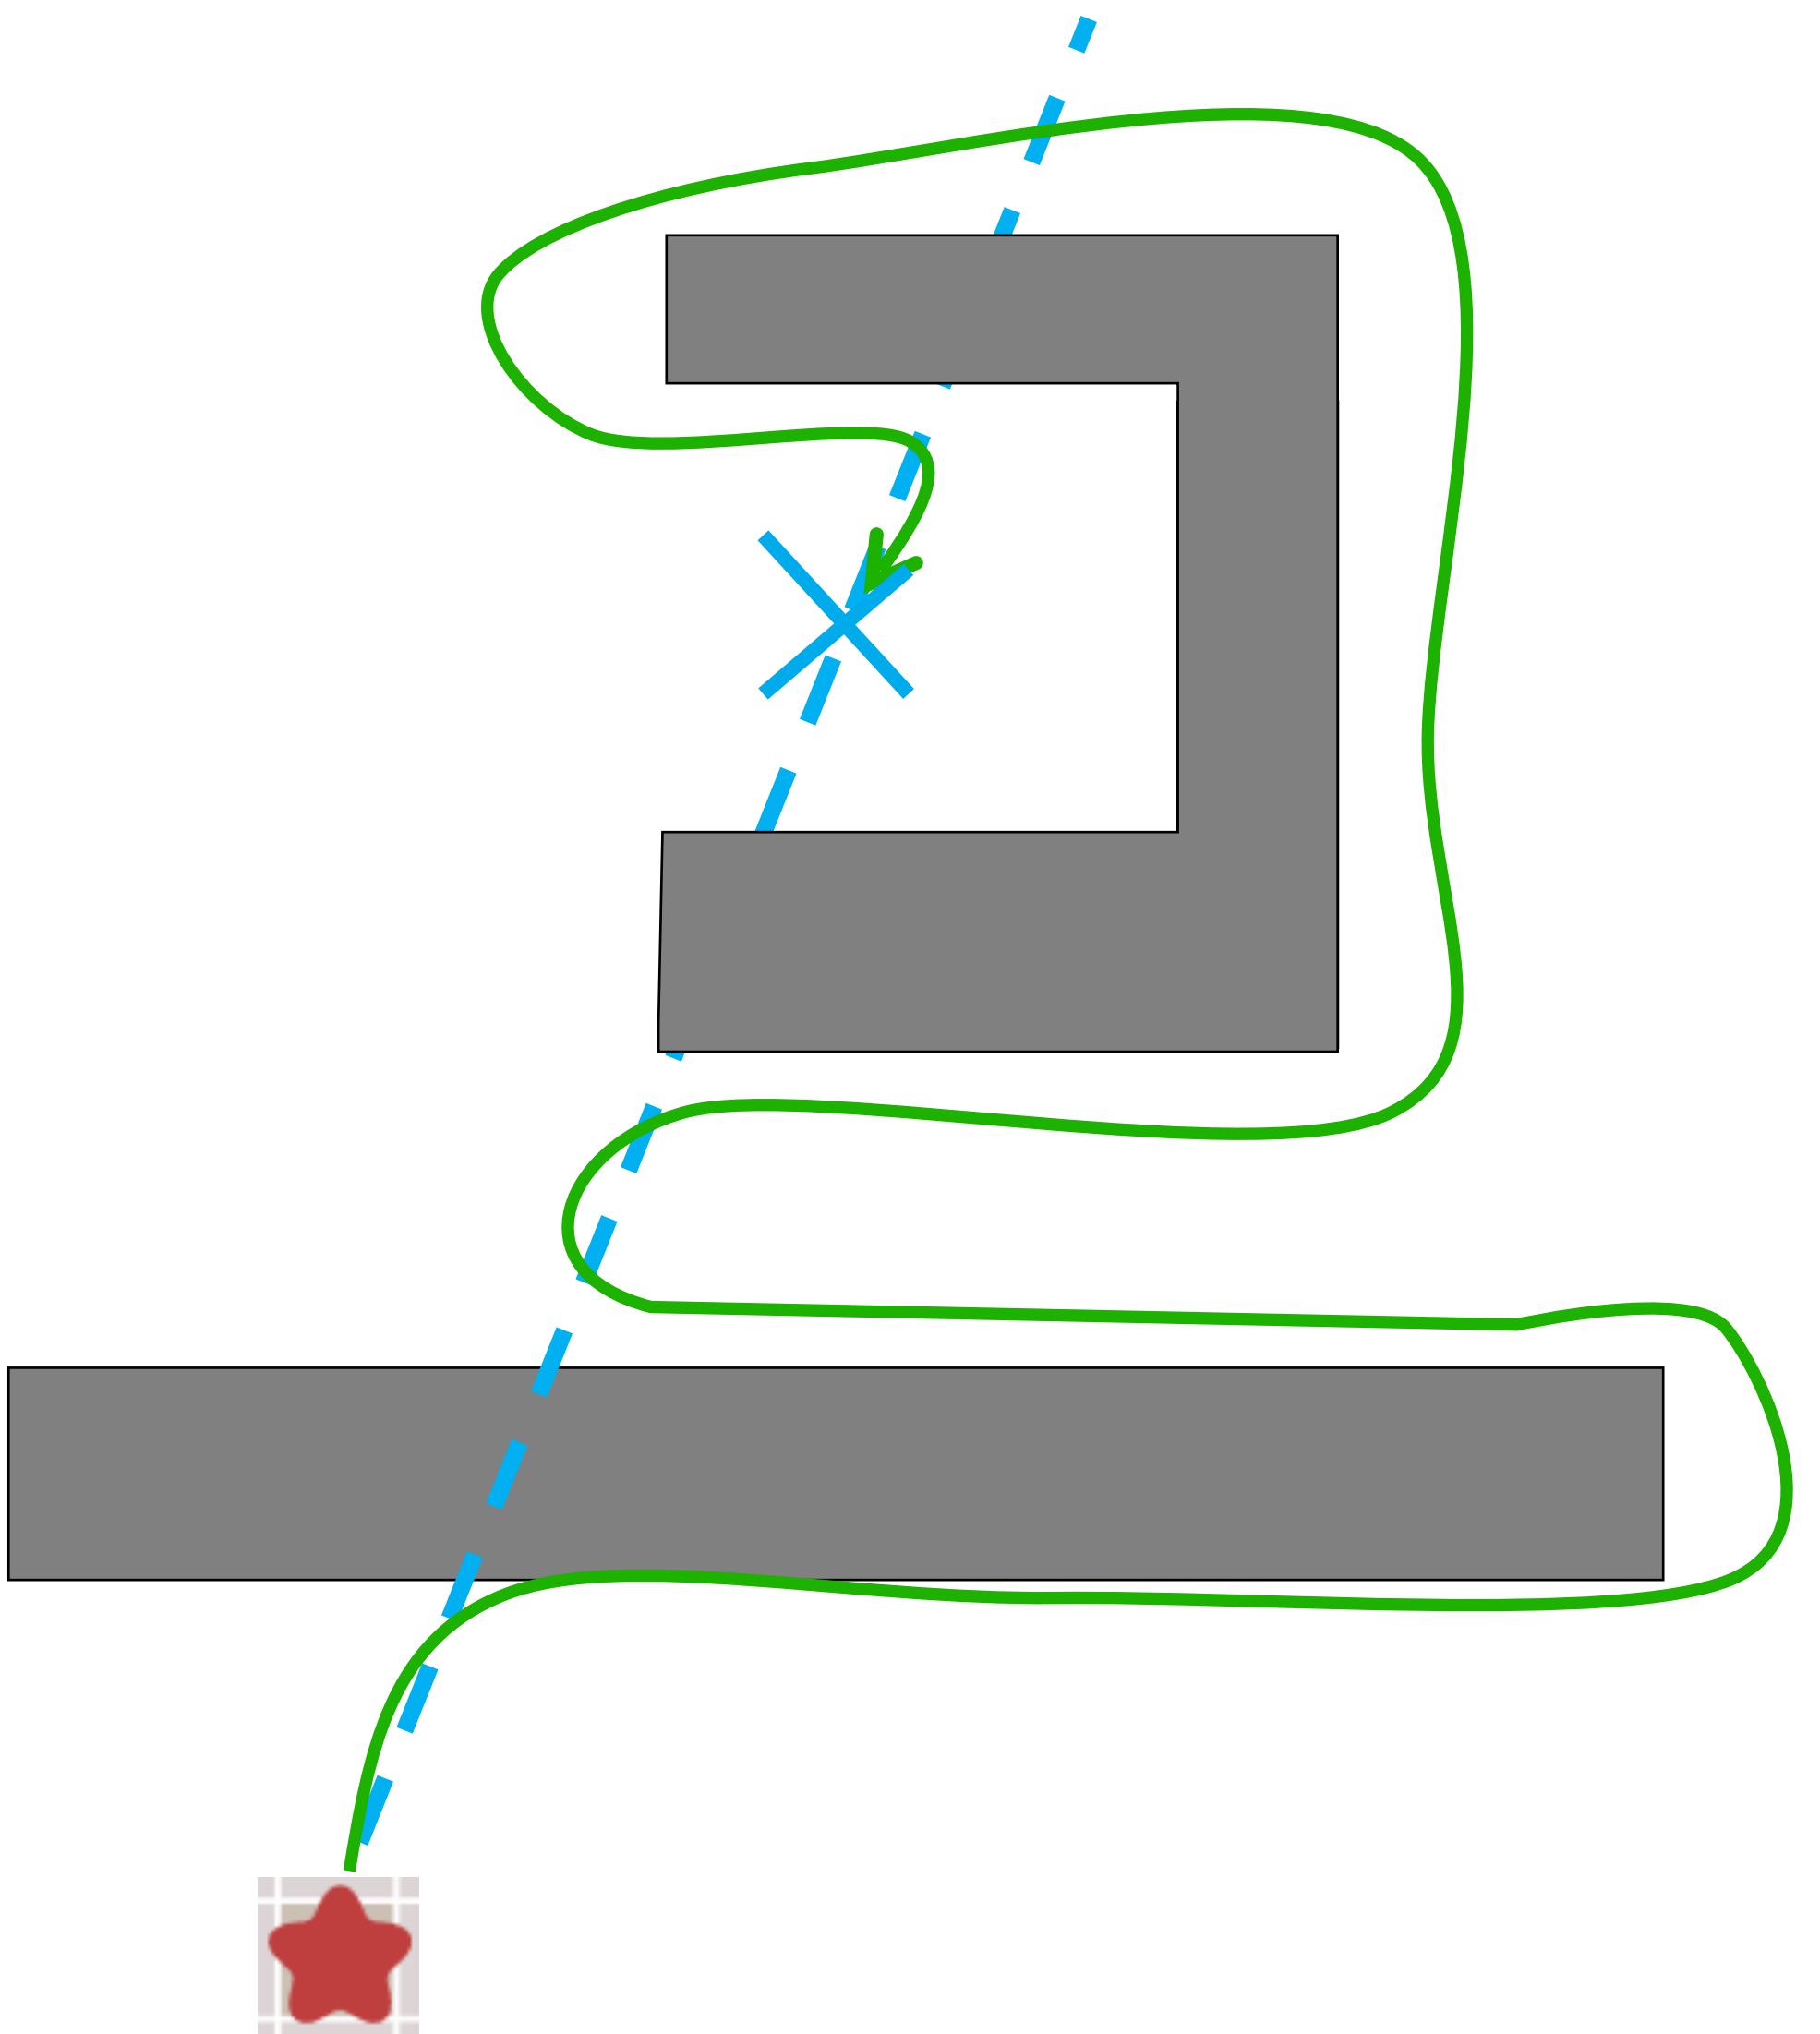
Bug 2

- Sensor Assumptions
 - Direction to the goal
 - Detect walls
 - Odometry
 - Original vector to the goal
- Algorithm
 - Go towards goal on the vector
 - Follow obstacles *until you are back on the vector (and closer to the goal)*
 - Loop



Bug 2

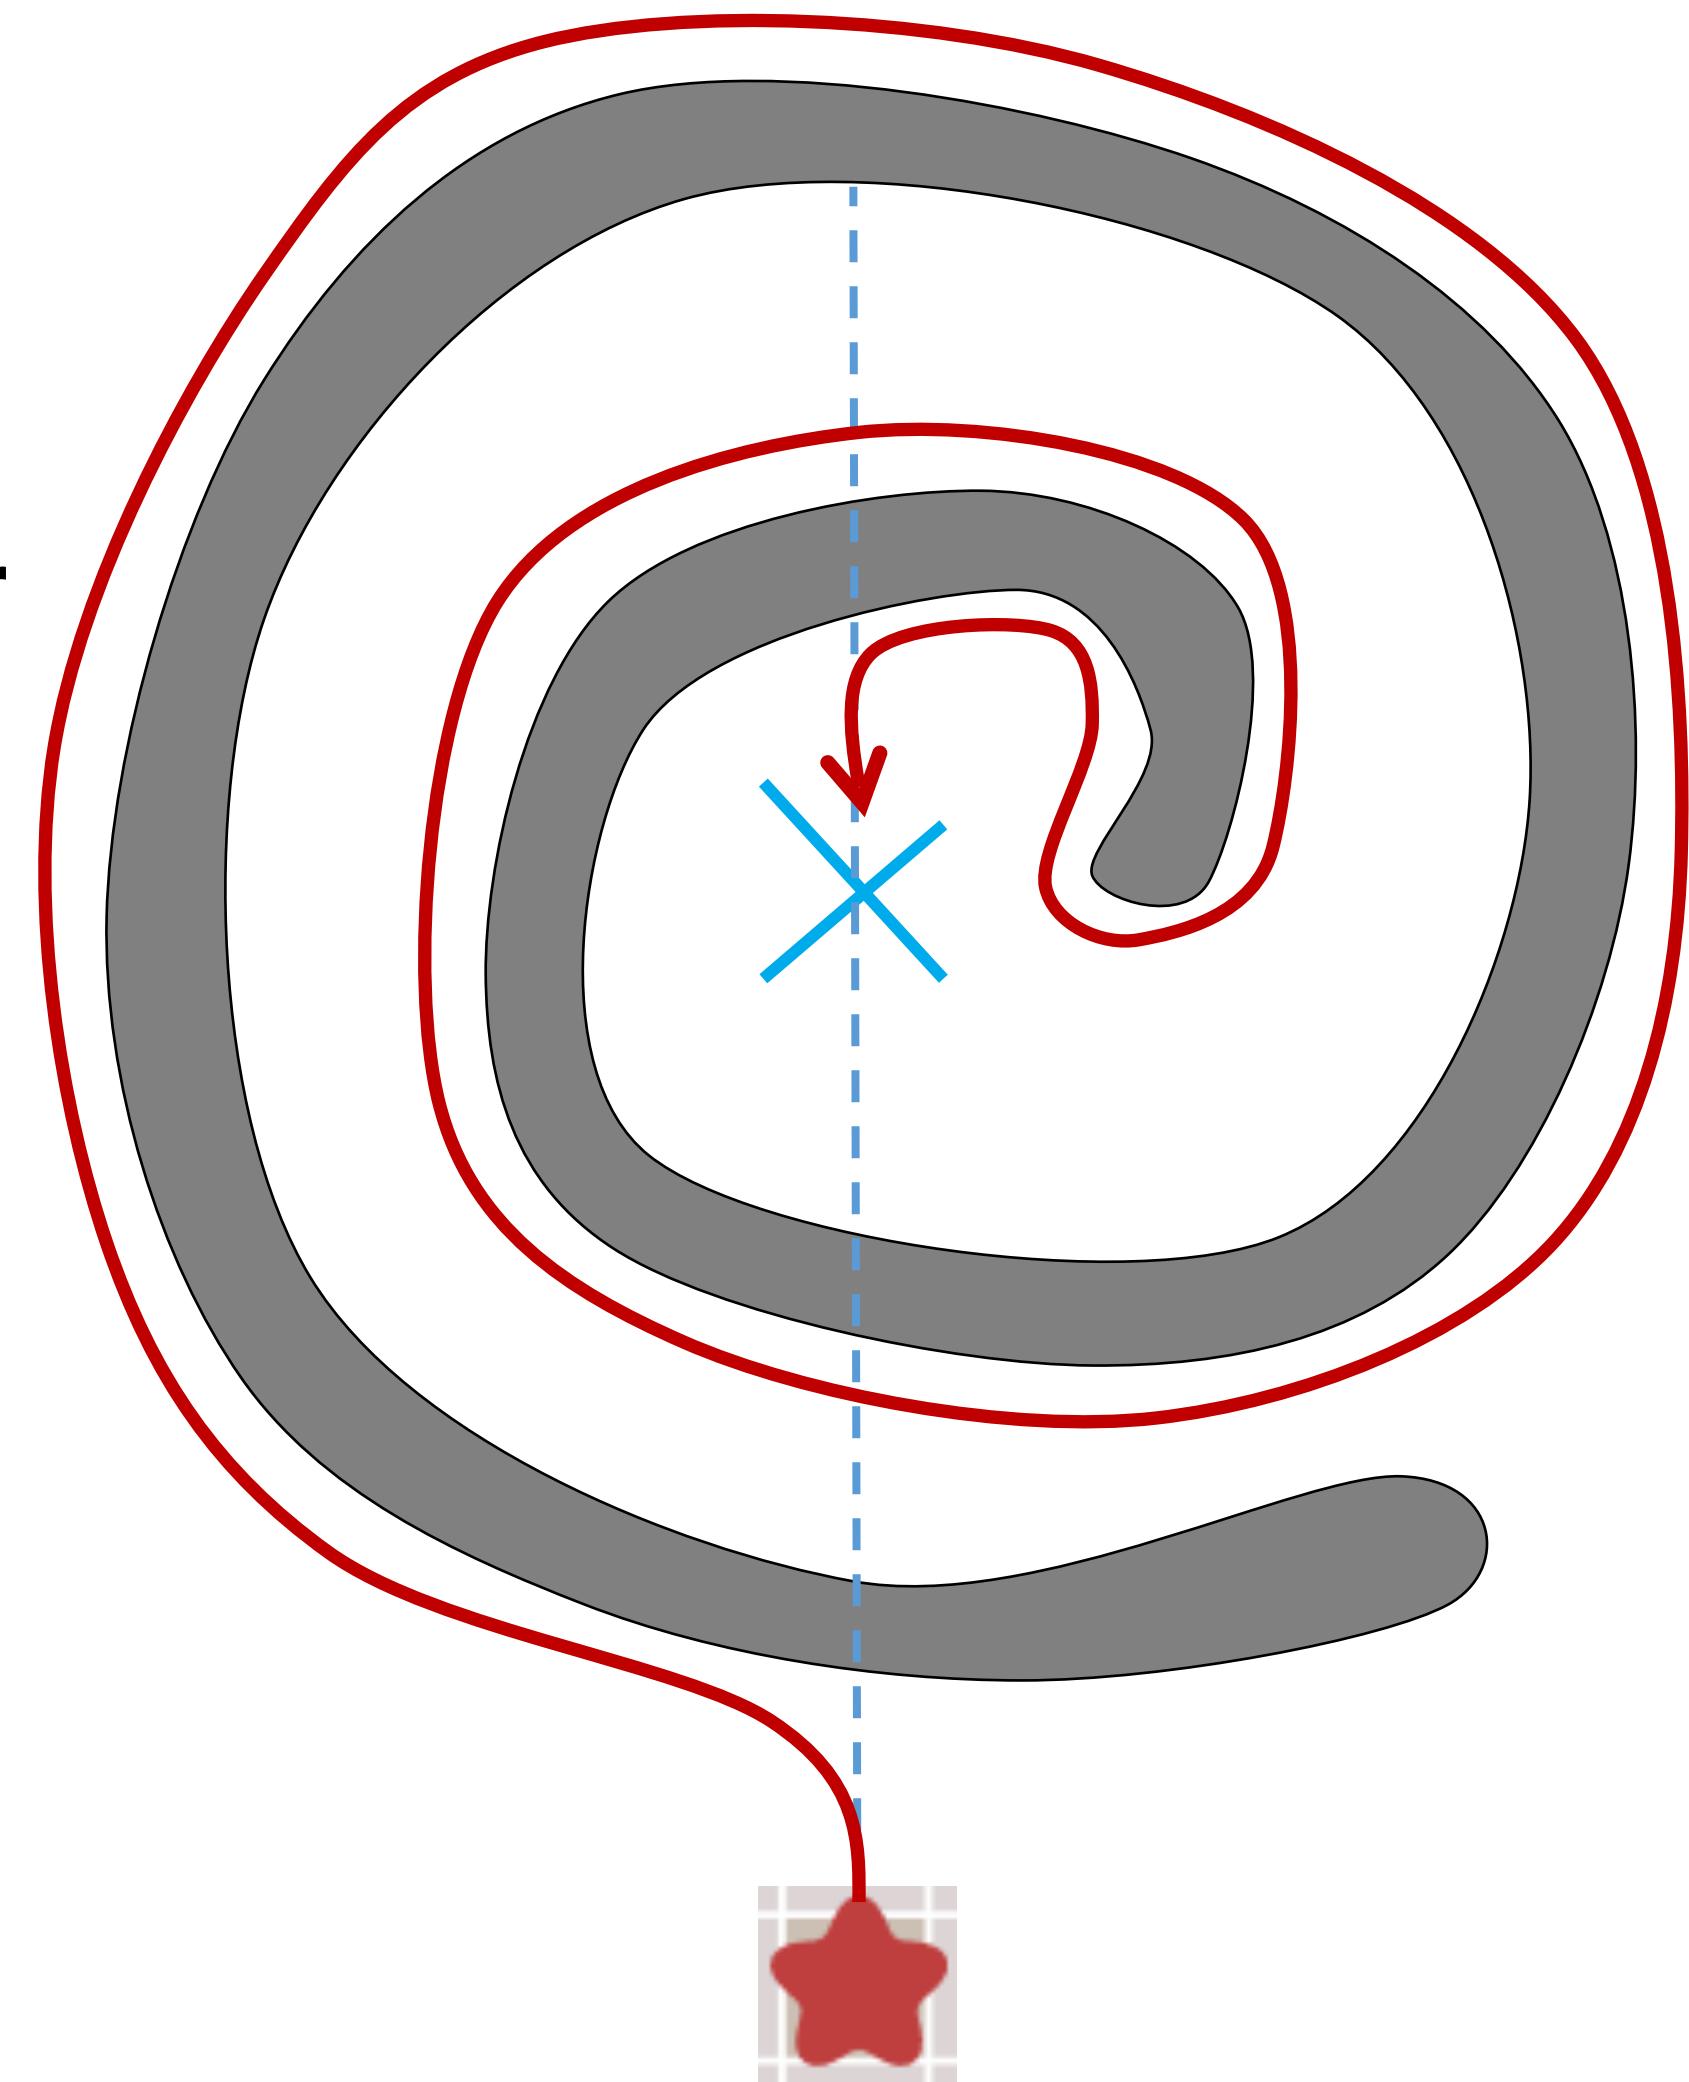
- Sensor Assumptions
 - Direction to the goal
 - Detect walls
 - Odometry
 - Original vector to the goal
- Algorithm
 - Go towards goal on the vector
 - Follow obstacles *until you are back on the vector (and closer to the goal)*
- Loop



Bug 2

- Sensor Assumptions
 - Direction to the goal
 - Detect walls
 - Odometry
 - Original vector to the goal
- Algorithm
 - Go towards goal on the vector
 - Follow obstacles *until you are back on the vector (and closer to the goal)*
- Loop

What is faster, right- or left- wall following?



Battle of the bugs (1 vs 2)

Bug 1
Layout 1

Bug 1
Layout 1

Battle of the bugs (1 vs 2)

Exhaustive search

Bug 1
Layout 2

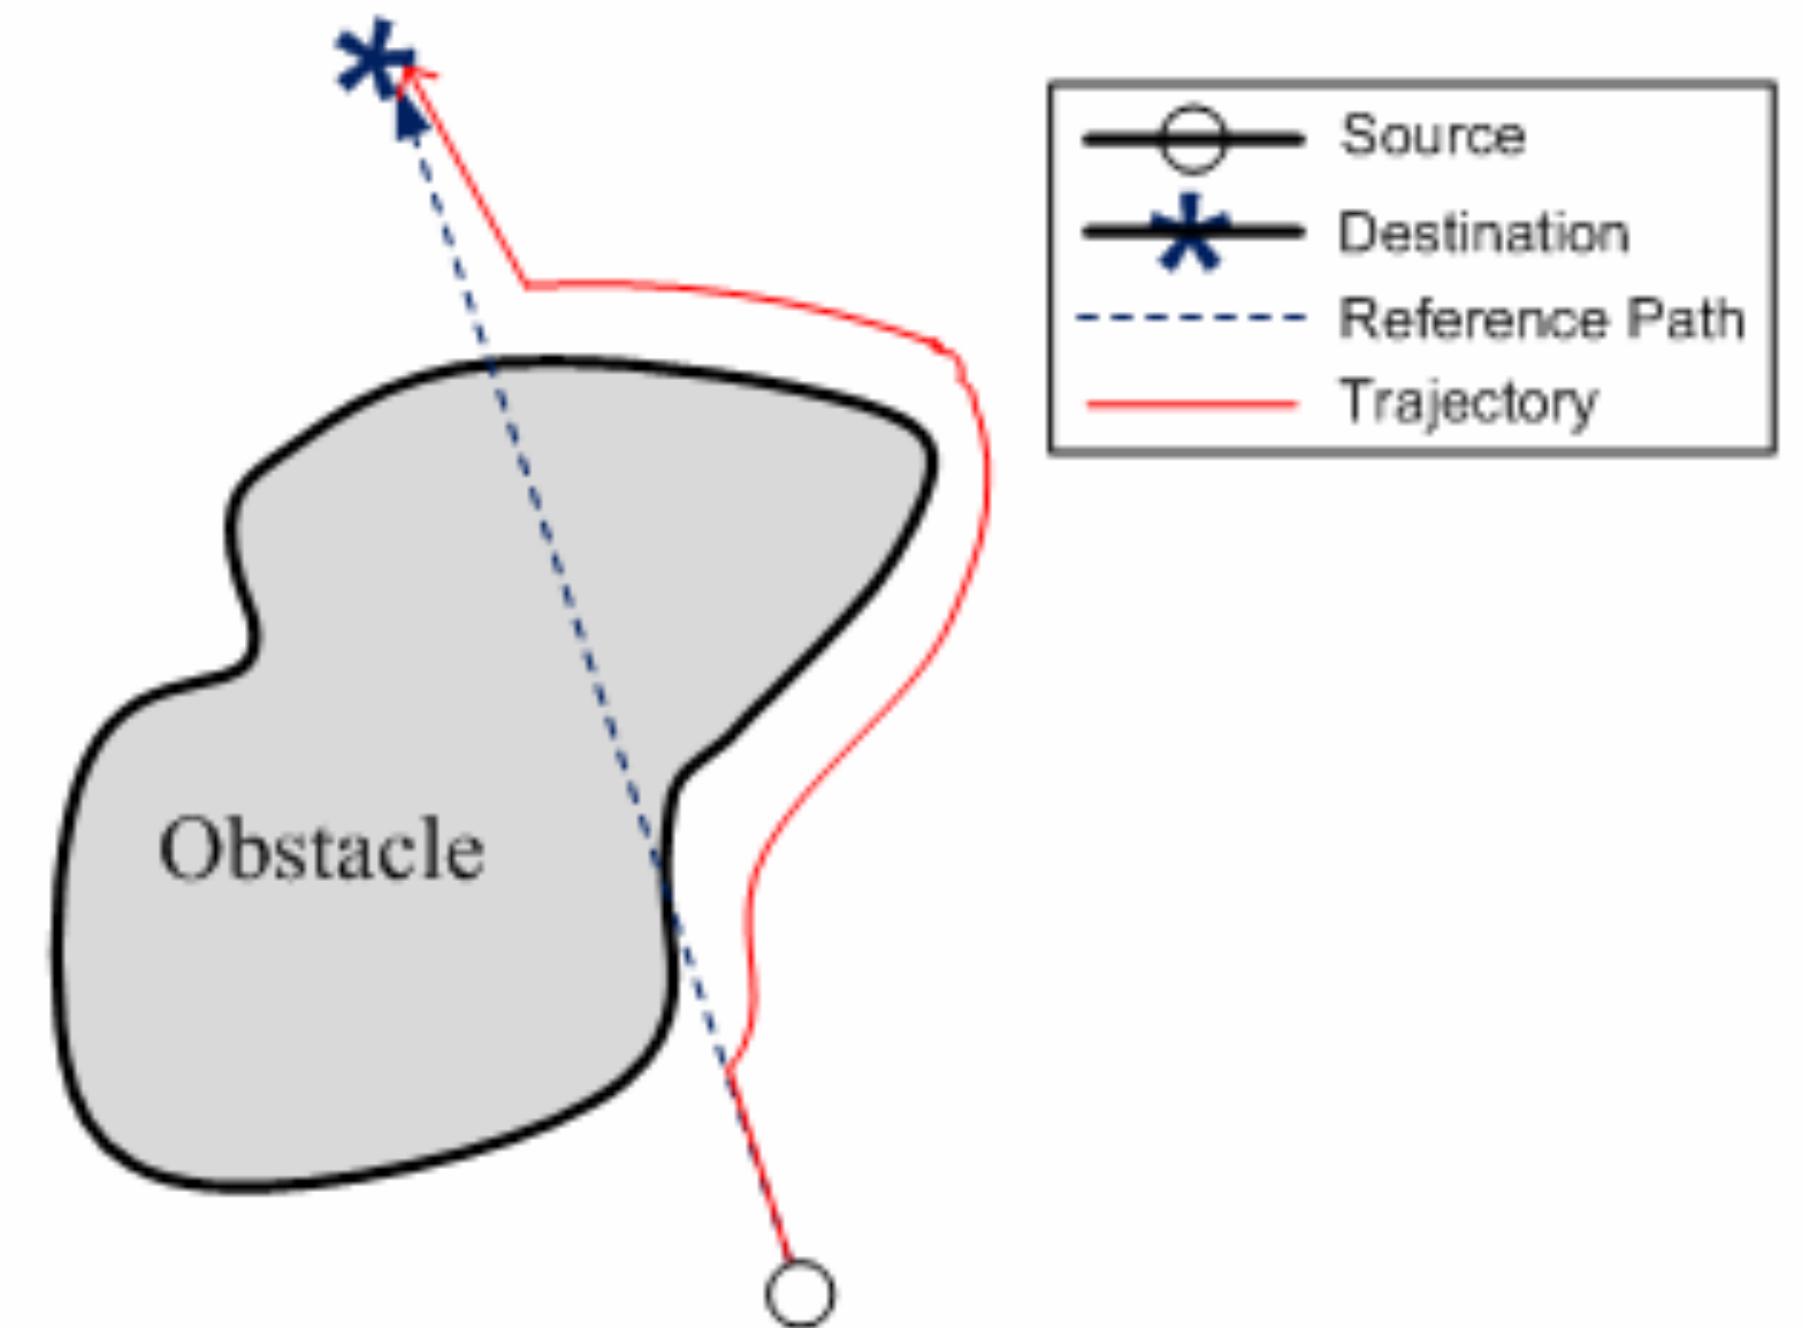
Greedy search

Bug 2
Layout 2

Bug algorithms

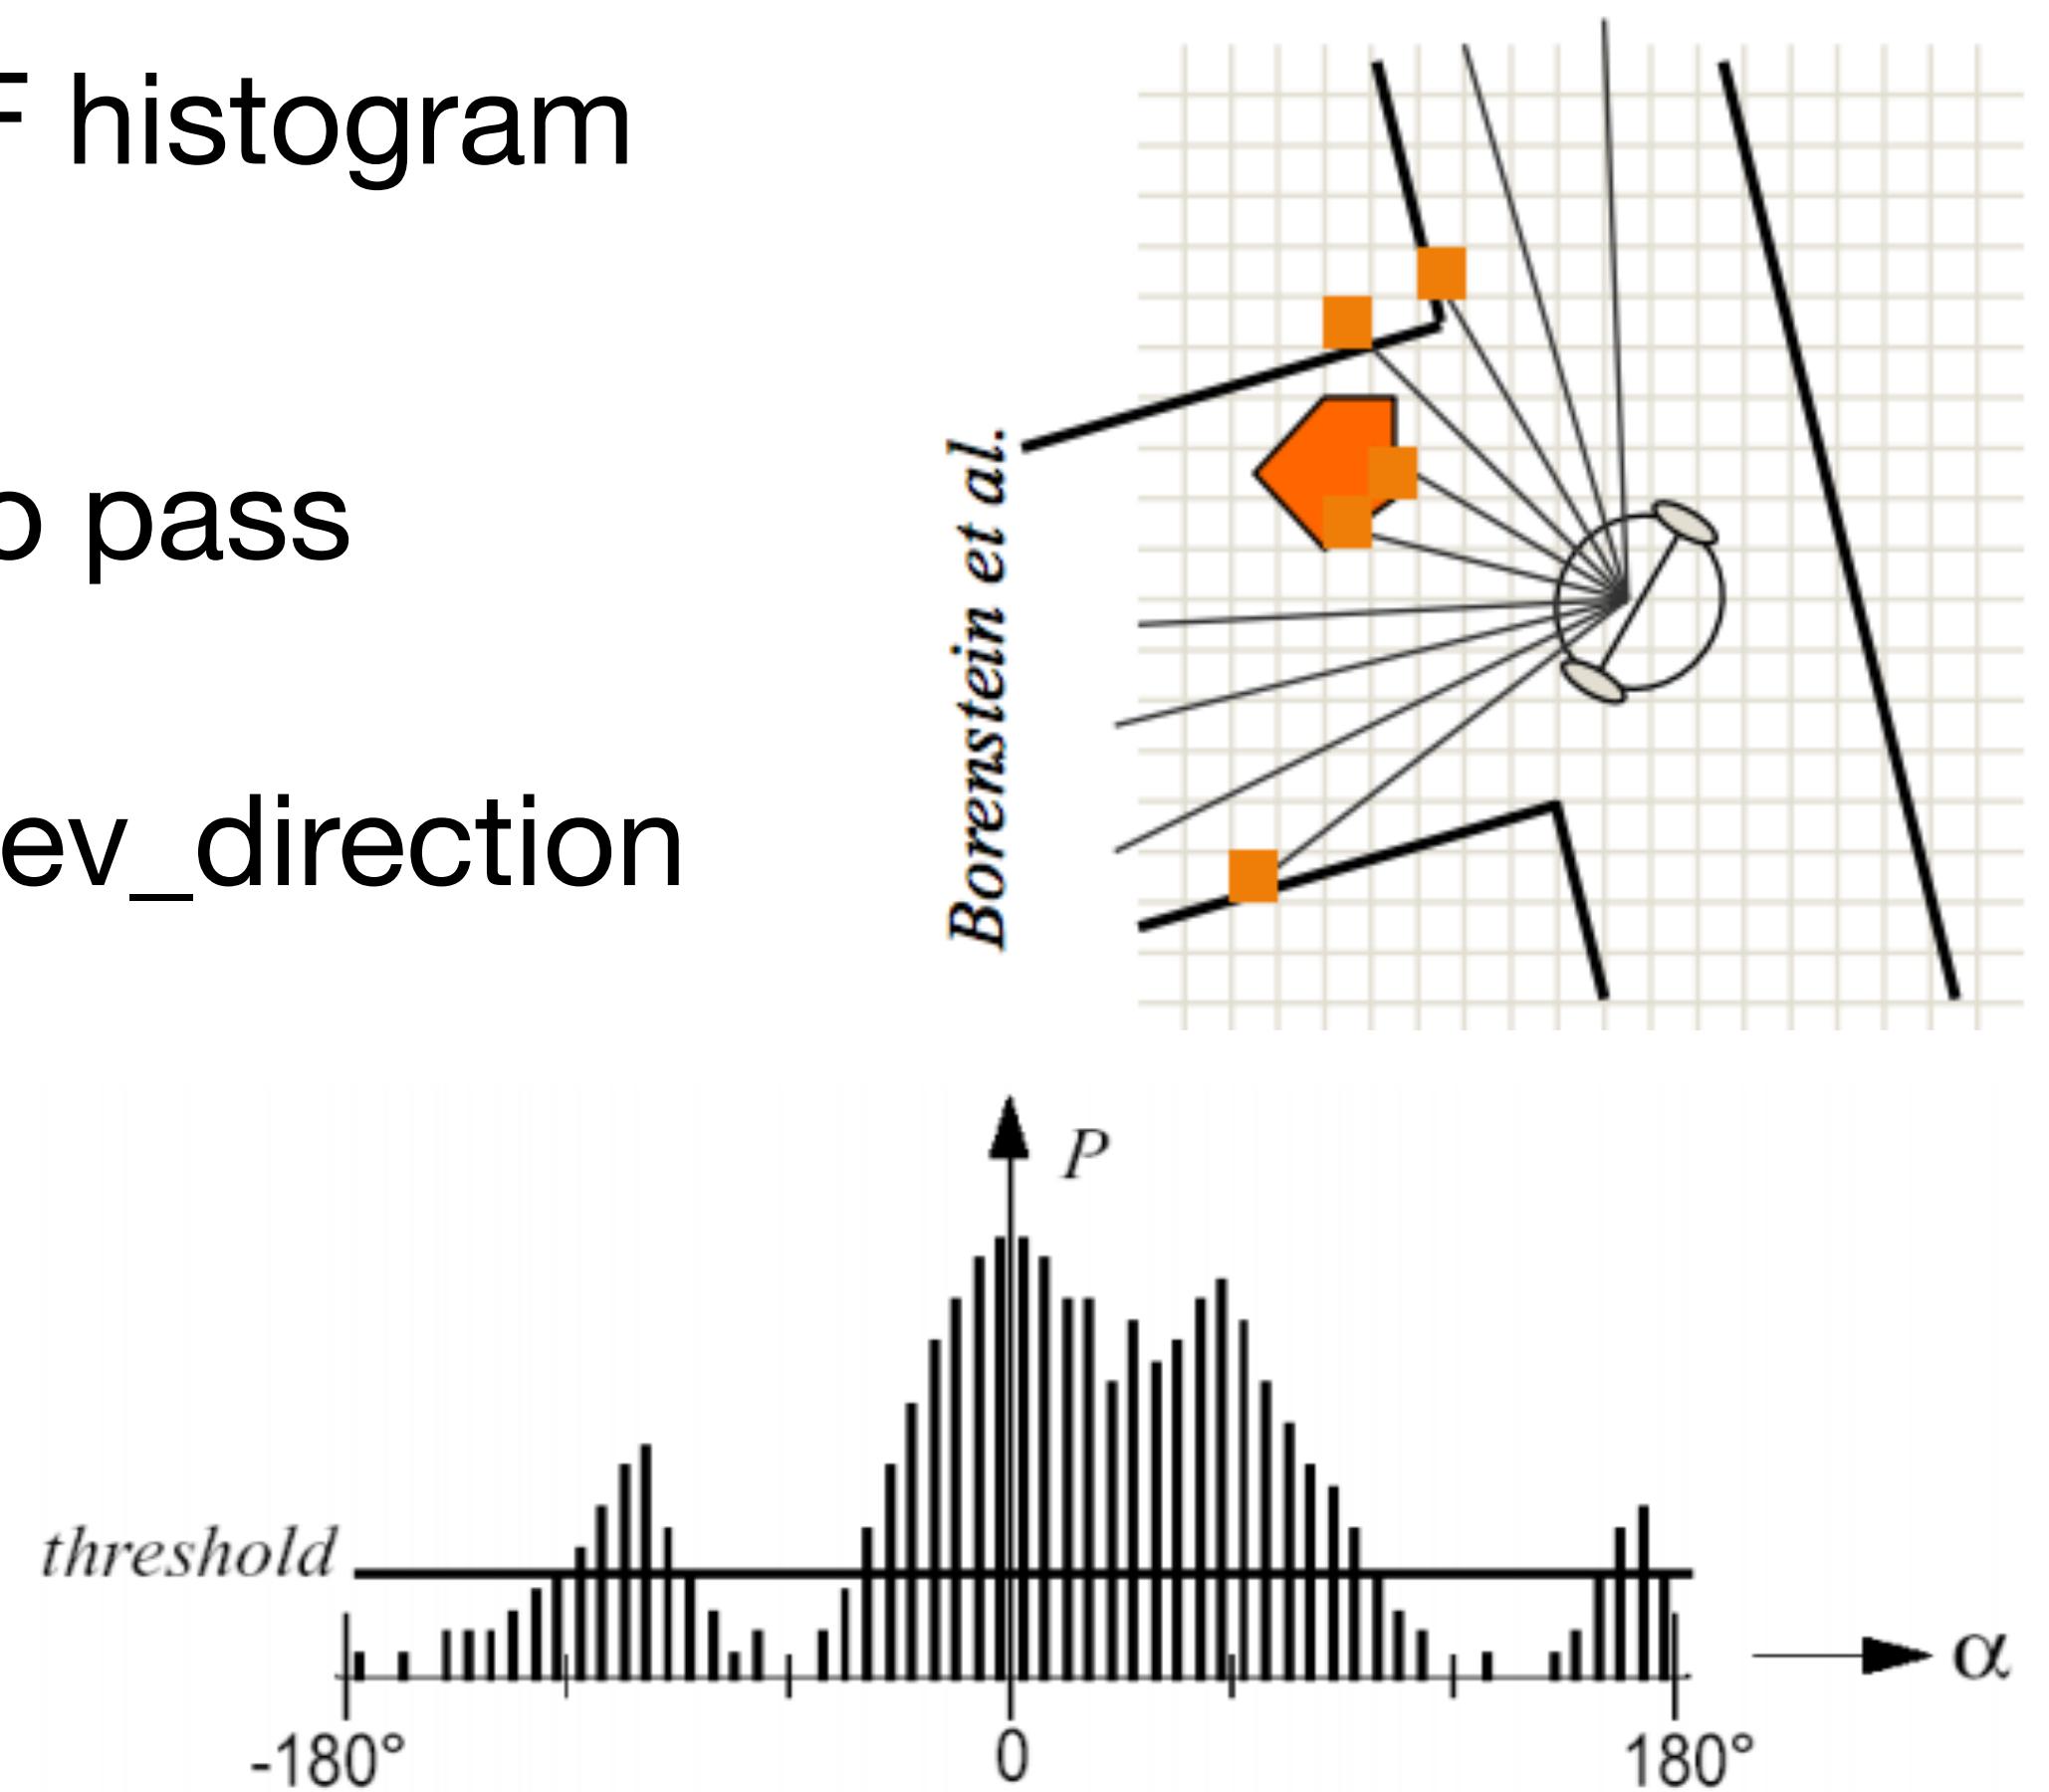
- Uses local knowledge and the direction and distance to the goal
- Basic idea
 - Follow the contour of obstacles until you see the goal
 - State 1: seek goal
 - State 2: follow wall
- Different Variants: Bug0, Bug1, Bug2

• The robot motion behavior is reactive
• Issues if the instantaneous sensor readings do not provide enough information or are noisy



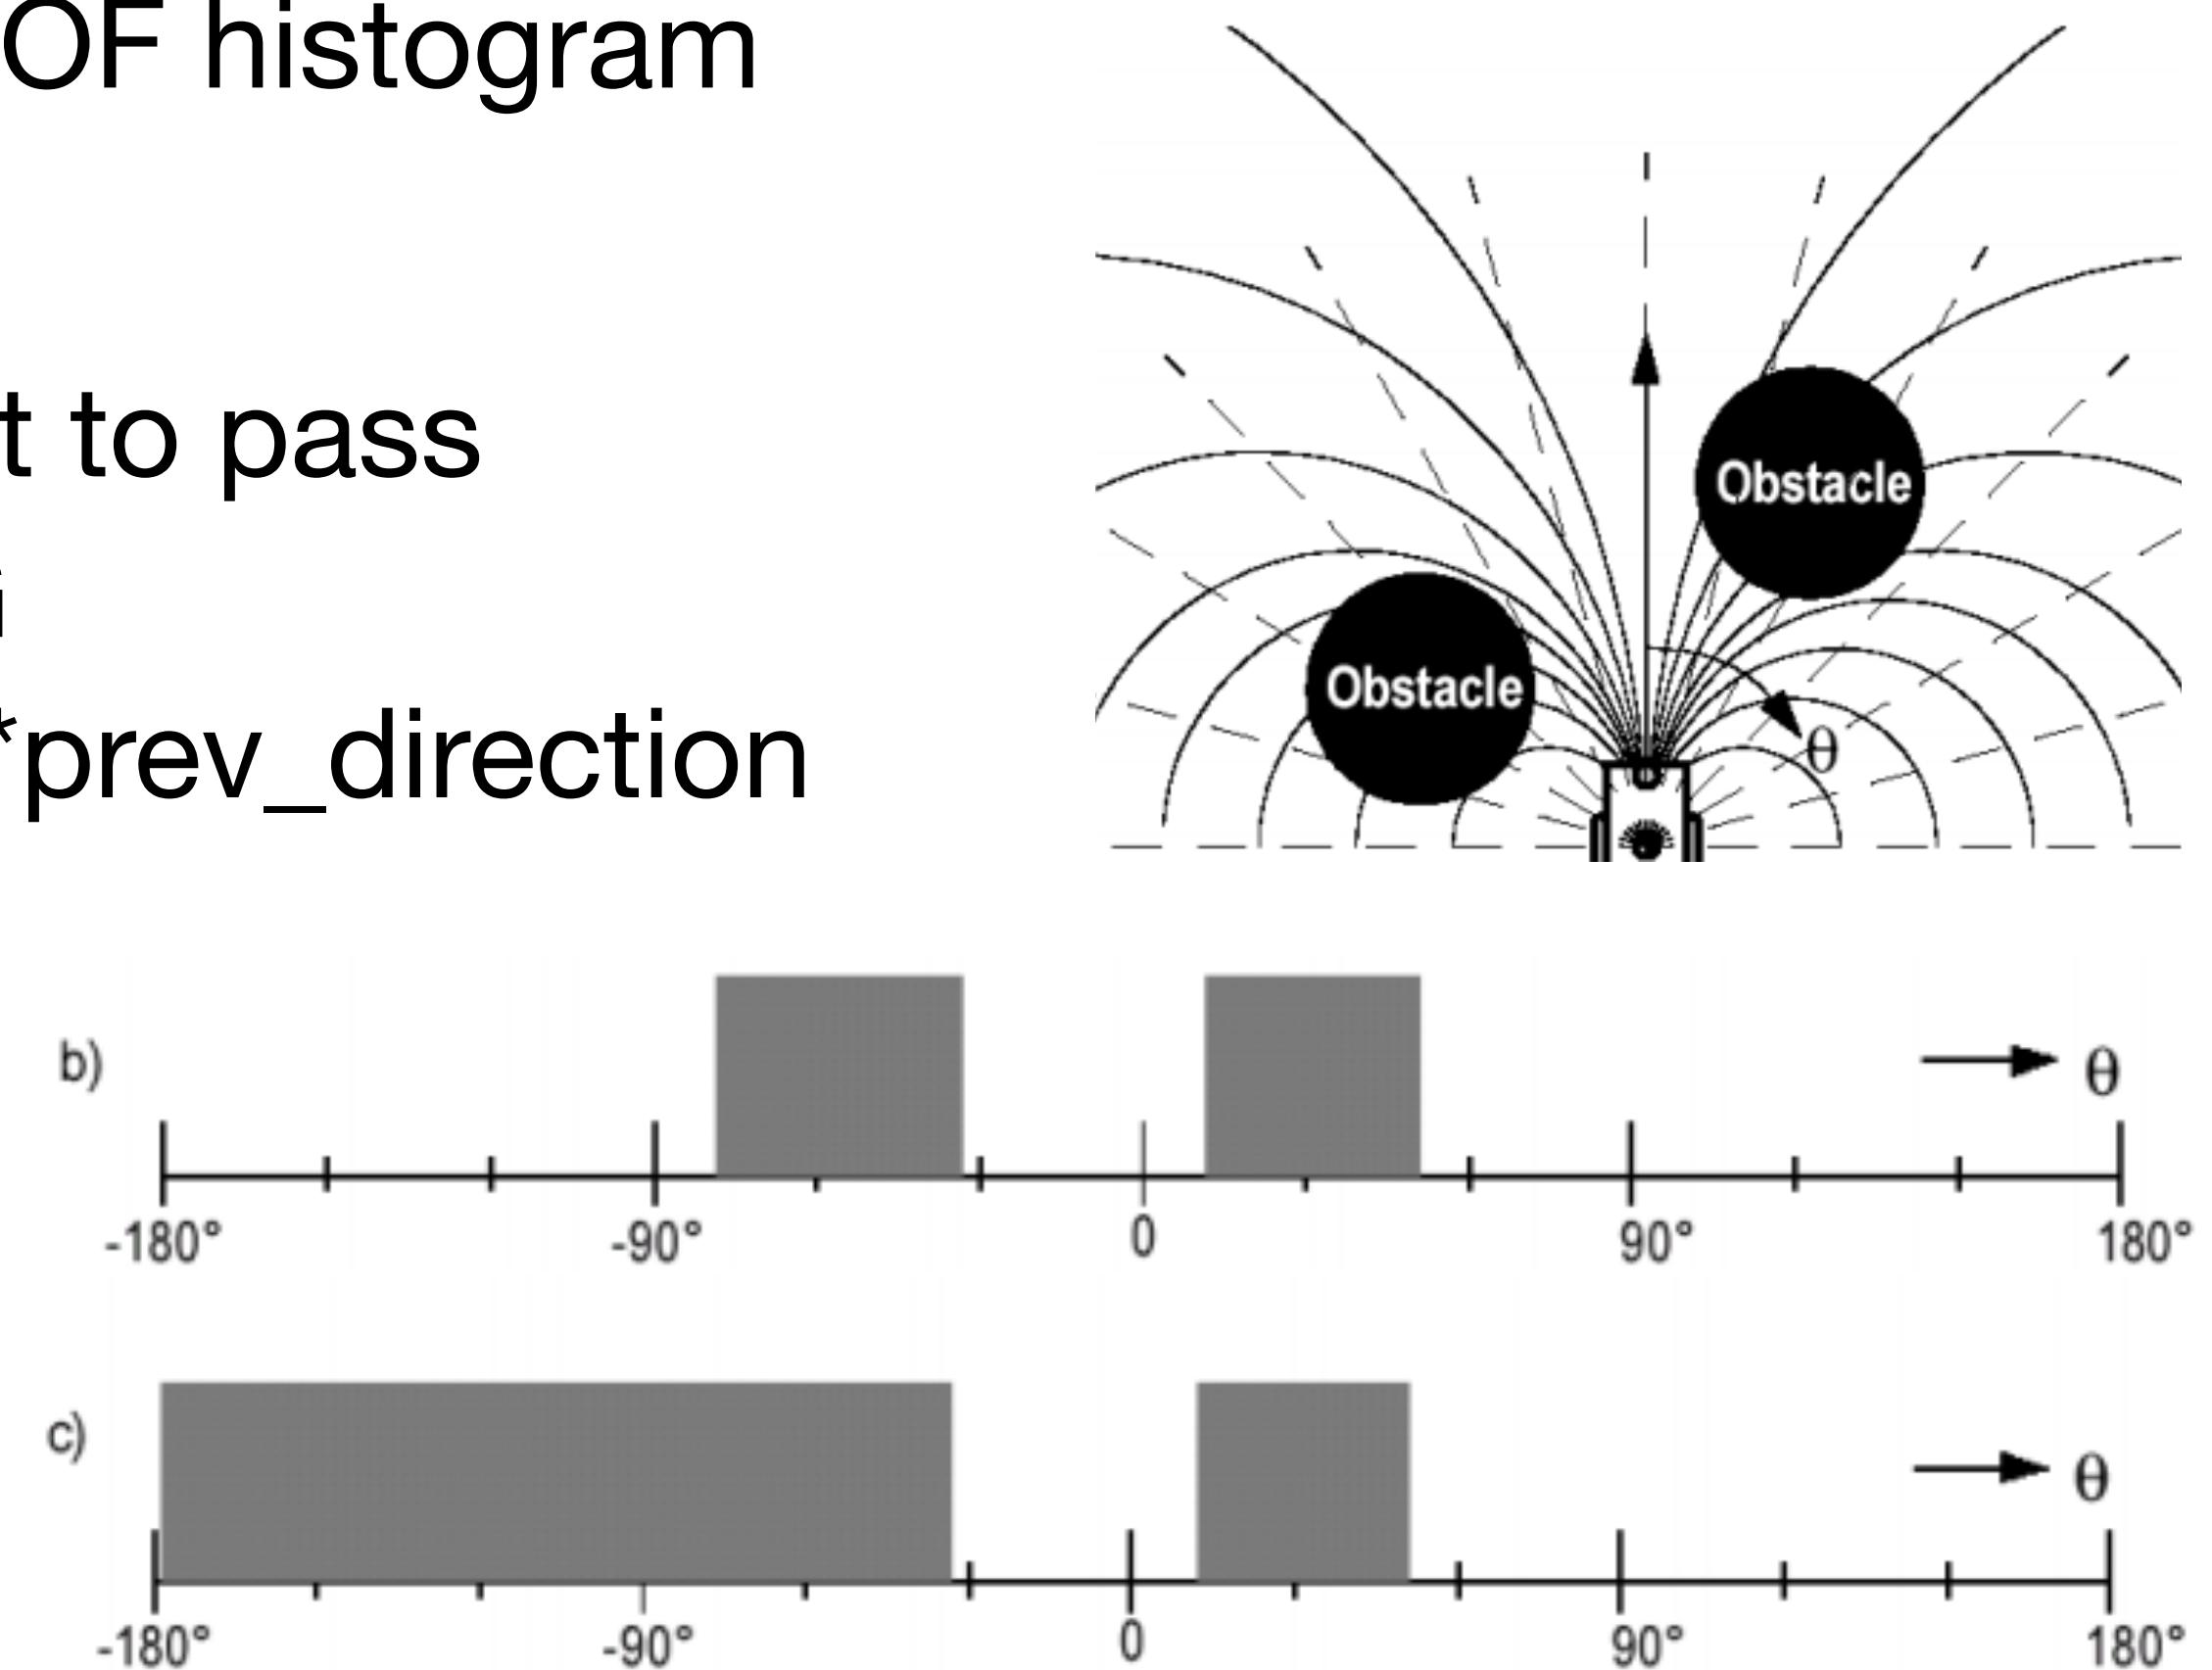
Vector Field Histograms

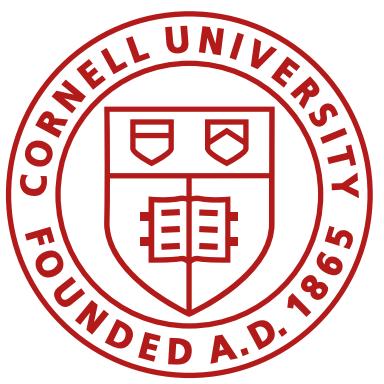
- VFH creates a local map of the environment around the robot populated by “relatively” recent sensor readings
- Build a local 3D grid map reduce to a 1-DOF histogram
- Planning
 - Find all openings large enough for robot to pass
 - Choose the one with the lowest cost, G
 - $G = a^*goal_direction + b^*orientation + c^*prev_direction$



Vector Field Histograms

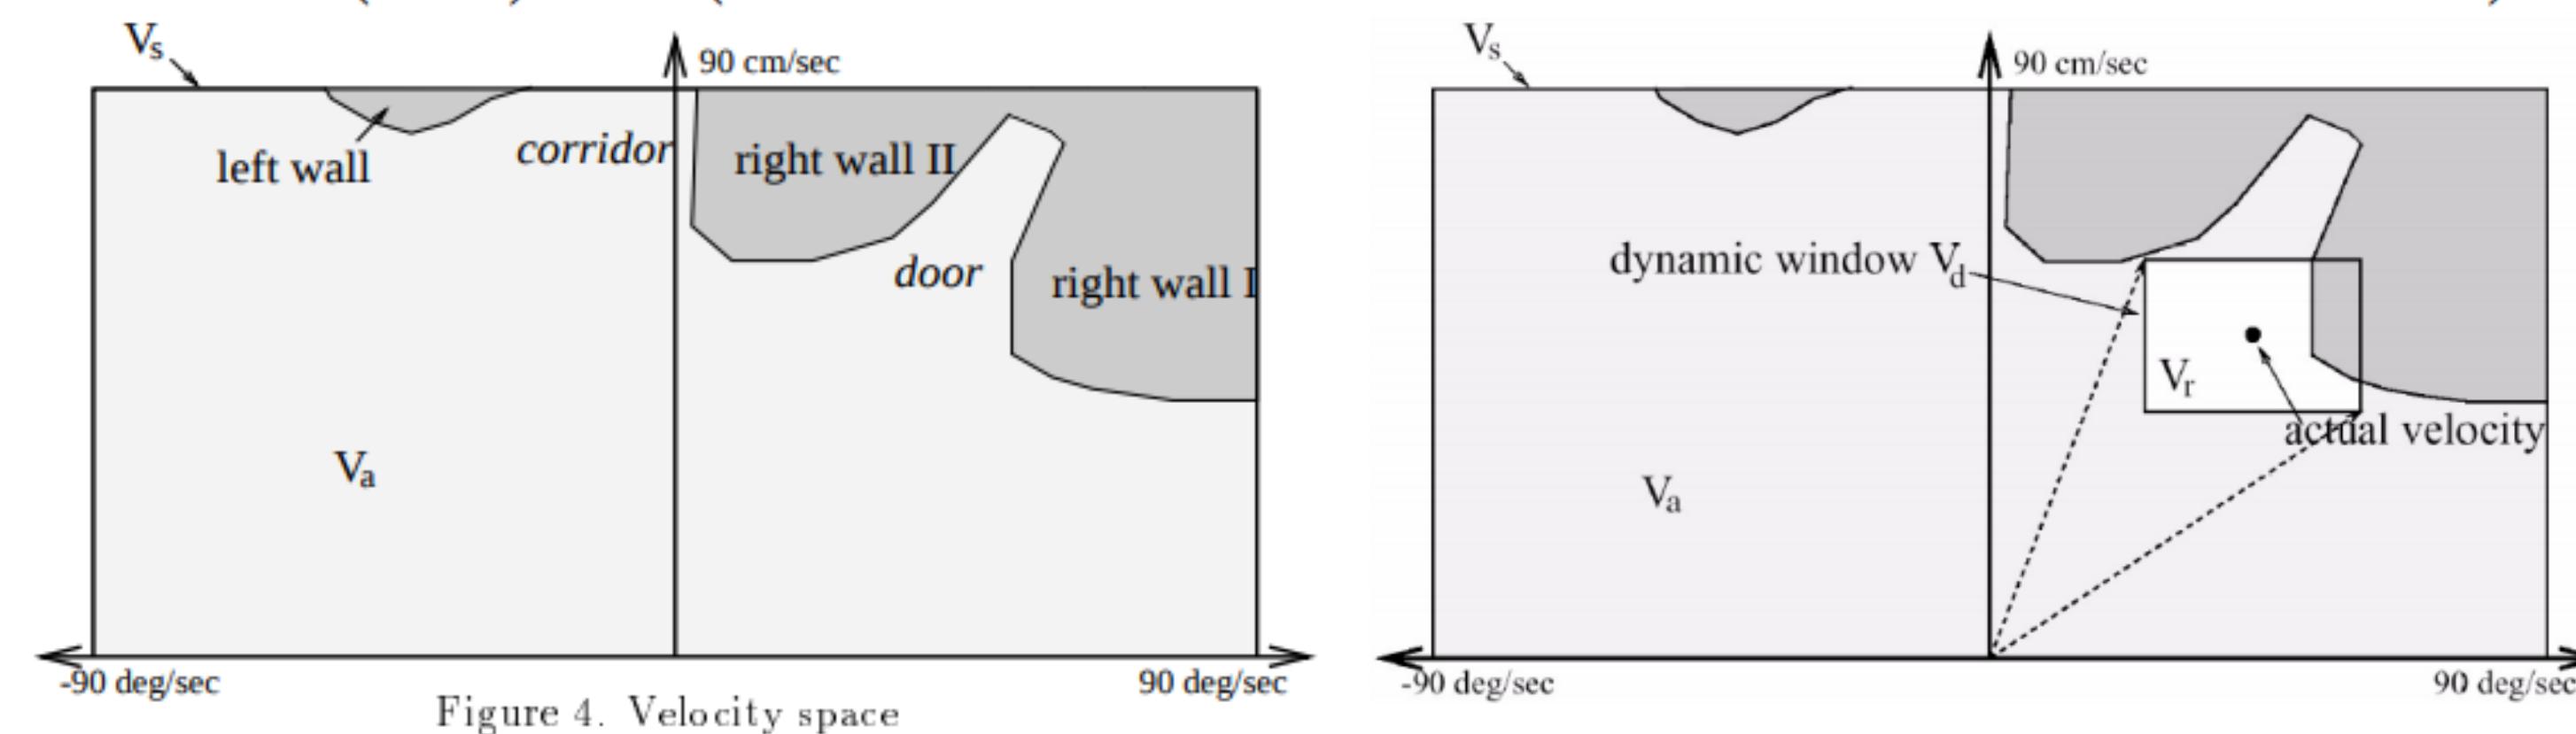
- VFH creates a local map of the environment around the robot populated by “relatively” recent sensor readings
- Build a local 3D grid map reduce to a 1-DOF histogram
- Planning
 - Find all openings large enough for robot to pass
 - Choose the one with the lowest cost, G
 - $G = a^*goal_direction + b^*orientation + c^*prev_direction$
 - VFH+: incorporate kinematics
- Limitations
 - Does not avoid local minima
 - Not guaranteed to reach goal

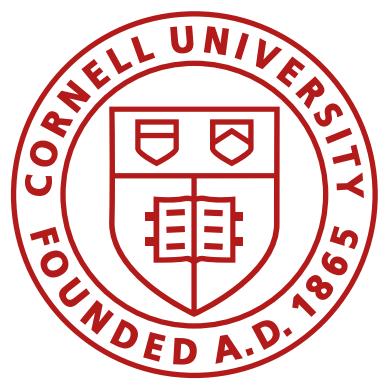




Dynamic Window Approach

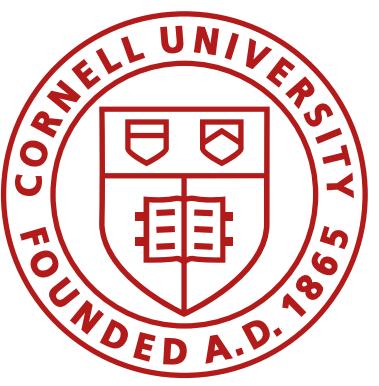
- Search in the velocity space (robot moves in circular arcs)
 - Takes into account robot acceleration and update rates
- A dynamic window, V_d , is the set of all tuples (v_d, ω_d) that can be reached
- Admissible velocities, V_a , include those where the robot can stop before collision
- The search space is then $V_r = V_s \cap V_a \cap V_d$
- Cost function: $G(v, \omega) = \sigma(\alpha \cdot \text{heading}(v, \omega) + \beta \cdot \text{dist}(v, \omega) + \gamma \cdot \text{velocity}(v, \omega))$





Local Planners

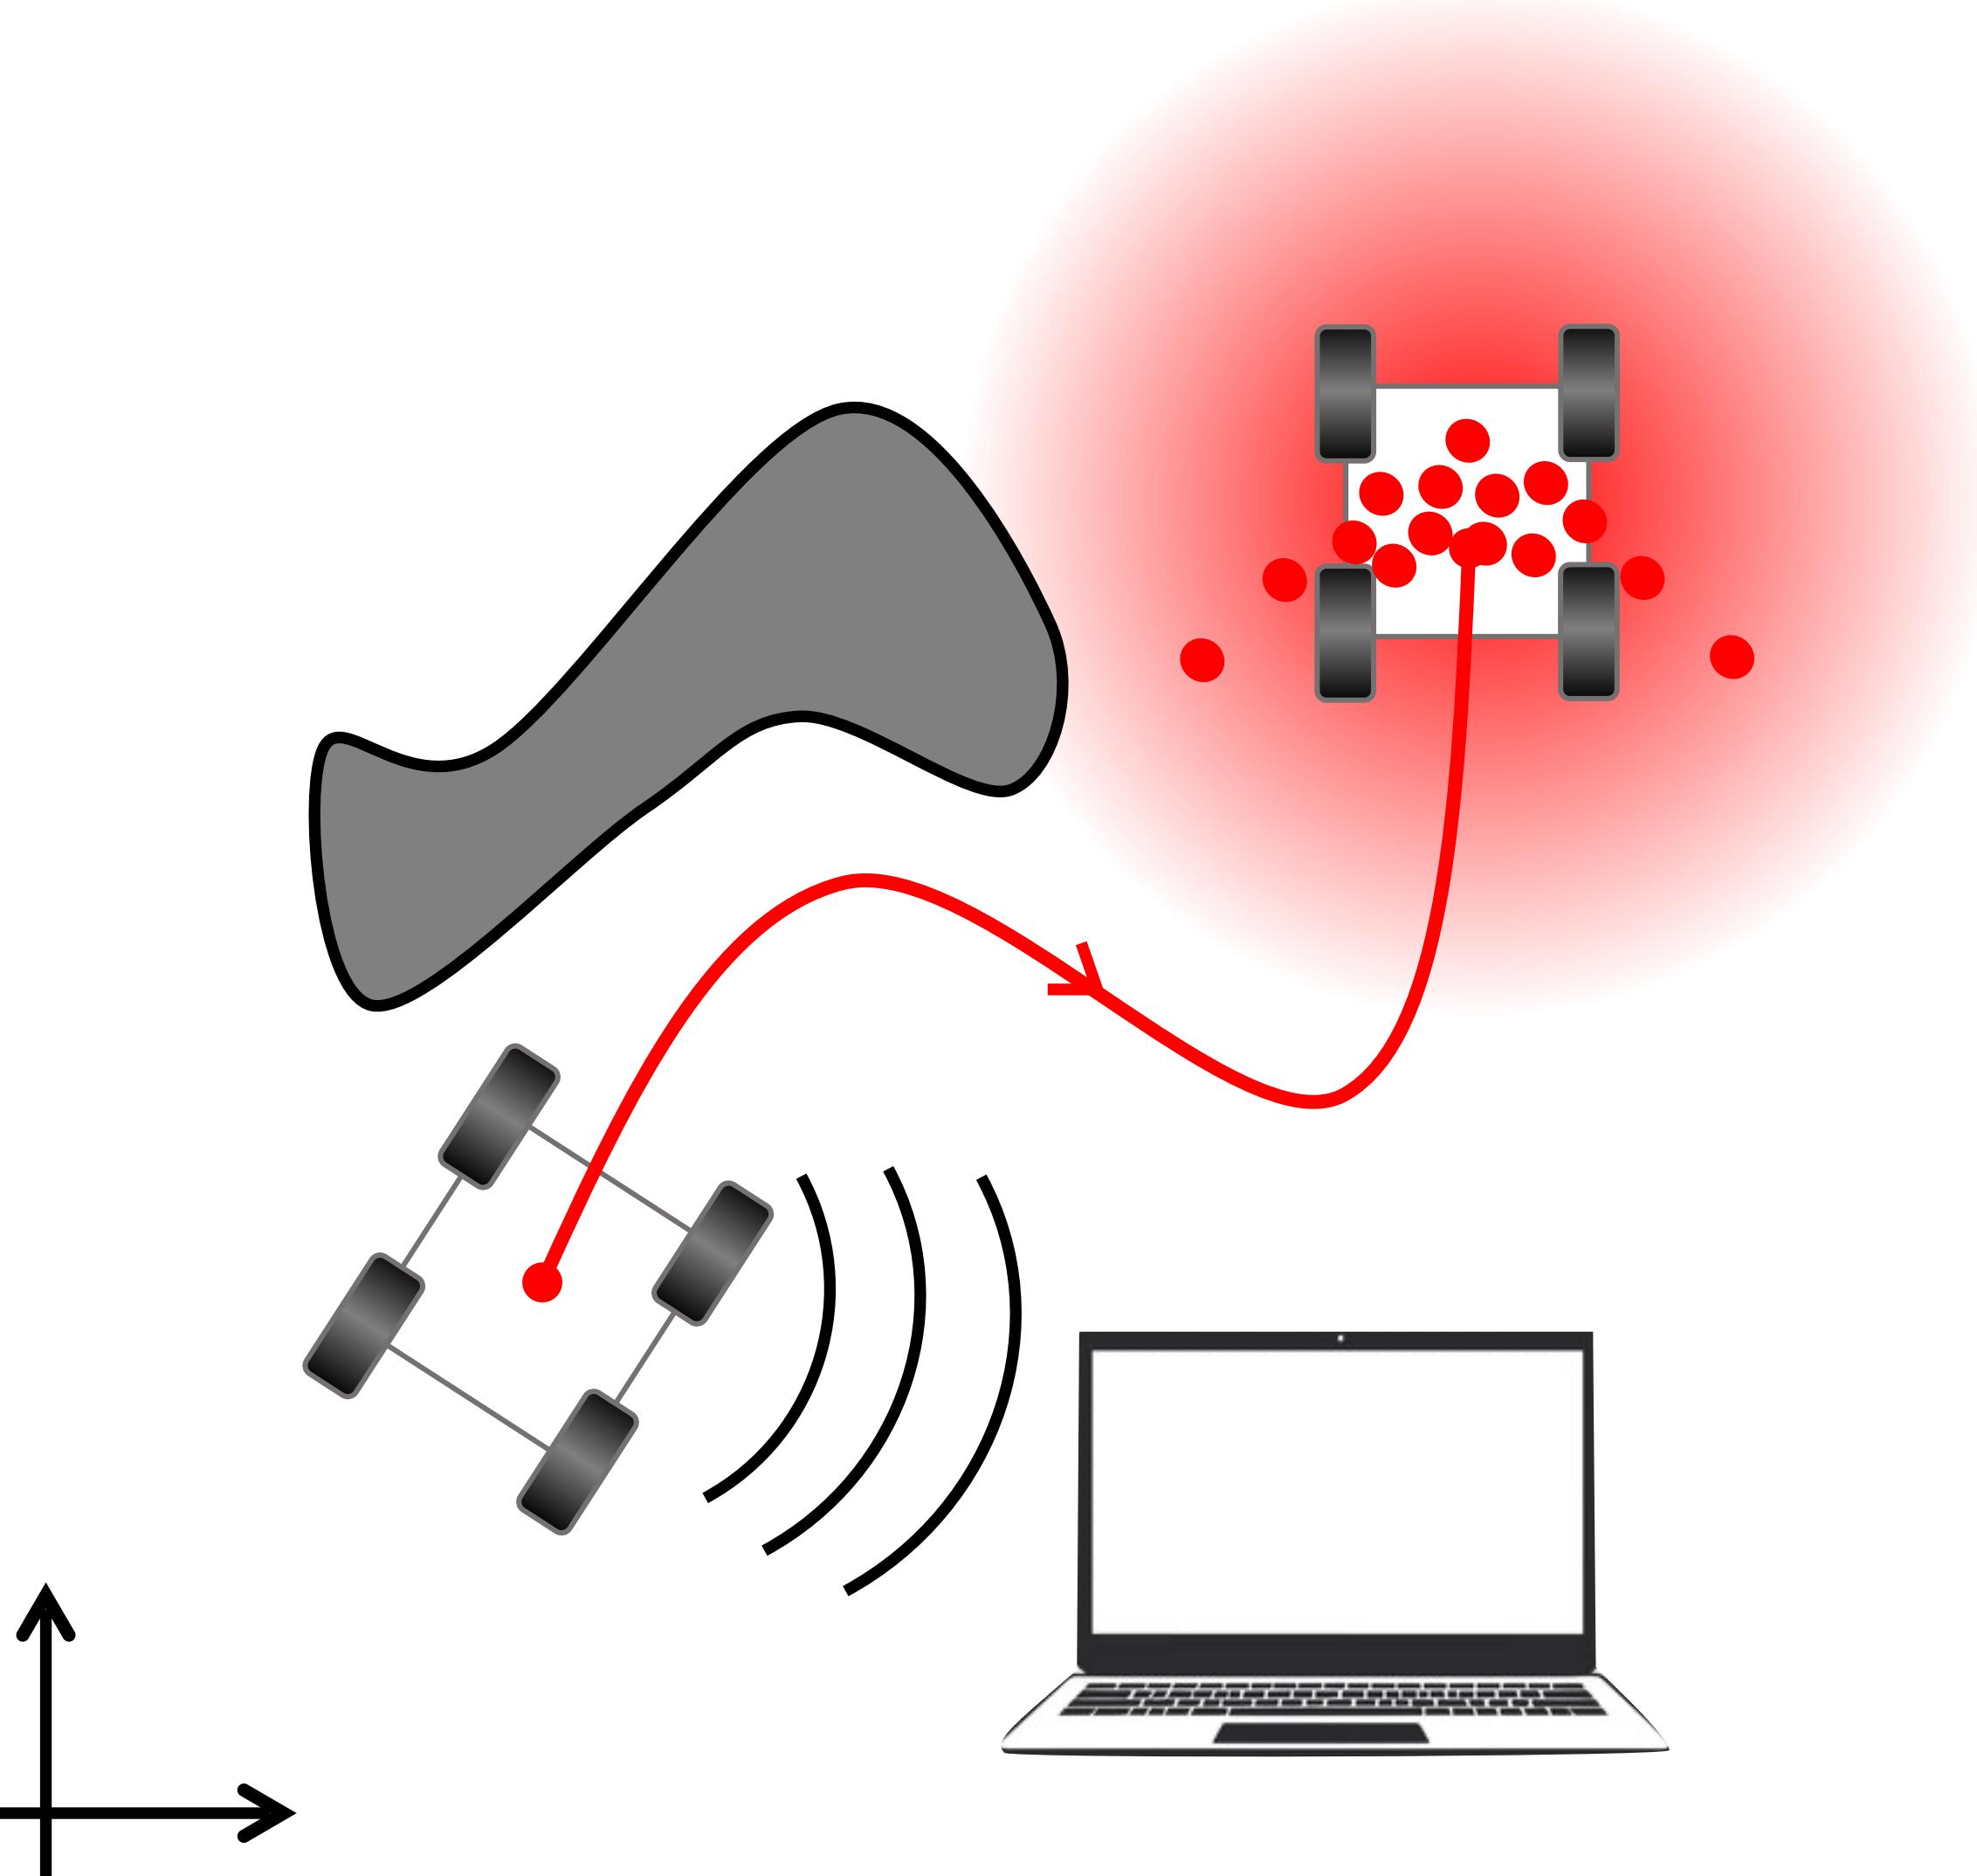
- Bug algorithms
 - Inefficient but can be exhaustive
- Vector Field Histograms
 - Takes into account probabilistic sensor measurements
- Vector Field Histograms+
 - Takes into account probabilistic sensor measurements and robot kinematics
- Dynamic window approach
 - Takes into account robot dynamics



Global localization

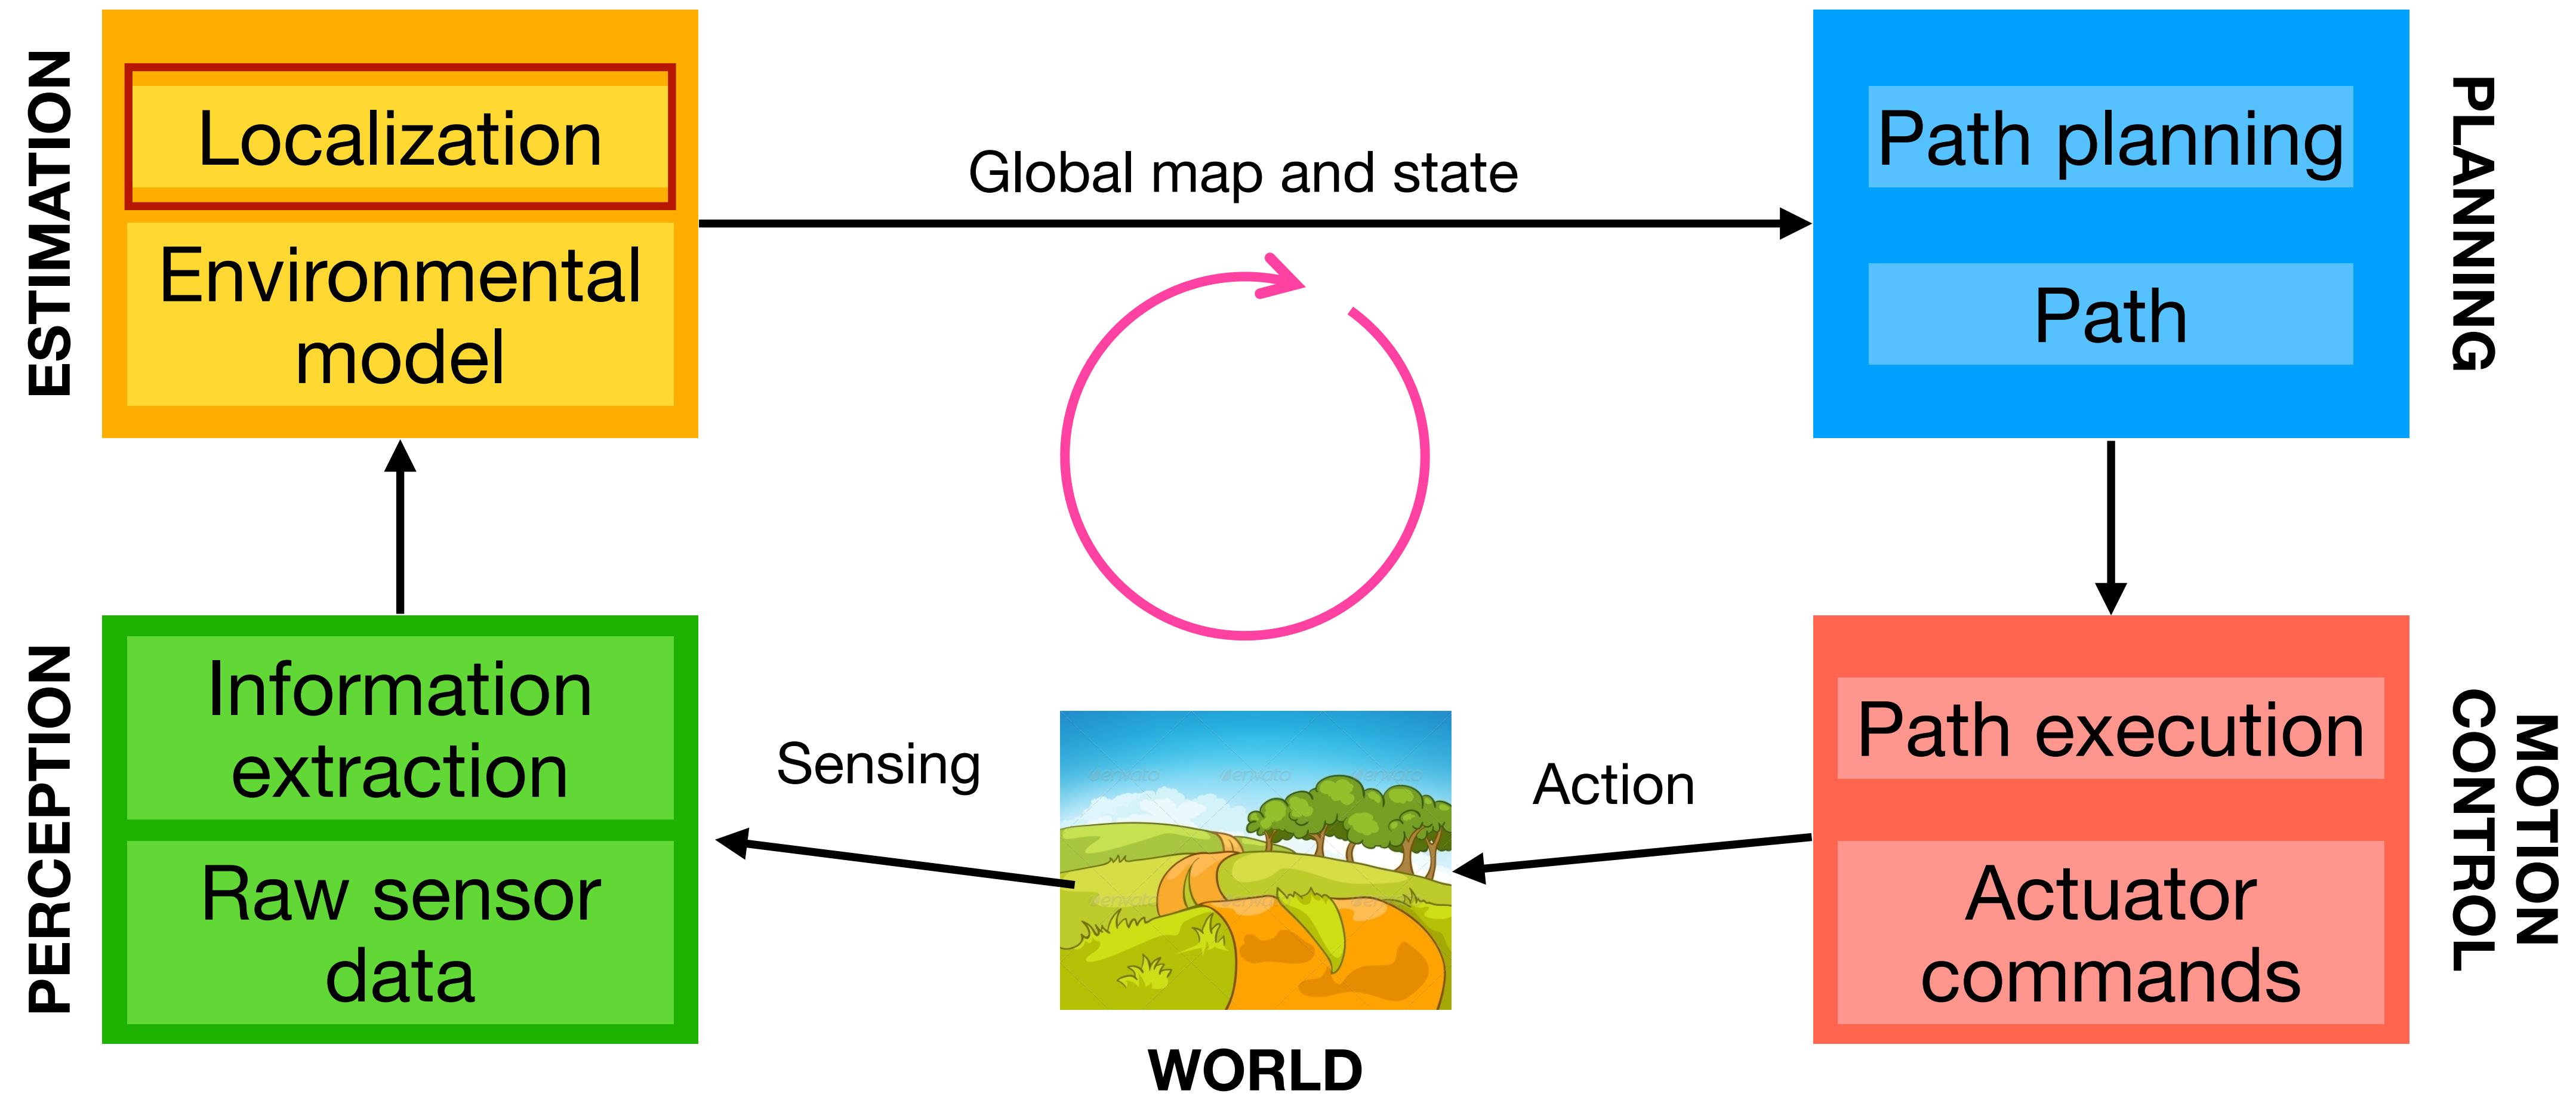
Next module on navigation

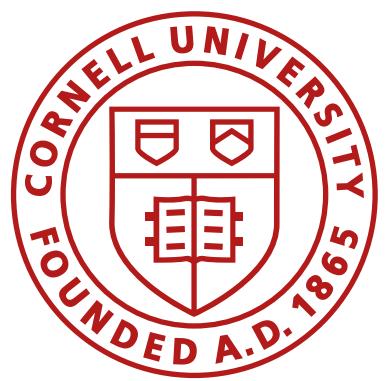
- Local planners
- Global localization and planning
 - Map representations
 - Continuous
 - Discrete
 - Topological
 - Maps as graphs
 - Graph search algorithms
 - Breadth first search
 - Depth first search
 - Dijkstras
 - A*



Navigation

- Break the problem down: localization, map building, path planning





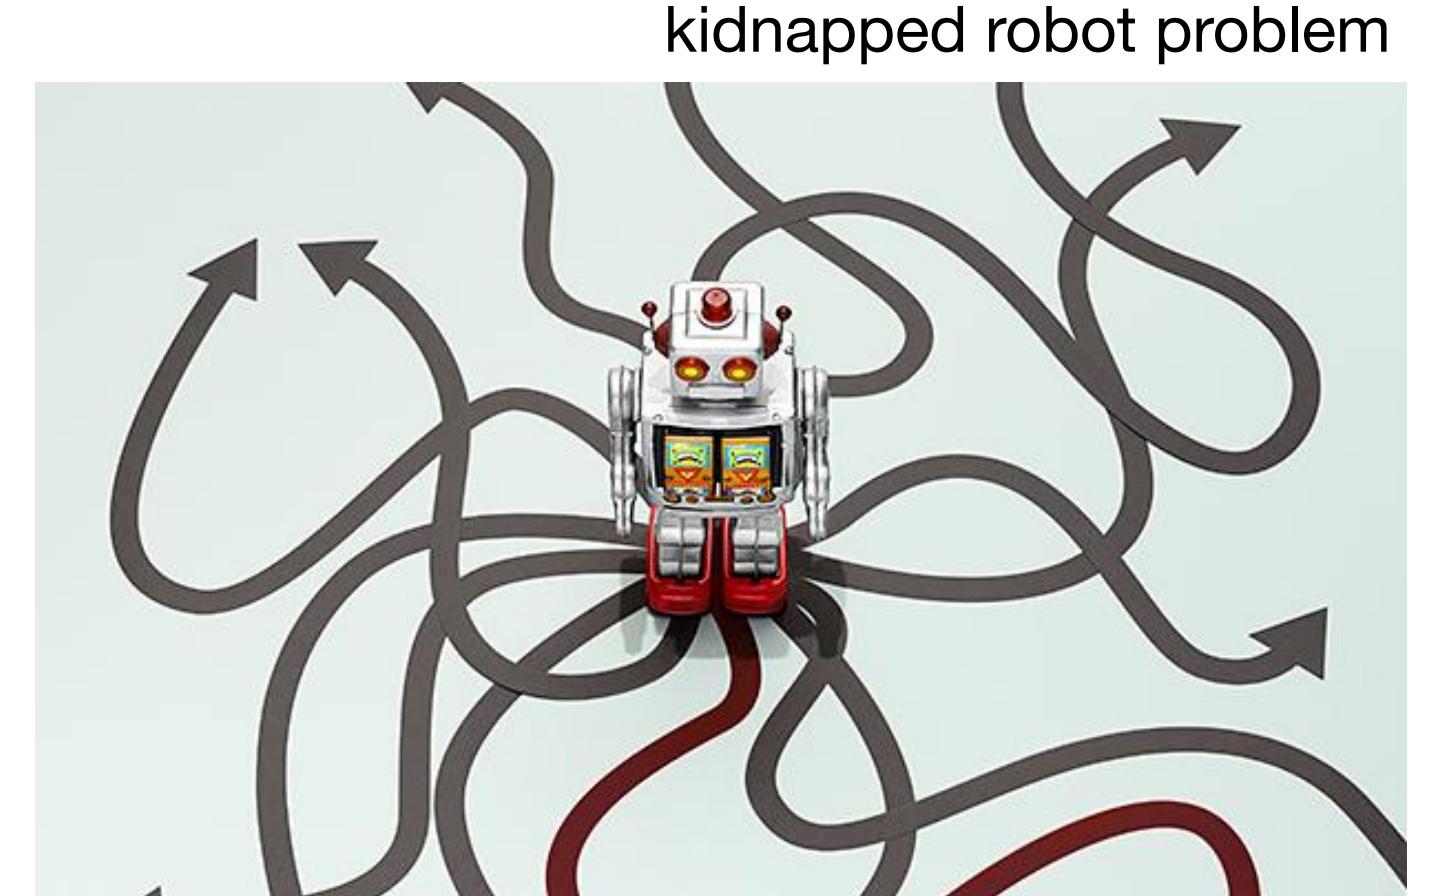
Localization Problem

Position Tracking

- Initial **robot pose is known**
- Either deterministically (odometry) or through Bayesian statistic (motion and sensor models)
- It is a “**local**” problem, as the uncertainty is local (often small) and confined to a region near the robot’s true pose

Global Localization

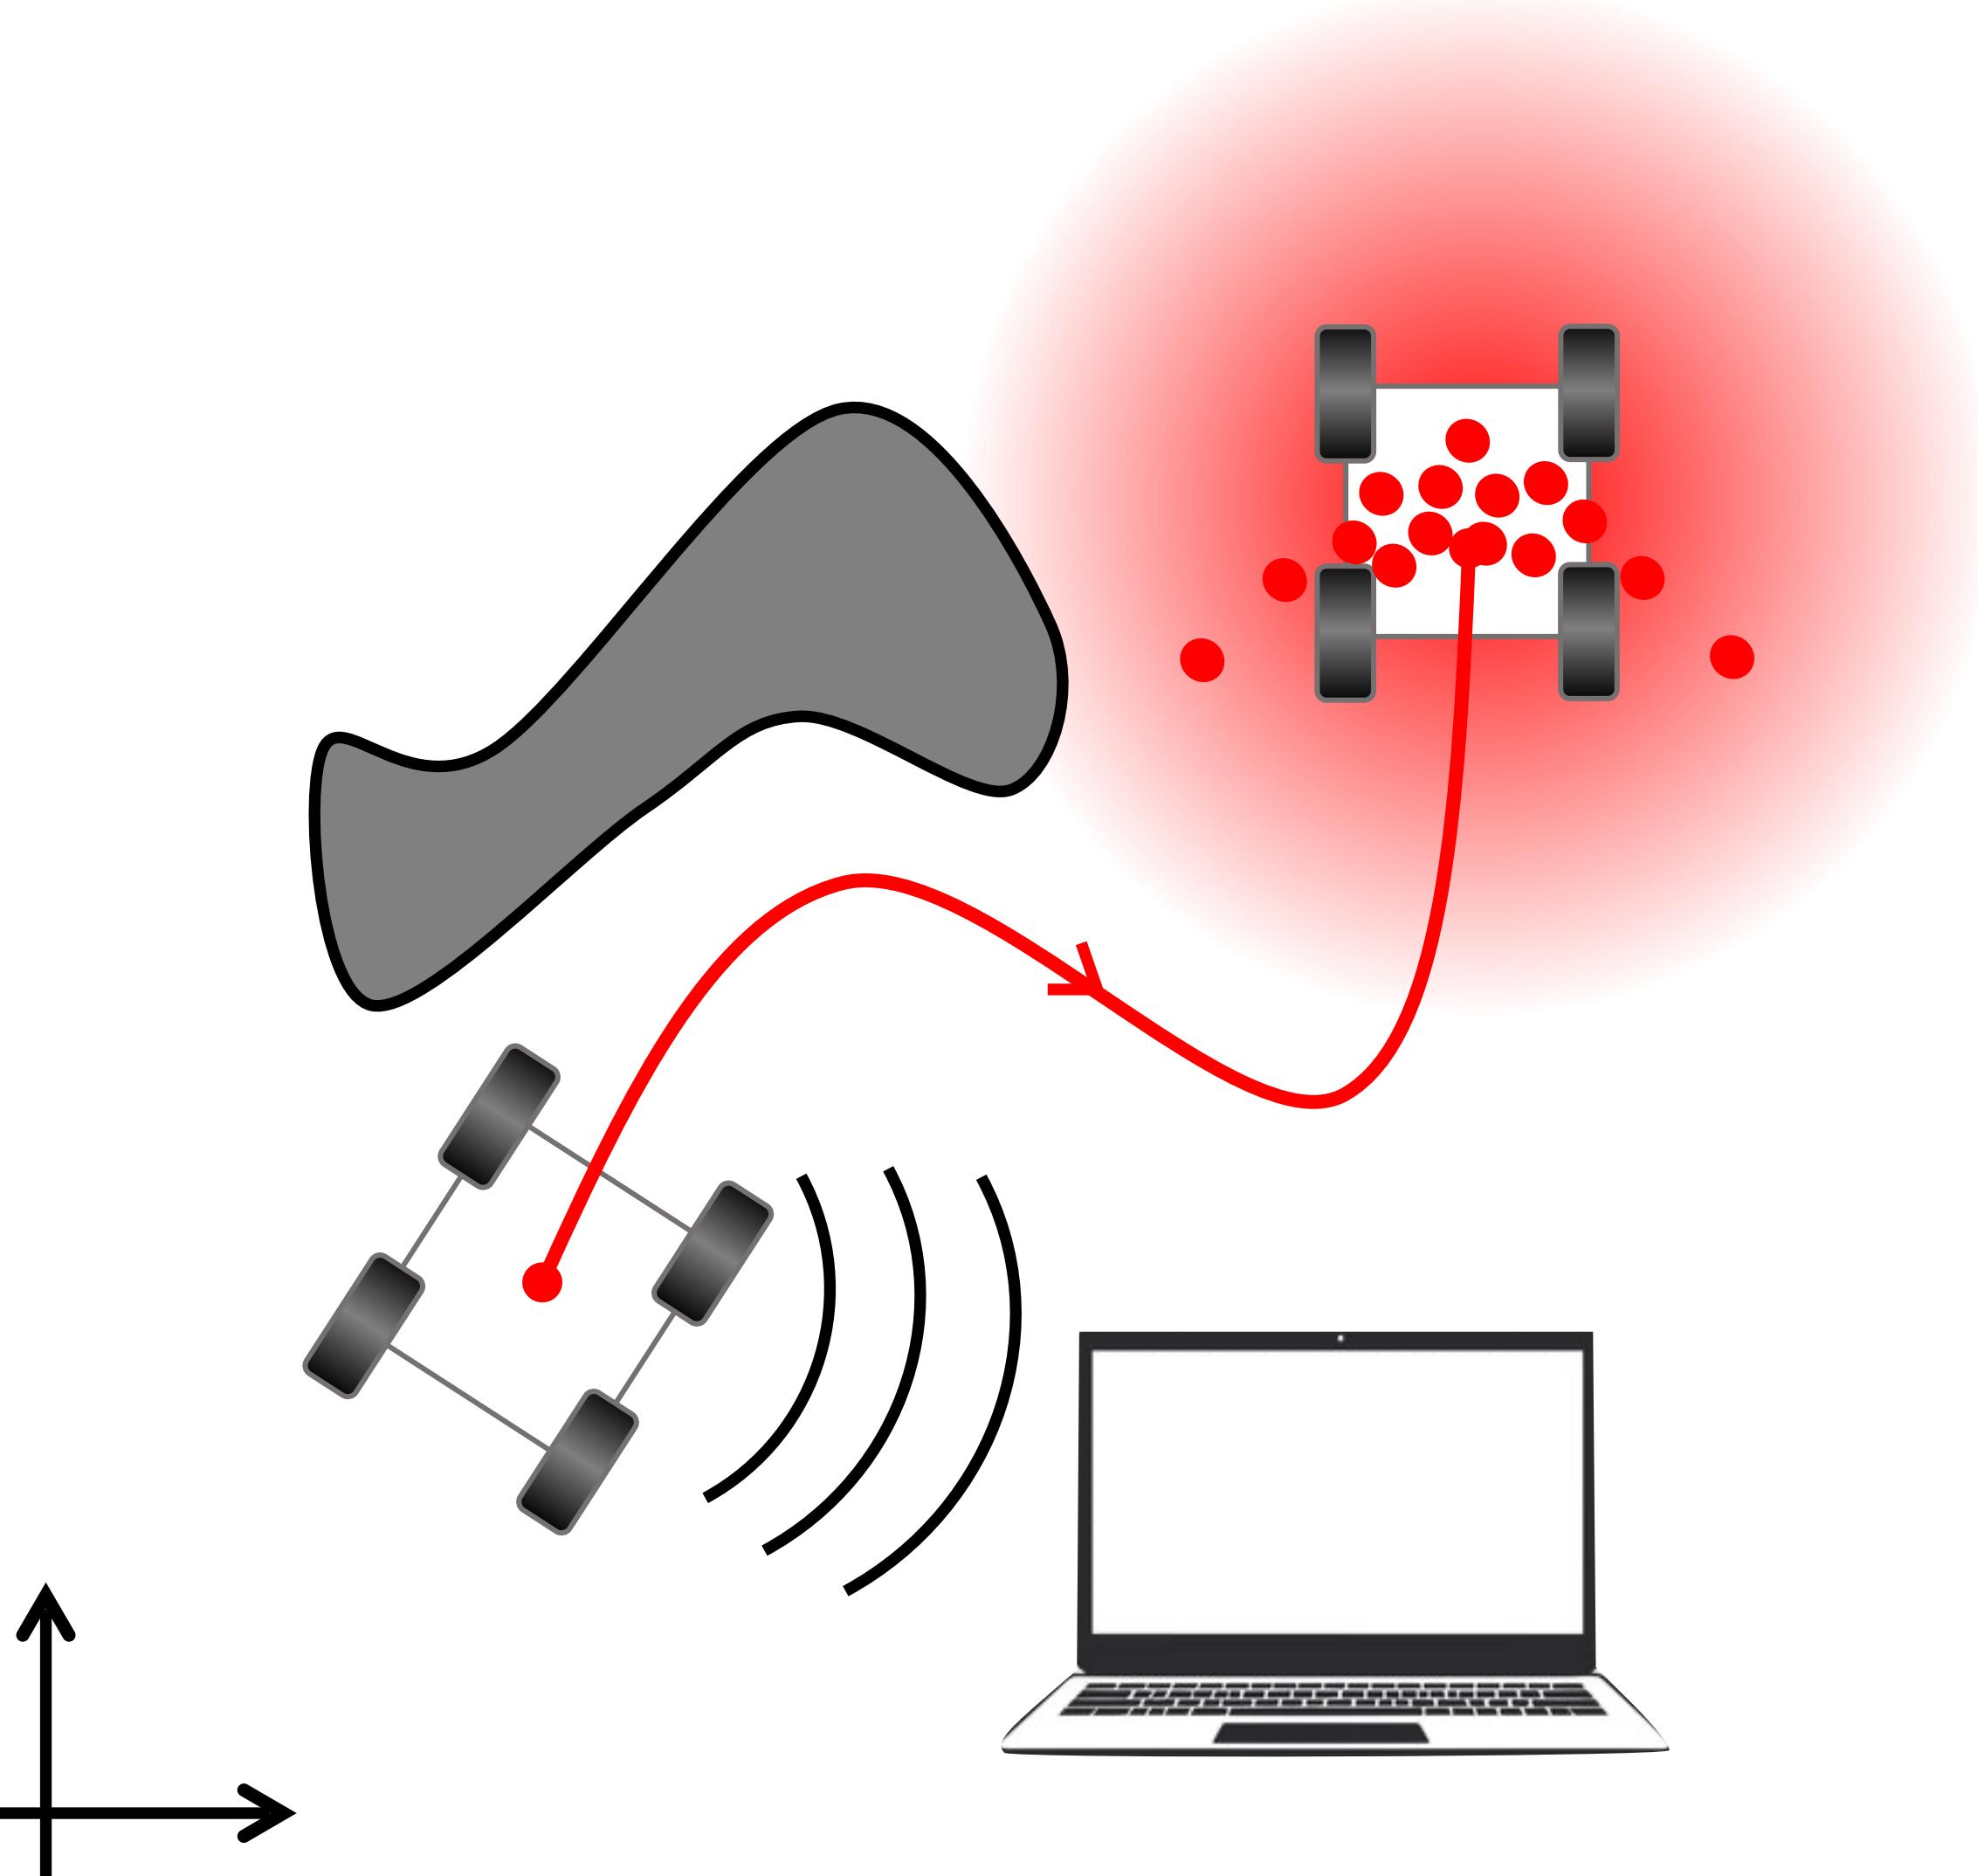
- Initial **robot pose is unknown**
- Need to estimate position from scratch
- A more difficult “**global**” problem, where you cannot assume boundedness in pose error



Next module on navigation

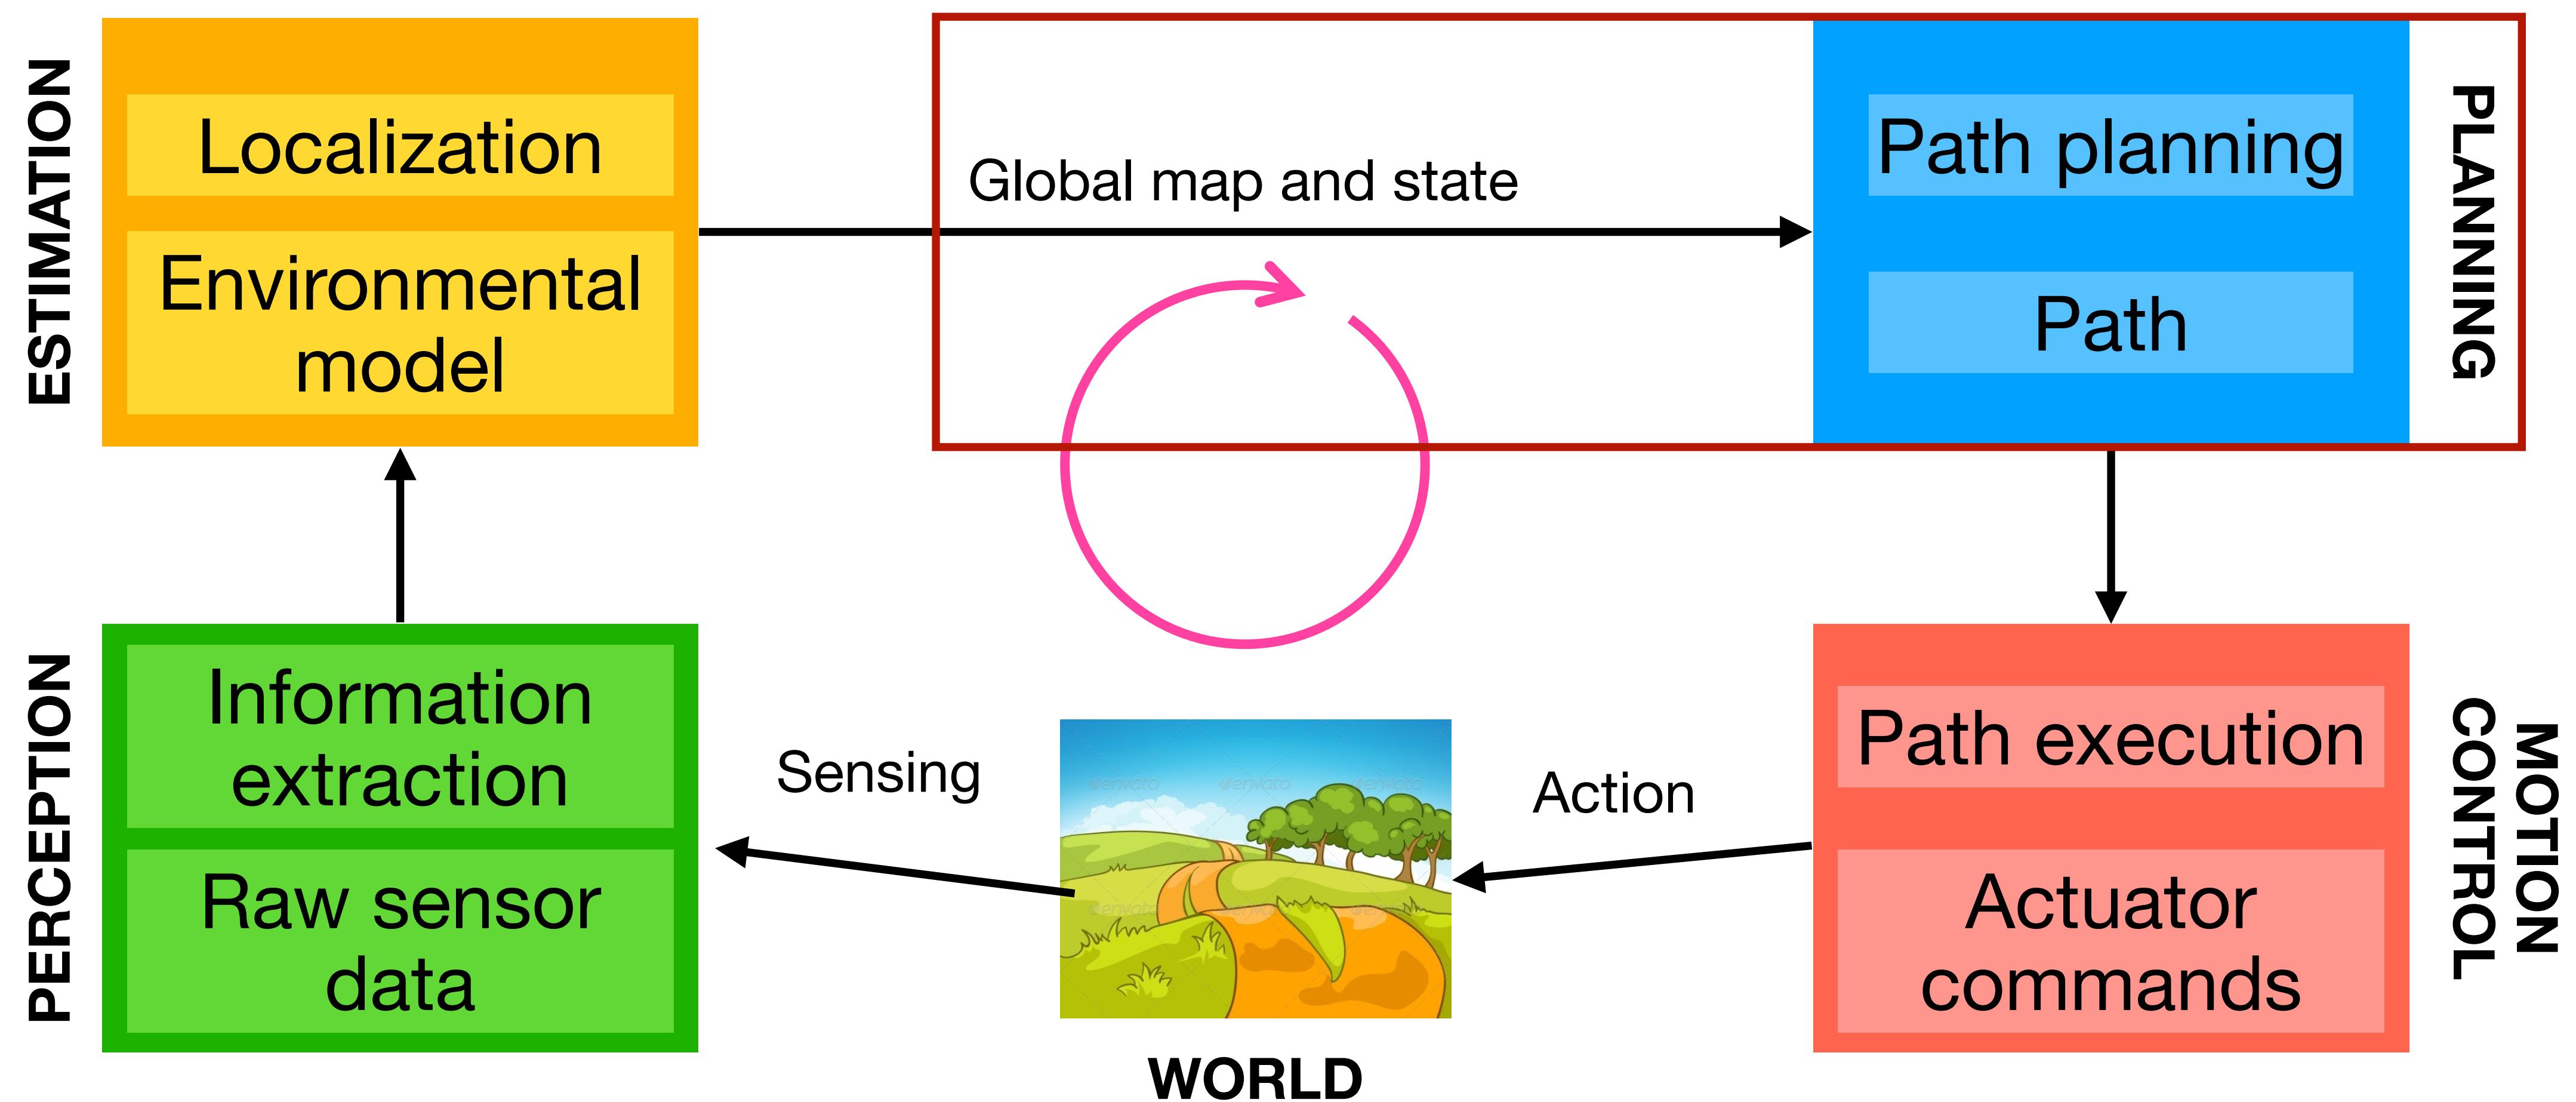
- Local planners
- Global localization and planning

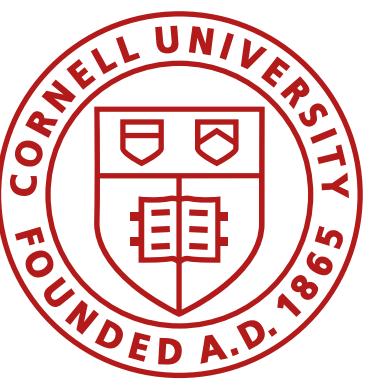
- Map representations
 - Continuous
 - Discrete
 - Topological
- Maps as graphs
 - Graph search algorithms
 - Breadth first search
 - Depth first search
 - Dijkstras
 - A*



Navigation

- Break the problem down: localization, map building, path planning



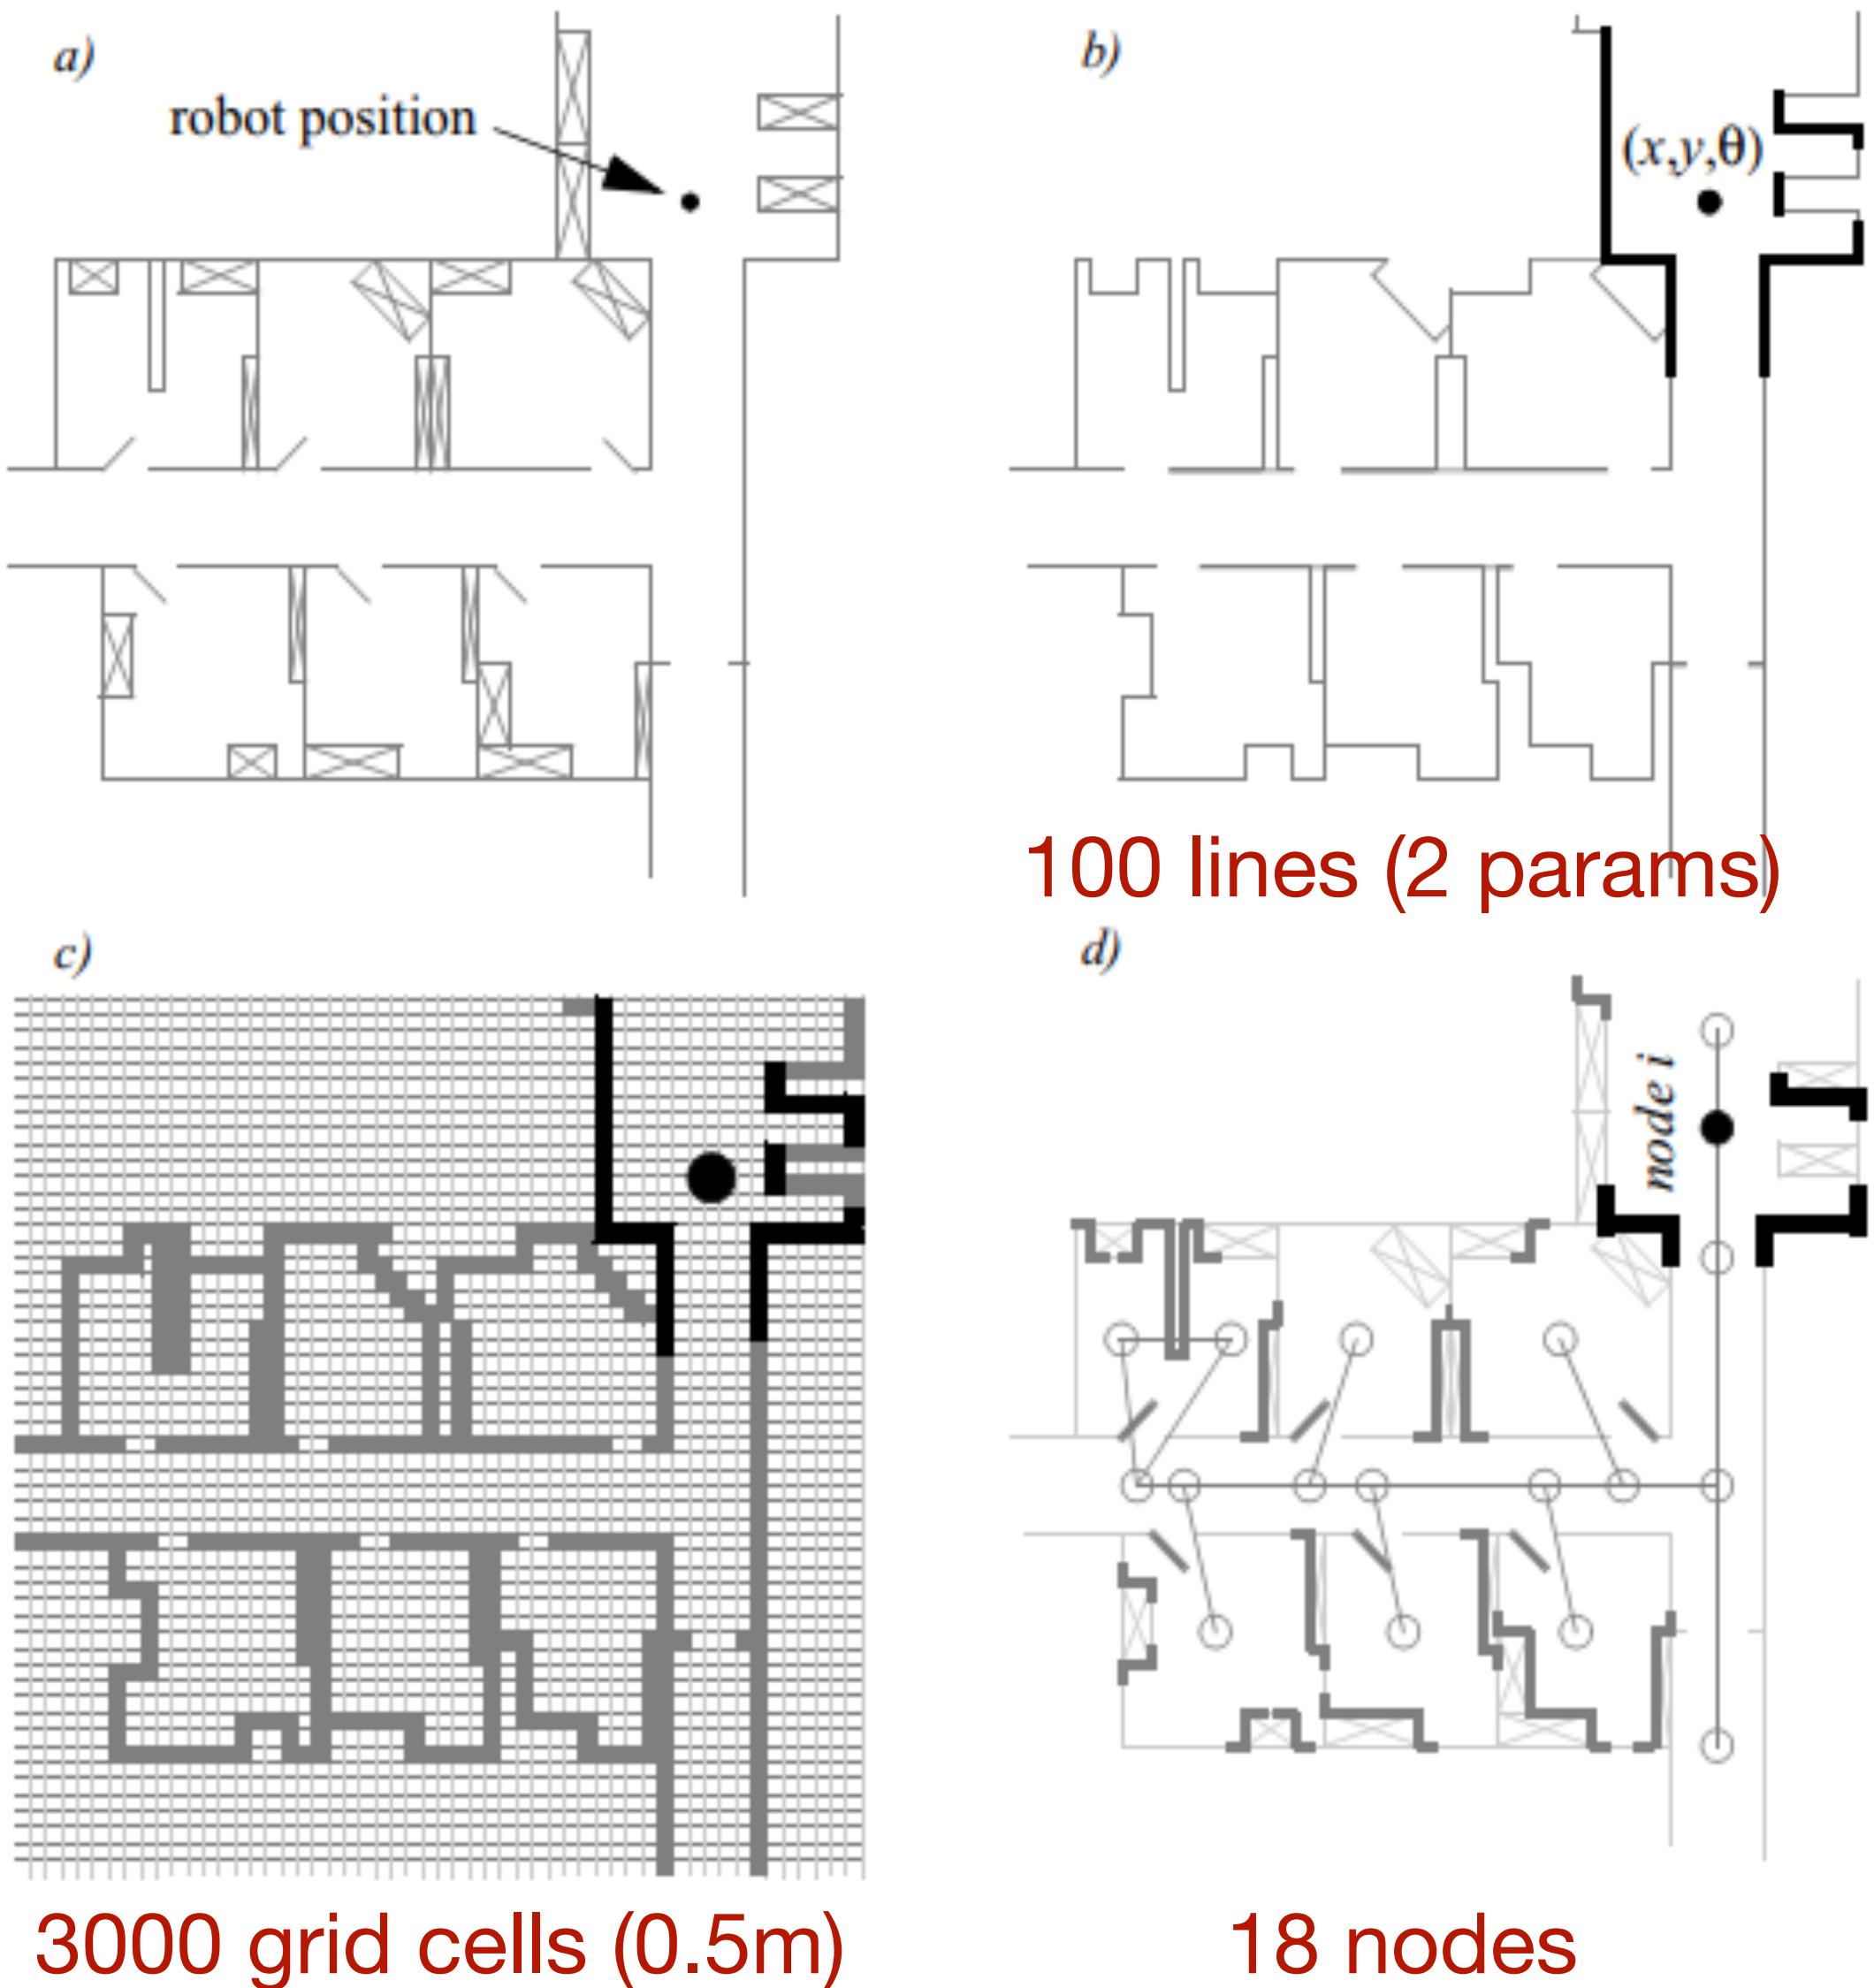


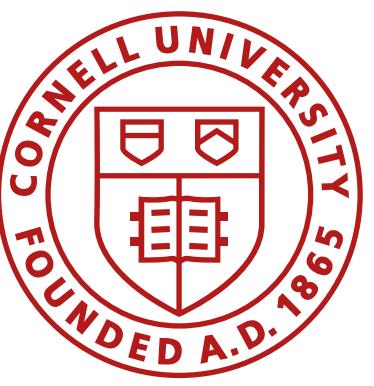
Map Representation

- a) Building Plan
- b) Line-based map
- c) Occupancy grid-based map
- d) Topological map

Important Properties

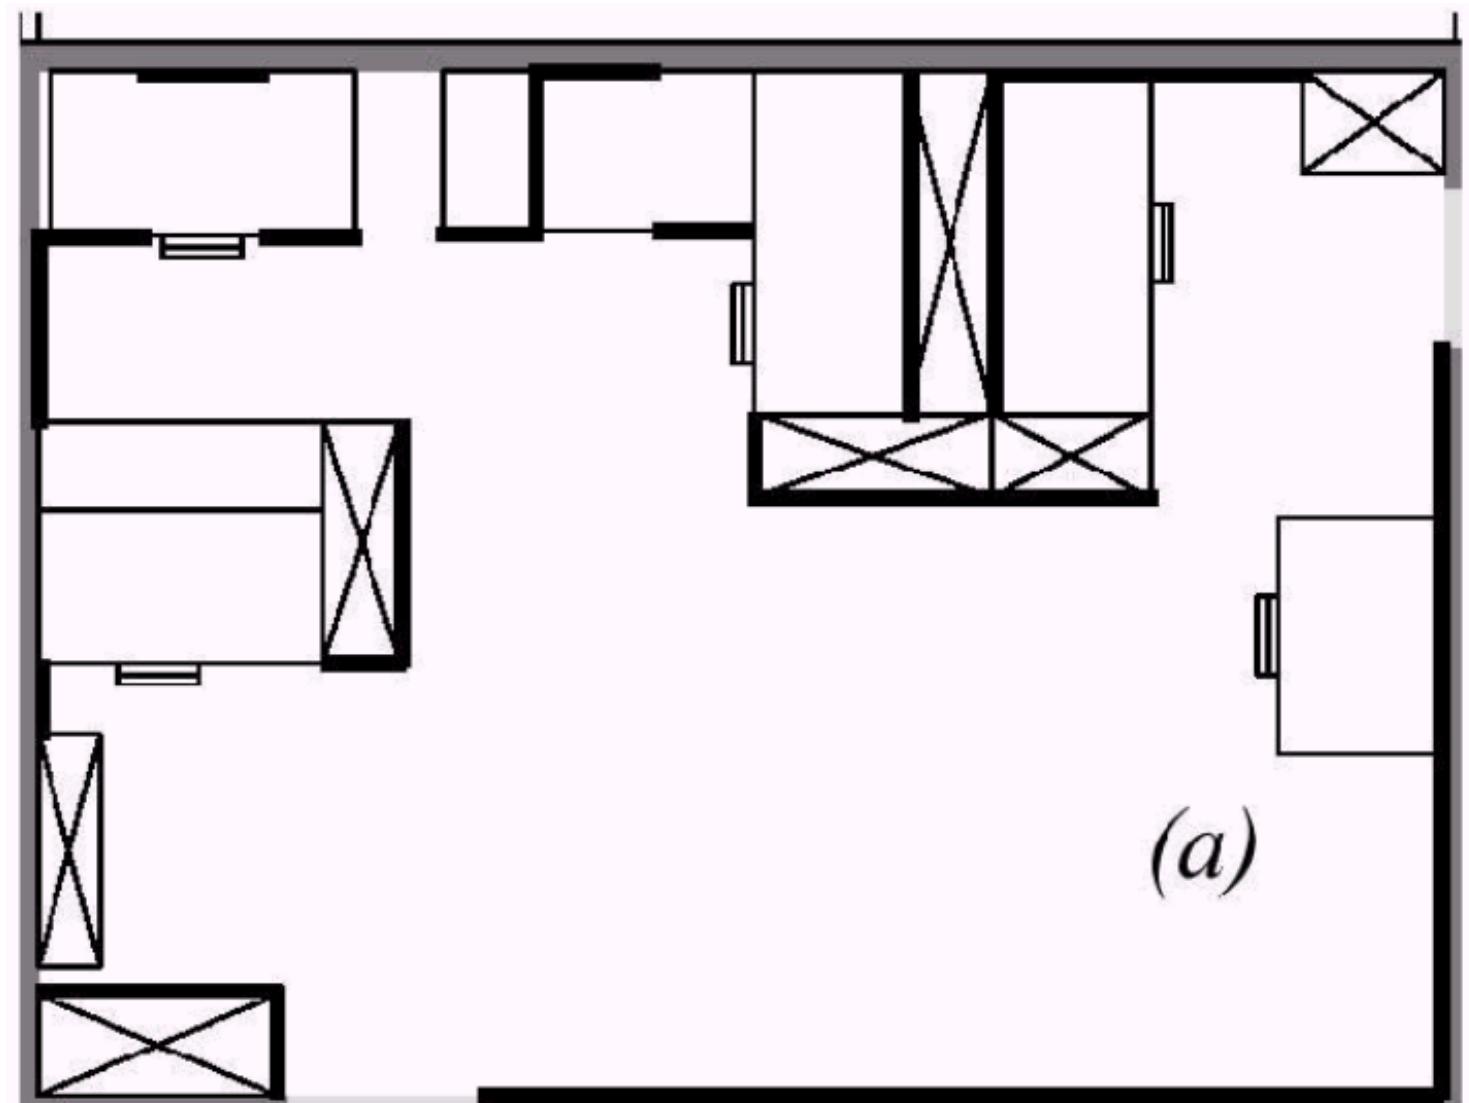
- Memory allocation
- Computation



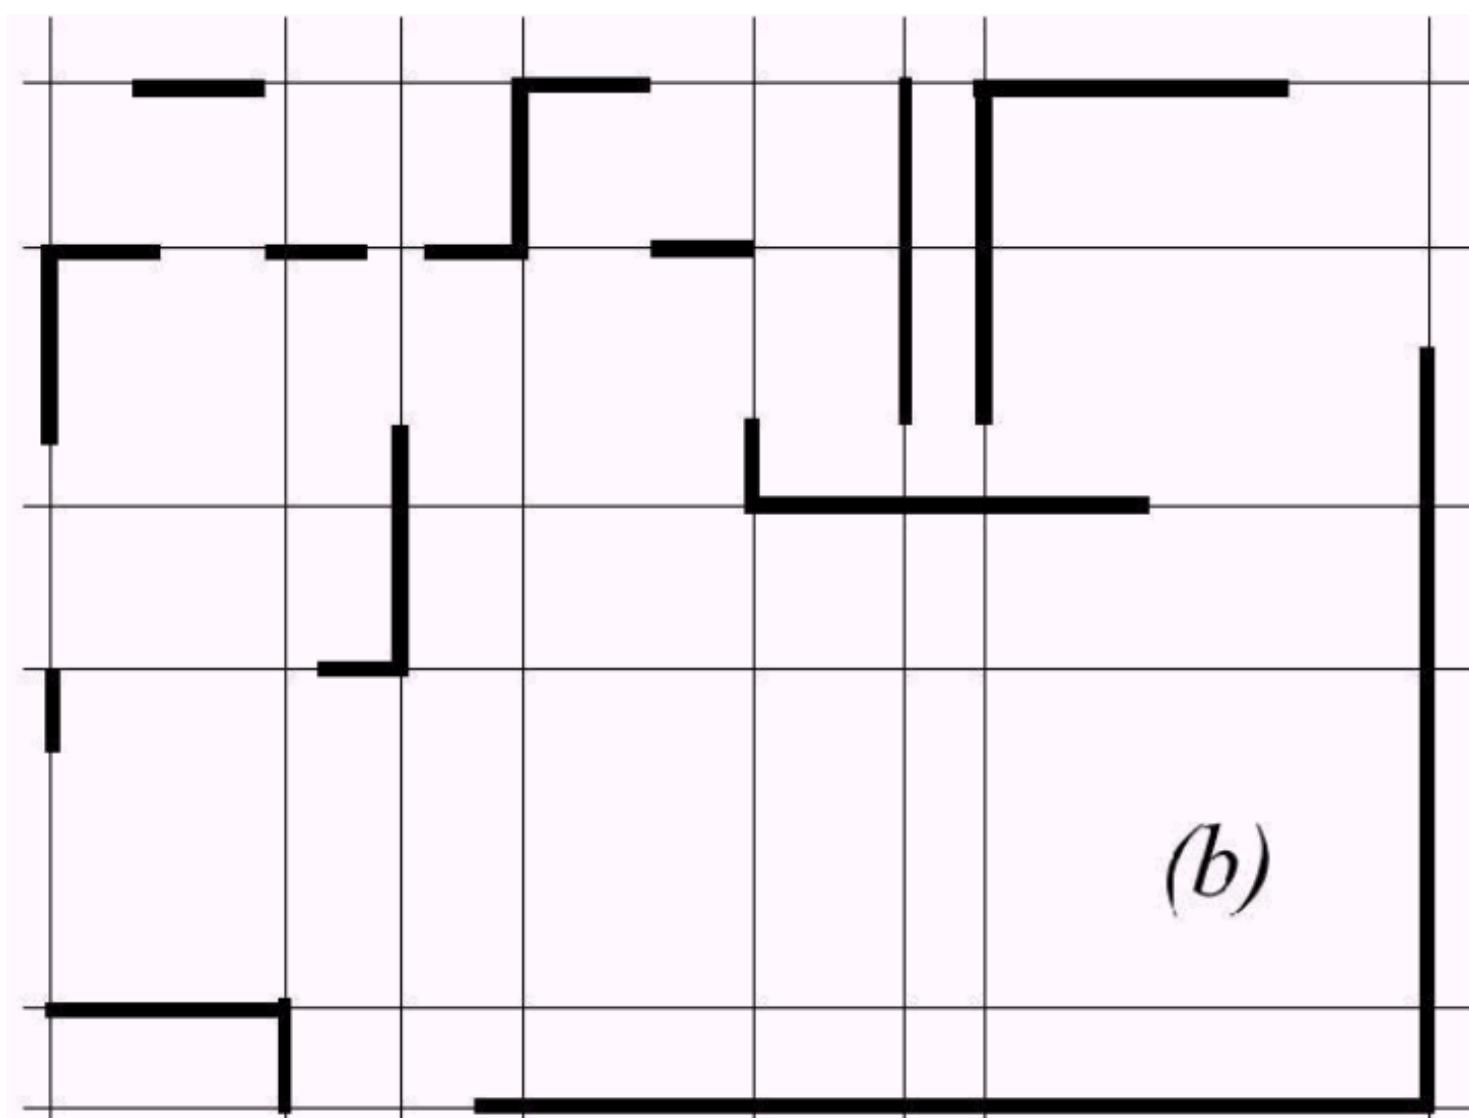


Continuous Representations

- Exact decomposition of the environment
- Used mainly in 2D representations
- Closed-world assumption
- Storage proportional to object density
- Example: Continuous line representations
 - Using range finders, we can extract lines/ line segments in the environment



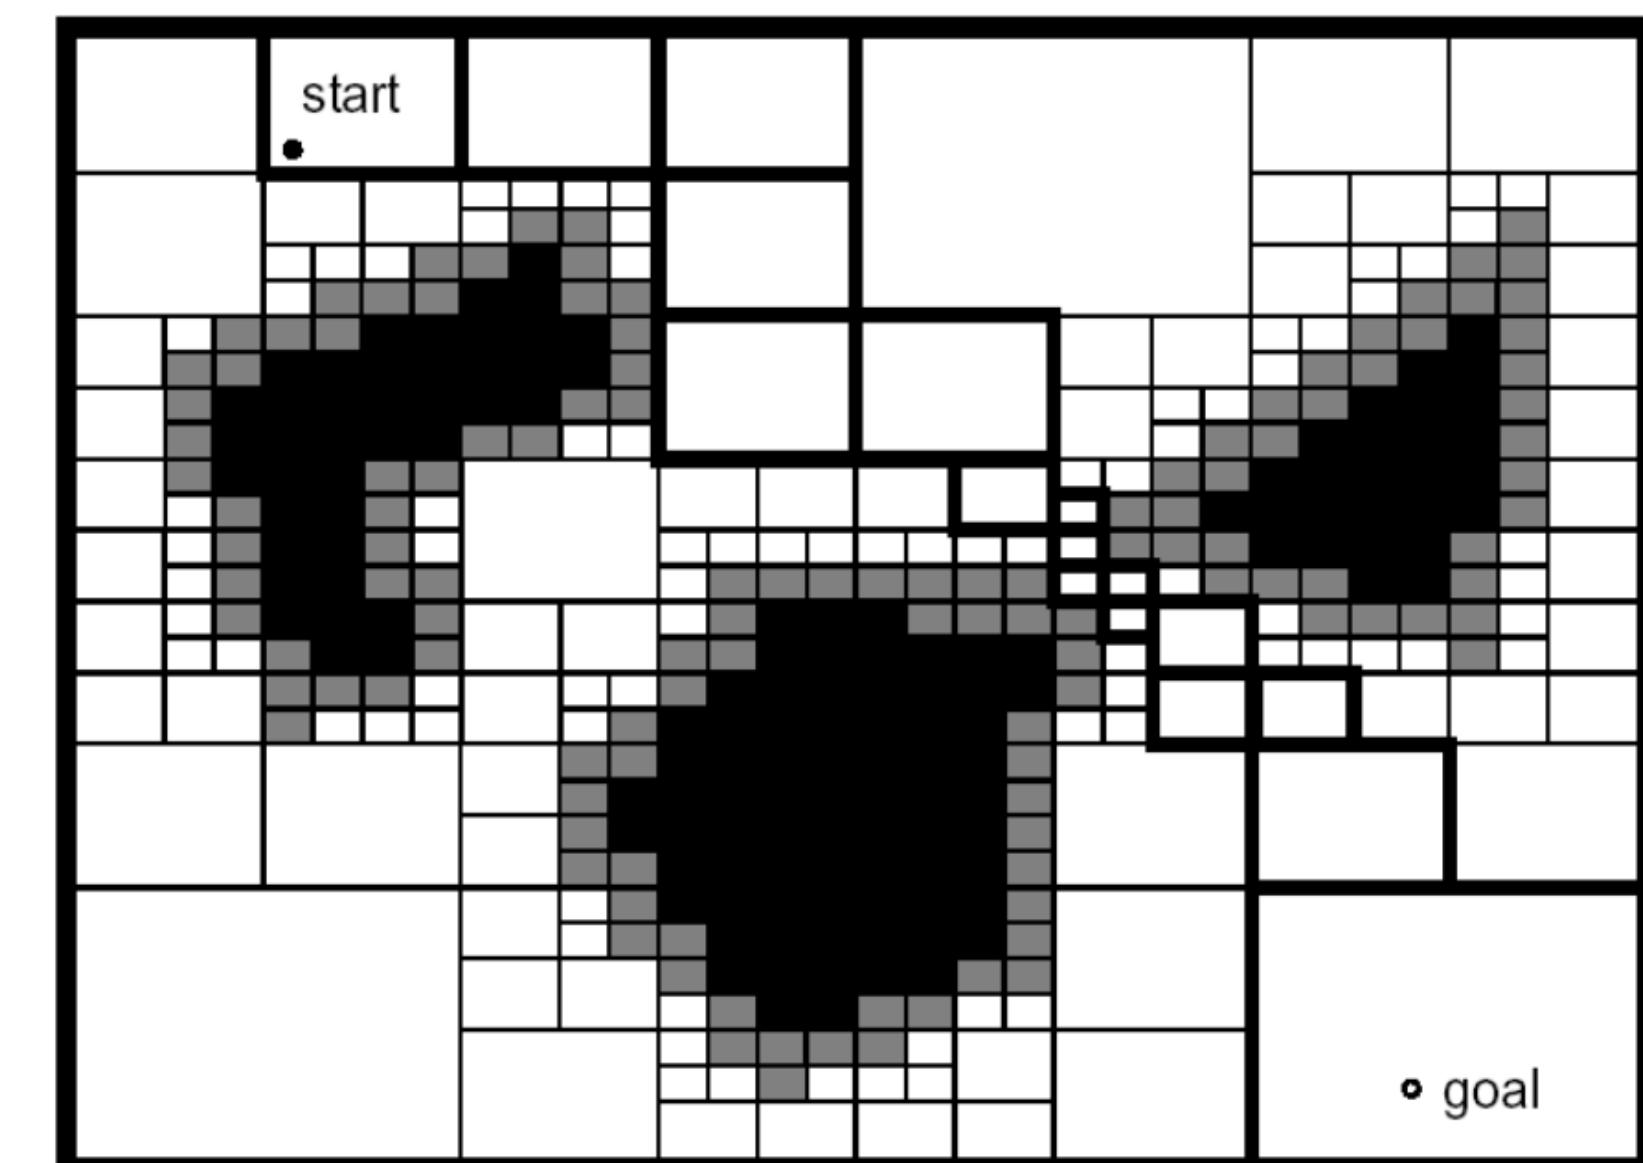
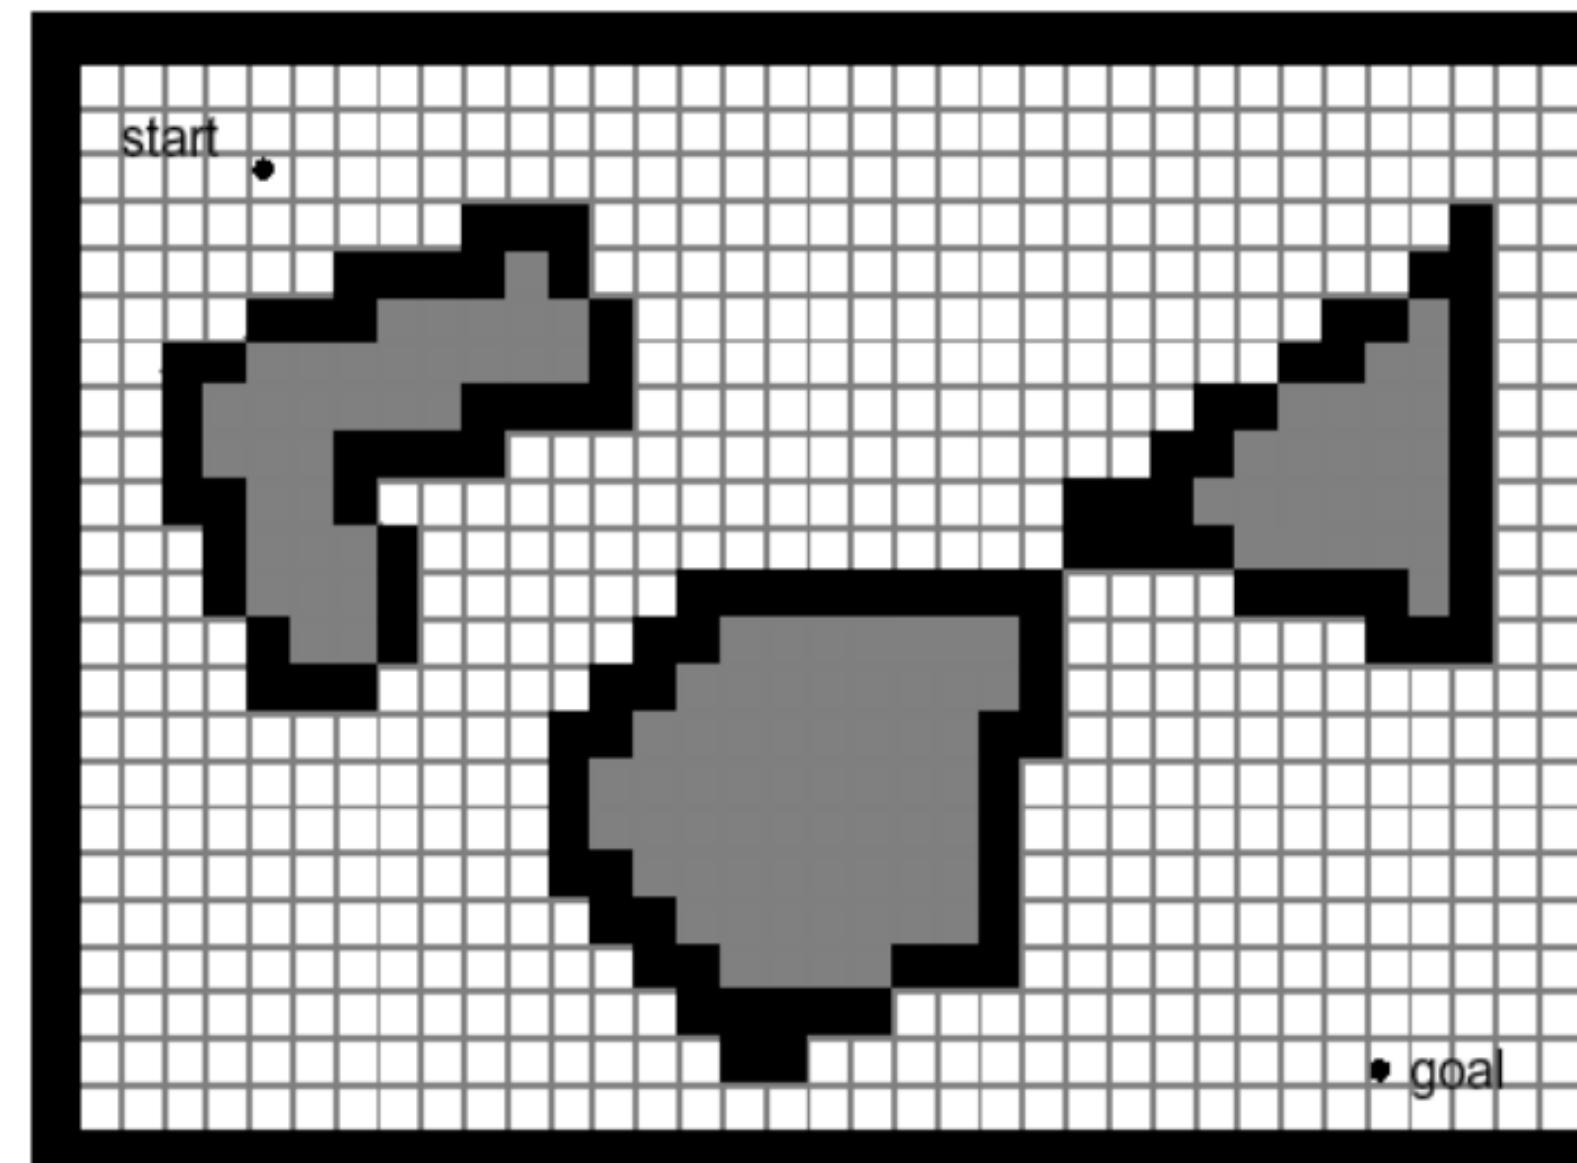
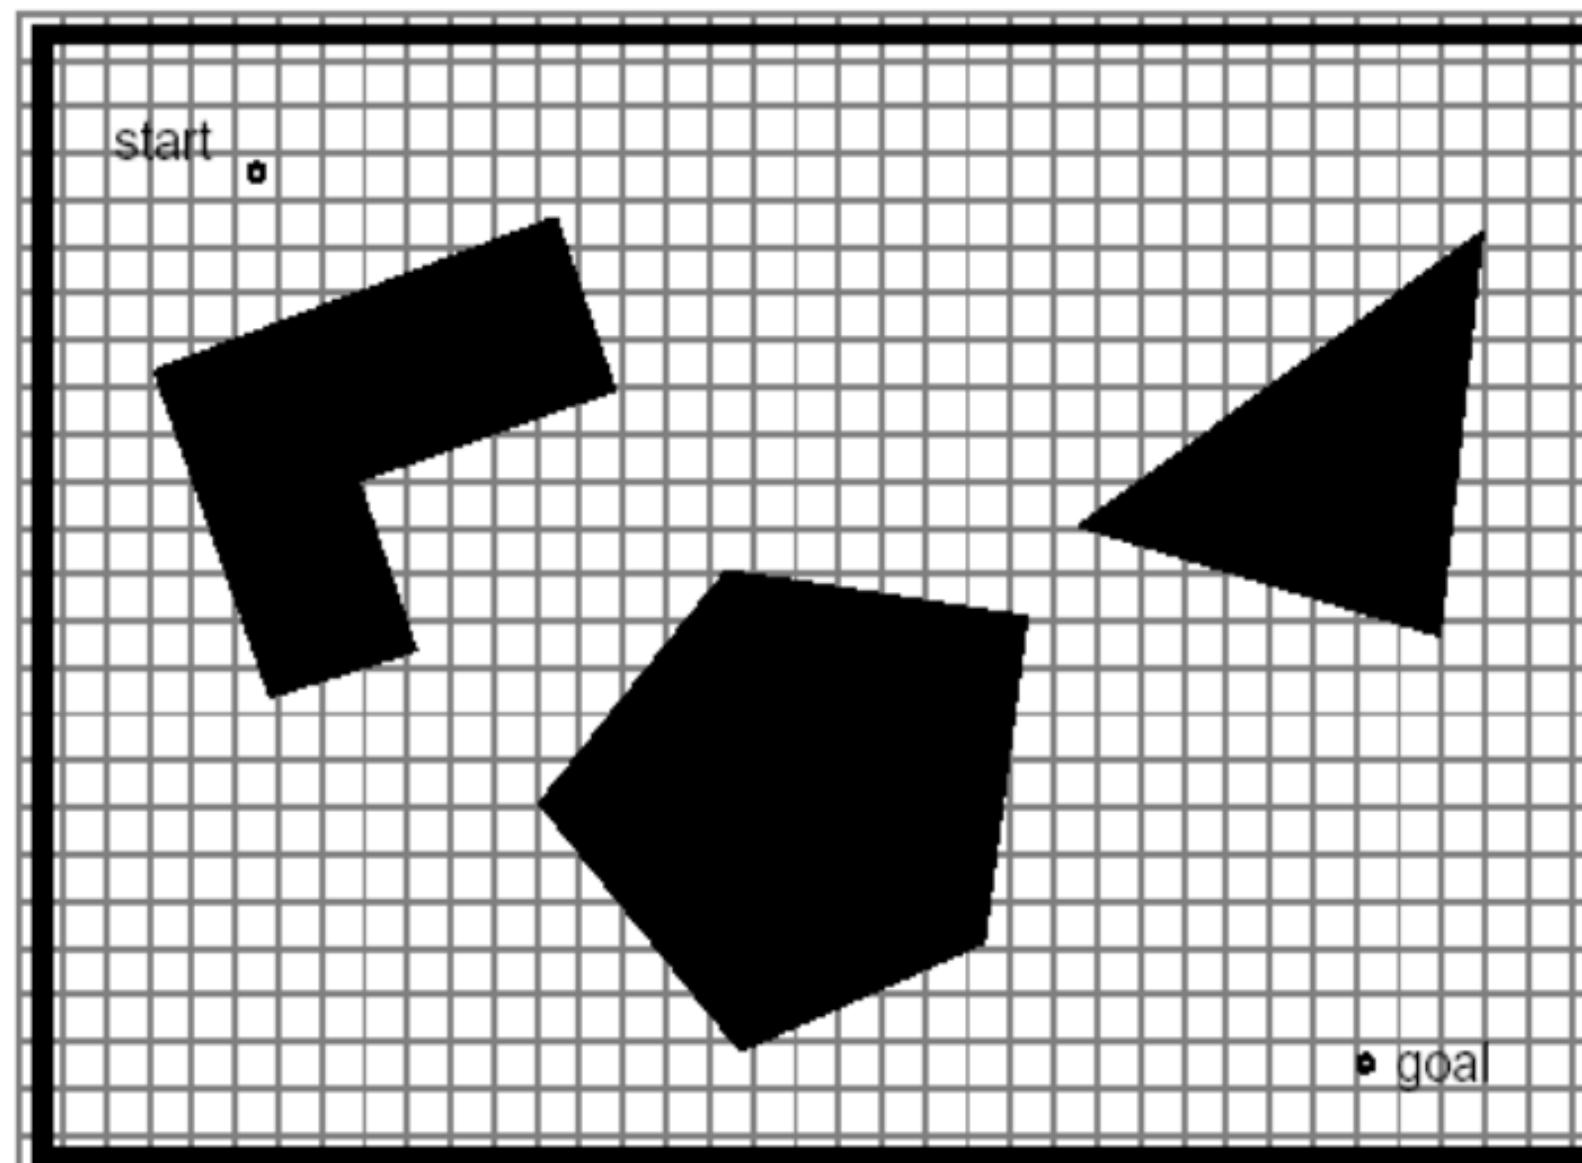
(a)

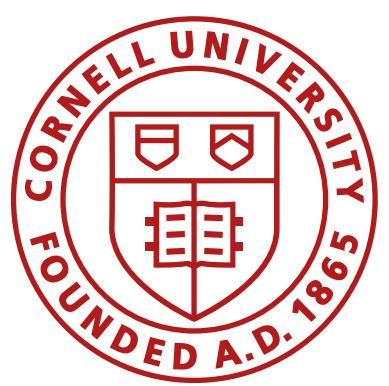


(b)

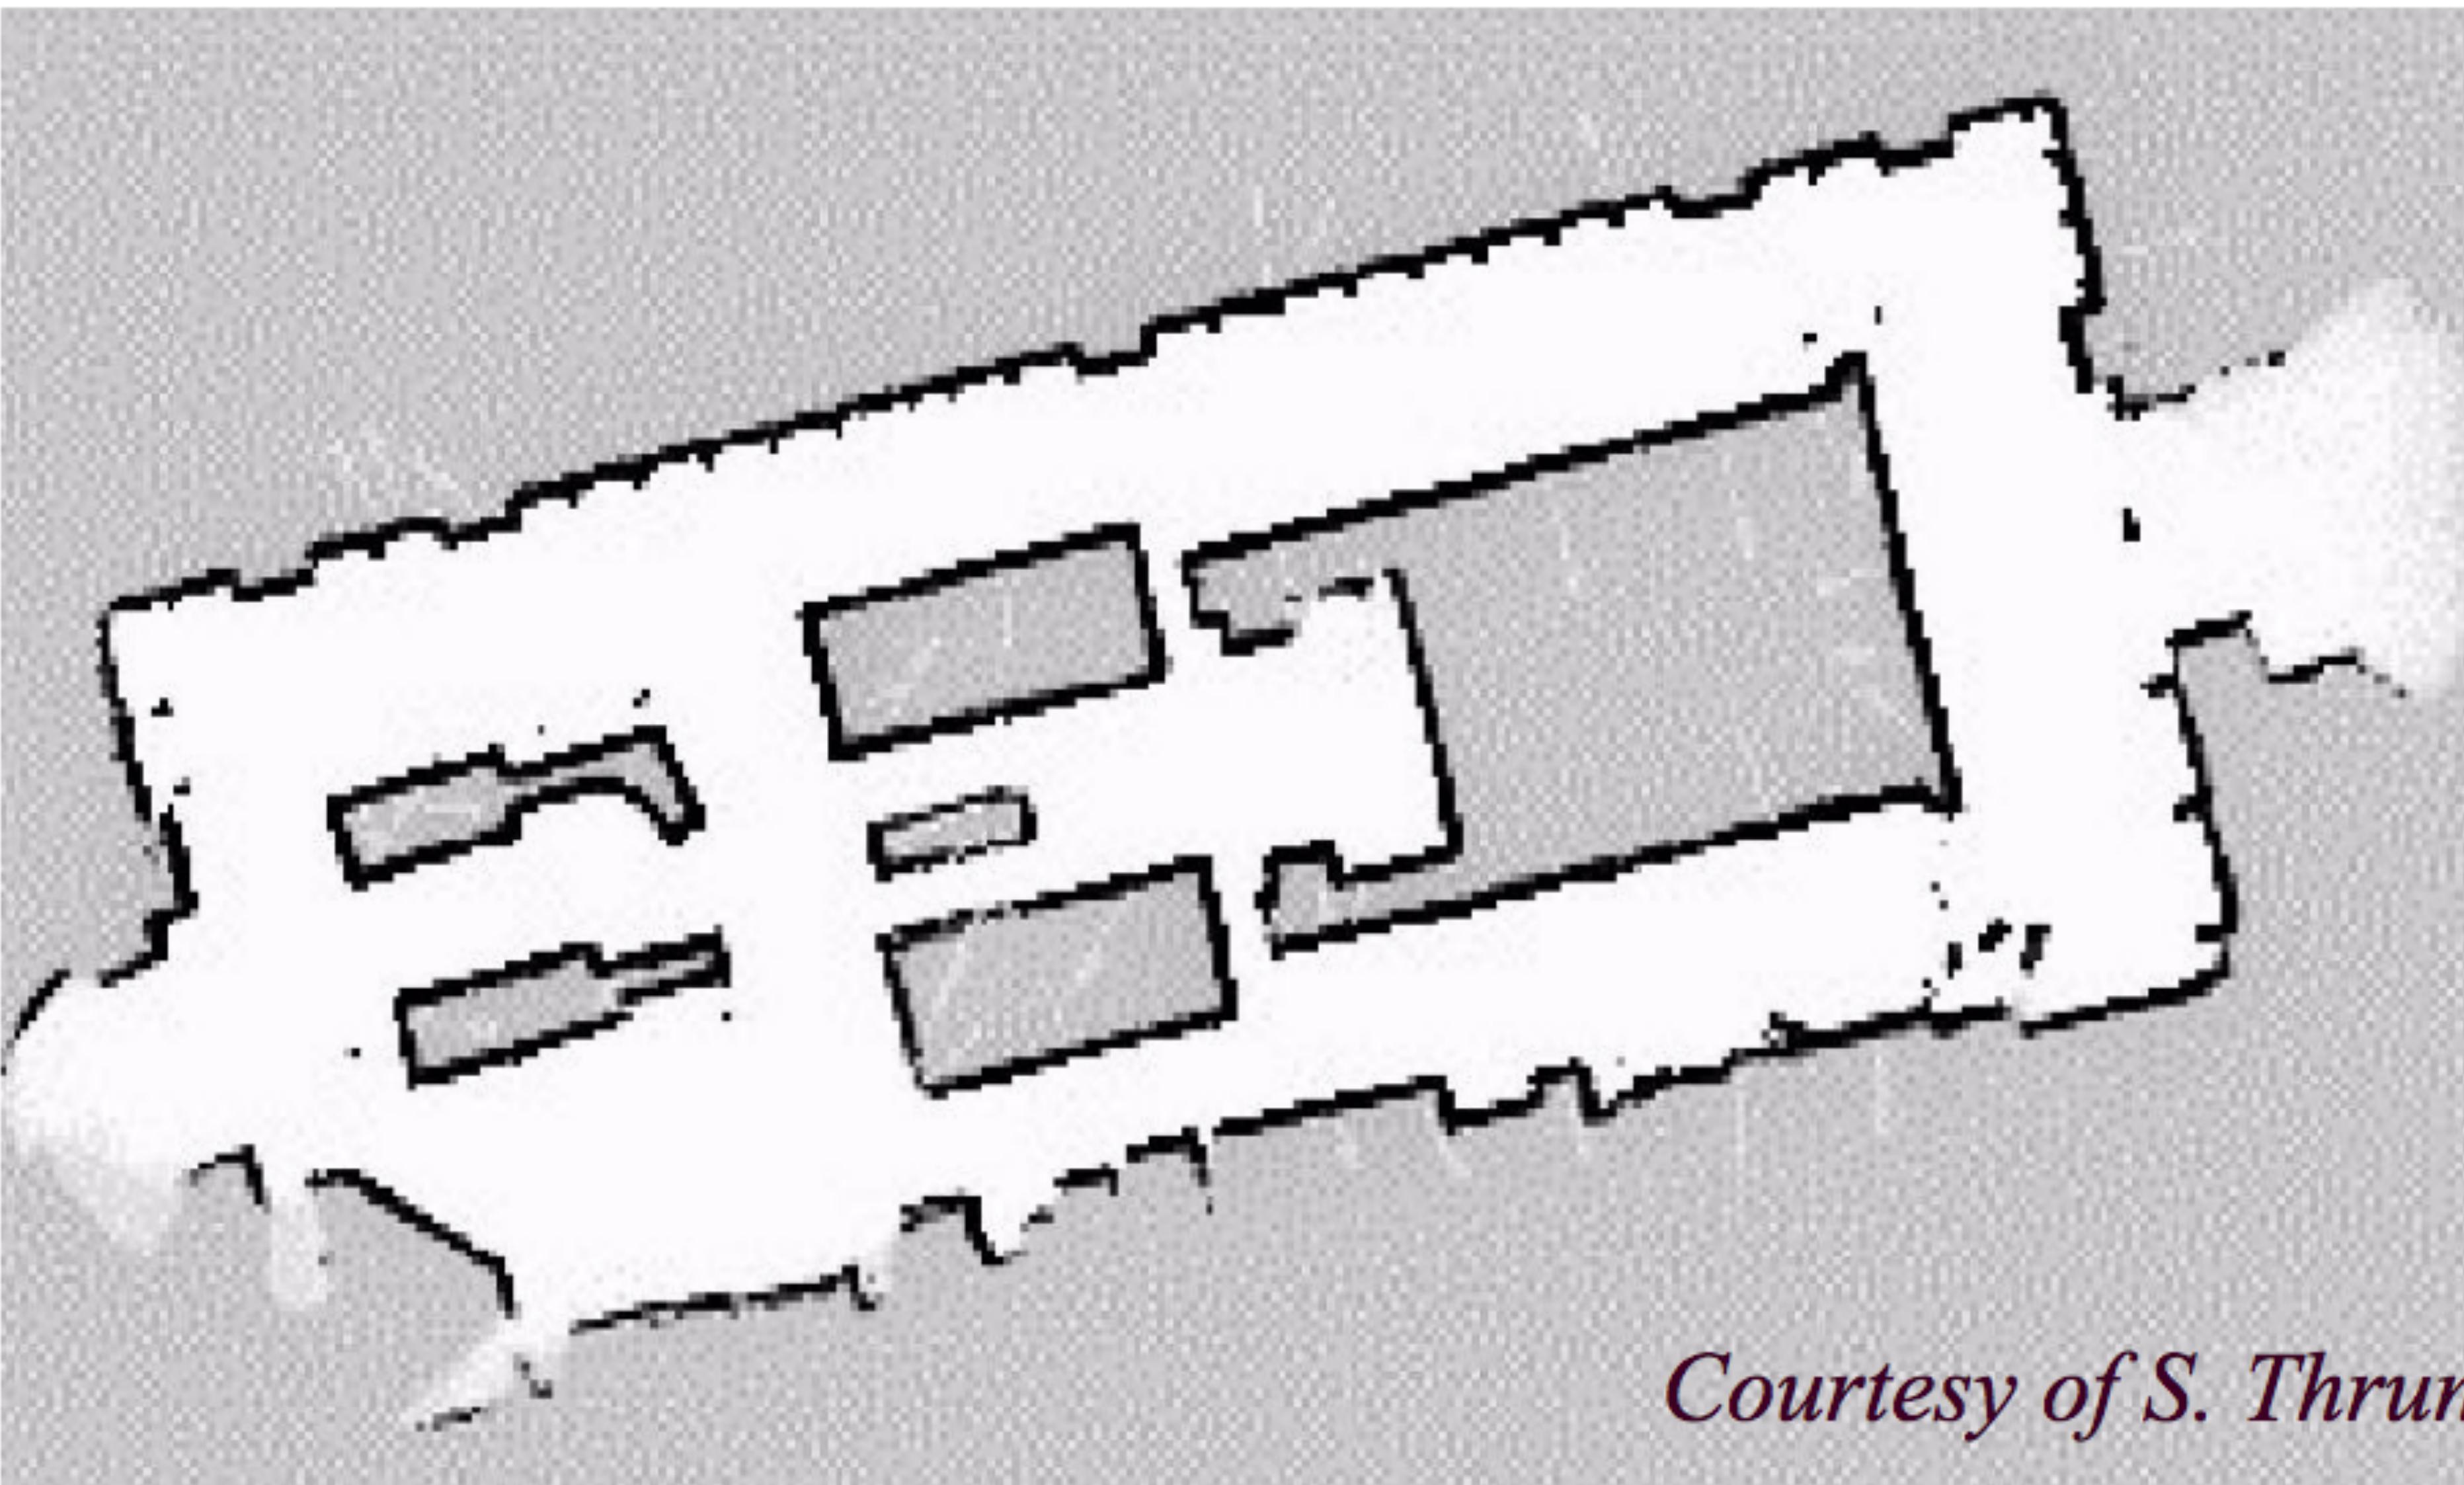
Cell Decomposition

- Fixed: Tesselate the world at a fixed resolution
- Approximate features given the resolution
- Most commonly used: occupancy grid
- Adaptive: Tesselate the world at varying resolutions, finer near objects





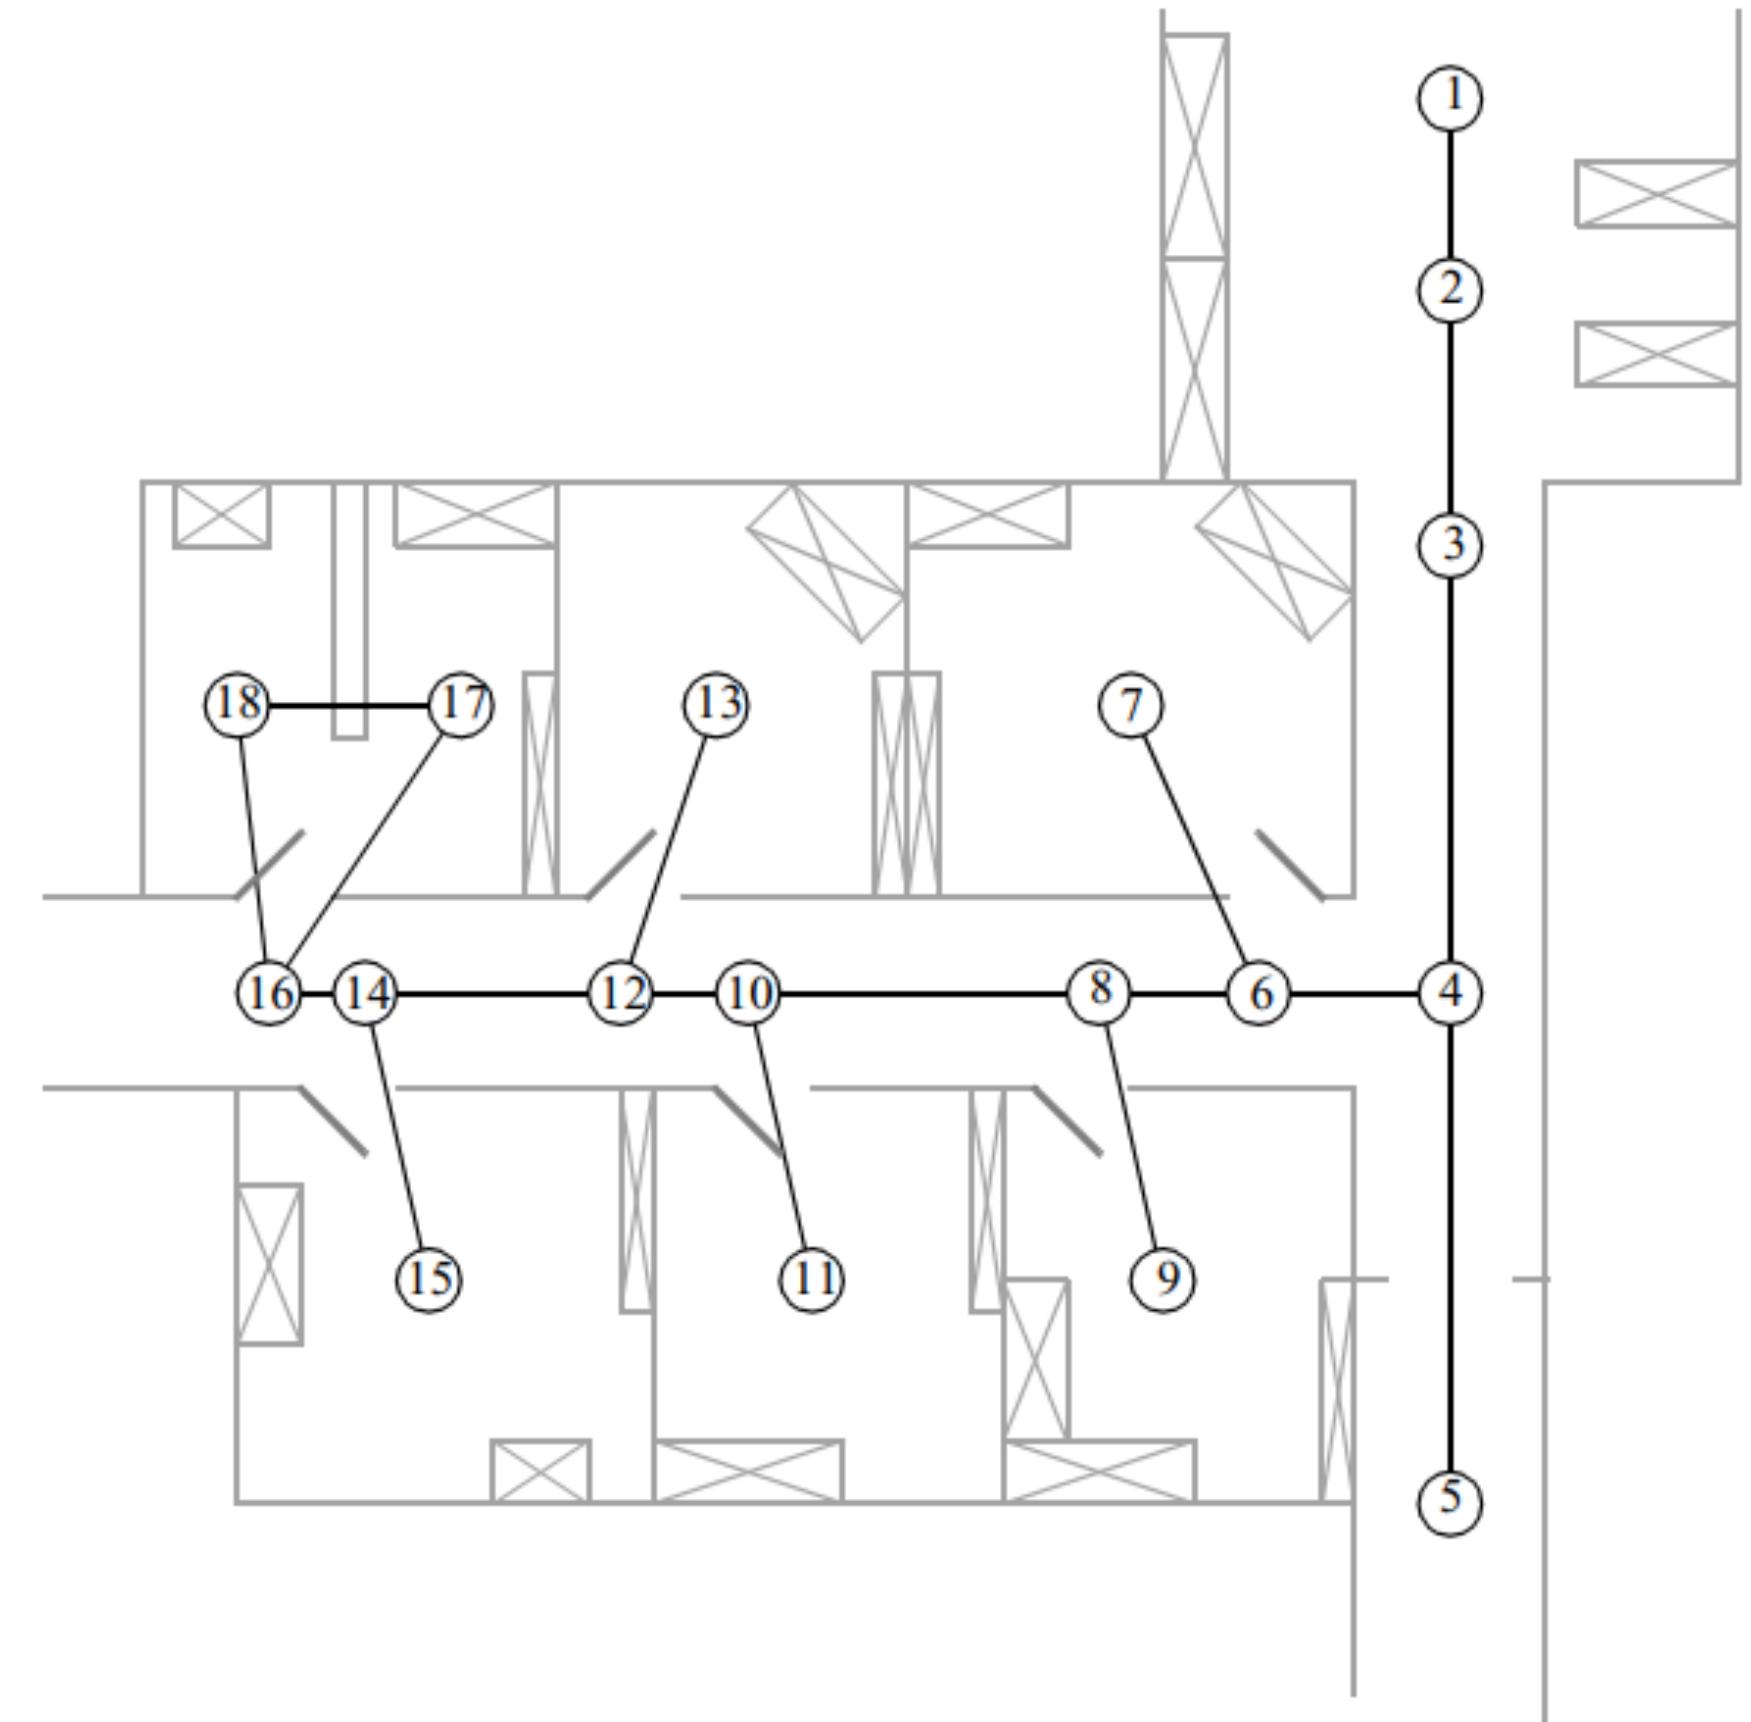
Fixed Decomposition



Courtesy of S. Thrun

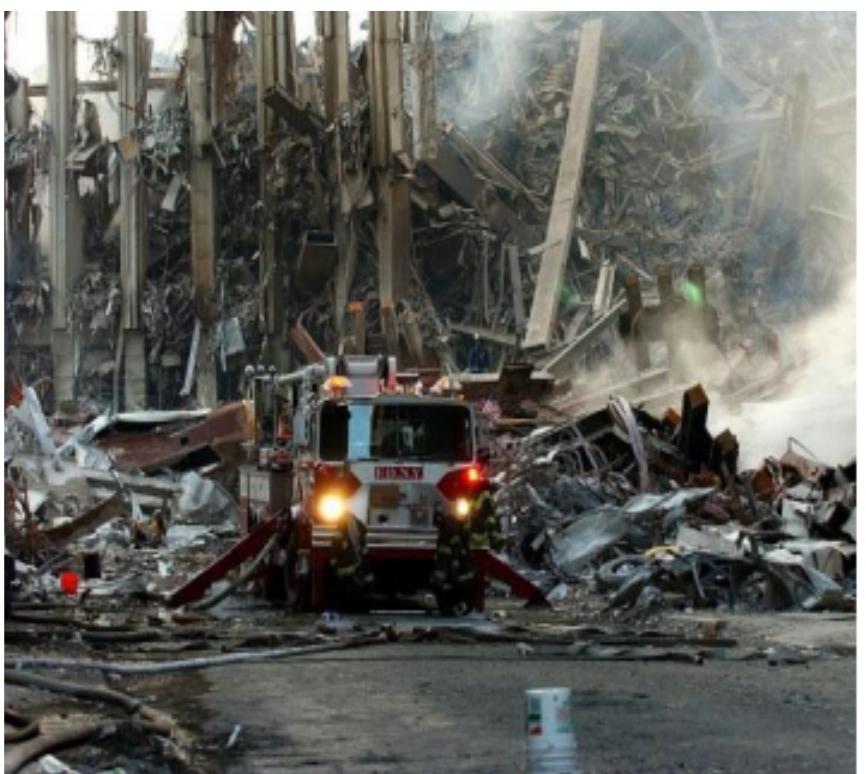
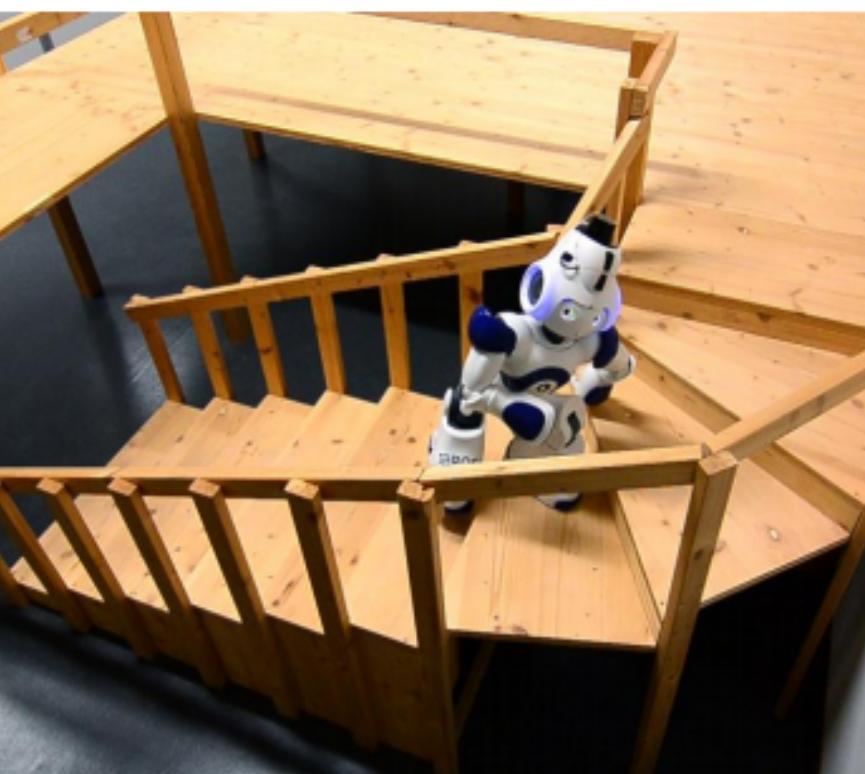
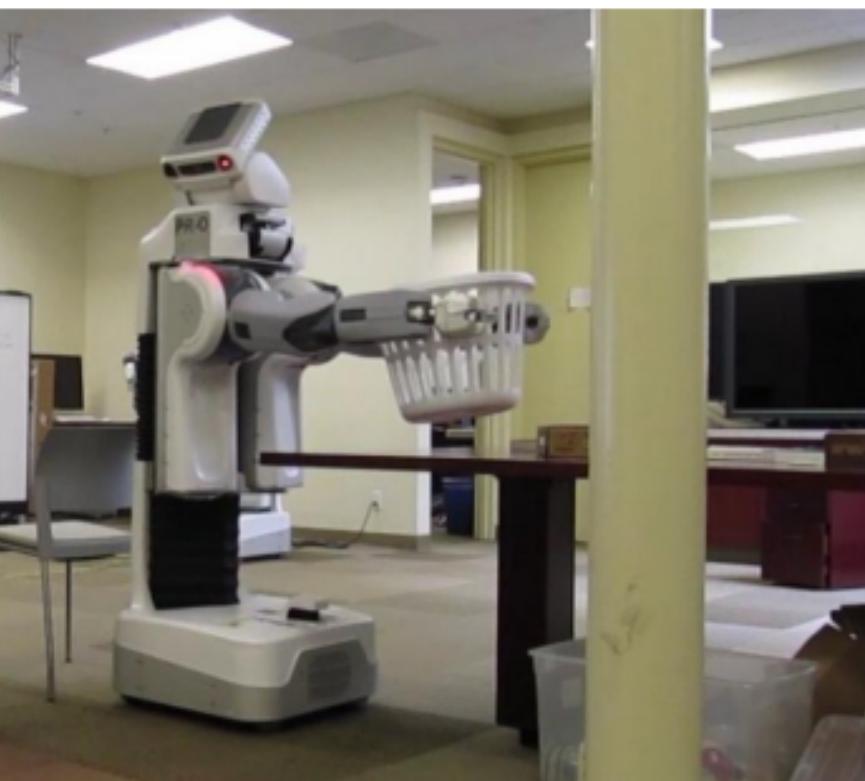
Topological decomposition

- A topological representation is a graph that specifies nodes and edges
 - Nodes denotes areas in the environment
 - Edges describe environment connectivity
- Robots can...
 - ...detect their current position in terms of the nodes of the topological graph
 - ...travel between nodes using robot motion



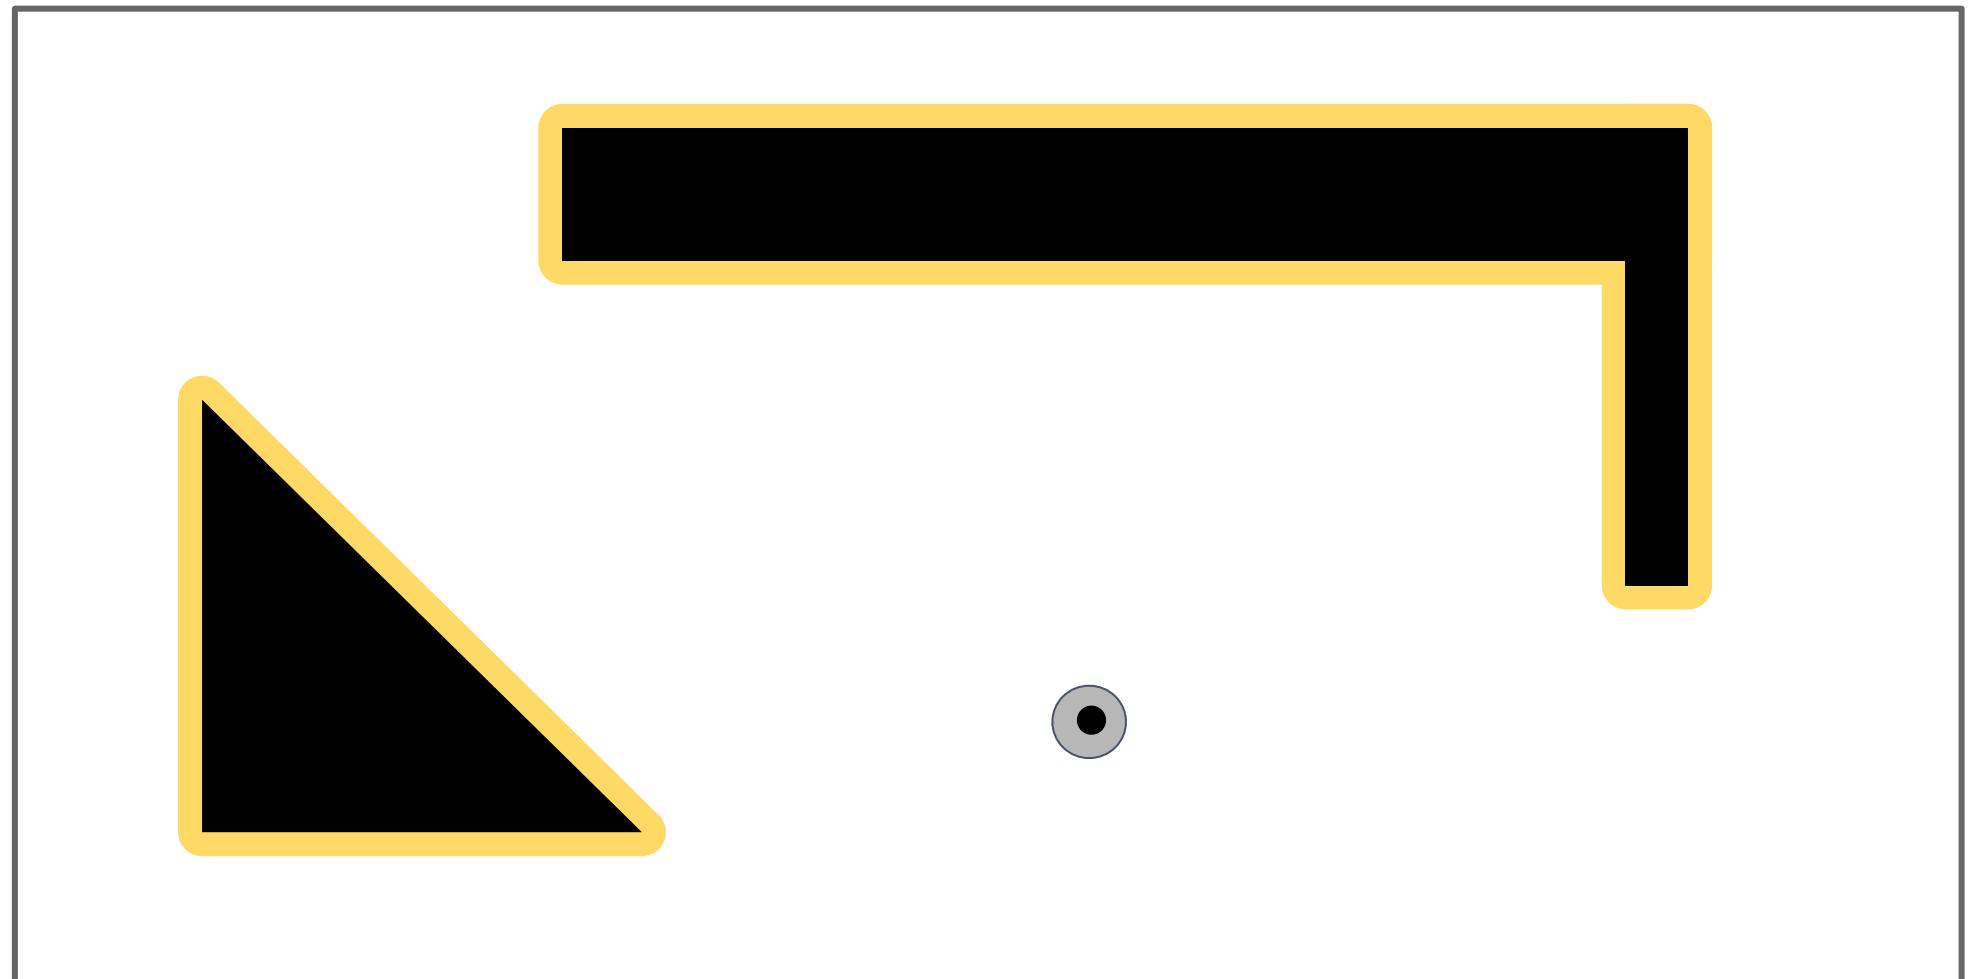
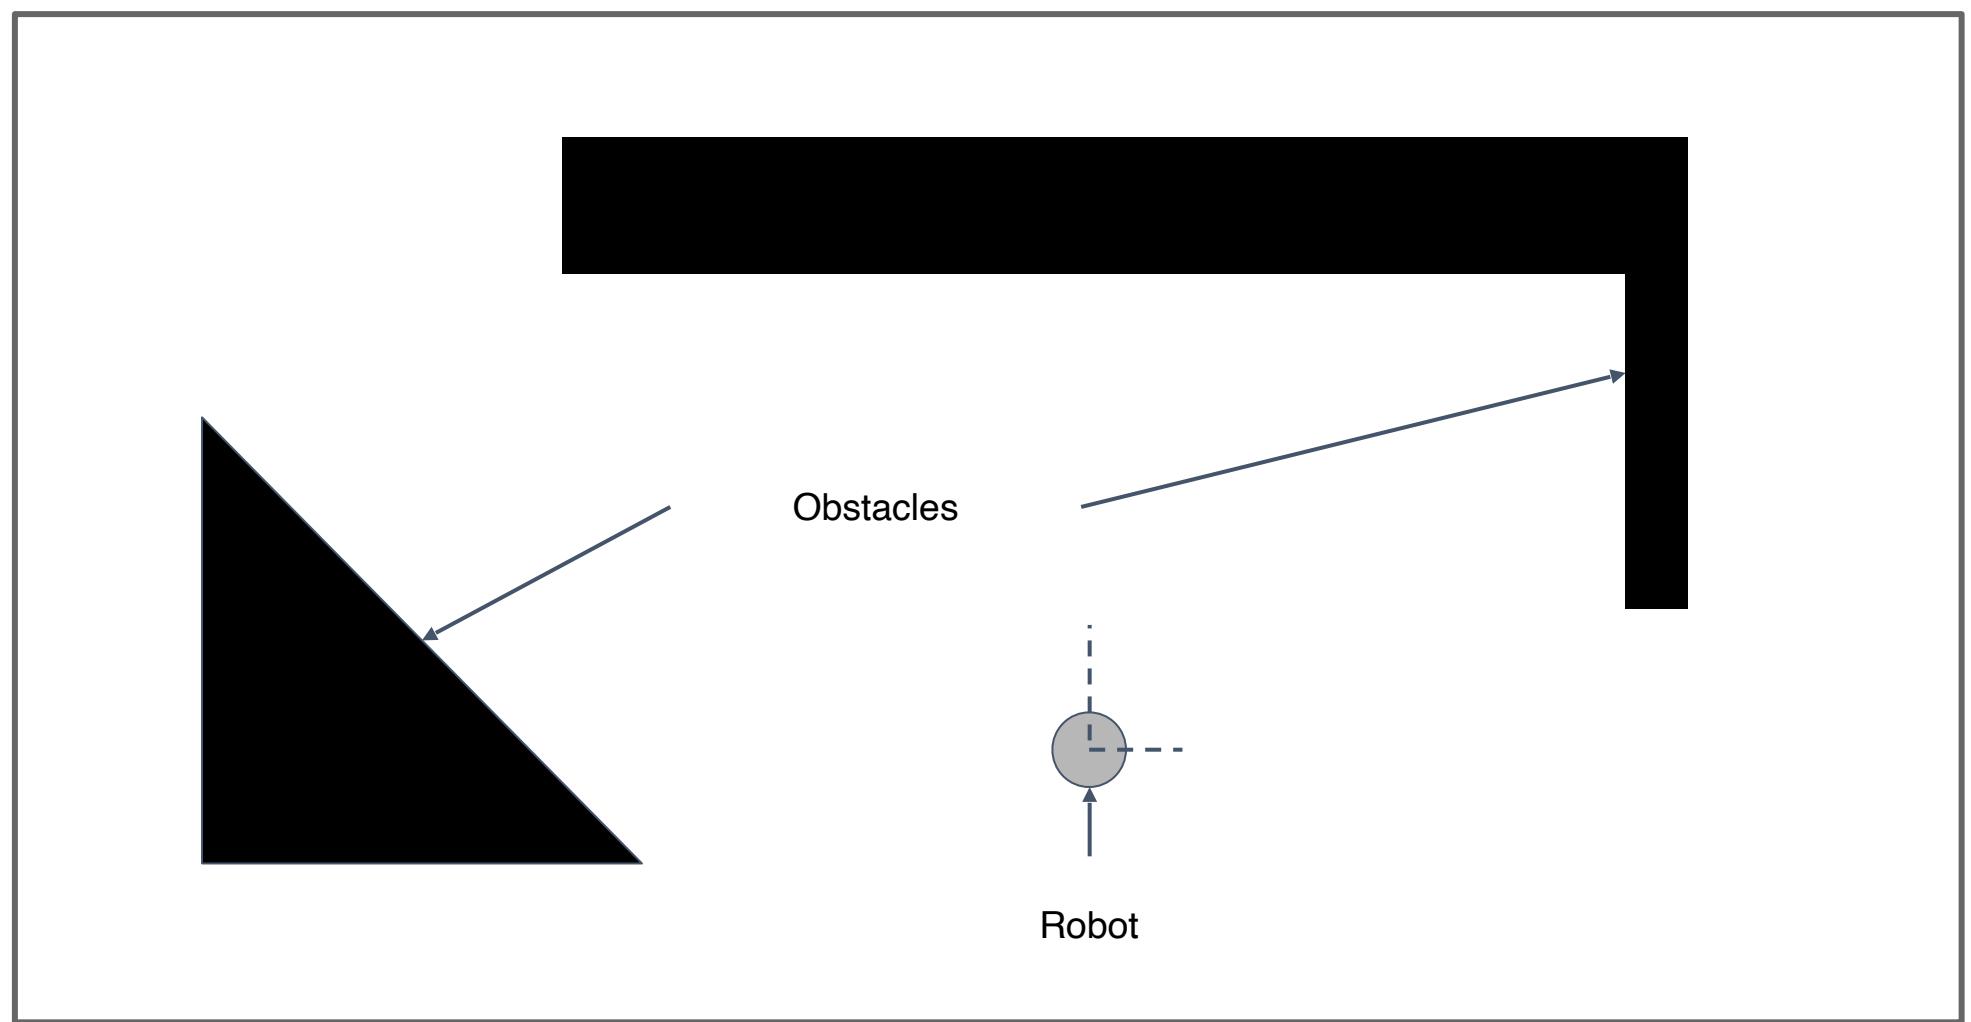
Topological decomposition

- A topological representation is a graph that specifies nodes and edges
 - Nodes denotes areas in the environment
 - Edges describe environment connectivity
- Robots can...
 - ...detect their current position in terms of the nodes of the topological graph
 - ...travel between nodes using robot motion
- Typical for 3D maps

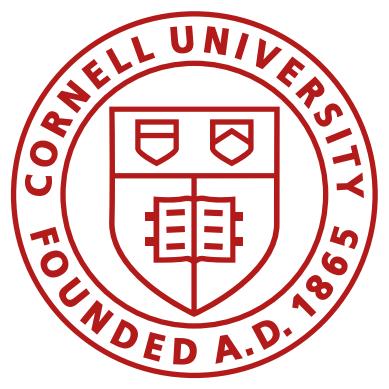


How to represent the robot pose?

- Physical robots take up space
- Expand obstacles
- Represent maps in configuration space instead of Euclidean space

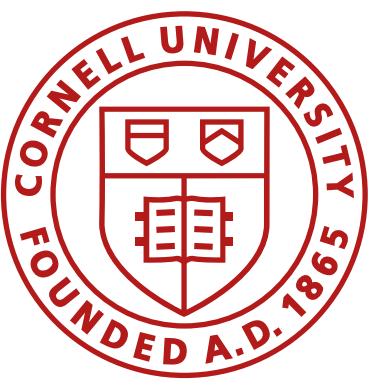


Robot can be treated as a point object



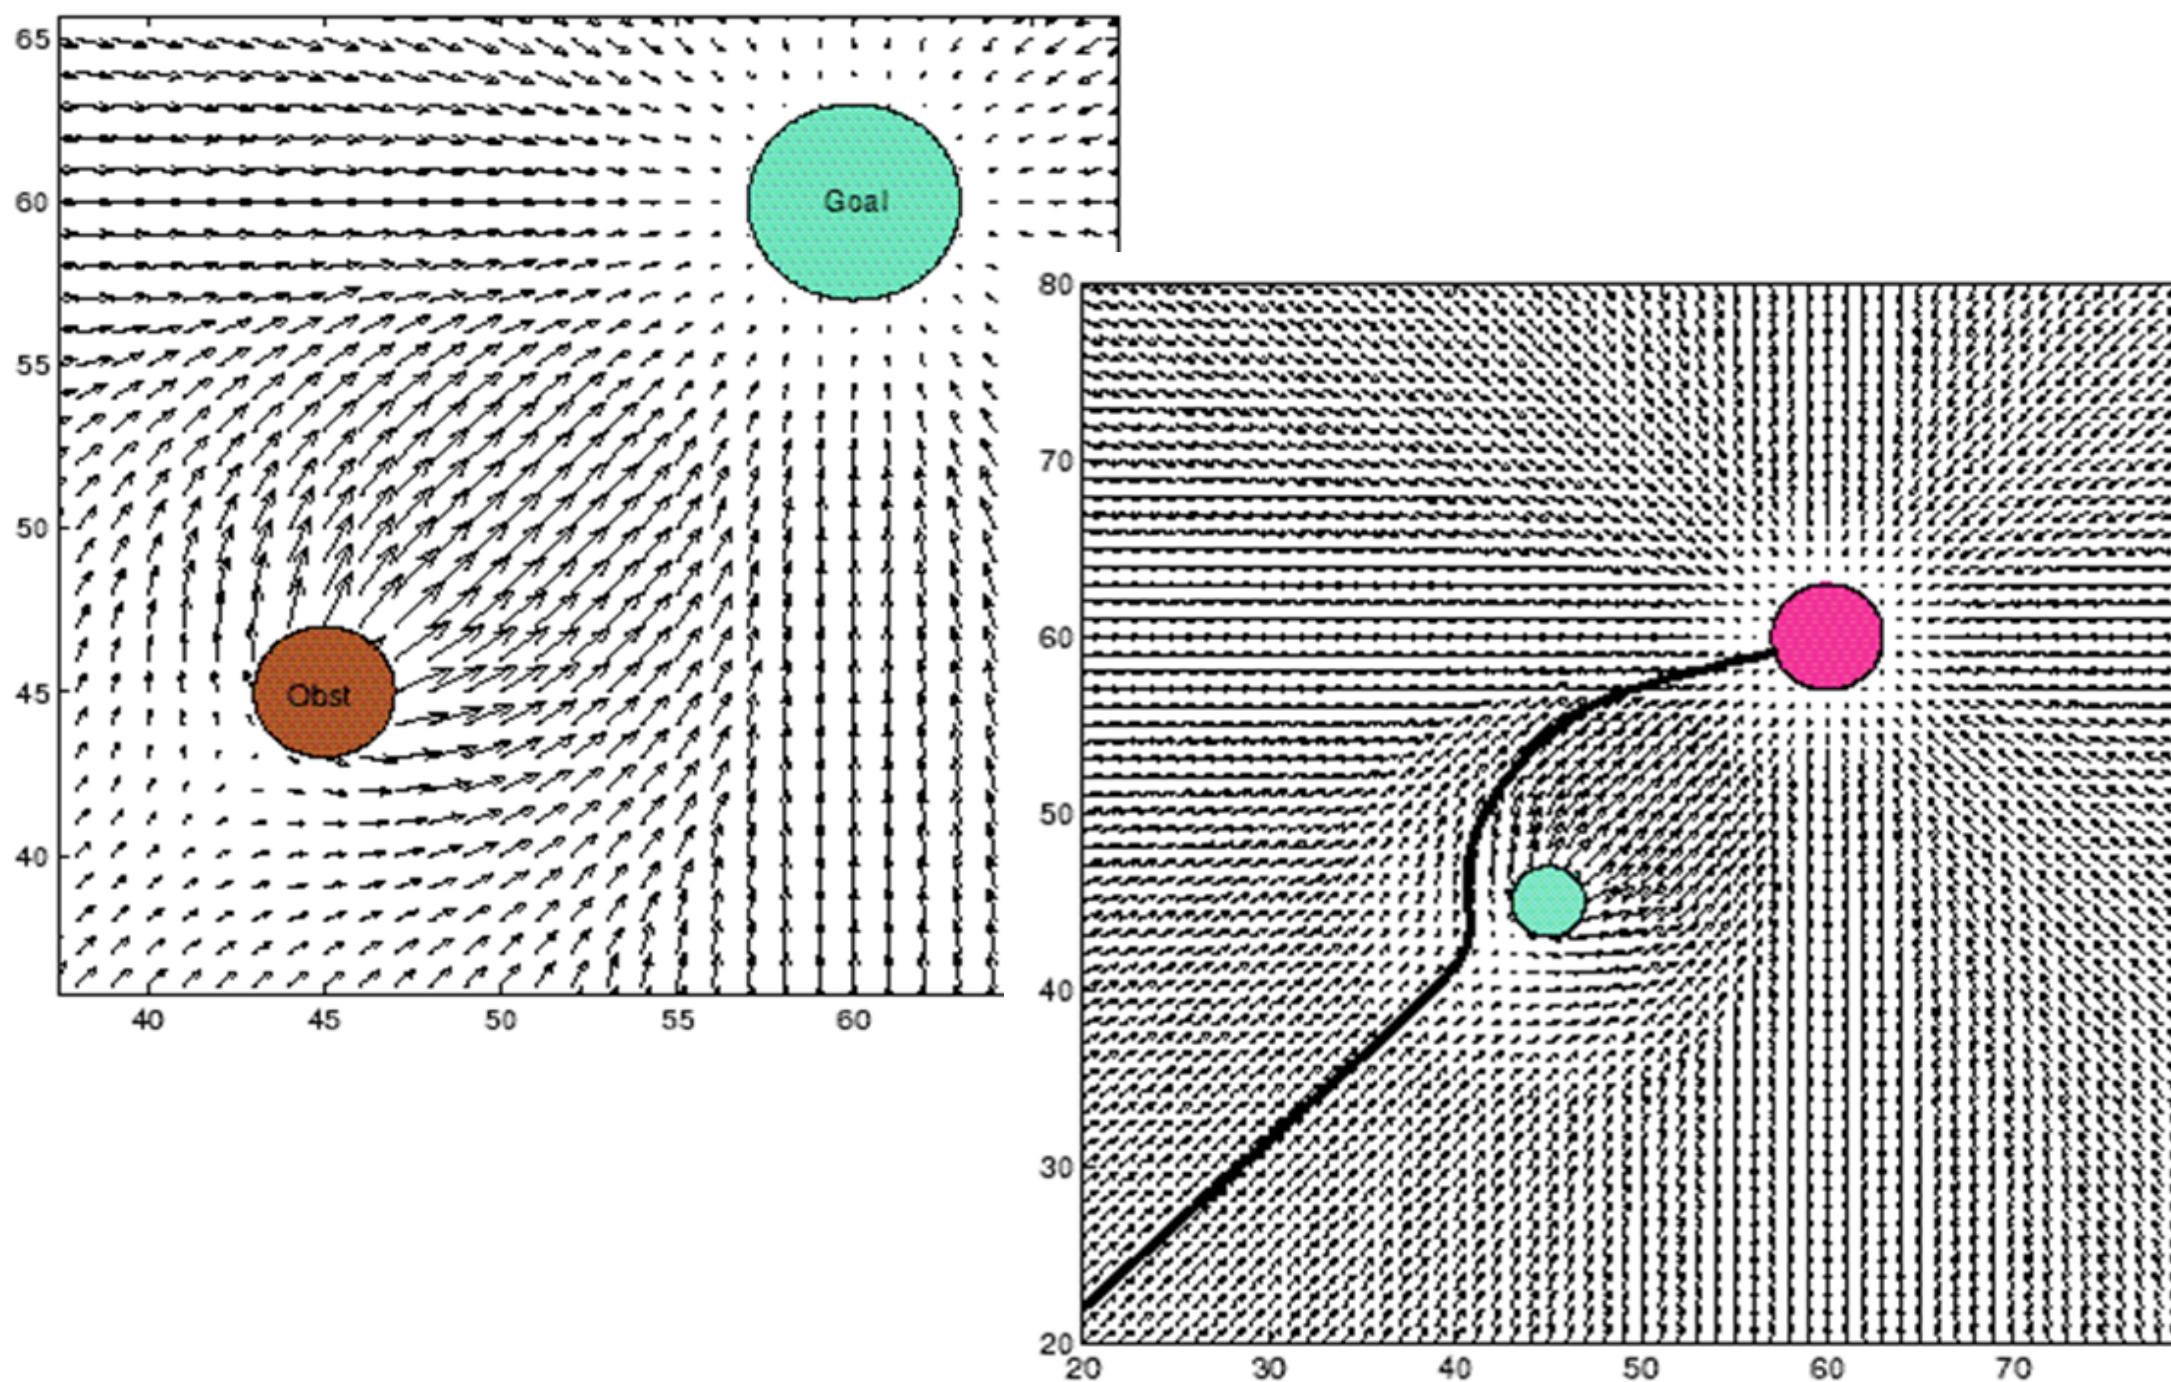
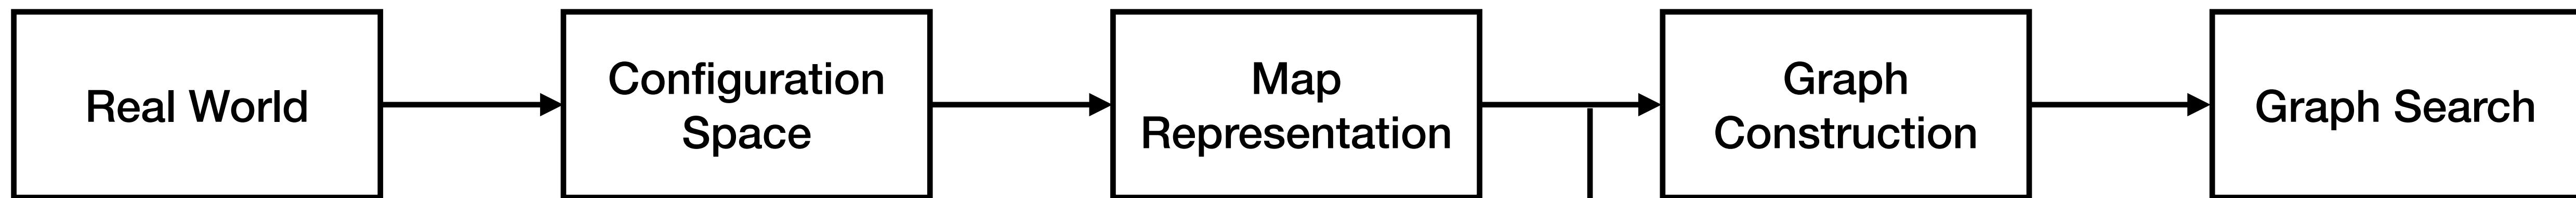
Map Representation Considerations

- The precision of the map must appropriately match the precision with which the robot needs to achieve its goals
- The precision of the map and the type of features represented must match the precision and data types returned by the robot's sensors
- The complexity of the map representation has direct impact on the computational complexity of reasoning about mapping, localization, and navigation



Constructing Graphs

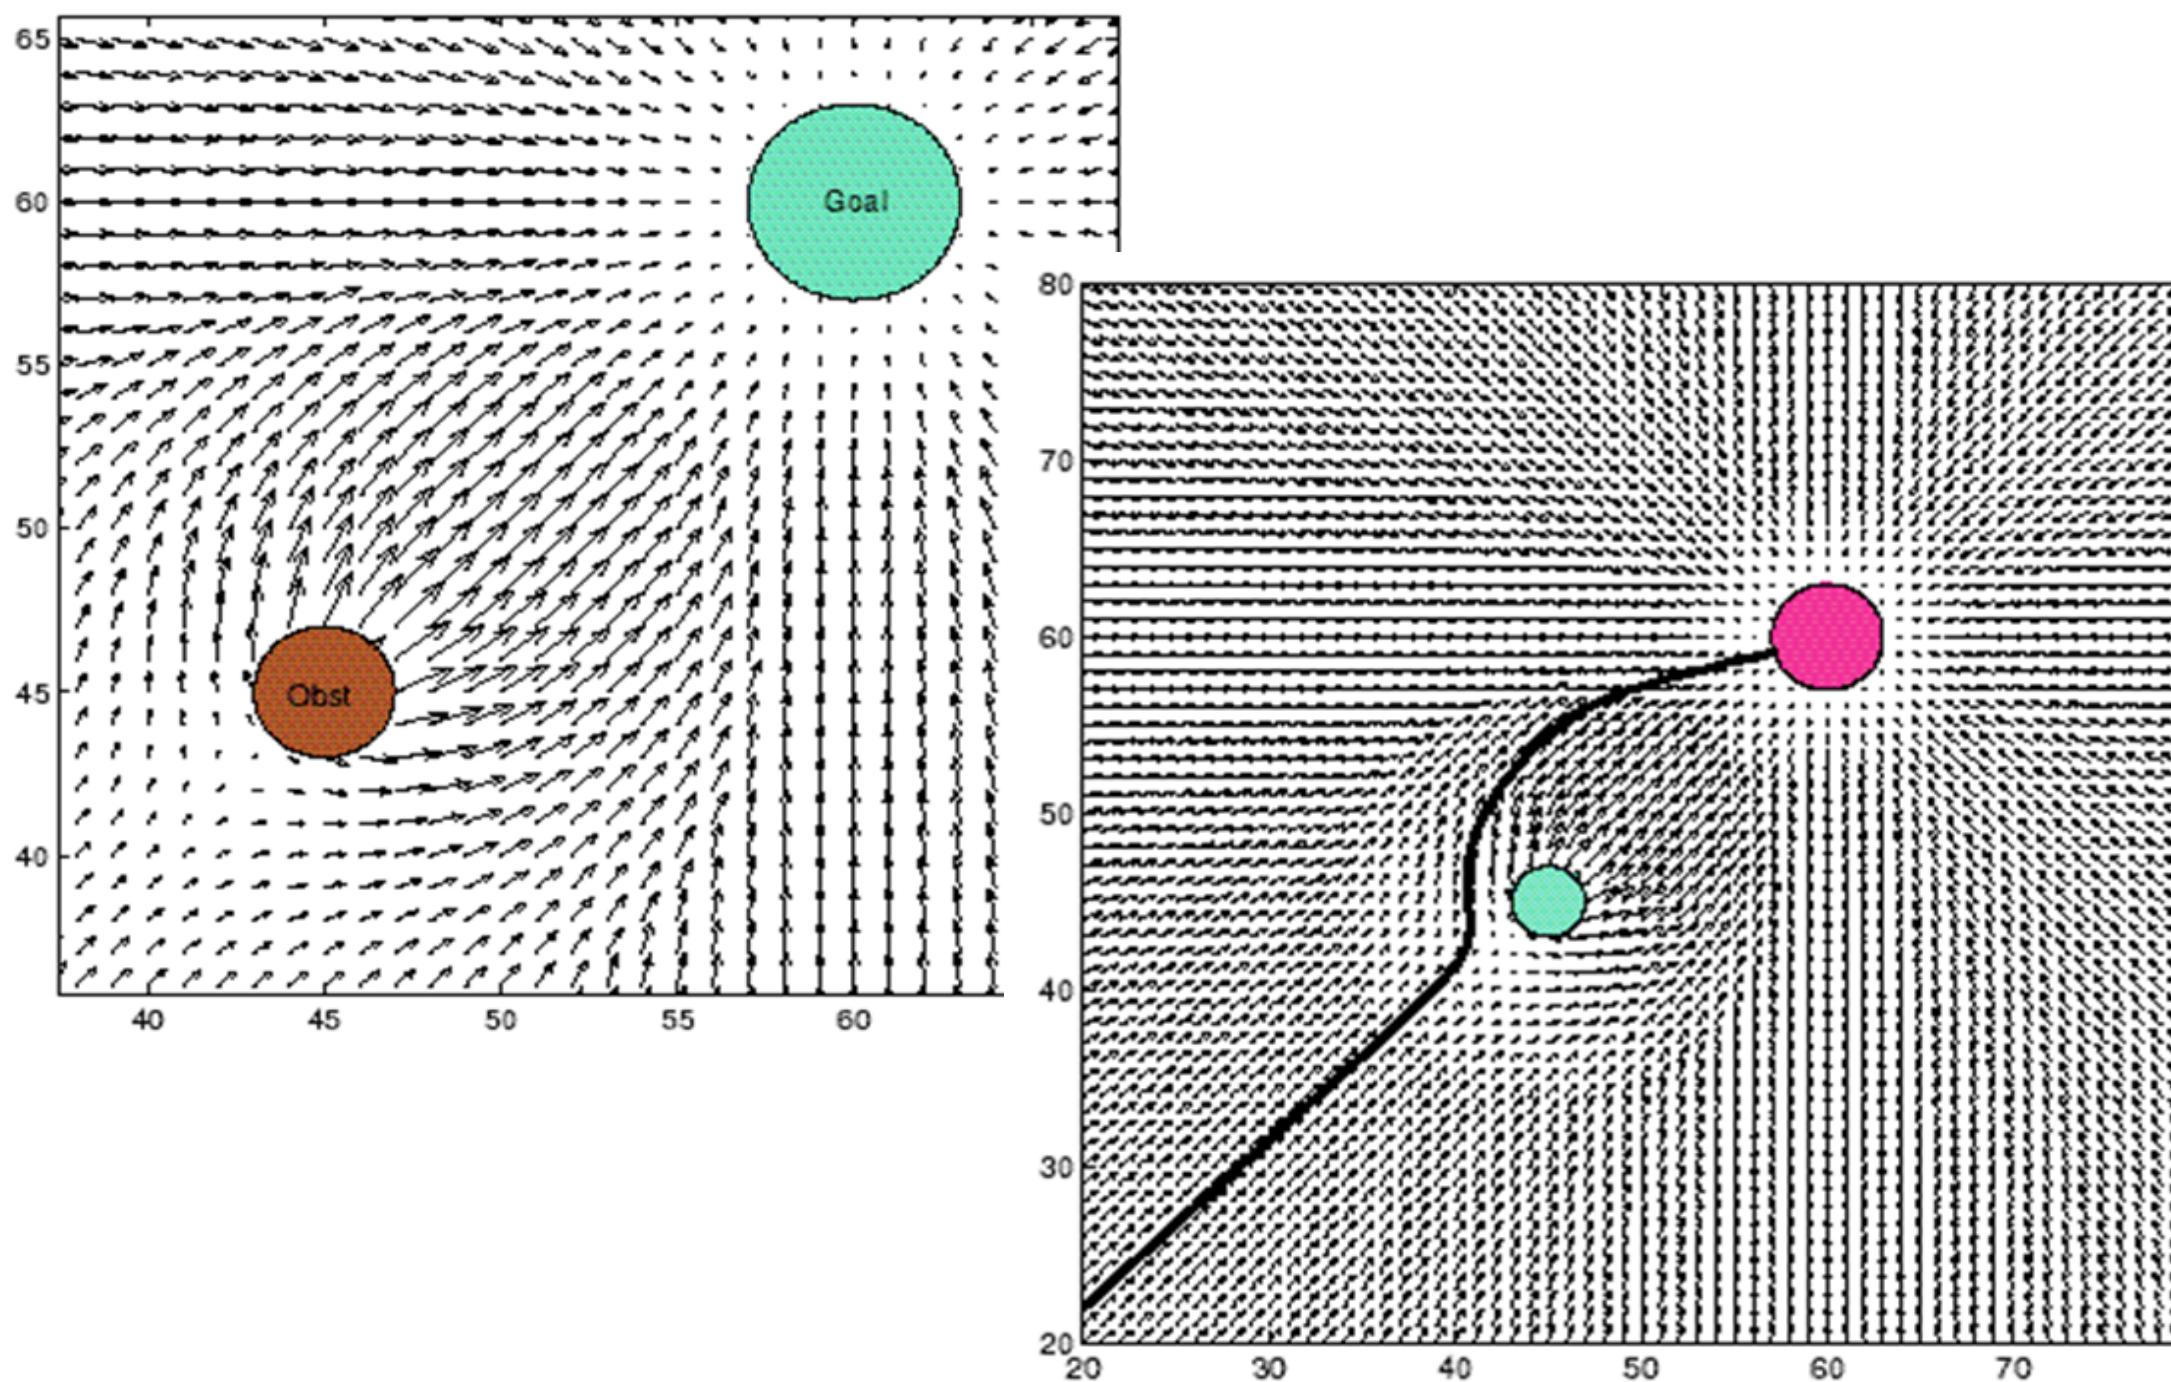
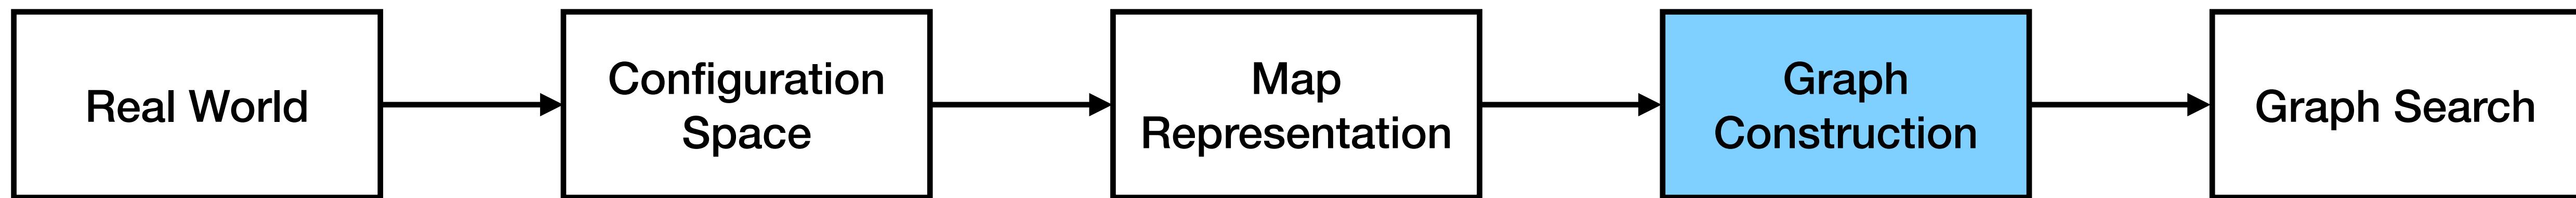
Modeling path planning as a graph search problem



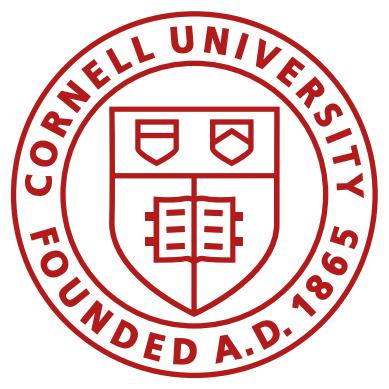
Common alternatives

- Optimal control
- Potential fields

Modeling path planning as a graph search problem



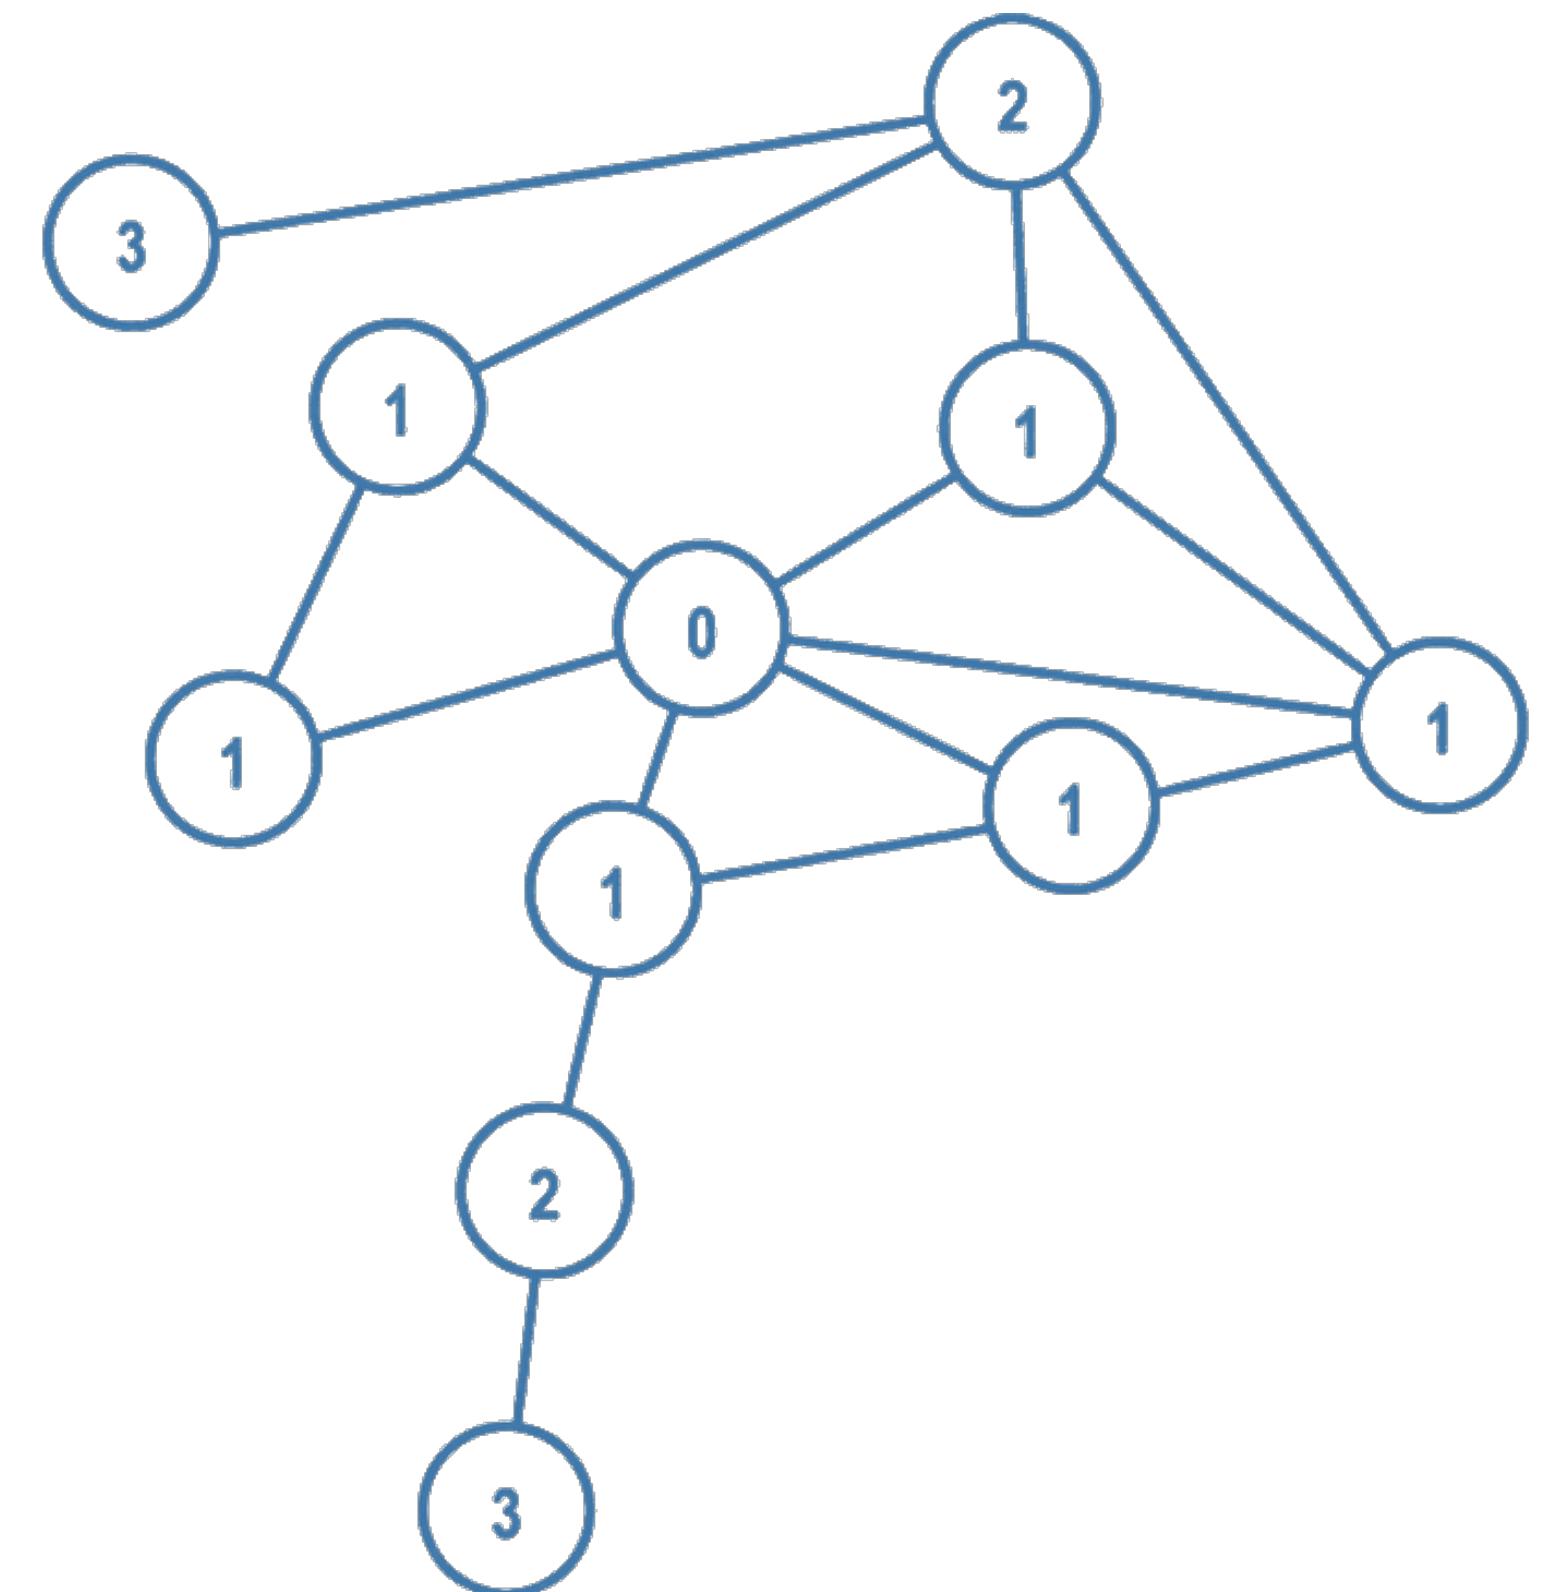
- Topological Graphs
- Cell decomposition
- Visibility Graphs
- RRT
- PRM

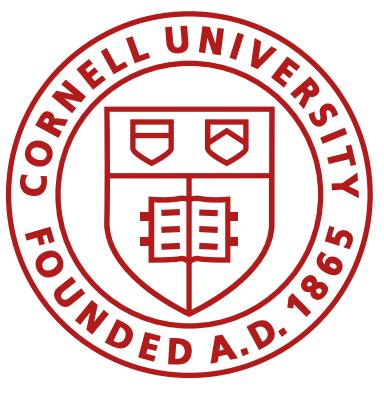


Graph Construction

- Transform continuous/ discrete/ topological maps to a discrete graph
- Why?
 - Model the path planning problem as a search problem
 - Graph theory has lots of tools
 - Real-time capable algorithms
 - Can accommodate for evolving maps

1. Divide space into simple, connected regions, or “cells”
2. Determine adjacency of open cells
3. Construct a connectivity graph
4. Find cells with initial and goal configuration
5. Search for a path in the connectivity graph to join them
6. From the sequence of cells, compute a path within each cell
 - e.g. passing through the midpoints of cell boundaries or by sequence of wall following movements

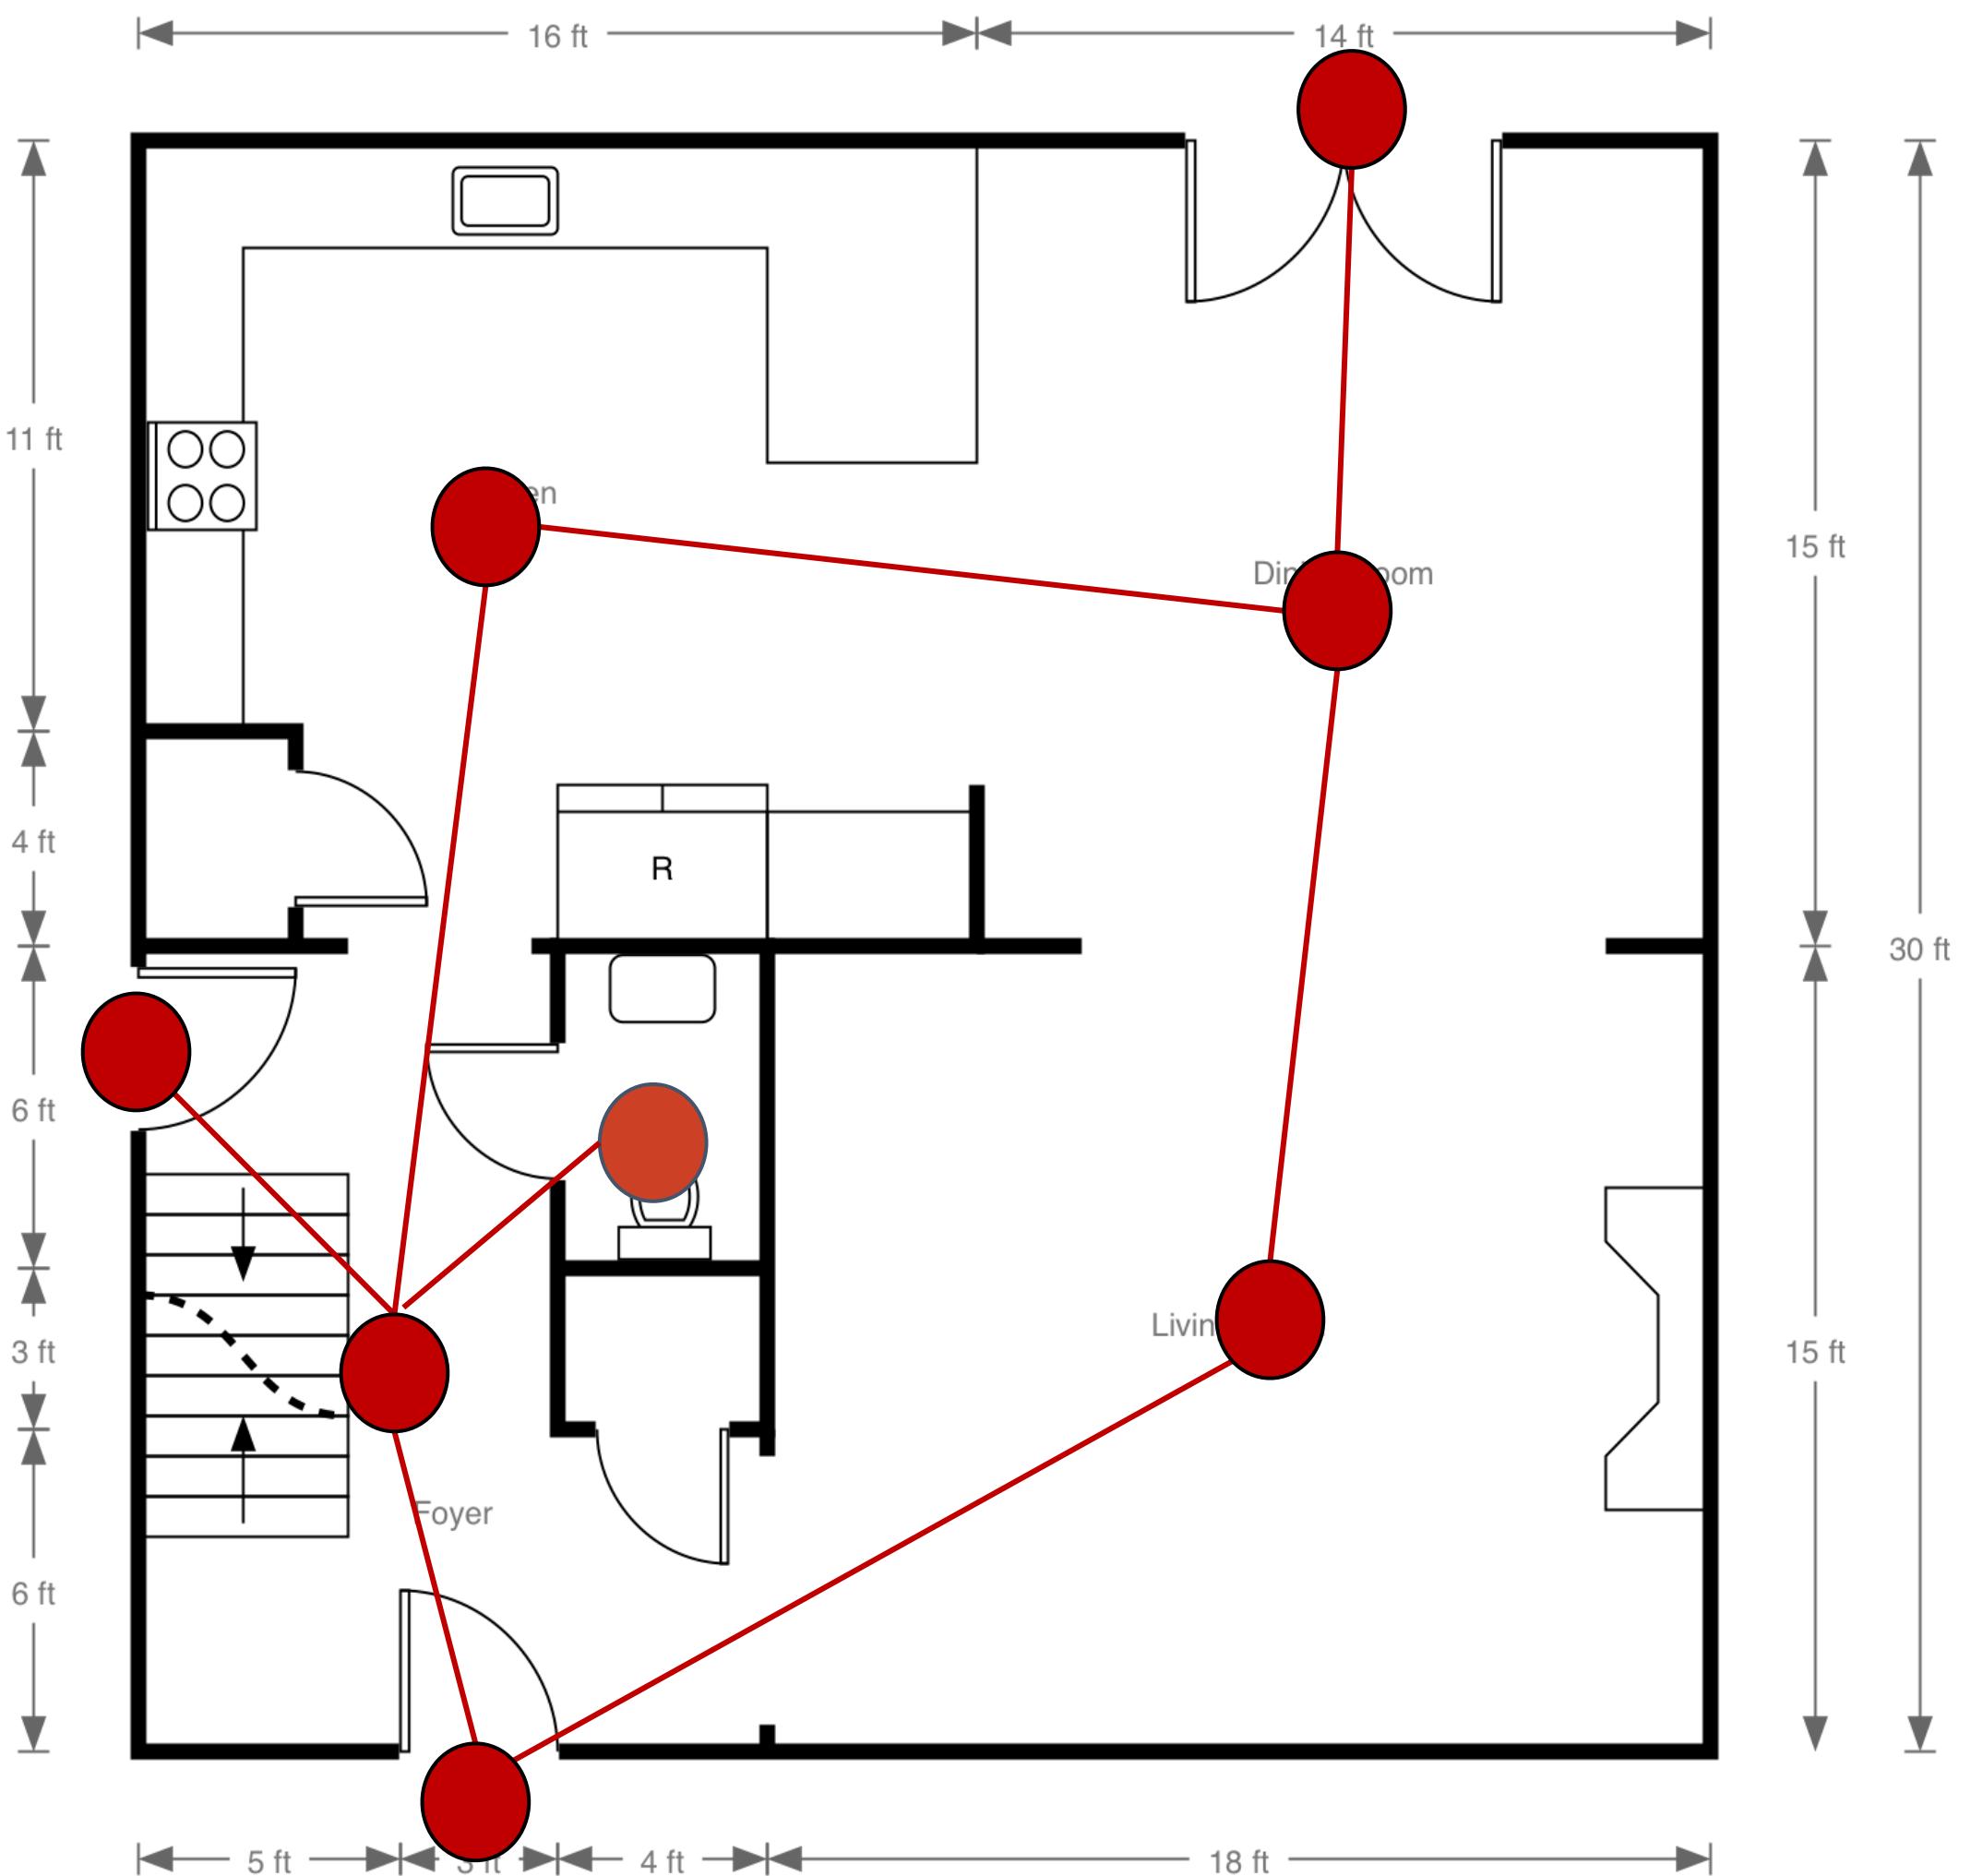


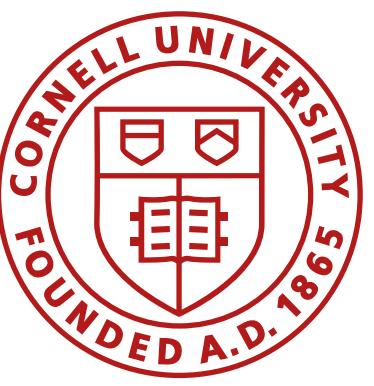


Geometry-based planners

Topological Maps

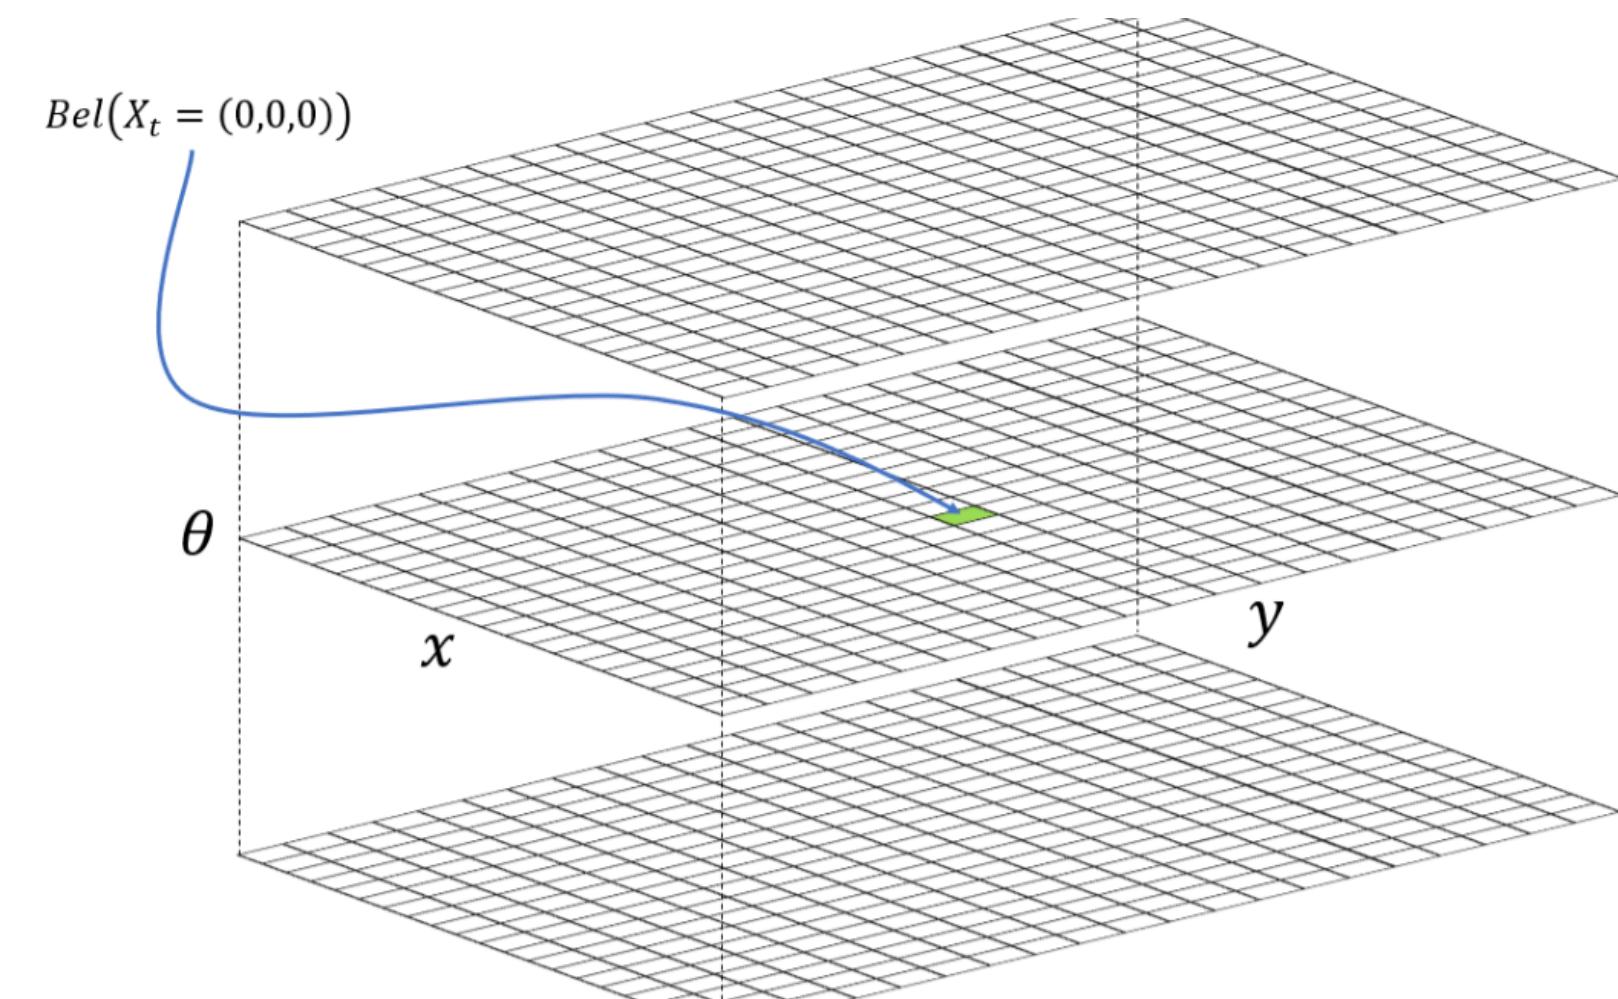
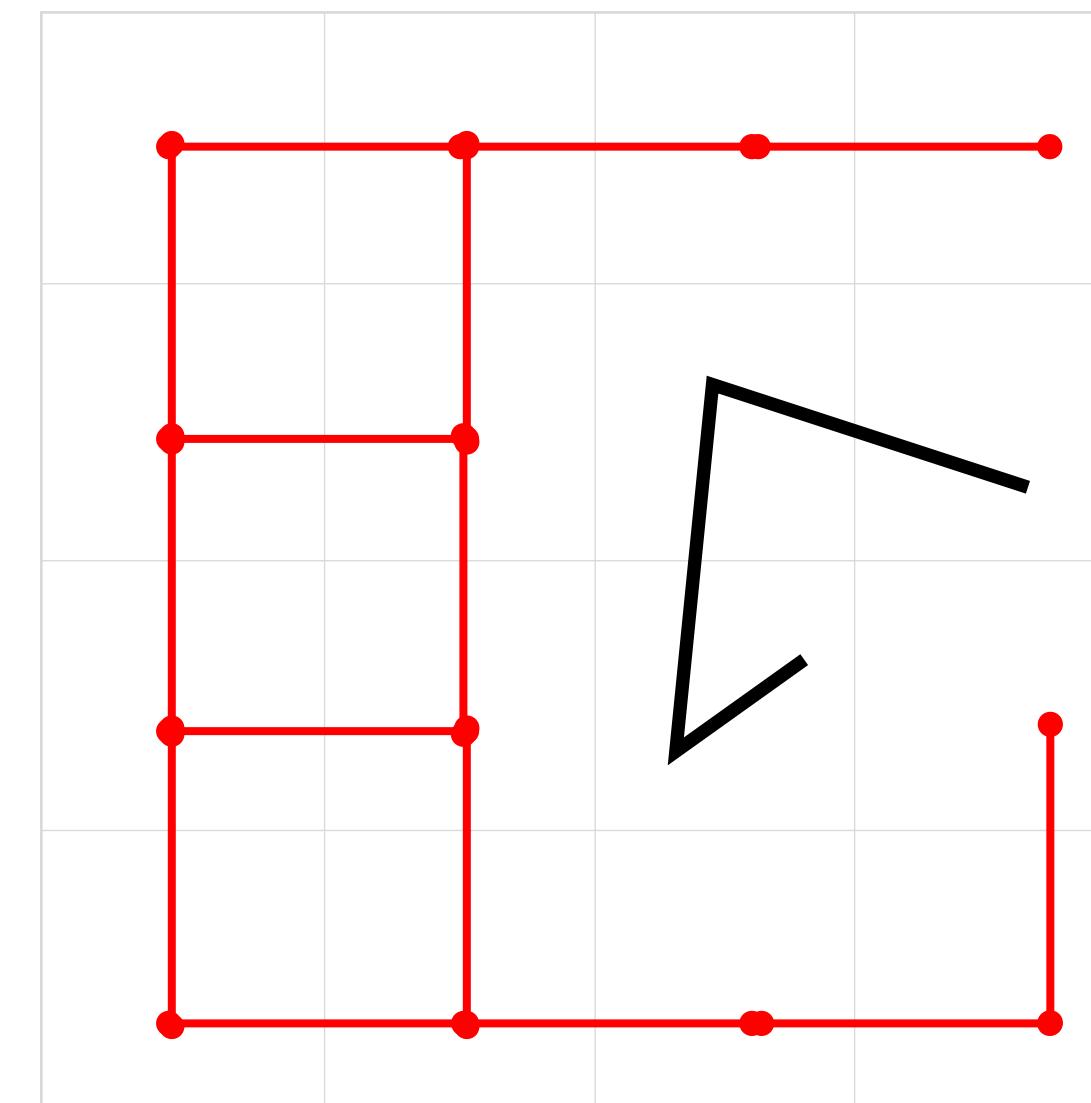
- Good abstract representation
 - Tradeoff in # of nodes
 - Complexity vs. accuracy
 - Efficient in large, sparse environments
 - Loss in geometric precision
 - Edges can carry weight
 - Con: limited information



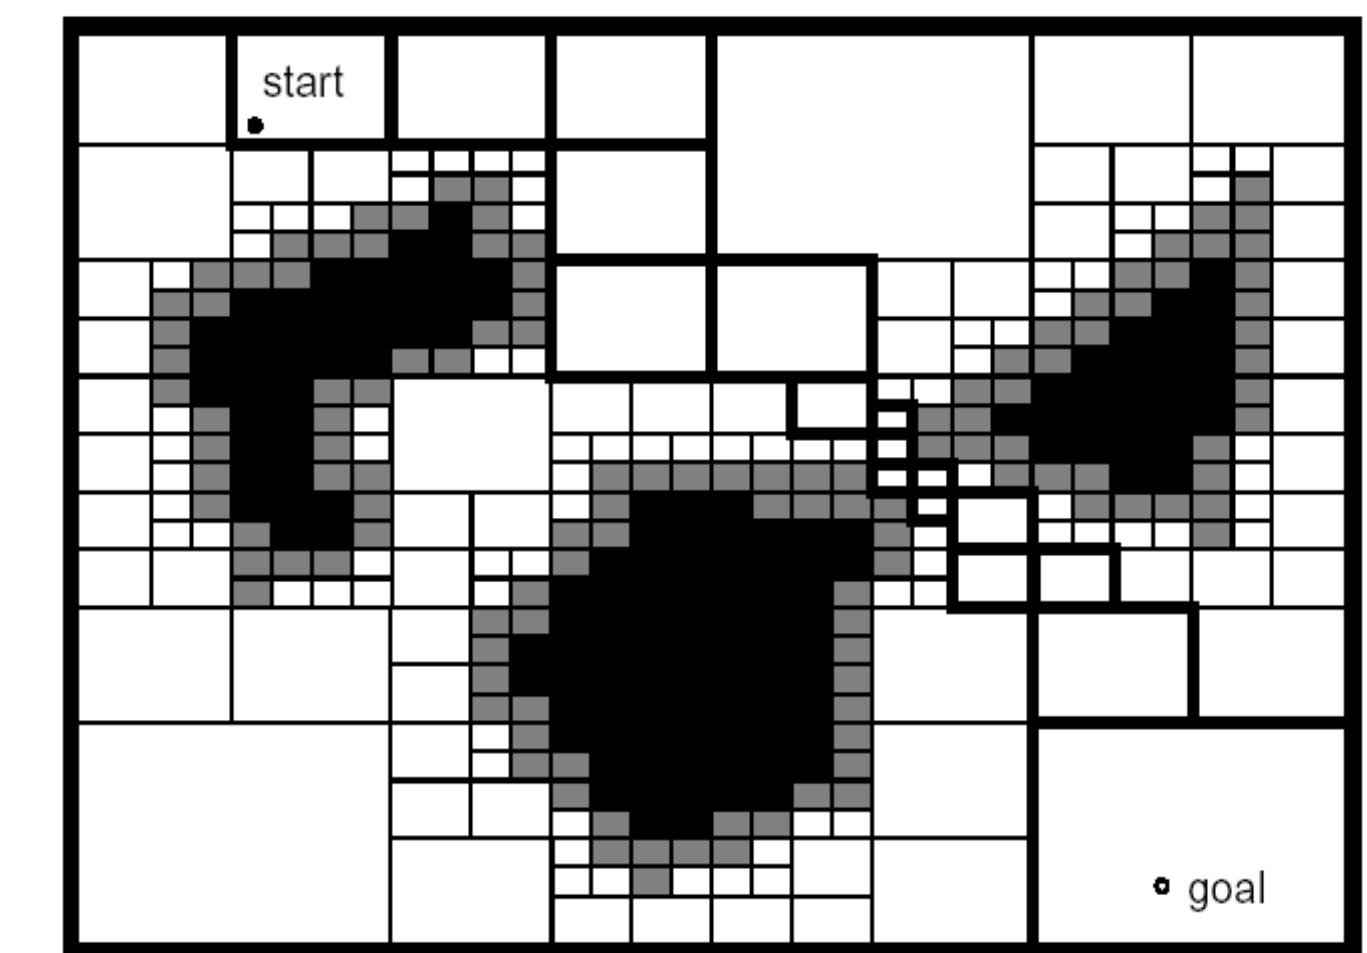


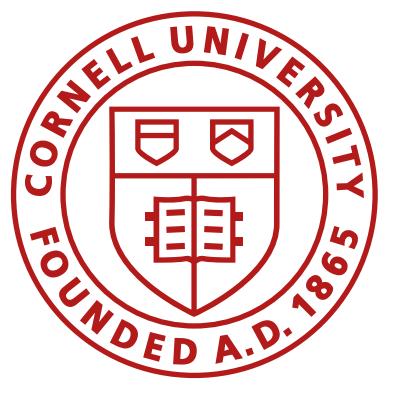
Fixed Cell Decomposition

(Lab 9-12)

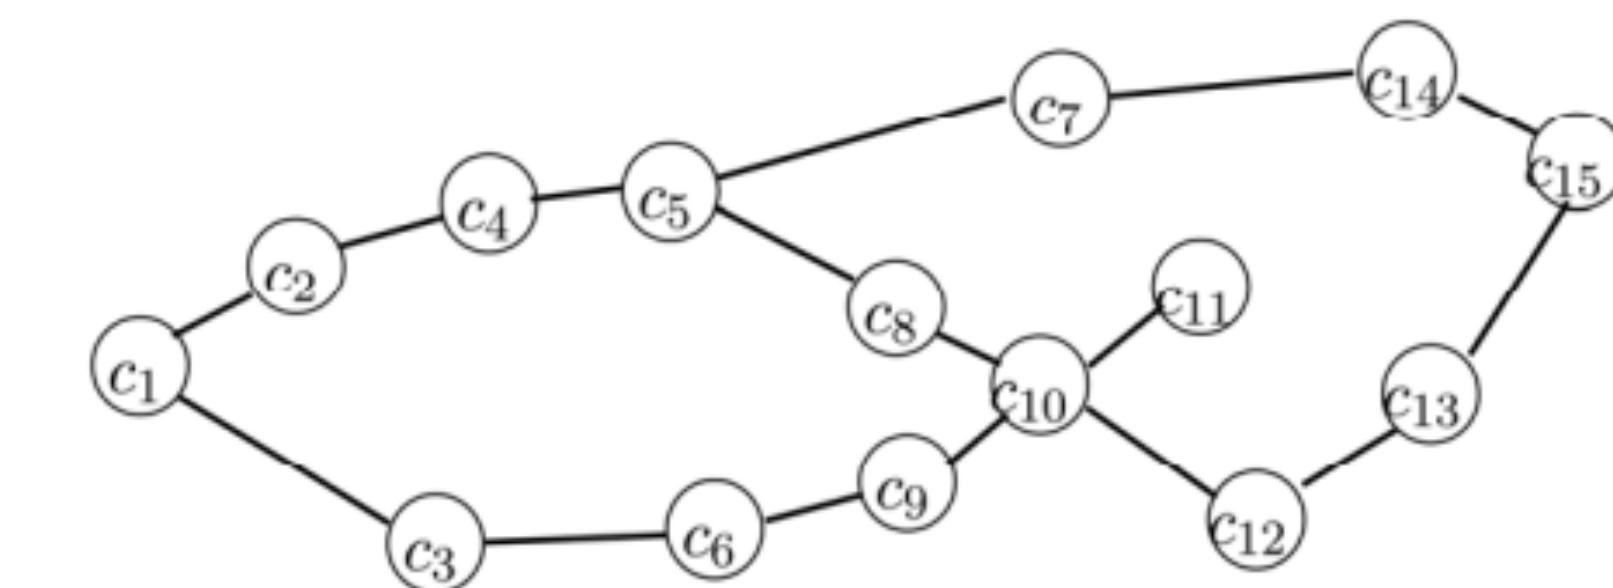
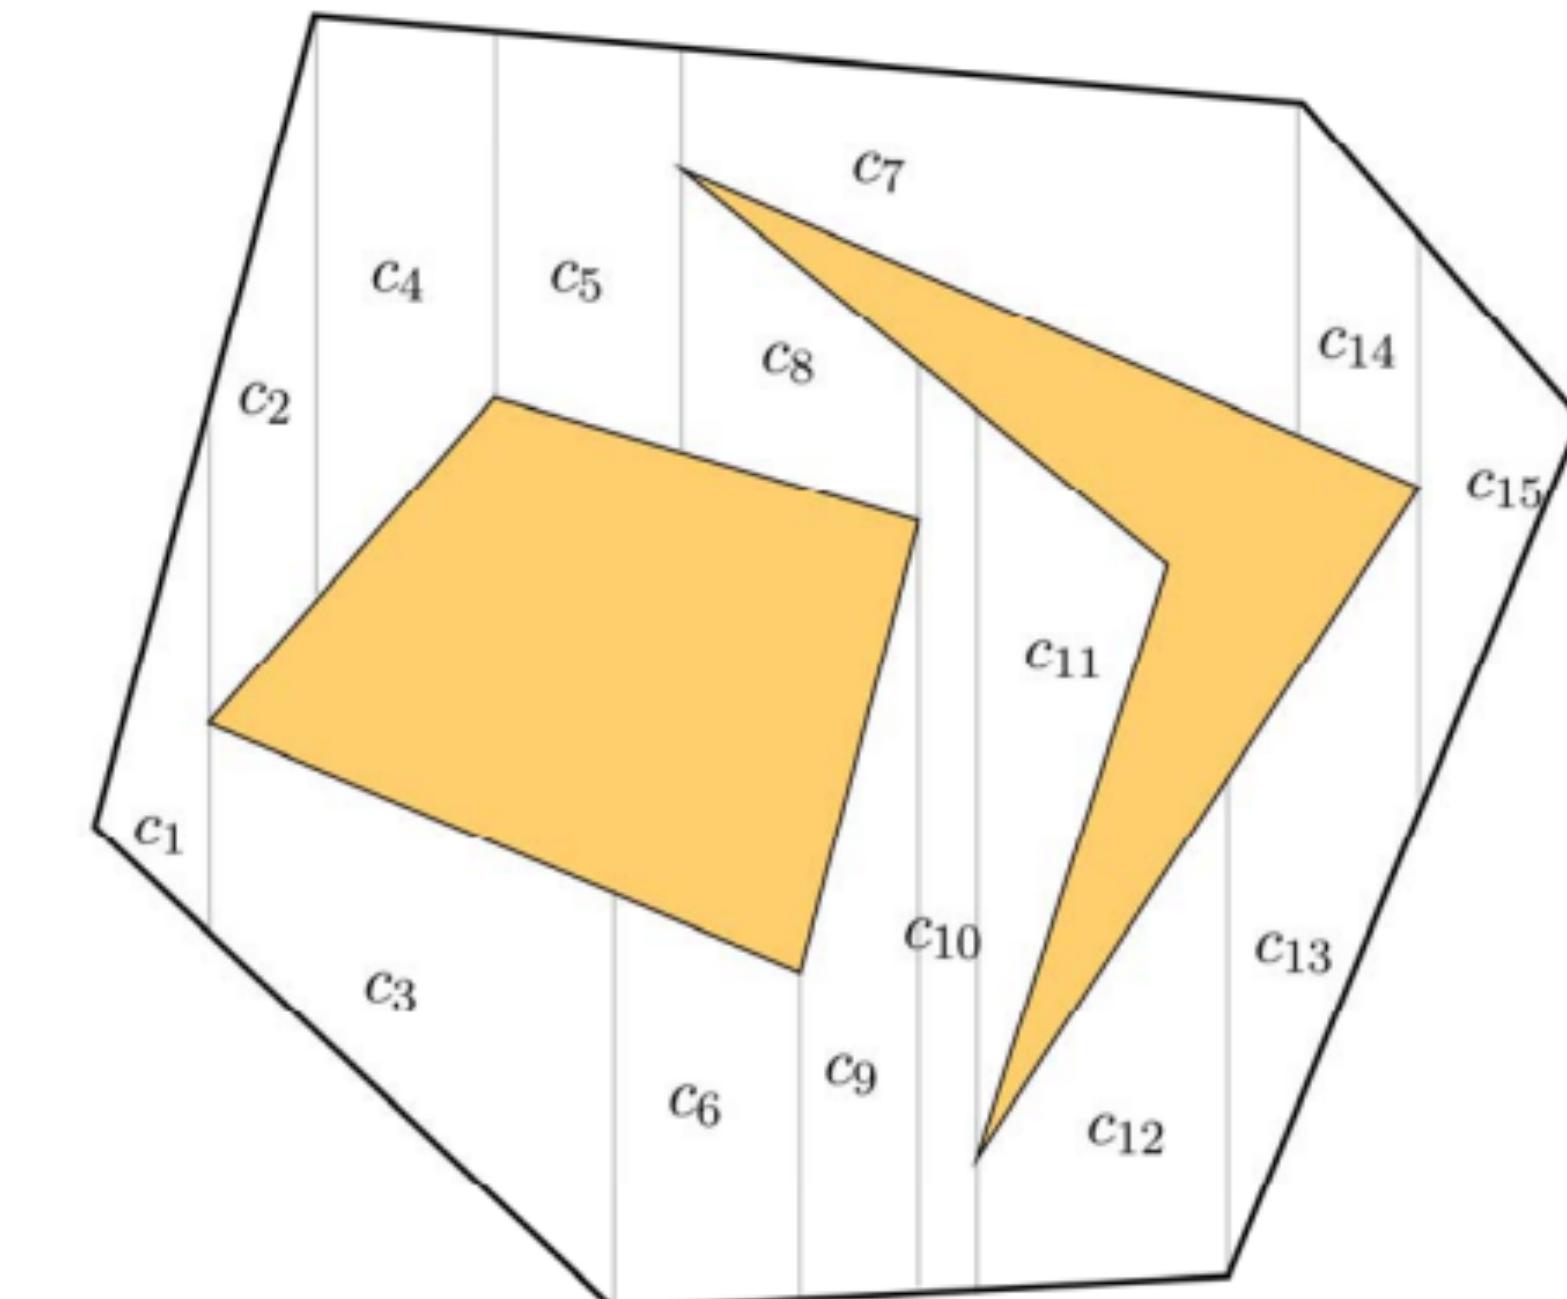
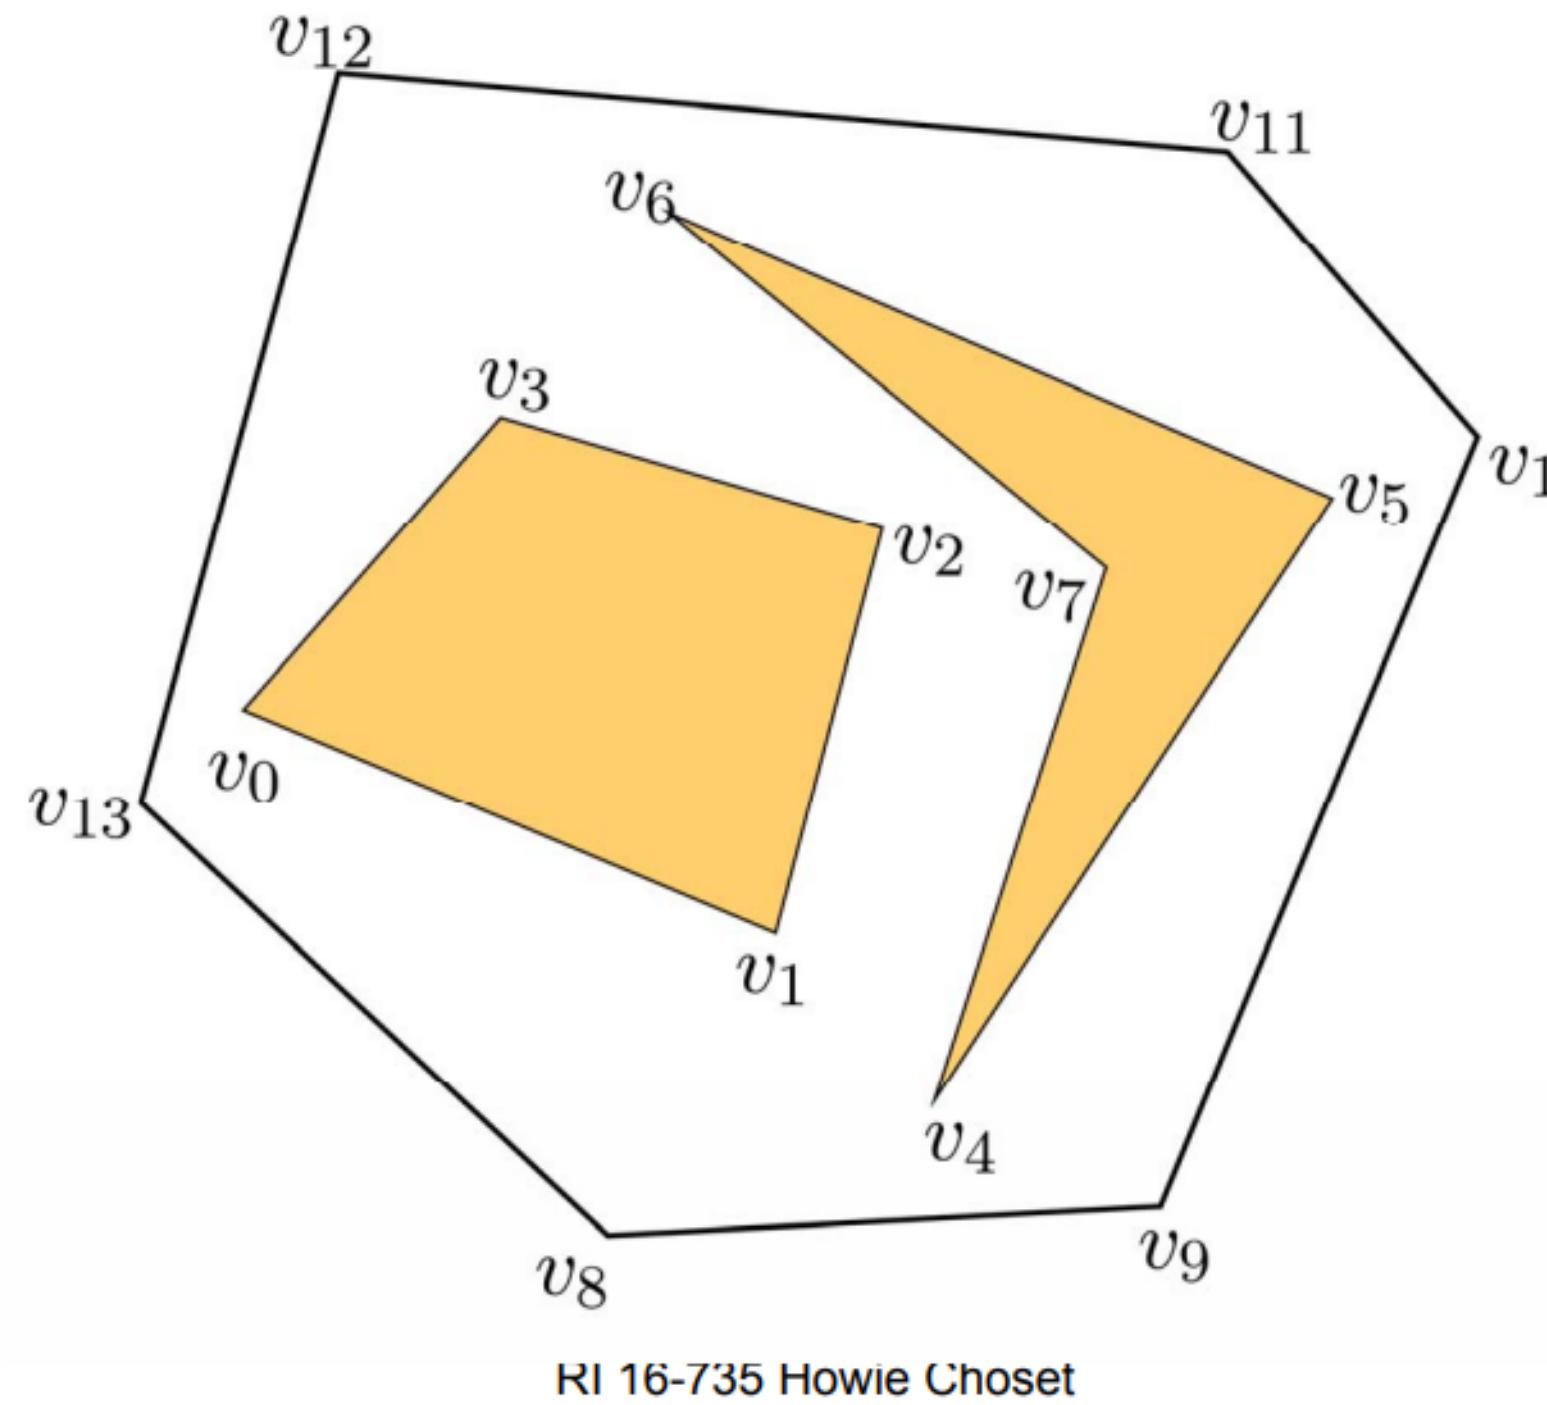


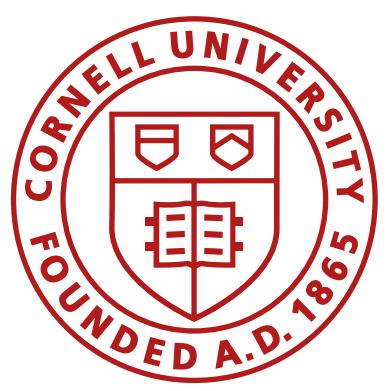
Adaptive Cell Decomposition





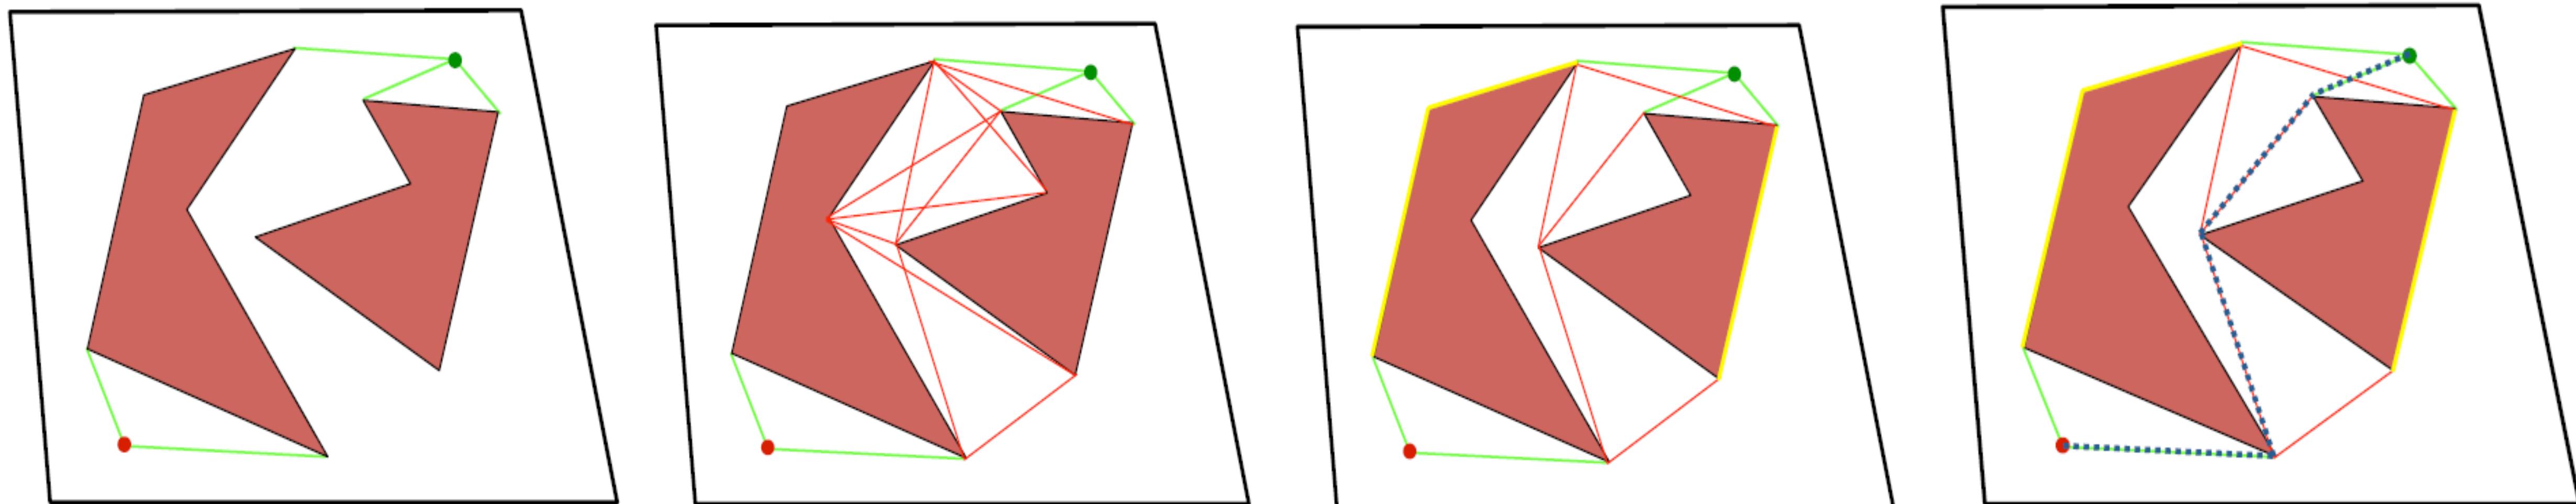
Trapezoidal Cell Decomposition

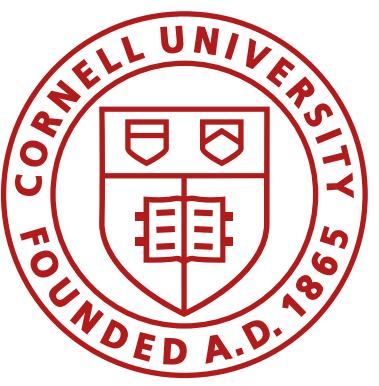




Visibility Graphs

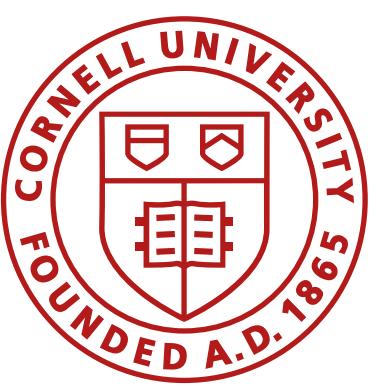
- Connect initial and goal locations with all visible vertices
- Connect each obstacle vertex to every visible obstacle vertex
- Remove edges that intersect the interior of an obstacle
- Plan on the resulting graph





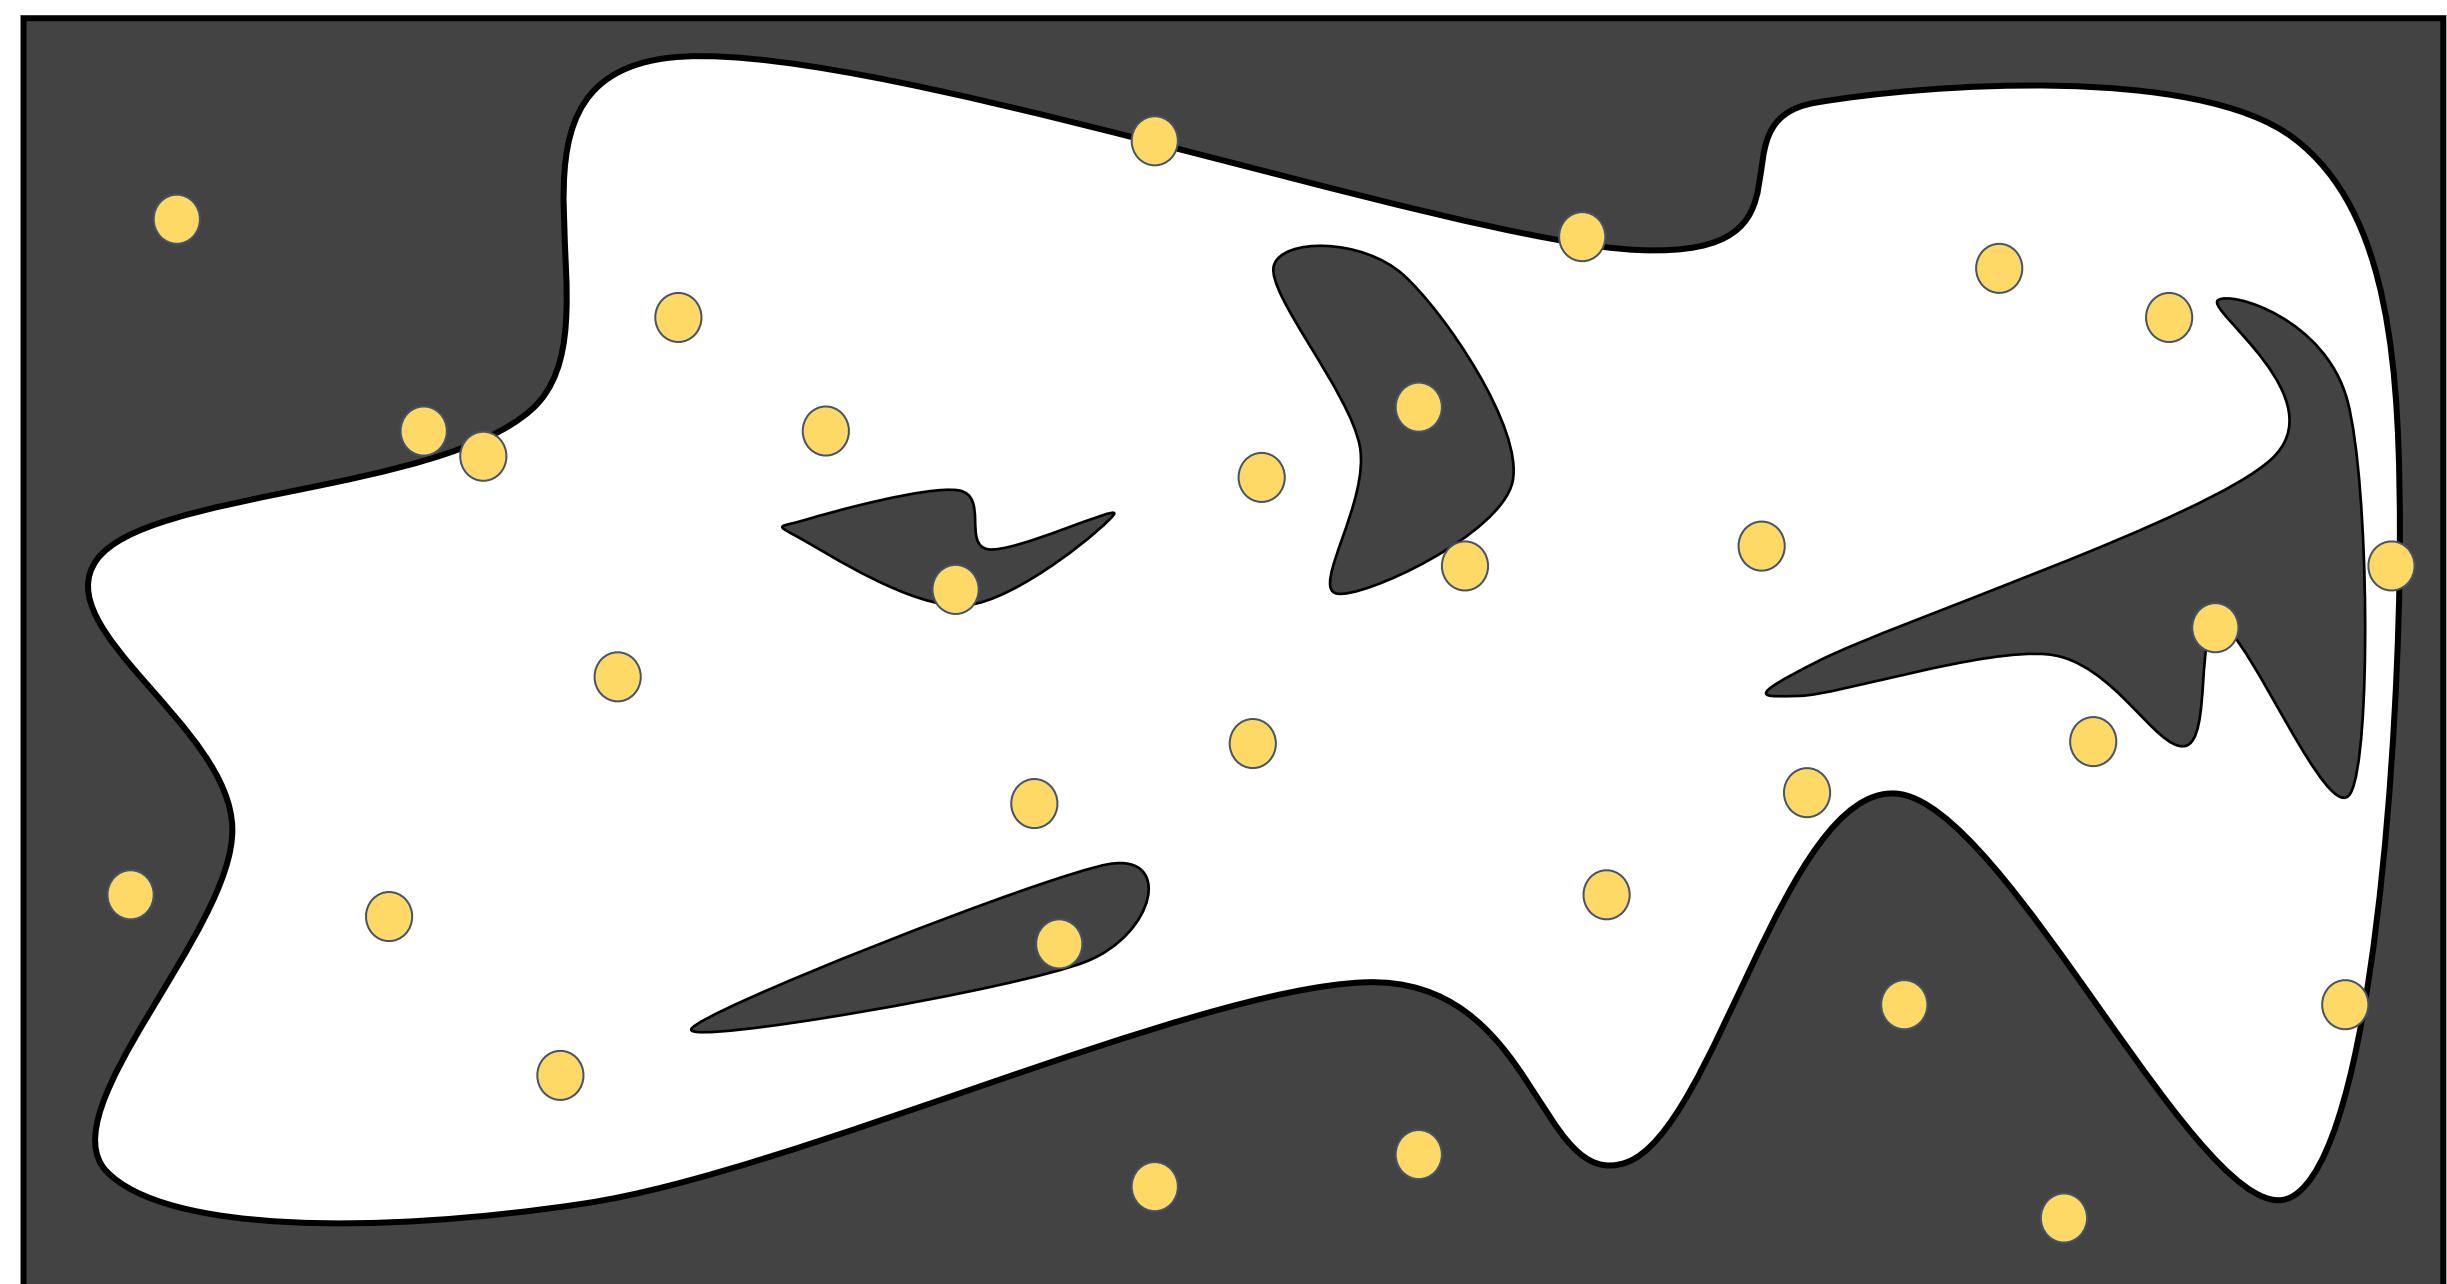
Sampling-based planners

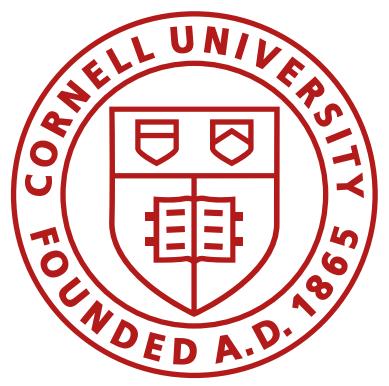
- Rather than computing the C-Space explicitly, we sample it
- Often efficient in high dimensional spaces
- Compute if a robot configuration has collisions
 - Just requires forward kinematics
 - (Local path plans between configurations)
- Examples
 - Probabilistic Roadmaps (PRM)
 - Rapidly Exploring Random Trees (RRT)



Probabilistic Roadmaps

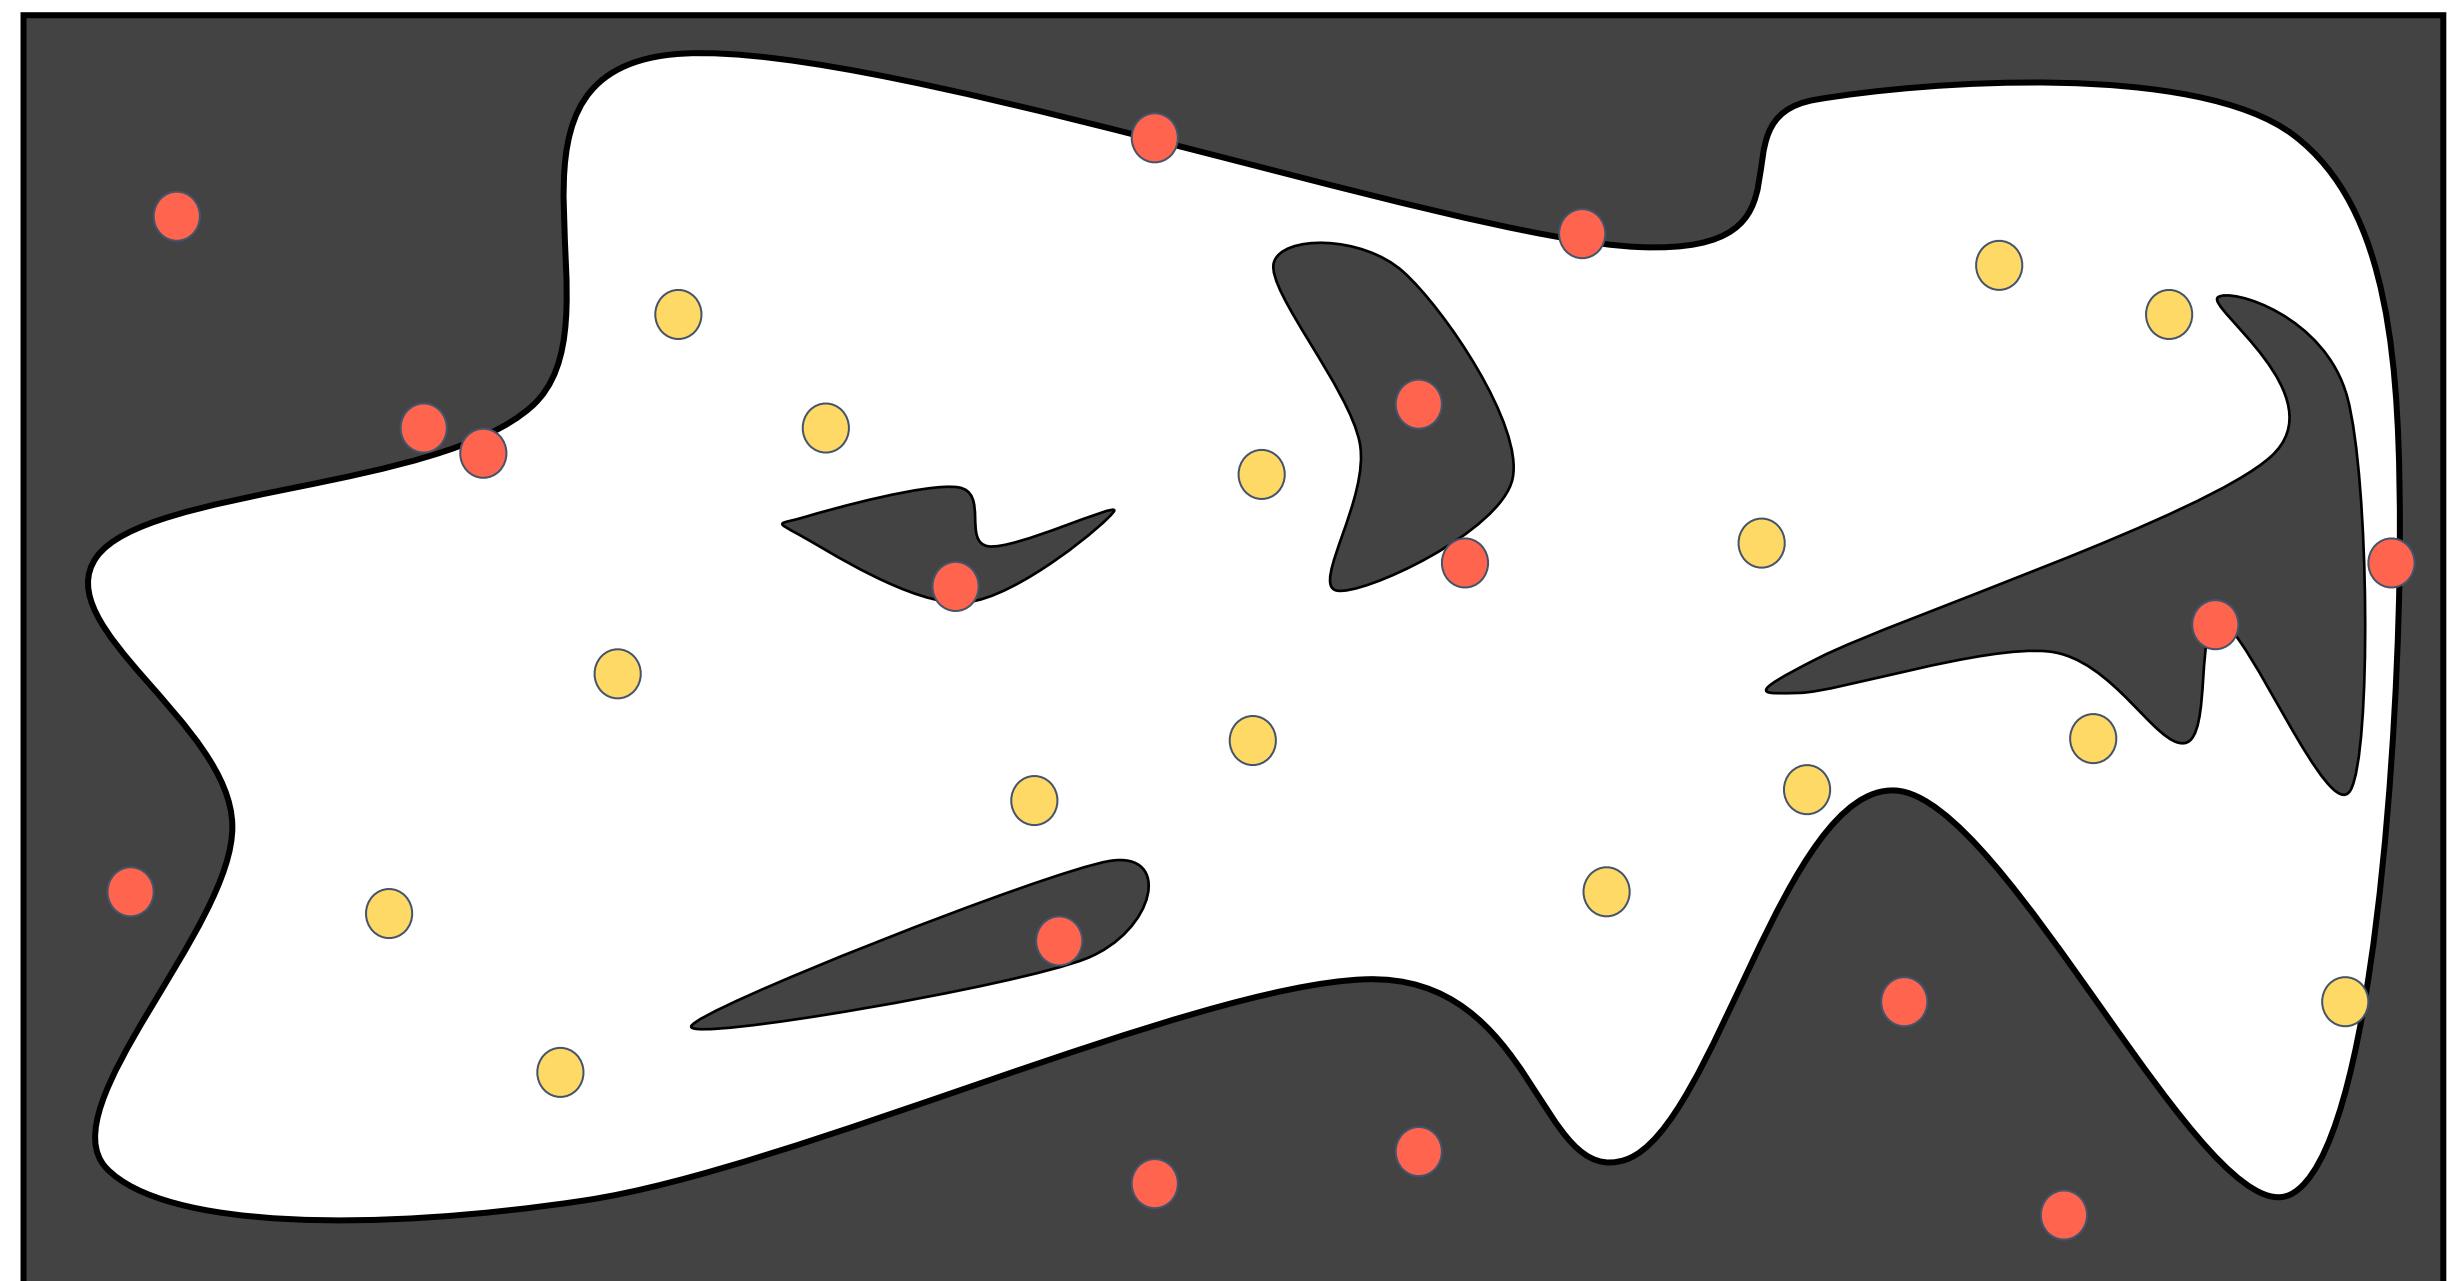
- Configurations are sampled by picking coordinates at random





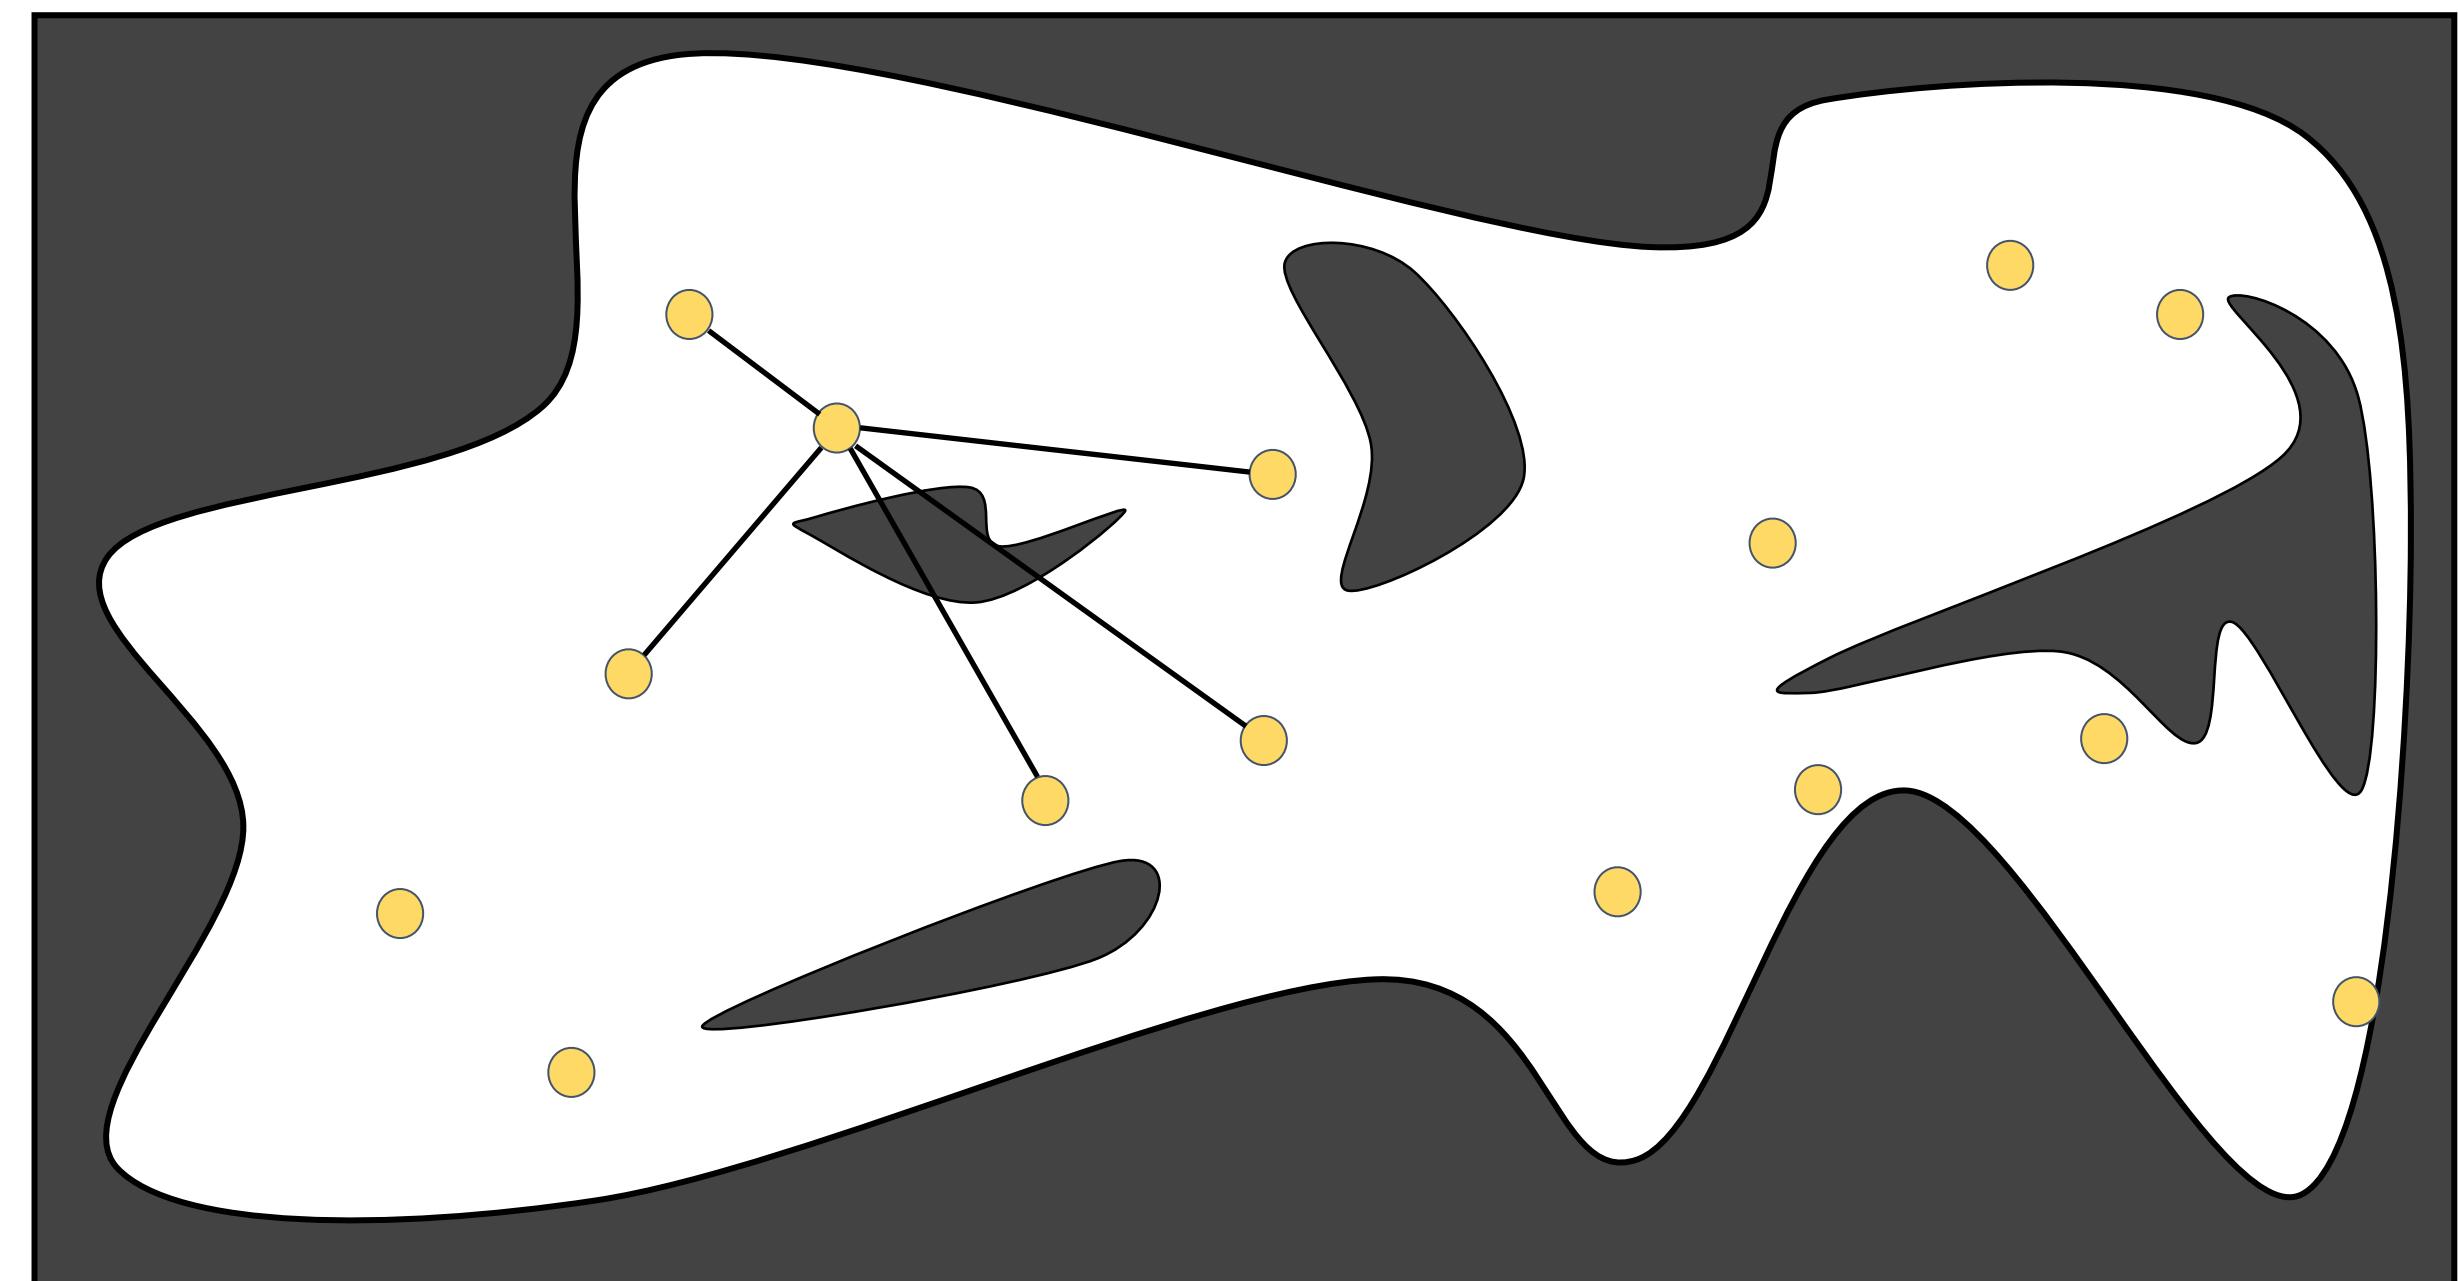
Probabilistic Roadmaps

- Configurations are sampled by picking coordinates at random
- Sampled configurations are tested for collision



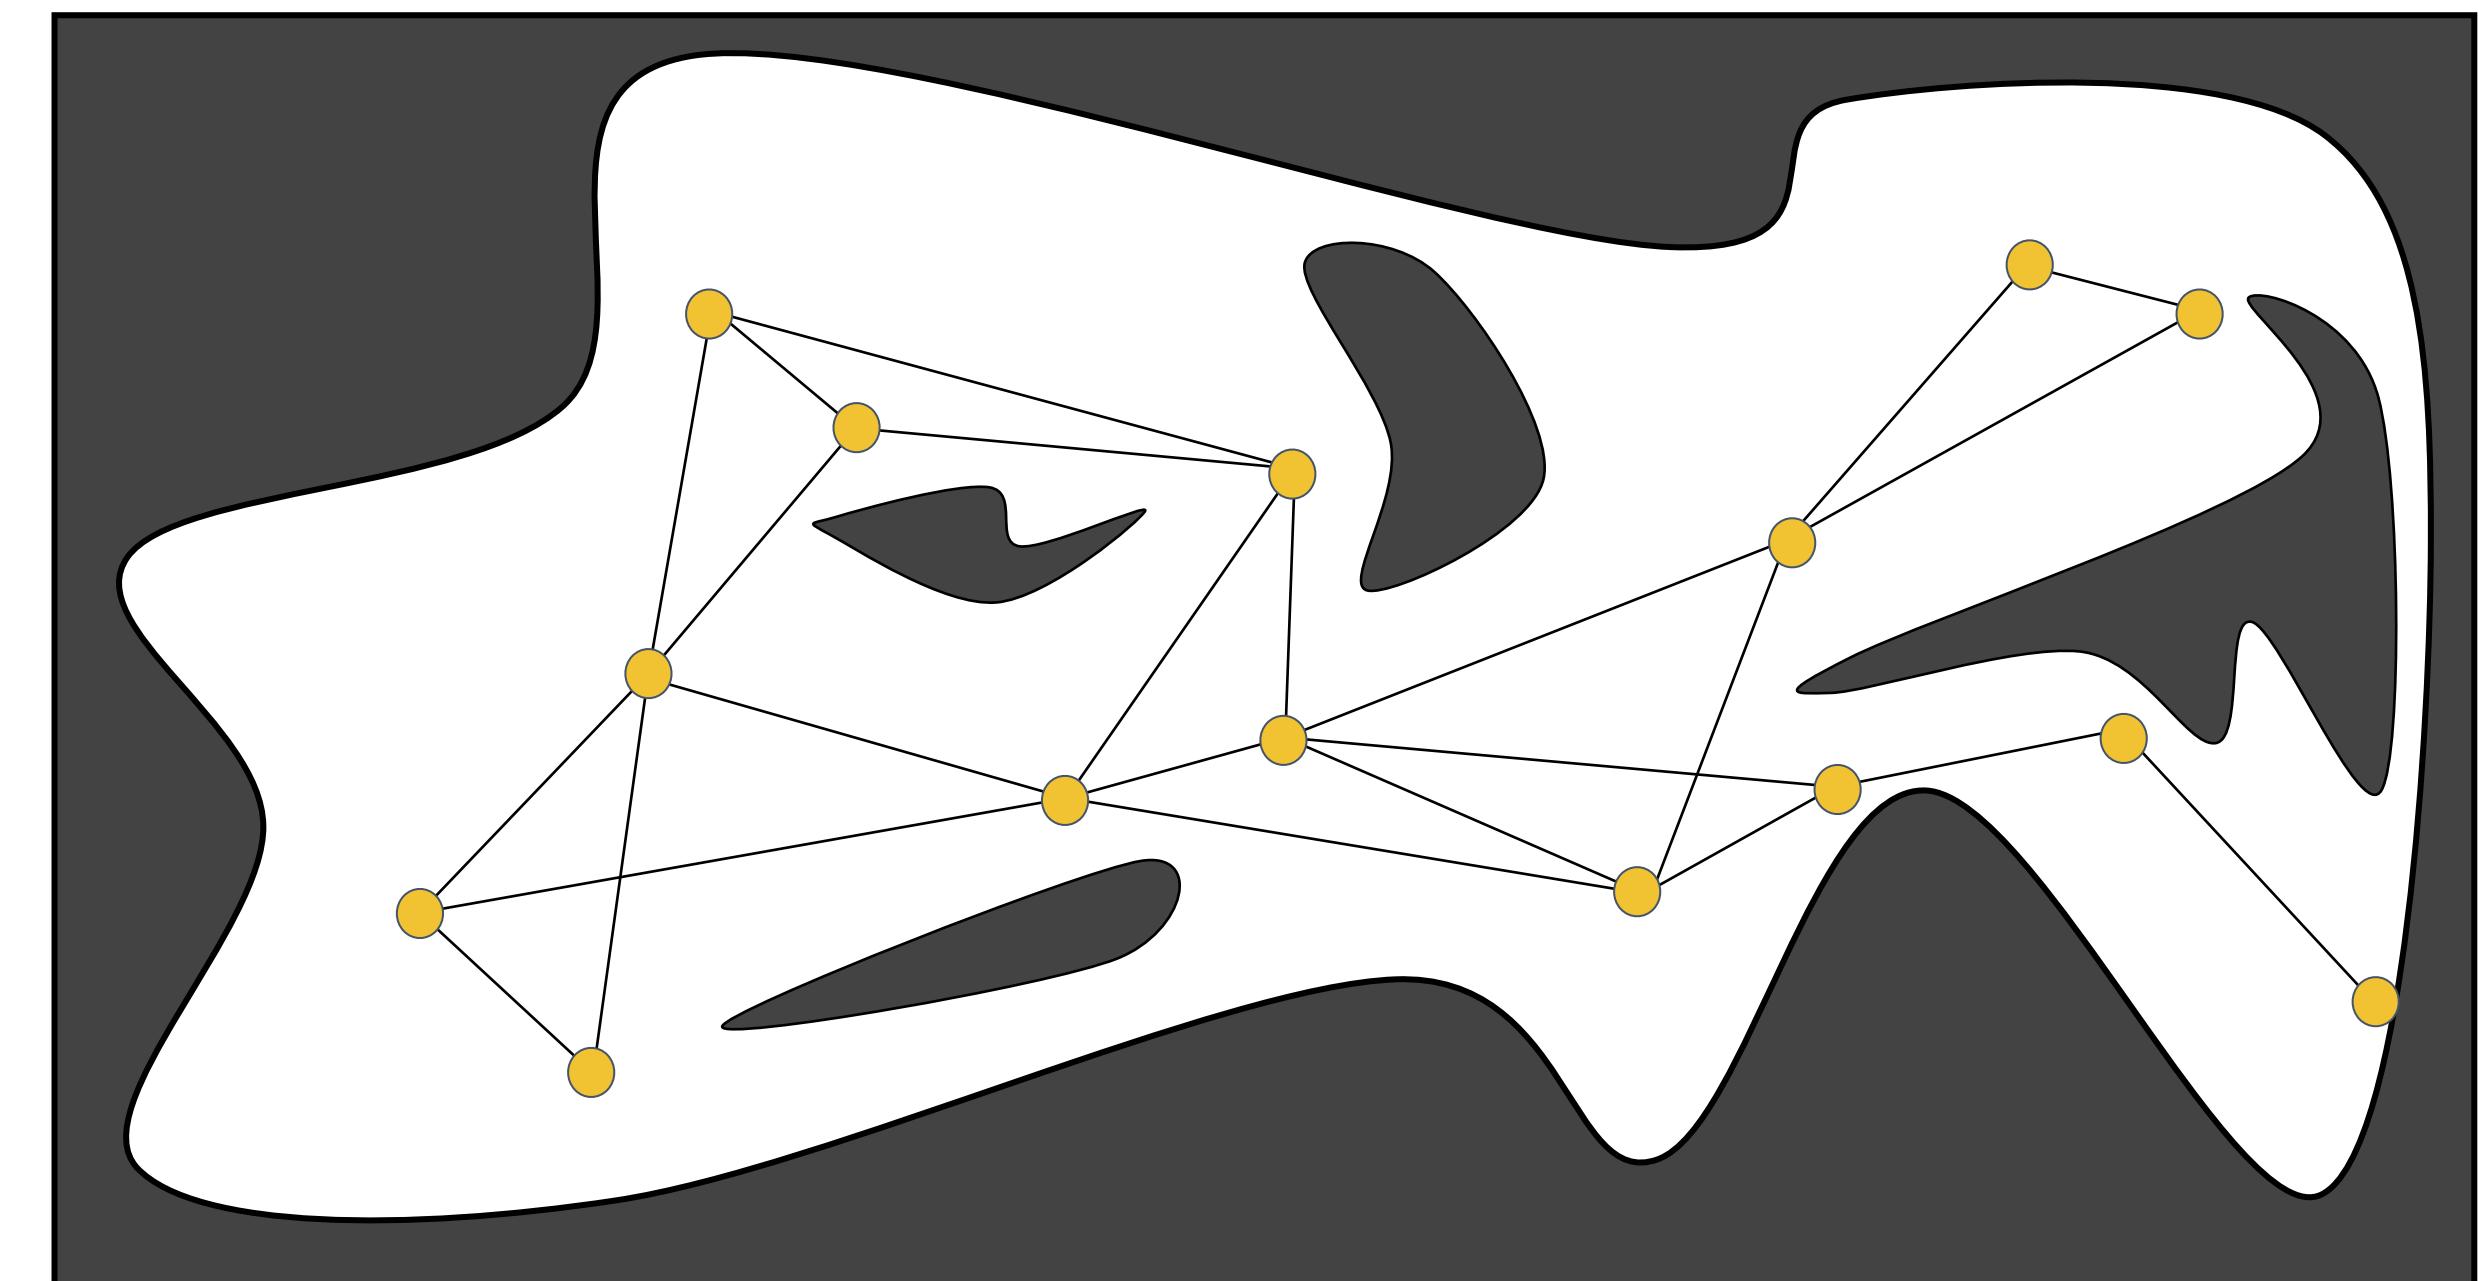
Probabilistic Roadmaps

- Configurations are sampled by picking coordinates at random
- Sampled configurations are tested for collision
- Each configuration is linked by straight paths to its nearest neighbors



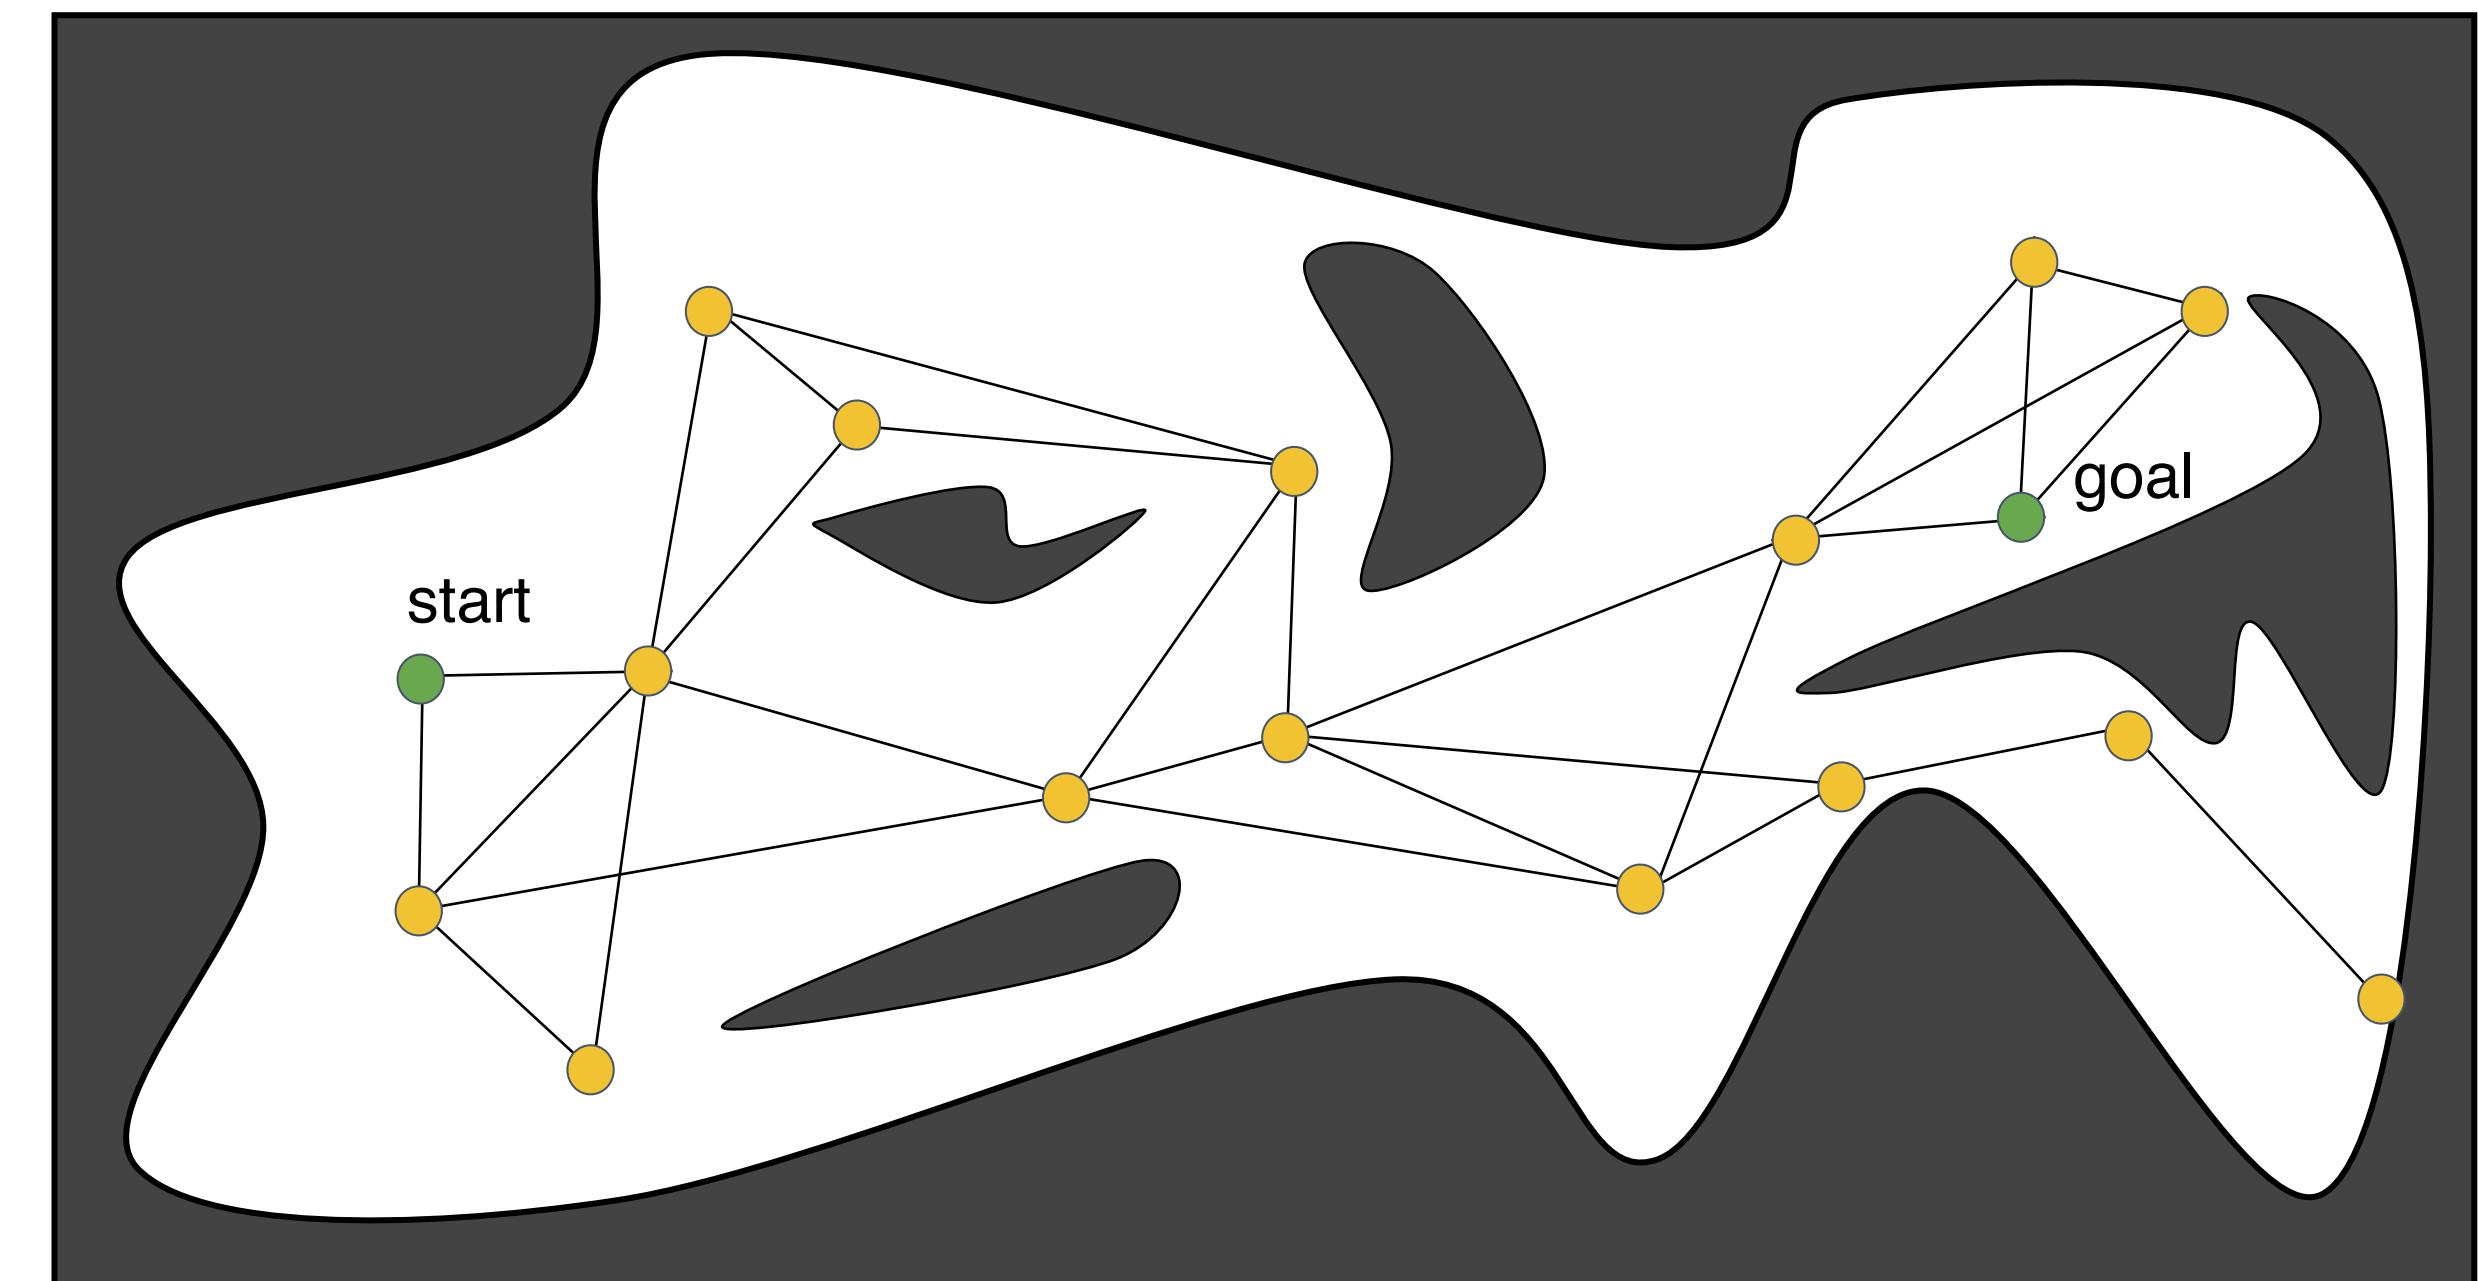
Probabilistic Roadmaps

- Configurations are sampled by picking coordinates at random
- Sampled configurations are tested for collision
- Each configuration is linked by straight paths to its nearest neighbors
- The collision-free links are retained as local paths to form the PRM



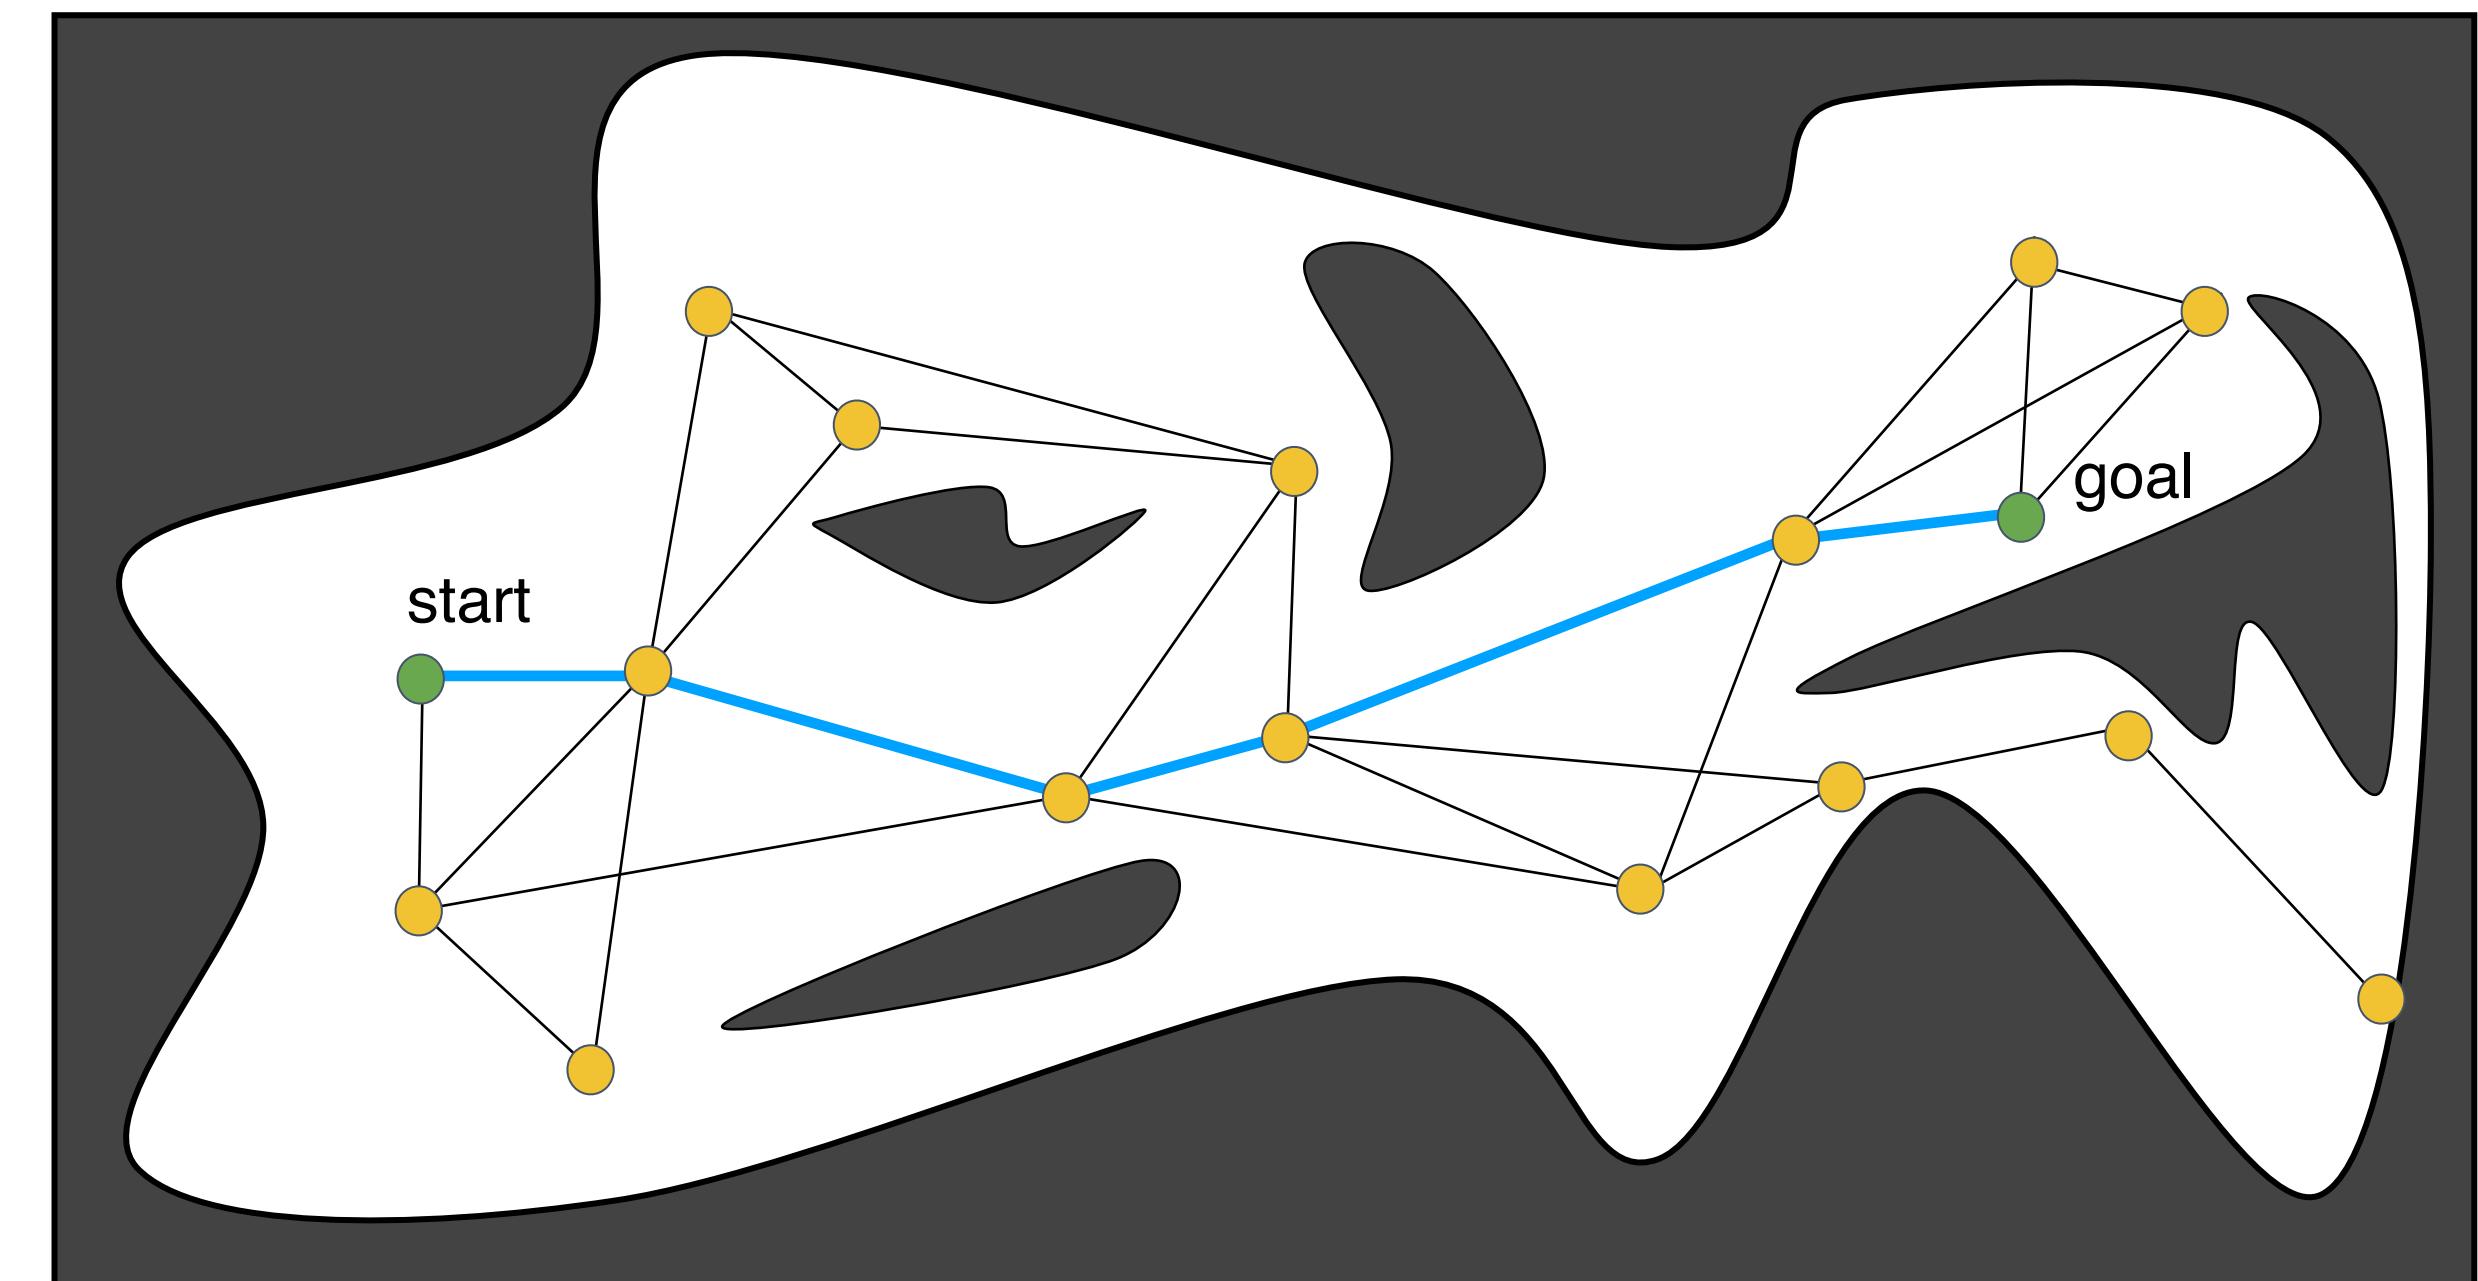
Probabilistic Roadmaps

- Configurations are sampled by picking coordinates at random
- Sampled configurations are tested for collision
- Each configuration is linked by straight paths to its nearest neighbors
- The collision-free links are retained as local paths to form the PRM
- The start and goal configurations are included as milestones



Probabilistic Roadmaps

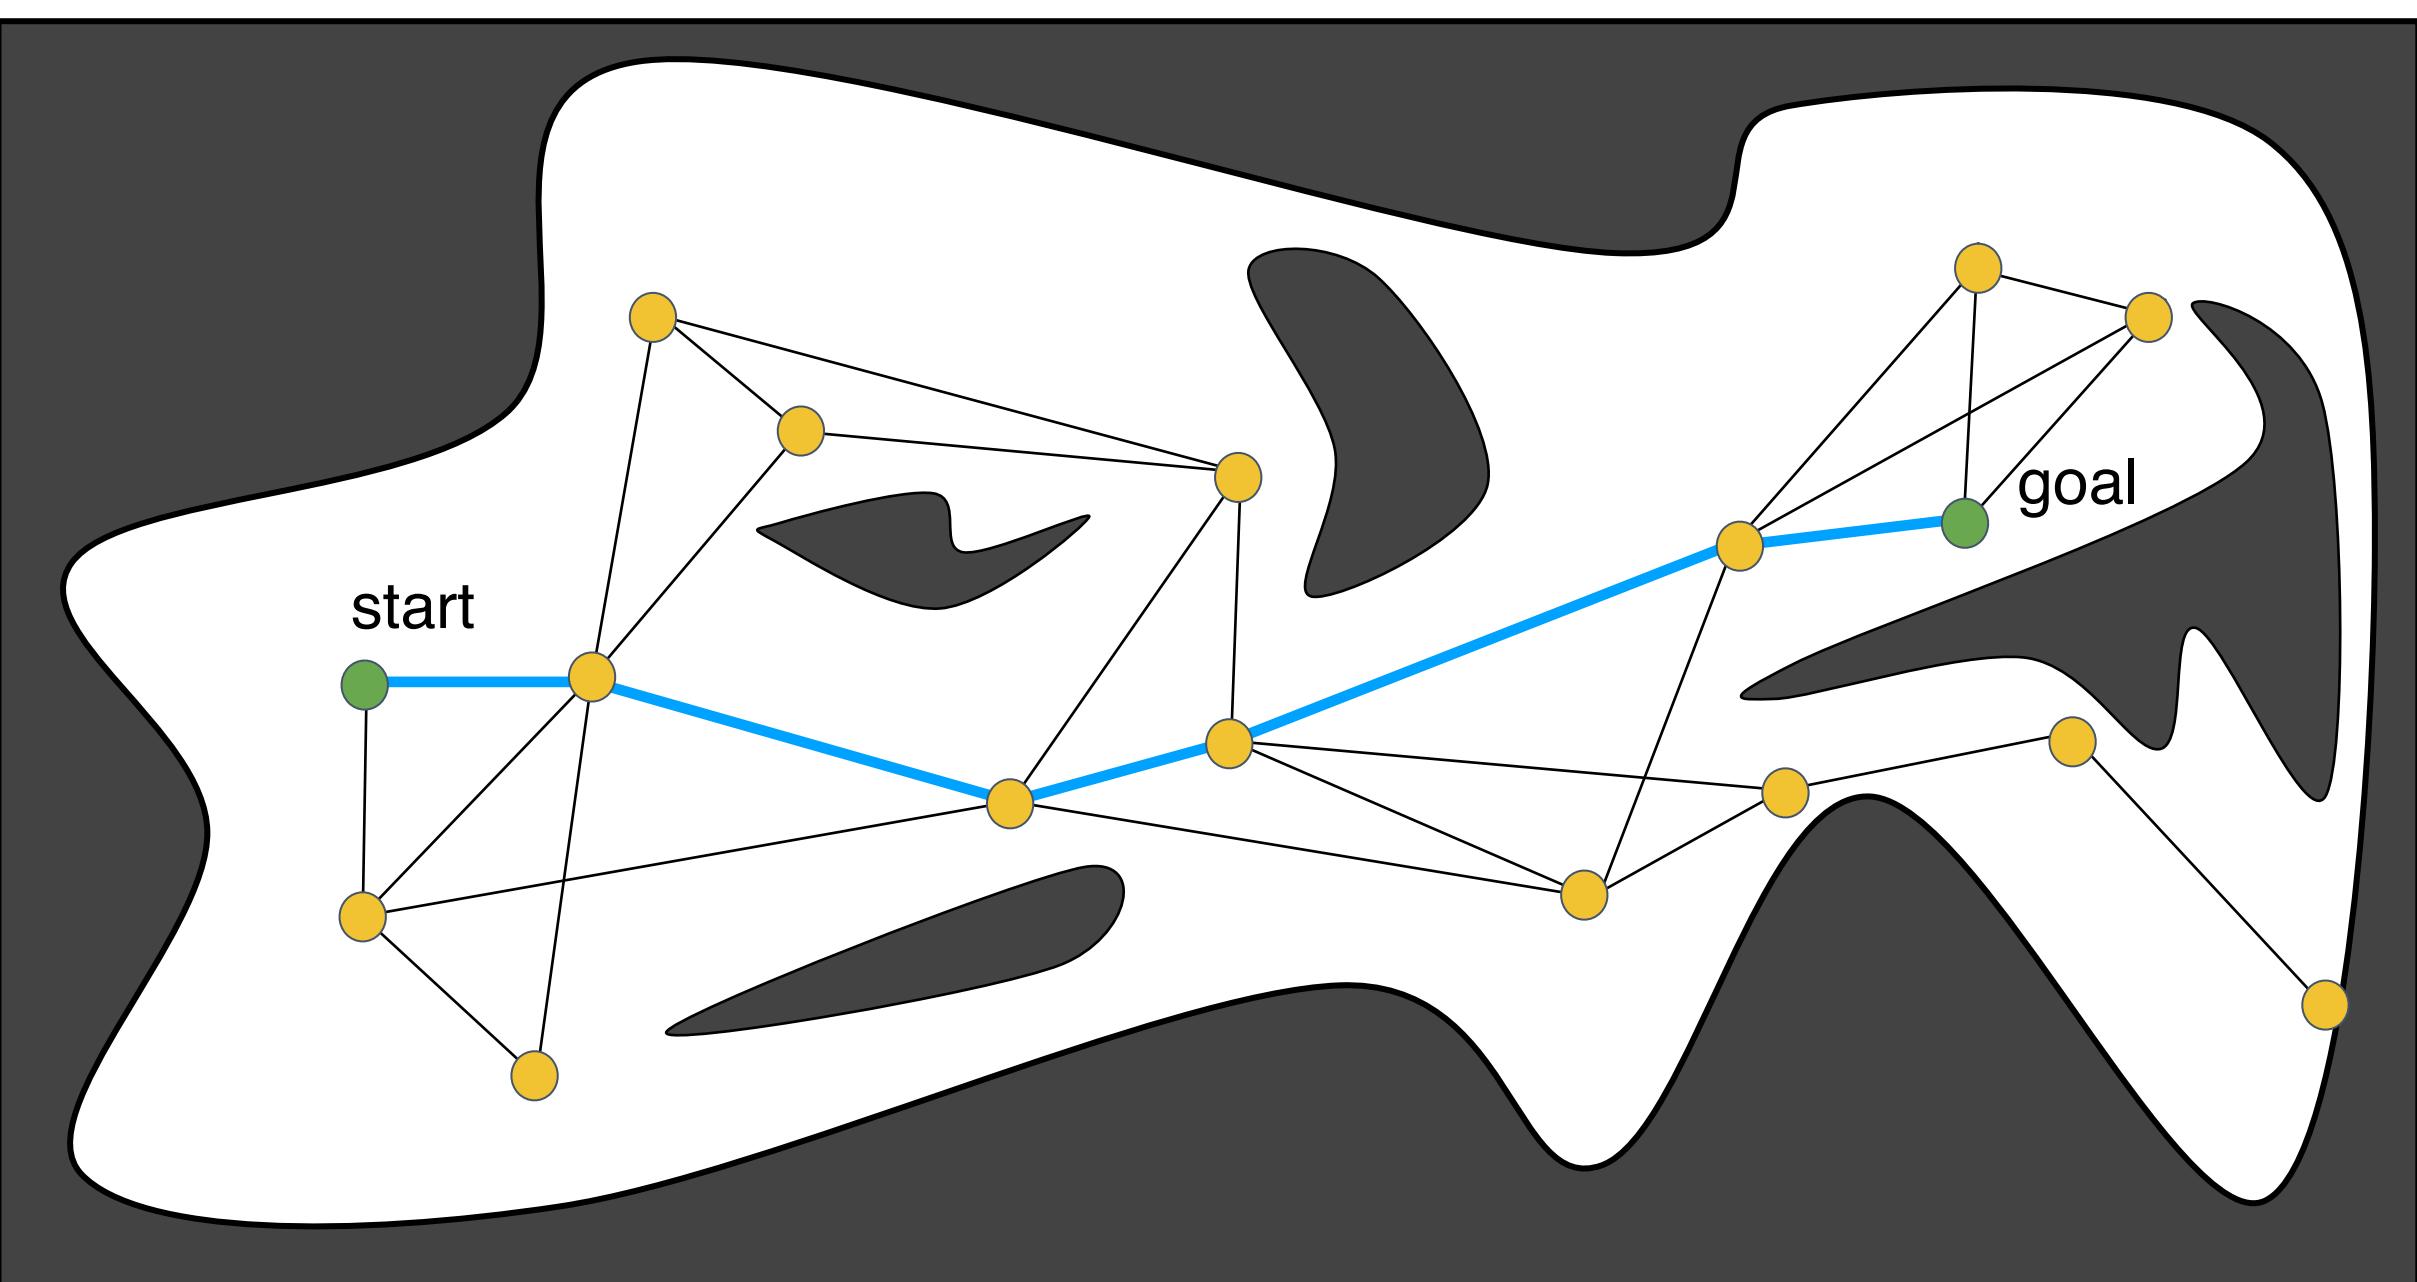
- Configurations are sampled by picking coordinates at random
- Sampled configurations are tested for collision
- Each configuration is linked by straight paths to its nearest neighbors
- The collision-free links are retained as local paths to form the PRM
- The start and goal configurations are included as milestones
- The PRM is searched for a path from start to goal



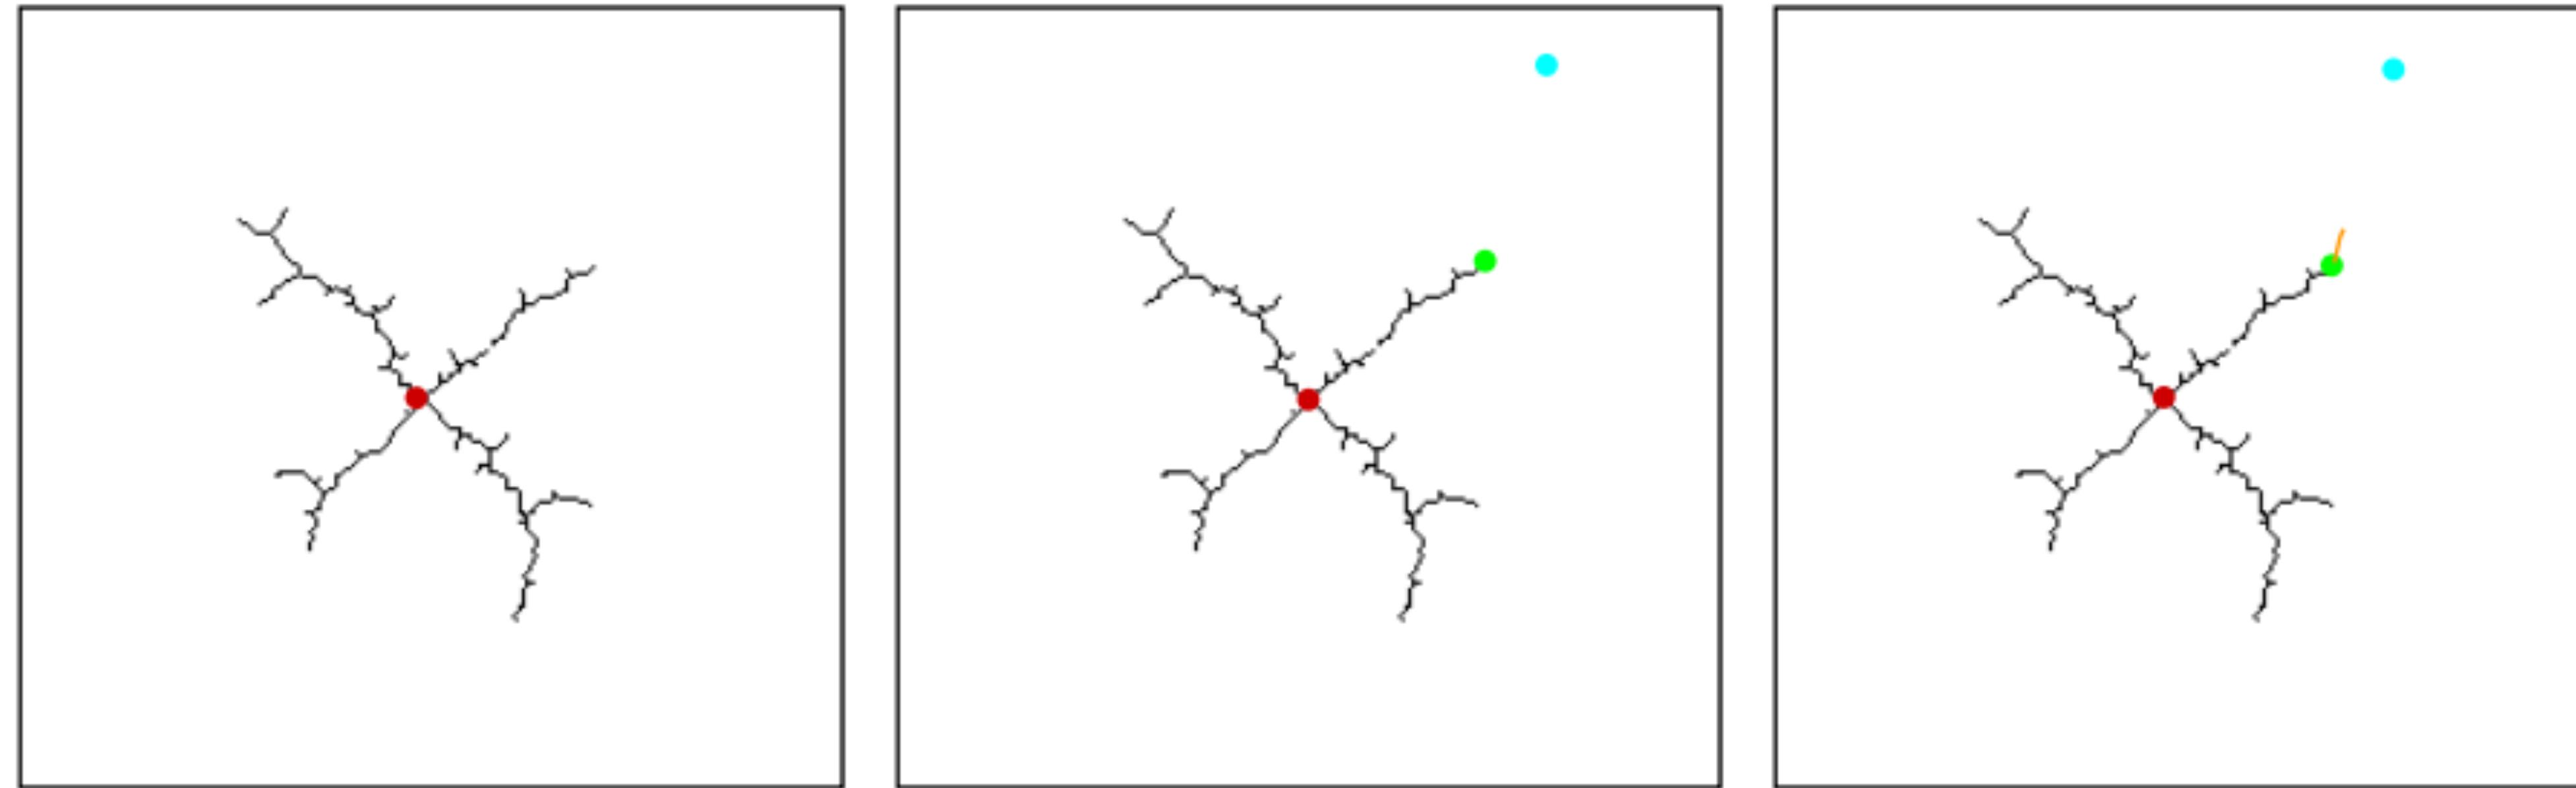
Probabilistic Roadmaps

Considerations

- Single query/ multi query
- How are nodes placed?
 - Uniform sampling strategies
 - Non-uniform sampling strategies
- How are local neighbors found?
- How is collision detection performed?
 - Dominates time consumption in PRMs



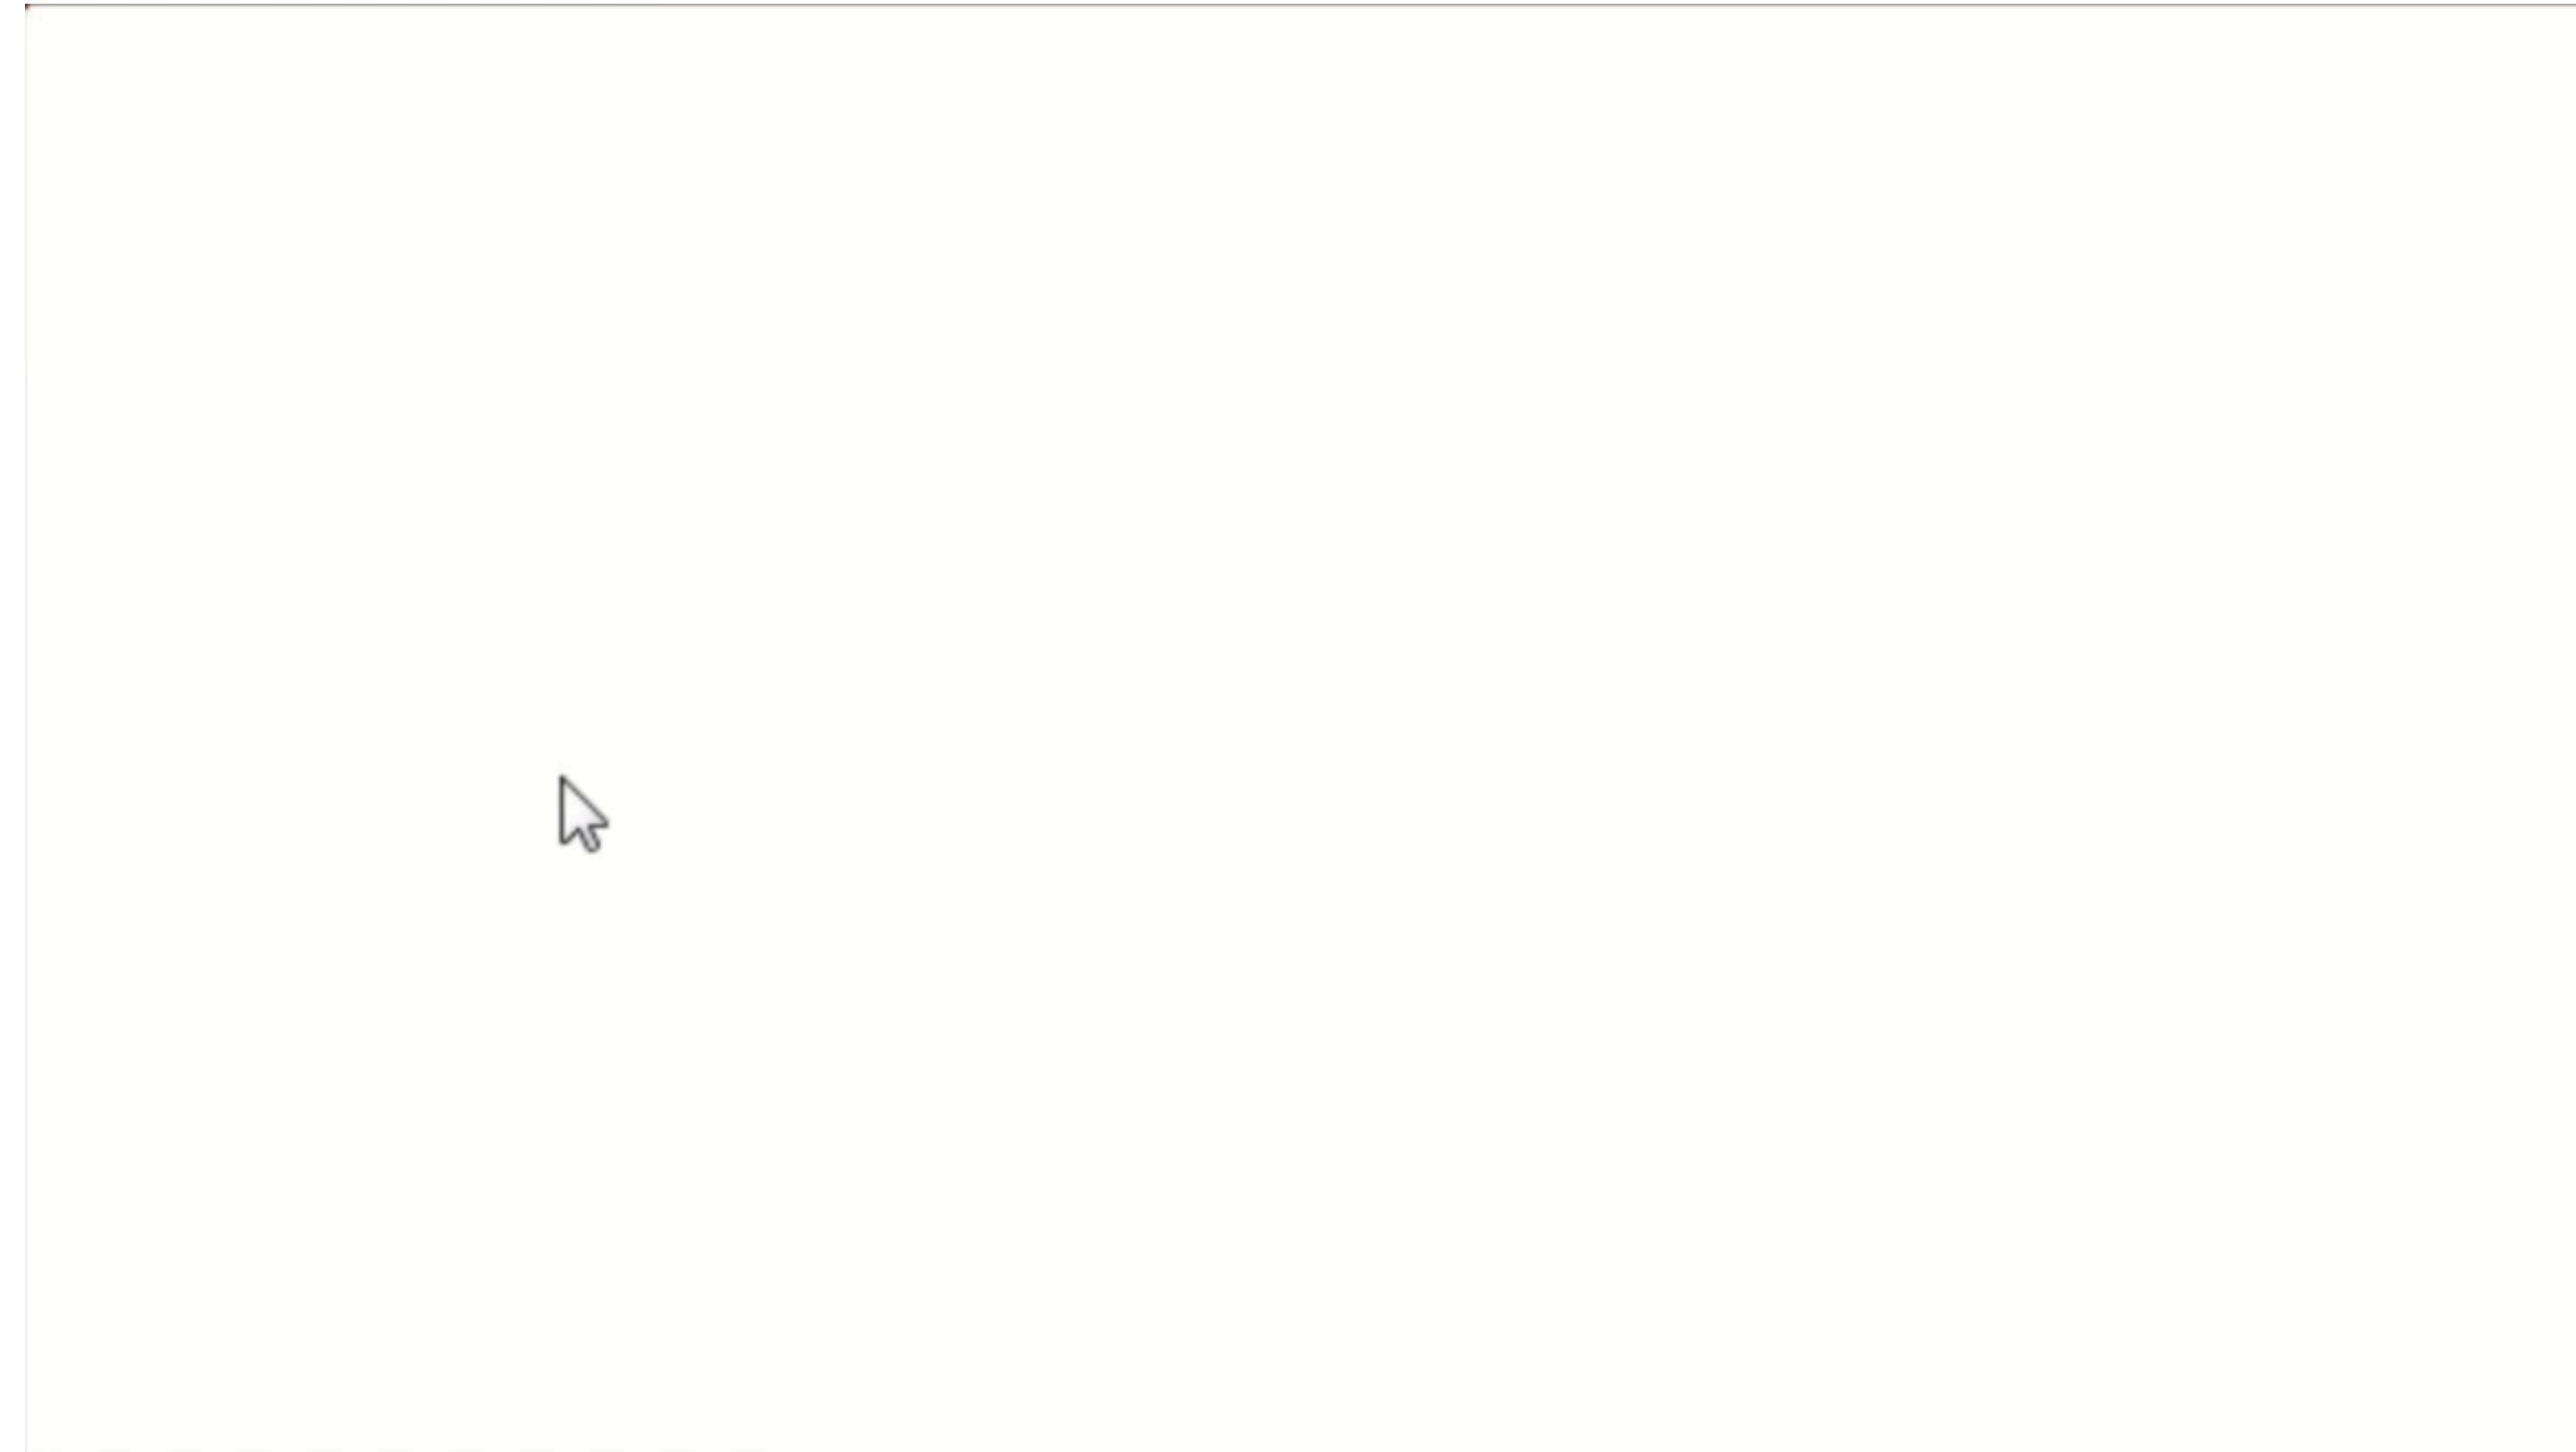
Rapidly Exploring Random Trees (RRT)



1. Maintain a tree rooted at the starting point ●
2. Choose a point at random from free space ●
3. Find the closest configuration already in the tree ●
4. Extend the tree in the direction of the new configuration /

Rapidly Exploring Random Trees (RRT)

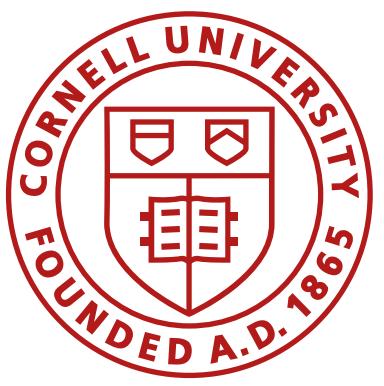
Uniform/ biased sampling



Rapidly Exploring Random Trees (RRT)

Considerations

- Sensitive to step-size (Δq)
 - Small: many nodes, closely spaced, slowing down nearest neighbor computation
 - Large: Increased risk of suboptimal plans / not finding a solution
- How are samples chosen?
 - Uniform sampling may need too many samples to find the goal
 - Biased sampling towards goal can ease this problem
- How are closest neighbors found?
- How are local paths generated?
- Variations
 - RRT Connect, A*-RRT, Informed RRT*, Real-Time RRT*, Theta*-RRT, etc.



**Next class... graph search...
see you Tuesday!**