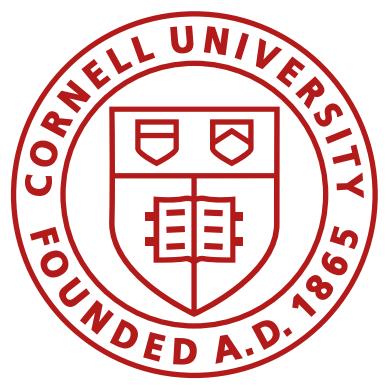


# **Particle Filters and SLAM**

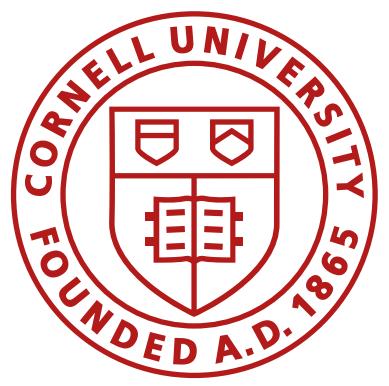
**Fast Robots, ECE4160/5160, MAE 4190/5190**

**E. Farrell Helbling, 4/15/25**




# Class Action Items

- Lab 10: localization (sim) starts today!
  - The lab is graded S/U (all of the code is already posted on previous websites)
  - If you do get an unsatisfactory, you will need to redo it.
- About to send a bunch of google forms your way
  - ECE Robotics Day availability (if you are unavailable it's because of another course conflict)
  - Lab 8 votes on stunts and bloopers
- This is the last technical content lecture of the course!




# Lab 10 Localization

- This lab is graded S/U
- Tasks
  - Read the full lab and the notebook (before you show up to lab)
  - Perform grid localization for the sample trajectory
  - Video demo
  - Discuss:
    - Control
    - Motion model and the prediction step
    - Sensor model and the update step
    - Choosing parameters, effect of changing parameters
    - Ways to mitigate computational load
    - Evaluate the Bayes Filter
    - Evaluate how well this will work on your robot



# Grid-based localization

- Simple
- ... but it is computationally expensive for large workspaces




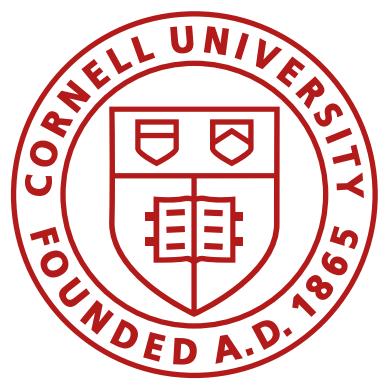
```

1. Algorithm Bayes_Filter ( $bel(x_{t-1})$ ,  $u_t$ ,  $z_t$ ) :
2.   for all  $x_t$  do
3.      $\overline{bel}(x_t) = \sum_{x_{t-1}} p(x_t | u_t, x_{t-1}) bel(x_{t-1})$ 
4.      $bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$ 
5.   end for
6.   return  $bel(x_t)$ 
  
```

# Monte Carlo Localization

- Non-parametric approach based on Particle Filters
- Model the distribution by samples
  - Prediction step
    - Draw from the samples
    - Move forward based on motion model
  - Update step
    - Weigh samples by their importance
    - Sensor model
  - Resample based on their weight
- The more samples we use, the better the estimate!

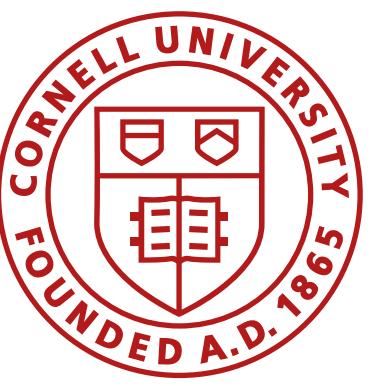




# Monte Carlo Localization



- Non-parametric approach based on Particle Filters
- Model the distribution by samples

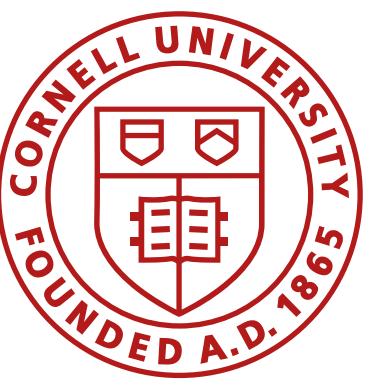
```
1. Algorithm Bayes_Filter ( $bel(x_{t-1})$ ,  $u_t$ ,  $z_t$ ) :  
2.   for all  $x_t$  do           Draw  $x_t^i$  from  $p(x_t | u_t, x_{t-1}^i)$   
3.      $\overline{bel}(x_t) = \sum_{x_{t-1}} p(x_t | u_t, x_{t-1}) bel(x_{t-1})$  ← Prior samples  
4.      $bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$   
5.   end for                Importance factor  $w_t^i \propto p(z_t | x_t)$   
6.   return  $bel(x_t)$ 
```



# Monte Carlo Localization



- How do you obtain samples from an arbitrary distribution?
  - Closed form solution for a uniform distribution
  - Closed form solution for a Gaussian distribution






# Monte Carlo Localization

- How do you obtain samples from an arbitrary distribution?
  - Closed form solution for a uniform distribution
  - Closed form solution for a Gaussian distribution
- Use a proposal distribution to generate samples from the target distribution
- Account for differences using a weight
  - $w = \text{target} / \text{proposal}$





# Monte Carlo Localization


- Each particle,  $j$ , is a pose hypothesis
- Proposal distribution from the motion model
  - $x_t^{[j]} \sim p(x_t | x_{t-1}, u_t)$
- Correction via the observation model
  - $w_t^{[j]} = \frac{\text{target}(x_t^{[j]})}{\text{proposal}(x_t^{[j]})} = p(z_t | x_t^{[j]})$
- Resample
  - Draw sample  $i$  with probably  $w_t^{[j]}$  and repeat  $J$  times


```

Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):
1:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
2:   for  $j = 1$  to  $J$  do
3:     sample  $x_t^{[j]} \sim p(x_t | u_t, x_{t-1}^{[j]})$ 
4:      $w_t^{[j]} = p(z_t | x_t^{[j]})$ 
5:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[j]}, w_t^{[j]} \rangle$ 
6:   endfor
7:   for  $j = 1$  to  $J$  do
8:     draw  $i \in 1, \dots, J$  with probability  $\propto w_t^{[i]}$ 
9:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
10:  endfor
11:  return  $\mathcal{X}_t$ 

```


# Monte Carlo Localization





# Monte Carlo Localization

- How would you deal with a kidnapped robot?
  - Randomly insert samples proportional to the average likelihood of the particles



- Pros
  - Works well for high-uncertainty scenarios
  - Much more efficient than grid cells
- Cons
  - Scales poorly with higher dimensional workspaces




# **Brief intro to SLAM**

# Related terms

- State estimation
- Localization **Inferring a location given a map**
- Mapping **Inferring a map given a location**
- SLAM **Learning a map and locating the robot simultaneously**
- Navigation
- Motion planning

# Related terms


- State estimation
- Localization
- Mapping
- **SLAM**
- Navigation
- Motion planning



# Given all we have learned...


# Given all we have learned...

- Markov localization in a grid
  - Localization: estimate your cell pose within the map
  - Mapping: estimate if cells are occupied or not
    - Every grid cell is a random variable
  - SLAM: estimate pose *and* if cells are occupied or not
    - 100x100 grid cells (pretty small map)
    - Localization:  $(x, y, \text{theta}) = 100 \times 100 \times 100$  states
    - Map:  $(x, y) = 10,000$  states
    - SLAM  $100 \times 100 \times 100 \times 10,000$  states
- Same issue for particle filters...
  - Balance parametric and non-parametric approaches




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?

- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated




# Why is SLAM hard?


- Robot pose/path and map are both unknown (not independent)
- Map and pose estimates are correlated
- Good data association is key




# Why is SLAM hard?


- The mapping between observations and the map is unknown
- Picking the wrong data association can cause map divergence







# Trilateration using range measurements





# Trilateration using range measurements




# Related terms

- State estimation
- Localization
- Mapping
- **SLAM**
- Navigation
- Motion planning

- Given
  - Control inputs  
 $U_{0:k} = \{u_1, u_2, \dots, u_k\}$
  - Relative observations  
 $Z = \{z_1, z_2, \dots, z_n\}$
- Compute
  - Map of the environment  
 $m = \{m_1, m_2, \dots, m_n\}$
  - Robot path (seq. of poses)  
 $X_{0:k} = \{x_0, x_1, \dots, x_k\}$

- Error in pose
- Error in observation
- Error in mapping
- Errors accumulate




Landmarks are considered motionless

# Simultaneous Localization and Mapping




## Graphical model

- Nodes are random variables
- Directed edges are variable dependencies
- Gray nodes: observed or directly measured variables
- White nodes: inferred latent variables




# SLAM Representations

- Grid maps or scans





[Lu & Milios, 97; Gutmann, 98; Thrun 98; Burgard, 99; Konolige & Gutmann, 00; Thrun, 00; Arras, 99; Haehnel, 01;...]

- Landmark-based




[Leonard et al., 98; Castelanos et al., 99; Dissanayake et al., 2001; Montemerlo et al., 2002;...]



# Simultaneous Localization and Mapping


- Nodes are random variables
- Directed edges are variable dependencies
- Gray nodes: observed or directly measured variables
- White nodes: inferred latent variables
- Full SLAM: compute a joint posterior over the whole path of the robot and the map



$$p(x_{1:t}, m | z_{1:t}, u_{1:t}, x_0)$$

# Simultaneous Localization and Mapping

- Nodes are random variables
- Directed edges are variable dependencies
- Gray nodes: observed or directly measured variables
- White nodes: inferred latent variables
- Full SLAM: compute a joint posterior over the whole path of the robot and the map
- Online SLAM: compute a posterior over the current pose along with the map



$$p(x_t, m | z_{1:t}, u_{1:t})$$

# Simultaneous Localization and Mapping

- Prediction (prediction step):

- $p(x_t, m | z_{0:t}, u_{1:t}, x_0) = \sum_{t-1} P(x_t | x_{t-1}, u_{1:t})P(x_{t-1}, m | Z_{0:t-1}, U_{1:t}, x_0)$

- Correction (update step):

- $p(x_t, m | z_{0:t-1}, u_{0:t}, x_0) = \eta P(z_t | x_t, m)P(x_t, m | Z_{0:t}, U_{1:t}, x_0)$

- We can solve the localization problem with the assumption that we know the map





- $P(x_t | Z_{0:t}, U_{0:t}, m)$

- We can solve the mapping problem with the assumption that we know the location

- $P(m | X_{0:t}, Z_{0:t}, U_{0:t})$

# Interesting Features


- Robot observations of the **relative landmark locations can be considered nearly independent**, because the relative landmark locations are independent from the robot's coordinate frame
- Robot observations of the **absolute landmark locations are less certain**, because the absolute landmark location is strongly related to the robot's coordinate frame
- Because landmarks are correlated even unobserved landmarks can be updated, such that correlations are increased for every observation we make
- The accuracy of the relative map increases for more observations



# Simultaneous Localization and Mapping

## Why is it hard?

- Map size
  - The larger the environment relative to the robot's perceptual range, the more difficult it is to acquire the map
- Perceptual ambiguity
  - The more different places look alike, the more difficult it is to establish correspondence between different locations traversed at different points in time
- Cycles
  - Motion-cycles are particularly difficult to map




# SLAM solutions

- The trick is to find an appropriate representation for the observation and the motion problem
    - Graph SLAM      Global optimization: outputs the most likely map and trajectory
    - EKF SLAM
    - Fast SLAM
- Probability distribution over landmarks and the most recent pose (online SLAM)

# **Graph SLAM**


# Graph SLAM

- Graph represents a set of objects where pairs of objects are connected by links encoding relations between them
- Create an edge if...
  - ... the robot moves from  $x_i$  to  $x_{i+1}$ 
    - edge corresponds to odometry measurement
  - ... the robot observes the same part of the environment from  $x_i$  and from  $x_j$
- Edges represent constraints
- Nodes represent the state (poses and landmarks)
  - Given a state, we can compute predicted observations
  - Find a configuration of the nodes so that the real and predicted constraints are as similar as possible
  - Minimize the Least Square Error over all constraints



# Graph SLAM


- Treat constraints (generated by motions and observations) as elastic springs
- Minimize the energy in all the springs
- Any modern SLAM implementation has some version of this
  - Pro: globally optimal
  - Con: BIG optimization problem, only one output
- Tricks
  - Combine poses over many time steps into single nodes to make the graph smaller
  - If you see the same landmark from several poses, you can get rid of the pose and add a stronger constraint between those landmarks



# **EKF SLAM**

# EKF SLAM

- Goal: Estimate  $p(x_k, m \mid u_{1:k}, z_{1:N})$
- Assume all noise is Gaussian
- Track a Gaussian belief of the current state and landmarks
- Apply the Kalman Filter...



# Kalman Filter

Kalman Filter ( $\mu(t - 1)$ ,  $\Sigma(t - 1)$ ,  $u(t)$ ,  $z(t)$ )

$$1. \mu_p(t) = A\mu(t - 1) + Bu(t)$$

$$2. \Sigma_p(t) = A\Sigma(t - 1)A^T + \Sigma_u$$

$$3. K_{KF} = \Sigma_p(t)C^T(C\Sigma_p(t)C^T + \Sigma_z)^{-1}$$


$$4. \mu(t) = \mu_p(t) + K_{KF}(z(t) - C\mu_p(t))$$

$$5. \Sigma(t) = (I - K_{KF}C)\Sigma_p(t)$$


6. Return  $\mu(t)$  and  $\Sigma(t)$

**prediction**

**update**



State estimate:  $\mu(t)$   
 State uncertainty:  $\Sigma(t)$   
 Process noise:  $\Sigma_u$   
 Kalman filter gain:  $K_{KF}$   
 Measurement noise:  $\Sigma_z$



# EKF SLAM


- Goal: Estimate  $p(x_k, m \mid u_{1:k}, z_{1:N})$
- Assume all noise is Gaussian
- Track a Gaussian belief of the current state and landmarks
- Linearize around every state and run the Kalman Filter

$$x = \begin{bmatrix} \bar{\mathcal{O}} \\ \mathcal{L}_1 \\ \vdots \\ \mathcal{L}_n \end{bmatrix} \quad P = \begin{bmatrix} P_{\mathcal{O}\mathcal{O}} & P_{\mathcal{O}\mathcal{M}} \\ P_{\mathcal{M}\mathcal{O}} & P_{\mathcal{M}\mathcal{M}} \end{bmatrix} = \begin{bmatrix} P_{\mathcal{O}\mathcal{O}} & P_{\mathcal{O}\mathcal{L}_1} & \cdots & P_{\mathcal{O}\mathcal{L}_n} \\ P_{\mathcal{L}_1\mathcal{O}} & \boxed{P_{\mathcal{L}_1\mathcal{L}_1} & \cdots & P_{\mathcal{L}_1\mathcal{L}_n}} \\ \vdots & \ddots & \ddots & \vdots \\ P_{\mathcal{L}_n\mathcal{O}} & P_{\mathcal{L}_n\mathcal{L}_1} & \cdots & P_{\mathcal{L}_n\mathcal{L}_n} \end{bmatrix}$$

- Landmark matrix grows, making the inversion step costly!
- In Full SLAM the trajectory matrix grows even faster

# EKF SLAM

- Goal: Estimate  $p(x_k, m \mid u_{1:k}, z_{1:N})$
- Assume all noise is Gaussian
- Track a Gaussian belief of the current state and landmarks
- Linearize around every state and run the Kalman Filter
- Pros
  - Super easy, well understood, runs online
  - Works well for low-uncertainty problems
- Cons
  - Works poorly for high-uncertainty problems
  - States must be well-approximated by Gaussians



# Fast SLAM

# Fast SLAM

- Half sample-based solution
  - Particle filter
    - Every particle has its own version of the map with a given trajectory
- Half analytical solution
  - Landmark-based
    - Each pose and map of independent features is updated analytically through EKF
  - Grid-based map
    - Occupancy of each grid cell is estimated by Bayes Filter

*Victoria Park dataset*  
University of Sydney

**GPS**  
**FastSLAM**  
4km traverse  
100 particles  
<5m RMS position error




# Fast SLAM

- Key idea: factorize the posterior
  - $p(x_{1:k}, m | z_{1:k}) = p(m | x_{1:k}, z_{1:k})p(x_{1:k} | z_{1:k})$
- $p(x_{1:k} | z_{1:k})$ : pose estimation is approximated by the Particle Filter (can represent multiple hypotheses)
- $p(m | x_{1:k}, z_{1:k})$ : classic mapping problem, approx using EKF (efficient at representing belief in high dimensions)
- Outcome is a Marginalized Particle Filter (MPF)
  - Each particle is a pose trajectory with an attached map corresponding to mean and covariance of each landmark

*Victoria Park dataset*  
University of Sydney

GPS  
**FastSLAM**  
4km traverse  
100 particles  
<5m RMS position error



# Fast SLAM

- Distribution is estimated by a fixed number of particles
  - Each particle,  $k$ , contains an estimate of robot path and the mean and covariance of each of the  $n$  features
  - $P^{[k]}(x_t^{[k]}; \mu^{[k]}, \Sigma_1^{[k]}, \dots, \mu^{[k]}, \Sigma_n^{[k]})$
- **Step 1:** Update particle trajectory (motion model)
- **Step 2:** Update particle landmarks with EKF (sensor model)
  - Linearize the observation model at  $(x_t^{[k]}, m)$
  - Only update associated landmarks
- **Step 3:** Update weights based on  $p(z_t | x_t^{[k]}, m^{[k]})$
- **Step 4:** Resample distribution


*Victoria Park dataset*  
University of Sydney

GPS  
**FastSLAM**  
4km traverse  
100 particles  
<5m RMS position error



# **SLAM State of the Art**

# State of the Art in SLAM



# State of the Art in SLAM



# State of the Art in SLAM

- Robotics
  - 3D cameras with depth maps and high frame rates and resolution
  - Dense 3D models of the world
  - Uses ROS and deep learning to recognize features
  - Come built-in in a range of robots
  - Inherent to e.g. the RealSense tracking cameras
- 3D scanning/ reconstruction
- Virtual and augmented reality



# State of the Art in SLAM

